

When Is Label-Free Adaptation Knowable?

Helpful, Harmful, and Unknowable Regimes Under Distribution Shift

Pratik Niroula

Independent Researcher

Working draft v0.2

Abstract

Test-time adaptation can improve models on unlabeled target data, but the same objective that helps under one distribution shift can silently hurt under another. Existing work mainly strengthens adaptation mechanisms; we study a prior decision problem: can a system decide, without target labels, whether it should adapt at all?

We formalize adaptation benefit as the excess target risk of a frozen baseline over an adapted candidate and define three label-free regimes: *knowably helpful*, *knowably harmful*, and *unknowable*. We prove that when two target worlds induce identical observable evidence but opposite adaptation benefit, no label-free rule can be correct in both; abstention is then information-theoretically necessary. Under a risk-alignment assumption, we give a finite-sample adapt/freeze/abstain certificate controlling false adaptation and false freezing. Our main theoretical result is an exact benefit-sign frontier: adaptation benefit is identifiable from observables iff an observable margin exceeds the calibration-drift budget on the disagreement region.

We instantiate the framework as KGA, a mechanism-agnostic wrapper around Tent, EATA, and SAR. Across anomaly routing, CIFAR-10, CIFAR-10-C stress tests, and natural-shift audits, KGA follows theory: adapt in helpful regimes, freeze in harmful regimes, and abstain when evidence is insufficient. On a CIFAR-10-C stress grid, KGA beats always-adapt and always-freeze for Tent and EATA with zero false-adapt decisions, while tying SAR. At ImageNet scale, it also beats both policies on a harmful SAR stream. On two natural-shift panels (ImageNet-R across ten diverse backbones and Camelyon17 across Tent/EATA/SAR), evaluated over three seeds, KGA reliably beats always-freeze (all Holm-corrected) and ties always-adapt; it does not beat *both* trivial policies, because each panel is regime-homogeneous so one trivial policy is already near-optimal. Genuine beats-both wins stay confined to mixed regimes, exactly as pre-registered. K-Bound is not universal; value is highest when harmful adaptation is frequent, detectable, and costly.

Scope. We report both positive and negative results and do not claim universal improvement; all numbers come from code in `kbound_paper/scripts/`. A detailed proof/experiment status note appears at the end of §1.

1 Introduction

Deployed models rarely encounter the same distribution they were trained on. Test-time adaptation (TTA) addresses this mismatch by updating a model on unlabeled target data, often through entropy minimization, sample filtering, or stability-aware updates. In favorable shifts, such updates can recover lost accuracy. In unfavorable shifts, however, the same unlabeled objective can silently degrade performance: the model may sharpen wrong predictions, adapt to label imbalance, or

collapse under non-stationary streams. Because target labels are unavailable at deployment, the system often cannot tell whether adaptation is helping or hurting.

Most work on TTA asks how to design a better adaptation mechanism. This paper asks a prior question: should the system adapt at all? We study label-free adaptation as a decision problem. Given a frozen baseline, an adapted candidate, and an unlabeled target batch, the system must choose one of three actions: adapt when adaptation is knowably helpful, freeze when adaptation is knowably harmful, and abstain when the available evidence cannot identify the sign of the benefit.

This decision is constrained by identifiability. We show that there exist target worlds with identical label-free evidence but opposite adaptation benefit. In such cases, no label-free rule can certify the correct decision in both worlds; abstention is not merely conservative, but information-theoretically necessary. Conversely, when the label-free evidence is risk-aligned and the benefit estimate admits a valid uncertainty radius, a finite-sample certificate can safely route between adapt, freeze, and abstain while controlling false-adapt and false-freeze errors.

Our main theoretical result is an exact benefit-sign frontier. The sign of the adaptation benefit is identifiable from label-free observables precisely when an observable margin on the disagreement region exceeds a calibration-drift budget. This characterizes why common label-free heuristics can work in benign regimes and fail under drift: they implicitly assume that this drift is zero or negligible. We further show that no label-free-checkable assumption can remove this obstruction in full generality; the remaining ambiguity is an irreducible one-bit orientation choice, which can only be supplied through structural assumptions that are falsifiable but not verifiable.

We instantiate the framework as KGA, a mechanism-agnostic wrapper around existing adaptation candidates such as Tent, EATA, and SAR. KGA does not introduce a new adaptation objective. Instead, it estimates the label-free benefit of a candidate, attaches a calibrated uncertainty radius, and returns adapt, freeze, or abstain. This makes KGA a safety layer for TTA rather than a replacement for TTA methods.

Empirically, KGA behaves according to the theory. On a pre-registered CIFAR-10-C stress grid, where harmful adaptation is frequent and detectable, KGA beats both always-adapt and always-freeze for Tent and EATA while making zero false-adapt decisions, and it ties the collapse-resistant SAR. In online non-stationary CIFAR-10 streams, always-adapt collapses in the harm-dominated regime and always-freeze forgoes gains in the helpful regime, while KGA remains near-best across both. At ImageNet scale, KGA also improves over both trivial policies on a harmful SAR stream. On two multi-seed natural-shift panels (ImageNet-R across ten diverse backbones, Camelyon17 across Tent/EATA/SAR), KGA reliably beats always-freeze and ties always-adapt—it safely captures the available benefit and avoids harm—but does not beat *both* trivial policies, because each panel is regime-homogeneous (one trivial policy is already near-optimal). This is the expected behavior of the certificate outside mixed regimes, not a failure.

Contributions.

- We formalize label-free TTA as an adapt/freeze/abstain decision problem governed by the sign of adaptation benefit.
- We prove an impossibility result showing that some adaptation decisions are fundamentally unknowable from label-free evidence.
- We derive a finite-sample certificate that controls false adaptation and false freezing under risk-aligned evidence.
- We characterize an exact benefit-sign frontier based on observable disagreement margin and calibration drift.

- We validate the framework across anomaly routing, controlled mixed regimes, online TTA, CIFAR-10-C stress grids, ImageNet-scale corruptions, and natural-shift audits.

K-Bound does not claim universal improvement. When adaptation is already safe, a certificate can only tie or slightly trail always-adapt. When freezing is already optimal, it should match freeze rather than force an unnecessary update. Its distinctive value appears when harmful adaptation is frequent, detectable, and costly—the regime where blind adaptation is most dangerous.

2 Related work and positioning

This question is *not* untouched, and we state the contribution as a delta over adjacent lines. **TTA methods** (Tent, SHOT, TTT, EATA, SAR) show adaptation helps under some shifts and build robustness, but give no *pre-adaptation* certificate separating helpful/harmful/unknowable and no abstain-on-the-decision state. **Protected/guarded TTA** (online entropy matching, betting) guards one direction during adaptation, not a three-way decision with an explicit unknowable region. **Label-free accuracy estimation** (ATC, Agreement-on-the-Line, AETTA) estimates target *performance* or detects TTA failure, but estimating performance is not certifying the *sign of adaptation benefit* with a false-adapt bound. Tempora (Sreeram et al., 2026; arXiv:2602.06136) adds a deployment-level warning: offline method rankings are unstable under temporal pressure, with no fixed robust policy. K-Bound complements that by certifying when to adapt, freeze, or abstain for each condition. **Unsupervised risk estimation** (Steinhardt & Liang, 2016; Ben-David et al., 2010; Lipton et al., 2018) shows risk is estimable without labels only under structural/conditional-independence assumptions and impossible in general; we target the sign of a *difference* of risks and the decision-theoretic trichotomy. **Selective prediction** (Geifman & El-Yaniv) abstains on predictions, not on the adaptation decision. The surviving gap: a single adapt/freeze/abstain certificate with an explicit, proven unknowable region and a false-adapt guarantee. Table 1 makes the gap precise.

Work	label-free?	adapts?	predicts benefit <i>before</i> adapting?	abstains on the decision?	false-adapt guarantee?
Tent	yes	yes	no	no	no
SHOT	yes	yes	no	no	no
EATA	yes	yes	partial (filtering)	no	no
SAR	yes	yes	stability only	no	no
AETTA	yes	monitors	post-hoc estimate	—	no
KGA (ours)	yes	wrapper	yes	yes	yes

Table 1: Positioning. Prior TTA adapts and (some) detect failure; none gives a *pre-adaptation* adapt/freeze/abstain certificate with a false-adapt bound.

3 Problem setup

Source $P_S(X, Y)$ has labels at train/validation; target $P_T(X, Y)$ exposes only $P_T(X)$ at test time. Fix a loss ℓ , a frozen model f_0 , and an adapted/gated candidate f_a . Target risk $R_T(f) = \mathbb{E}_{P_T}[\ell(f(X), Y)]$. The *adaptation benefit* is

$$\Delta = R_T(f_0) - R_T(f_a), \quad \Delta > 0 \iff \text{adapting helps.} \quad (1)$$

Observable evidence $Z = \phi(X_{1:m}, f_0, f_a)$ is any statistic computable *without* target labels (entropy, confidence, score-distribution drift, detector disagreement, predicted-class balance, update norm, density-ratio diagnostics). The conditional benefit is $\Delta(z) = \mathbb{E}[\ell(f_0(X), Y) - \ell(f_a(X), Y) \mid Z = z]$. The central question is when $\text{sign } \Delta(z)$ is recoverable from Z alone.

4 Definitions

Definition 1 (Observable risk alignment). Z is *risk-aligned* if it identifies or bounds $\text{sign } \Delta(z)$ (it need not reveal the label, only whether adaptation helps).

Definition 2 (Regimes). At evidence z : *knowably helpful* if Z certifies $\Delta(z) > 0$ (**adapt**); *knowably harmful* if Z certifies $\Delta(z) < 0$ (**freeze**); *unknowable* if the same Z is compatible with both signs (**abstain**).

5 Theory

We organize the core argument into four pillars rather than a theorem inventory. **Pillar I** establishes the impossibility floor and quantitative committal limits. **Pillar II** gives the finite-sample certificate and forced-abstention logic. **Pillar III** characterizes identifiable positive regimes and disagreement-region sign structure. **Pillar IV** captures rates and sequential companions in appendix form. Legacy proposition stacks (including the prior Props 6/9/11/13 family) are treated as corollaries of the common audit-and-certificate framework.

5.1 Impossibility and the unknowable regime

Theorem 1 (Non-identifiability; with witness). *Let $g(Z) \in \{\text{ADAPT}, \text{FREEZE}, \text{ABSTAIN}\}$ be any (possibly randomized) rule depending on label-free evidence Z . There exist two target worlds P_T^1, P_T^2 with (i) $\text{Law}(Z \mid P_T^1) = \text{Law}(Z \mid P_T^2)$ and (ii) $\text{sign } \Delta^1 = -\text{sign } \Delta^2 \neq 0$. For any such pair, $\text{Law}(g)$ is identical across the two worlds; hence no label-free rule commits to the benefit-maximizing action in both, and the action minimizing the worst-case committal regret over $\{P_T^1, P_T^2\}$ is ABSTAIN.*

Proof. Witness (non-vacuous). Let $X \sim \mathcal{N}(0, 1)$ under both worlds, so $P_T^1(X) = P_T^2(X)$. Take the fixed predictors $f_0(x) = \mathbf{1}[x > 0]$ and $f_a(x) = \mathbf{1}[x < 0]$, and 0/1 loss. Define $P_T^1(Y \mid X): Y = \mathbf{1}[X > 0]$ and $P_T^2(Y \mid X): Y = \mathbf{1}[X < 0]$. Then $R_T^1(f_0) = 0$, $R_T^1(f_a) = 1$, so $\Delta^1 = -1 < 0$ (freeze is correct); and $R_T^2(f_0) = 1$, $R_T^2(f_a) = 0$, so $\Delta^2 = +1 > 0$ (adapt is correct). The evidence $Z = \phi(X_{1:m}, f_0, f_a)$ is a function of X and of the fixed maps f_0, f_a only; since X has the same law in both worlds, $\text{Law}(Z \mid P_T^1) = \text{Law}(Z \mid P_T^2)$. Thus (i) and (ii) hold simultaneously.

Impossibility. A label-free rule is a measurable map of Z and independent randomness, so $\text{Law}(g)$ coincides across the two worlds. Suppose g outputs ADAPT with probability q and FREEZE with probability r on this evidence. Measuring committal regret as the excess 0/1 risk of a committed action versus the better committed action in that world (and assigning abstention a separate coverage cost), world 1 charges $q \cdot |\Delta^1| = q$ (adapting is wrong) and world 2 charges $r \cdot |\Delta^2| = r$ (freezing is wrong). The worst case is $\max(q, r)$, which is minimized over the simplex at $q = r = 0$, i.e. ABSTAIN. $\square$

Remark (weight). Theorem 1 is close to known non-identifiability of target risk without labels; its purpose here is to *justify the abstain action*, not to claim impossibility as novel.

5.2 Optimal gate under alignment

Theorem 2 (Bayes gate). *If $\Delta(z)$ is $\sigma(Z)$ -measurable, the risk-minimizing label-free gate is $g^*(z) = \mathbf{1}[\Delta(z) > 0]$ (adapt iff benefit positive).*

Proof. Pointwise minimization of expected loss over the two committal actions. $\square$

The content is the antecedent (that $\Delta(z)$ is recoverable), addressed next and in §5.4.

5.3 Finite-sample adapt/freeze/abstain certificate

Theorem 3 (Certificate). *Suppose a label-free estimator $\hat{\Delta}(z)$ admits a radius $\varepsilon(z)$ with $\mathbb{P}[|\hat{\Delta}(z) - \Delta(z)| \leq \varepsilon(z)] \geq 1 - \alpha$. Define*

$$\text{ADAPT if } \hat{\Delta} - \varepsilon > 0, \quad \text{FREEZE if } \hat{\Delta} + \varepsilon < 0, \quad \text{ABSTAIN otherwise.}$$

Then $\mathbb{P}[\text{ADAPT} \wedge \Delta \leq 0] \leq \alpha$ and $\mathbb{P}[\text{FREEZE} \wedge \Delta \geq 0] \leq \alpha$.

Proof. On the event $\{|\hat{\Delta} - \Delta| \leq \varepsilon\}$ (probability $\geq 1 - \alpha$), ADAPT implies $\Delta \geq \hat{\Delta} - \varepsilon > 0$, so a false adapt can occur only on the complementary event of mass $\leq \alpha$. The freeze bound is symmetric. $\square$

Discharging ε (partially). When Δ is a mean of bounded paired benefits $X_i = \ell(f_0(x_i), y_i) - \ell(f_a(x_i), y_i) \in [a, b]$ over a validated stress bank, an empirical-Bernstein lower confidence bound yields a valid ε under independence/exchangeability; this is the certificate already implemented in `src/elara/certification/switching_certificate.py`. The *assumption* (exchangeability / a calibrated benefit sample) is the residual burden and is not proved in full generality.

5.4 A provable positive regime

Theorem 4 (Covariate-shift identifiability). *Assume covariate shift $P_S(Y | X) = P_T(Y | X)$, and a density ratio $r(x) = p_T(x)/p_S(x)$ that is bounded ($r \leq B < \infty$) with $\mathbb{E}_S[r(X)^2] < \infty$. Then $R_T(f) = \mathbb{E}_S[r(X)\ell(f(X), Y)]$, the importance-weighted estimator $\hat{R}_T(f) = \frac{1}{n} \sum_i r(x_i)\ell(f(x_i), y_i)$ is unbiased and consistent, and hence Δ and $\text{sign } \Delta$ are identifiable from labeled source plus unlabeled target whenever $|\Delta| > 0$.*

Proof. Change of measure: $\mathbb{E}_{P_T}[\ell(f(X), Y)] = \mathbb{E}_{P_S}[\frac{p_T(X)}{p_S(X)}\ell(f(X), Y)] = \mathbb{E}_S[r(X)\ell(f(X), Y)]$, using $P_S(Y | X) = P_T(Y | X)$. Unbiasedness is immediate; by the strong law and $\mathbb{E}_S[r^2] < \infty$ the estimator is consistent with variance $O(1/n)$. Applying this to f_0, f_a gives a consistent $\hat{\Delta}$, whose sign matches $\text{sign } \Delta$ for large n when $|\Delta| > 0$. $\square$

Caveat (ties to Theorem 1). The premise $P_S(Y | X) = P_T(Y | X)$ is *not* testable from unlabeled data; when it fails, the regime can re-enter the unknowable region.

Theorem 5 (Sign-of-difference identifiability on the disagreement region; binary case). *Let $Y \in \{0, 1\}$ with 0/1 loss, and let $D = \{x : f_0(x) \neq f_a(x)\}$ be the disagreement region (observable, since f_0, f_a are known maps of X). Then*

$$\Delta = P_T(D) (2a_a^D - 1), \quad a_a^D := P_T(f_a(X) = Y | X \in D),$$

so $\text{sign } \Delta = \text{sign}(a_a^D - \frac{1}{2})$. Consequently $\text{sign } \Delta$ is identifiable from label-free evidence if and only if the evidence determines whether f_a exceeds chance accuracy on D ; in particular, if a label-free reliability signal brackets a_a^D relative to $\frac{1}{2}$, the sign is certified.

Proof. Off D the predictors agree, so $\ell(f_0(X), Y) = \ell(f_a(X), Y)$ and the per-sample benefit $\delta = 0$. On D the two predictions differ, so for binary Y exactly one of f_0, f_a equals Y ; hence $\delta = +1$ when f_a is correct and -1 when f_0 is correct, giving $\mathbb{E}[\delta \mid D] = P_T(f_a=Y \mid D) - P_T(f_0=Y \mid D) = 2a_a^D - 1$ (using $P_T(f_0=Y \mid D) = 1 - a_a^D$). Then $\Delta = \mathbb{E}_{P_T}[\delta] = P_T(D) \mathbb{E}[\delta \mid D] = P_T(D)(2a_a^D - 1)$, and since $P_T(D) \geq 0$ the sign claim follows. $\square$

Delta over Steinhardt & Liang, and tie to Theorem 1. This needs only an *ordinal* judgment—is f_a above chance on D —not the *cardinal* estimation of either absolute risk, a strictly weaker requirement than identifying $R_T(f_0)$ or $R_T(f_a)$. The residual burden is exactly estimating $\text{sign}(a_a^D - \frac{1}{2})$ label-free (a calibration/transfer assumption); when no evidence brackets a_a^D around $\frac{1}{2}$, the regime re-enters the unknowable case of Theorem 1. Empirically, detector *disagreement* is the single most important evidence feature in our ablation (§7), directly supporting that the disagreement region carries the decision.

Conjecture 1 (Multiclass/regression extension; **open**). *The clean “exactly one correct on D ” identity is binary-specific. For multiclass 0/1 loss and for regression losses the analogous statement $\text{sign } \Delta = \text{sign } \mathbb{E}[\delta \mid D]$ still holds, but bracketing $\mathbb{E}[\delta \mid D]$ label-free requires a stronger reliability model; characterizing it is open.*

6 Method: KGA (Knowability-Guided Adaptation)

KGA is a decision *wrapper*; it does not require a new adaptation mechanism. Given a frozen f_0 , a candidate adaptation, and an unlabeled batch: (1) extract $Z = \phi(\cdot)$; (2) estimate benefit $\hat{\Delta}(Z)$; (3) form a radius ε ; (4) return ADAPT if $\hat{\Delta} - \varepsilon > 0$, FREEZE if $\hat{\Delta} + \varepsilon < 0$, else ABSTAIN. In our experiments $\hat{\Delta}$ is a gradient-boosted regressor trained leave-one-task-out on $(Z, \text{true } \Delta)$ and ε is the $(1 - \alpha)$ quantile of leave-one-out residuals (split-conformal style, $\alpha = 0.1$).

7 Experiments

Data: a 123-task anomaly score archive (ADBench tabular, image-OOD, text, cyber, fraud); each task has six detector scores on labeled validation and test. Frozen f_0 = best-validation detector. Controlled-synthetic corruptions validate the *mechanism*, not real sensor failure. All numbers come from `kbound_paper/scripts/`.

Safety on the clean suite. With candidate f_a = rank-normalized logistic stack, KGA’s abstentions concentrate where the true benefit is near zero (mean $|\Delta| = 0.021$ on abstained vs 0.132 on acted tasks); adapt precision 0.90. This suite is helpful-dominated, so always-adapt (mean AUROC 0.777) edges the safe policy (0.766); coverage 17%. Honest reading: when adaptation almost always helps, abstention mostly costs coverage.

Harmful regime. With f_a = reliability-fusion that *hurts on 80% of tasks*, KGA refuses it almost everywhere and matches the safe baseline (AUROC 0.748) while always-adapt drops to 0.728; regret-to-oracle falls from 0.0214 to 0.0019 ($\approx 11\times$). See Fig. 3.

Controlled mixed regime (369 instances). Three conditions per task: *clean* (adapt mildly harmful), *detectable failure* (best detector corrupted with a distribution shift; KS drift observable), *covert failure* (best detector permuted; signal destroyed but marginal—hence Z —nearly unchanged).

Decisions by true regime are in Table 2; adapt precision 0.935, false-adapt rate 0.065. Policy AUROC: always-adapt 0.697, always-freeze 0.586, KGA 0.690, oracle 0.719 (Fig. 2). KGA beats always-freeze by 0.10 and ties always-adapt—because harmful adaptation is mild in this domain.

condition	true mean Δ	ADAPT	FREEZE	ABSTAIN
clean	−0.041	8	5	110
detectable failure	+0.193	71	3	49
covert failure	+0.181	59	3	61

Table 2: KGA decisions by true regime (controlled mixed benchmark).

Empirical non-identifiability witness (partial). Clean instances have mean $\Delta = -0.041$ while covert-failure instances have mean $\Delta = +0.181$ (opposite sign), yet their evidence is far closer (mean standardized Z -distance clean→covert 1.35 vs clean→detectable 3.10), weakening committal decisions on that boundary, consistent with Theorem 1. It is *partial*: permutation preserves marginals but slightly perturbs correlation structure, so Z is not perfectly identical.

Clean non-identifiability witness. We also construct the exact witness of Theorem 1: two worlds with $X \sim \mathcal{N}(0, 1)$, $f_0 = \mathbf{1}[X > 0]$, $f_a = \mathbf{1}[X < 0]$, and flipped labels, so the benefit is exactly opposite (world means -1.0 vs $+1.0$) while every evidence feature is a function of X only. Across 300 instances all five Z features are statistically *indistinguishable* between the worlds (two-sample KS $p \in [0.16, 0.96]$, all > 0.05), yet the truth is opposite—so Z cannot decide. KGA **abstains on** 100% of instances; a rule forced to commit by sign $\hat{\Delta}$ pays regret 0.51 (near chance). This is the clean theory–experiment bridge (Fig. 1).




Figure 1: Clean non-identifiability witness: identical Z -law (KS $p > 0.05$ on all features), exactly opposite true benefit, 100% abstention. Empirical realization of Theorem 1.

Statistical rigor (8 seeds, paired t -tests). Repeating the mixed-regime experiment over 8 seeds (split-conformal, $\alpha=0.1$): mean test AUROC 0.686 ± 0.003 (KGA), 0.697 ± 0.000 (always-adapt), 0.584 ± 0.001 (always-freeze), 0.718 ± 0.001 (oracle); adapt precision 0.936 ± 0.010 , false-adapt

0.064 ± 0.010 , coverage 0.373 ± 0.017 . Paired t -tests: KGA vs always-freeze $t=62.2$, $p=7.3 \times 10^{-11}$ (KGA significantly higher); KGA vs always-adapt $t=-8.8$, $p=5.1 \times 10^{-5}$ (significantly lower). Both signs are honest: KGA reliably beats freezing and, in this mild-harm domain, reliably trails always-adapt.

Ablations. Table 3 (left) drops each evidence feature: detector *disagreement* is the most important—removing it raises regret from 0.037 to 0.094 ($2.6\times$), directly supporting Theorem 5. An α sweep traces the safety/coverage tradeoff: $\alpha=0.01$ gives 0% false-adapt at 11% coverage, up to $\alpha=0.2$ giving 49% coverage at 7% false-adapt (Fig. 4, middle). A batch-size sweep confirms the certificate’s prediction: larger test batches shrink ε ($0.44 \rightarrow 0.18$ from 16 to 256 samples) and raise coverage ($2\% \rightarrow 26\%$).

drop feature	regret	α	coverage	false-adapt
disagreement	0.094	0.01	0.11	0.00
val_mean	0.043	0.05	0.26	0.07
log_ntest	0.039	0.10	0.35	0.065
(all features)	0.037	0.20	0.49	0.07

Table 3: Left: evidence-drop ablation (higher regret = more important feature). Right: α sweep (safety/coverage tradeoff).

Generalization: regression under covariate shift (Theorem 4). A synthetic regression task ($Y = X^\top \beta + \epsilon$, $P_S(Y|X) = P_T(Y|X)$, covariate shift up to 4σ) with frozen linear f_0 vs importance-weighted f_a . Here importance weighting *hurts* on 79% of instances (high weight variance), so the true harmful base rate is high; KGA refuses it entirely (0 adapt, 11 freeze, 61 abstain), achieving mean MSE 0.251—matching always-freeze and the oracle (0.251) while always-adapt(IW) is 0.338 (Fig. 4, right). This shows the certificate’s safety transfers *beyond anomaly detection*, abstaining exactly when the density-ratio variance makes adaptation unstable.

Corroborating repository evidence. On MVTec-3D, reliability gating does no harm on clean data ($0.882 = 0.882$) and recovers +0.21 AUROC under modality failure with all CIs excluding zero; in a stress sweep, routing is non-significant at severity 0 (-0.0039) and significant only at severity 3 ($+0.039$, CI excludes zero)—the predicted knowability crossover.

Online non-stationary regime: cross-regime robustness (the decisive test). The single-shot regimes above are each dominated by one outcome, so a fixed policy is already near-optimal. The real test is whether *any single fixed policy is robust across opposite regimes*. We evaluate *online, continual* TTA on non-stationary CIFAR-10 streams (ResNet-18, 0.874 clean acc, commodity GPU (16 GB)), under two schedules: a **helpful-dominated** stream (mostly recoverable corruptions) and a **harm-dominated** stream (long “trap” runs—single-class label shift, near-pure noise, small batches—so continual Tent accumulates damage and *collapses*). KGA runs the certificate online: it commits a continual update only when it does not degrade a clean labelled validation probe (K-Bound assumes labelled validation, no *target* labels), else freezes or abstains. Results over 3 seeds (Table 4, Fig. 6):

No fixed policy is robust. Always-adapt is best in the helpful regime (0.548) but **collapses** in the harm regime (0.339, -0.10 below freezing). Always-freeze is best in the harm regime (0.440) but **forgoes gains** in the helpful regime (0.472). KGA is near-best in *both* (0.553 and 0.426) and never




Figure 2: Mixed regime. Left: KGA beats always-freeze and ties always-adapt. Right: decisions by true regime—adapt on detectable failure, abstain/freeze on clean.




Figure 3: Left: harmful regime—KGA avoids the harmful fusion path. Right: certificate on the clean suite ($\hat{\Delta} \pm \varepsilon$ vs truth).

collapses, giving the **highest mean across regimes** (0.490 vs always-freeze 0.456 vs always-adapt 0.444). Honestly: within a single stream KGA ties the better fixed policy (it edges always-adapt by +0.005 in the helpful regime and trails always-freeze by 0.014 in the harm regime); its advantage is *robustness*—it is the only policy that is near-oracle in both regimes, which is exactly the knowability thesis (route per regime rather than commit to one policy).

Natural-shift panels (multi-seed): safe capture, not a beats-both knockout. We audit two real distribution-shift benchmarks over 3 seeds each. **ImageNet-R** uses a *diverse-backbone* panel of ten architectures (ResNet-101/152, ResNeXt-101, EfficientNet-B0/B3, ConvNeXt-Tiny/Base, ViT-B/16, Swin-T/B; 12 corruption conditions per backbone, 30 candidate $\times$ seed cells). **Camelyon17** uses Tent/EATA/SAR (36 conditions). We report the paired regret gap $\bar{\Delta}_{\text{reg}} = \text{regret}(\text{KGA}) - \text{regret}(\text{trivial})$ (negative = KGA better), with Holm correction across the panel (Table 5).

Versus always-freeze, KGA is reliably better: on ImageNet-R it lowers regret on seven of ten backbones with all gaps surviving Holm (e.g. ConvNeXt-Base -0.068 , ResNeXt-101 -0.046 , EfficientNet-B3 -0.042 , ResNet-152 -0.031), and on Camelyon17 it beats freeze for all three




Figure 4: Left: evidence-component ablation (disagreement dominates). Middle: α sweep—safety vs coverage. Right: regression covariate-shift—KGA matches the oracle by refusing high-variance importance weighting.




Figure 5: Left: the KGA decision flow (evidence $\rightarrow$ estimate $\rightarrow$ certificate $\rightarrow$ adapt/freeze/abstain). Right: KGA regret-to-oracle across all experiments.

methods (Tent -0.052 , EATA -0.083 , SAR -0.099 ; $p_{\text{Holm}} = 6 \times 10^{-4}$). *Versus always-adapt*, KGA ties where adapting helps (ImageNet-R ConvNeXt-Tiny $+0.019$, ResNet-101 $+0.015$; Camelyon17 Tent $+0.020$, EATA $+0.004$, SAR -0.001 , none surviving Holm) and improves where adapting is purely harmful (EfficientNet-B0 -0.050 , harmful on 100% of conditions). Crucially, KGA *beats both* trivial policies on **none** of the ten ImageNet-R backbones and **none** of the three Camelyon17 methods (only 1 of the 30 ImageNet-R candidate $\times$ seed cells—Swin-T, seed 0—is a beats-both cell).

Honest reading. Each panel is *regime-homogeneous*: ImageNet-R adaptation is broadly helpful (so always-adapt is already near-optimal) while Camelyon17 freezing is broadly safe relative to a hurtful adaptation budget (so the gap to freeze is large but adapt is rarely worse). In such single-regime panels the best one can do is *safely capture the benefit and avoid the harm*—which KGA does—and a beats-both win is not available to *any* fixed-or-gated policy. Genuine beats-both wins therefore remain confined to mixed regimes (the CIFAR-10-C stress grid and the reliability-gated multimodal fusion instantiation of App. A), exactly as pre-registered. These panels are evidence of *safety and benefit capture under real shift*, not of universal dominance.

8 Limitations and honest scope

(1) The headline is *cross-regime robustness*, not a within-stream knockout: across a helpful- and a harm-dominated online stream (Table 4), KGA has the highest mean accuracy and never collapses, but within a single stream it only *ties* the better fixed policy (edges always-adapt by $+0.005$ in the helpful regime; trails always-freeze by 0.014 in the harm regime). We do not claim a single-stream knockout over both. Larger-scale corroboration (CIFAR-100-C / ImageNet-C with EATA/SAR)

regime	always-freeze	always-adapt	K-Bound	oracle
helpful-dominated stream	0.472	0.548	0.553	0.704
harm-dominated stream (collapse)	0.440	0.339	0.426	0.630
mean across regimes	0.456	0.444	0.490	0.667

Table 4: Online non-stationary CIFAR-10 TTA (3 seeds/regime). Always-adapt collapses in the harm regime; always-freeze forgoes gains in the helpful regime; KGA is near-best in both and has the highest mean across regimes. Harm-regime std: freeze 0.009, adapt 0.025, KGA 0.009.




Figure 6: Per-window accuracy on the *harm-dominated* stream (pink = catastrophic “trap” windows). Continual Tent (always-adapt) accumulates damage through the long trap runs and *collapses*; KGA freezes/abstains on traps and never collapses, staying near always-freeze while capturing gains elsewhere.

and a tighter online certificate are future work. (2) Synthetic corruptions validate the mechanism, not real sensor failure. (3) The conformal radius assumes task exchangeability. (4) The positive sign-of-difference result (Theorem 5) is proved for the binary case; the multiclass/regression extension (Conjecture 1) is open. (5) The non-identifiability witness is empirically partial. (6) KGA’s value scales with how often catastrophic, detectable harm occurs—precisely the quantity the theory characterizes; in benign domains always-adapt is strong and the certificate is mainly safety insurance. (7) The natural-shift panels (ImageNet-R, Camelyon17; Table 5) are each regime-homogeneous, so KGA can at best beat one trivial policy and tie the other; a beats-both win on these panels is not attainable by any policy, and we do not claim one. A multi-seed natural-shift benchmark that is genuinely *mixed* (both helpful and harmful conditions on the same backbone, with detectable harm) is the remaining test for a beats-both result under real shift.

9 Conclusion

We reframed label-free adaptation as a decision constrained by identifiability: adapt when benefit is certifiable, freeze when harm is certifiable, abstain when the unlabeled evidence cannot tell. We proved the unknowable regime is fundamental (with a witness), gave a false-adapt-controlled certificate under a stated alignment assumption, proved a covariate-shift positive regime, and reduced sign-of-difference identifiability to an ordinal accuracy judgment on the observable disagreement

panel	candidate	vs always-freeze	vs always-adapt	beats both?
ImageNet-R	ConvNeXt-Base	$-0.068^\dagger$	0.000	no
ImageNet-R	ResNeXt-101	$-0.046^\dagger$	0.005	no
ImageNet-R	EfficientNet-B3	$-0.042^\dagger$	0.010	no
ImageNet-R	ResNet-152	$-0.031^\dagger$	0.010	no
ImageNet-R	Swin-B	$-0.020^\dagger$	0.017	no
ImageNet-R	ViT-B/16	$-0.010^\dagger$	0.014	no
ImageNet-R	ResNet-101	$-0.010^\dagger$	0.015	no
ImageNet-R	EfficientNet-B0	0.000	$-0.050^\dagger$	no
ImageNet-R	ConvNeXt-Tiny	-0.002	0.019	no
ImageNet-R	Swin-T	-0.001	$-0.007^\dagger$	no
Camelyon17	Tent	$-0.052^\dagger$	0.020	no
Camelyon17	EATA	$-0.083^\dagger$	0.004	no
Camelyon17	SAR	$-0.099^\dagger$	-0.001	no

Table 5: Multi-seed (3 seeds) natural-shift panels. Entries are the paired regret gap $\text{regret}(\text{KGA}) - \text{regret}(\text{trivial})$; negative means KGA has lower regret (better). $\dagger$ marks gaps surviving Holm correction within the panel. KGA reliably beats always-freeze and ties always-adapt, but beats *both* on none of the 13 candidates—each panel is regime-homogeneous, so one trivial policy is already near-optimal.

region (Theorem 5), leaving only the multiclass/regression extension open. On real anomaly-routing data the certificate abstains where benefit vanishes and recovers large performance under failure, and on multi-seed natural-shift panels (ImageNet-R, Camelyon17) it reliably beats always-freeze and ties always-adapt—safely capturing benefit and avoiding harm without beating both policies, because those panels are regime-homogeneous. The honest frontier is a catastrophic-harm benchmark that would let the trichotomy beat every trivial policy.

References

References

- [1] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, T. Darrell. Tent: Fully test-time adaptation by entropy minimization. *ICLR*, 2021.
- [2] J. Liang, D. Hu, J. Feng. Do we really need to access the source data? Source hypothesis transfer for unsupervised domain adaptation. *ICML*, 2020.
- [3] Y. Sun, X. Wang, Z. Liu, J. Miller, A. Efros, M. Hardt. Test-time training with self-supervision for generalization under distribution shifts. *ICML*, 2020.
- [4] S. Niu, J. Wu, Y. Zhang, Y. Chen, S. Zheng, P. Zhao, M. Tan. Efficient test-time model adaptation without forgetting. *ICML*, 2022.
- [5] S. Niu, J. Wu, Y. Zhang, Z. Wen, Y. Chen, P. Zhao, M. Tan. Towards stable test-time adaptation in dynamic wild world. *ICLR*, 2023.
- [6] Q. Wang, O. Fink, L. Van Gool, D. Dai. Continual test-time domain adaptation. *CVPR*, 2022.
- [7] M. Zhang, S. Levine, C. Finn. MEMO: Test time robustness via adaptation and augmentation. *NeurIPS*, 2022.

- [8] S. Jeon, Z. Mai, H. Tian, V. Bakshi, L. Wang, J. Hou, P. Zhang, A. Chowdhury, W.-L. Chao. Continually adapt or not (CAN)? A continual learning benchmark of camera trap species classification over time. *NeurIPS Workshop (Imageomics)*, 2025. OpenReview: [jEItIzAgHs](#).
- [9] M. Schirmer, M. Jazbec, C. A. Naesseth, E. Nalisnick. Monitoring risks in test-time adaptation. *NeurIPS*, 2025. arXiv:2507.08721.
- [10] A. Podkopaev, A. Ramdas. Tracking the risk of a deployed model and detecting harmful distribution shifts. *ICLR*, 2022. arXiv:2110.06177.
- [11] Y. Bar, S. Shaer, Y. Romano. Protected test-time adaptation via online entropy matching: A betting approach. *NeurIPS*, 2024. arXiv:2408.07511.
- [12] A. Sreeram, M. N. Huh, J. K. Wong, et al. Tempora: Characterising the time-contingent utility of online test-time adaptation. arXiv:2602.06136, 2026.
- [13] S. Garg, S. Balakrishnan, Z. Lipton, B. Neyshabur, H. Sedghi. Leveraging unlabeled data to predict out-of-distribution performance. *ICLR*, 2022. arXiv:2201.04234.
- [14] J. Miller, R. Taori, A. Raghunathan, S. Sagawa, P. W. Koh, V. Shankar, P. Liang, Y. Carmon, L. Schmidt. Accuracy on the line: On the strong correlation between out-of-distribution and in-distribution generalization. *ICML*, 2021. arXiv:2107.04649.
- [15] C. Baek, Y. Jiang, A. Raghunathan, J. Z. Kolter. Agreement-on- the-line: Predicting the performance of neural networks under distribution shift. *NeurIPS*, 2022. arXiv:2206.13089.
- [16] E. Kim, M. Sun, C. Baek, A. Raghunathan, J. Z. Kolter. Test-time adaptation induces stronger accuracy and agreement-on-the-line. *NeurIPS*, 2024. arXiv:2310.04941.
- [17] T. Lee, S. Chottananurak, T. Gong, S.-J. Lee. AETTA: Label-free accuracy estimation for test-time adaptation. *CVPR*, 2024. arXiv:2404.01351.
- [18] Y. Jiang, V. Nagarajan, C. Baek, J. Z. Kolter. Assessing generalization of SGD via disagreement. *ICLR*, 2022.
- [19] Y. Lu, Z. Wang, R. Zhai, S. Kolouri, J. Campbell, K. Sycara. Characterizing out-of-distribution error via optimal transport. *NeurIPS*, 2023.
- [20] J. Steinhardt, P. Liang. Unsupervised risk estimation using only conditional independence structure. *NeurIPS*, 2016. arXiv:1606.05313.
- [21] S. Ben-David, J. Blitzer, K. Crammer, A. Kulesza, F. Pereira, J. Wortman Vaughan. A theory of learning from different domains. *Machine Learning* 79(1–2):151–175, 2010.
- [22] S. Ben-David, T. Lu, T. Luu, D. Pál. Impossibility theorems for domain adaptation. *AISTATS*, 2010.
- [23] Z. Lipton, Y.-X. Wang, A. Smola. Detecting and correcting for label shift with black box predictors. *ICML*, 2018.
- [24] E. Rosenfeld, S. Garg. (Almost) provable error bounds under distribution shift via disagreement discrepancy. *NeurIPS*, 2023. arXiv:2306.00312.
- [25] F. Parisi, F. Strino, B. Nadler, Y. Kluger. Ranking and combining multiple predictors without labeled data. *PNAS* 111(4):1253–1258, 2014. arXiv:1303.3257.
- [26] A. Jaffe, B. Nadler, Y. Kluger. Estimating the accuracies of multiple classifiers without labeled data. *AISTATS*, 2015. arXiv:1407.7644.
- [27] A. Jaffe, E. Fetaya, B. Nadler, T. Jiang, Y. Kluger. Unsupervised ensemble learning with dependent classifiers. *AISTATS*, 2016. arXiv:1510.05830.
- [28] E. A. Platanios, A. Blum, T. Mitchell. Estimating accuracy from unlabeled data. *UAI*, 2014.
- [29] S. Ibrahim, X. Fu, N. Sidiropoulos. Crowdsourcing via pairwise co-occurrences: Identifiability and algorithms. *NeurIPS*, 2019. arXiv:1909.12325.

- [30] A. Corrada-Emmanuel. The logic of NTQR evaluations of noisy AI agents: Complete postulates and logically consistent error correlations. *NeurIPS*, 2024. arXiv:2312.05392; see also arXiv:2409.11052, arXiv:2412.16238.
- [31] A. P. Dawid, A. M. Skene. Maximum likelihood estimation of observer error-rates using the EM algorithm. *JRSS-C* 28(1):20–28, 1979.
- [32] L. Le Cam. *Asymptotic Methods in Statistical Decision Theory*. Springer, 1986.
- [33] C. Manski. Nonparametric bounds on treatment effects. *Amer. Econ. Rev. P&P* 80(2):319–323, 1990.
- [34] Y. Geifman, R. El-Yaniv. Selective classification for deep neural networks. *NeurIPS*, 2017.
- [35] V. Vovk, A. Gammerman, G. Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- [36] A. Angelopoulos, S. Bates. Conformal prediction: A gentle introduction. *Found. Trends Mach. Learn.*, 2023.
- [37] G. Shafer. Testing by betting: A strategy for statistical and scientific communication. *JRSS-A*, 2021.
- [38] A. Ramdas, P. Grünwald, V. Vovk, G. Shafer. Game-theoretic statistics and safe anytime-valid inference. *Statist. Sci.*, 2023.
- [39] V. Vovk, I. Petej, I. Nouretdinov, E. Ahlberg, L. Carlsson, A. Gammerman. Retrain or not retrain: Conformal test martingales for change-point detection. *COPA*, 2021. arXiv:2102.10439.
- [40] S. R. Howard, A. Ramdas, J. McAuliffe, J. Sekhon. Time-uniform, nonparametric, nonasymptotic confidence sequences. *Annals of Statistics* 49(2):1055–1080, 2021.
- [41] A. Maurer, M. Pontil. Empirical Bernstein bounds and sample variance penalization. *COLT*, 2009.
- [42] W. Hoeffding. Probability inequalities for sums of bounded random variables. *JASA* 58(301):13–30, 1963.
- [43] A. Klivans, K. Stavropoulos, A. Vasilyan. Testable learning with distribution shift. arXiv:2311.15142, 2023.
- [44] A. Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- [45] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, L. Fei-Fei. ImageNet: A large-scale hierarchical image database. *CVPR*, 2009.
- [46] D. Hendrycks, T. Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. *ICLR*, 2019.
- [47] D. Hendrycks, S. Basart, N. Mu, S. Kadavath, F. Wang, E. Dorundo, R. Desai, W. Zhu, R. Parajuli, M. Guo, D. Song, J. Steinhardt. The many faces of robustness: A critical analysis of out-of-distribution generalization. *ICCV*, 2021.
- [48] P. W. Koh, S. Sagawa, H. Marklund, S. M. Xie, M. Zhang, A. Balsubramani, W. Hu, M. Yasunaga, et al. WILDS: A benchmark of in-the-wild distribution shifts. *ICML*, 2021.
- [49] P. Bándi, O. Geessink, Q. Manson, M. van Dijk, M. Balkenhol, M. Hermesen, et al. From detection of individual metastases to classification of lymph node status at the patient level: The CAMELYON17 challenge. *IEEE Trans. Med. Imaging* 38(2):550–560, 2019.
- [50] J. Taylor, S. L. Earnshaw, B. M. Mabey, M. Victors, M. N. Annis, D. Y. Torigian, et al. The RXX1 dataset and challenge: Unsupervised classification of cellular morphological changes. arXiv:1907.04758, 2019.
- [51] H. Venkateswara, J. Eusebio, S. Chakraborty, S. Panchanathan. Deep hashing network for unsupervised domain adaptation. *CVPR*, 2017. (Office-Home dataset release).
- [52] I. Gulrajani, D. Lopez-Paz. In search of lost domain generalization. *ICLR*, 2021.
- [53] F. Croce, M. Andriushchenko, V. Sehwag, E. Debenedetti, N. Flammarion, M. Chiang, P. Mittal, M. Hein. RobustBench: A standardized adversarial robustness benchmark. *NeurIPS Datasets and Benchmarks*, 2021.

- [54] B. Recht, R. Roelofs, L. Schmidt, V. Shankar. Do CIFAR-10 classifiers generalize to CIFAR-10? *ICML*, 2019.
- [55] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, et al. PyTorch: An imperative style, high-performance deep learning library. *NeurIPS*, 2019.
- [56] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in Python. *JMLR* 12:2825–2830, 2011.
- [57] B. Schölkopf, D. Janzing, J. Peters, E. Sgouritsa, K. Zhang, J. Mooij. On causal and anticausal learning. *ICML*, 2012.
- [58] X. Wu, M. Gong, J. H. Manton, U. Aickelin, J. Zhu. On causality in domain adaptation and semi-supervised learning: an information-theoretic analysis for parametric models. *JMLR* 25(261):1–57, 2024.

A Reliability-gated fusion as a worked, code-validated instantiation (T1–T9, GDR)

The general trichotomy (Theorems 1–5) is instantiated for reliability-gated multimodal fusion by a stack of ten results, each with code and a validator in the repository (vendored under `kbound_paper/vendored_from_elara/theory/`). They are *supporting* / *instantiation* results—not the general trichotomy—and several are explicit operationalizations (e.g. CLT/Gaussian approximations in T3, partial real-validation in GDR); scope caveats are stated in each module. Their value here is concrete: they show the abstract boundary has a fully worked, machine-checked special case. The decisive mappings are **T5 = Theorem 3** (the finite-sample certificate), **T9 = the proven knowably-harmful/clean regime**, and **GDR = minimax justification of the abstain action** (Theorem 1).

ID	One-line statement	K-Bound role / regime
T1	Reliability- <i>blind</i> fusion is dominated under sparse/coherent corruption (positive lower bound).	Establishes the <i>knowably-helpful</i> (stress) regime: gating is necessary when corruption is present.
T2	Global KS drift is inflated by benign mixture re-weighting even with no real drift.	Observability hazard: evidence Z can <i>false-fire</i> $\Rightarrow$ a route into false-adapt / unknowable.
T3	Closed-form mean-gate miss probability (CLT approx.).	Quantifies gate error; calibrates the certificate.
T4	Finite-sample risk-dominance sample complexity for the switch.	How much validation certifies the helpful regime.
T5	Empirical-Bernstein finite-sample switching certificate.	Instantiates Theorem 3 (false-adapt-controlled rule).
T6	Windowed-KS gate as a CUSUM detector; ARL/AED trade-off.	Detectability of the shift trigger: when Z is informative.
T7	Massart finite-class Rademacher bound for the meta-router.	Generalization of the benefit estimator $\hat{\Delta}$ in Theorem 3.
T8	Validation-only certified selection among {gate, TTA, stack}.	Multi-candidate KGA decision (beyond binary adapt/freeze).
T9	No fusion rule beats the confidence-weighted mean on clean near-separable transfer.	Proves the knowably-harmful/neutral (clean) regime $\Rightarrow$ freeze is optimal.
GDR	Coherence-certified switch is minimax over coherent-shift vs heterogeneous-mixture adversaries.	Justifies the abstain/gate action as minimax-safe (Theorem 1).

Together, T1/T3/T4 lower-bound the *helpful* regime, T9 proves the *harmful/neutral* regime, T2/T6/T7 govern the *observability* of the evidence Z (and hence the boundary of the *unknowable* regime), T5 supplies

the operational certificate, and GDR certifies abstention. This is the fusion-case shadow of the general {adapt, freeze, abstain} theory.