

# K-Bound: When Is Label-Free Adaptation Knowable?

Helpful, Harmful, and Unknowable Regimes Under Distribution Shift

Full Companion Manuscript

Pratik Niroula

Minnesota State University, Mankato

pratik.niroula@mnsu.edu

## Abstract

Test-time adaptation updates a model without new target labels, but an update can help or harm. We ask when the available evidence supports replacing a frozen predictor with a fixed adapted candidate. K-Bound answers this question for a binary population model with a declared bound  $\beta$  on the target calibration residual. We construct target label laws that preserve all label-free observations and derive the exact feasible interval for paired population benefit. A strict decision is justified precisely when  $|M| > \beta$ , where  $M$  is the observable correctness margin on the disagreement region; inside the closed band  $|M| \leq \beta$ , no strict direction is uniformly justified. Furthermore, an evidence-fibre audit limit shows that the same label-free observations cannot audit away this unresolved residual bound: any uniformly valid audit must cover at least the declared budget  $\beta$ . Knowability-Guided Adaptation (KGA) is an empirical companion, not an implementation of this population frontier. It learns to predict observed evaluation-cell benefit from labeled historical outcomes, then uses a residual-based calibrated interval to choose ADAPT, FREEZE, or ABSTAIN without new evaluation labels. On an exploratory, cross-fitted CIFAR-10-C panel, Tent and EATA have lower point regret than either fixed policy, while SAR favors always-adapt. On CCT-20, KGA retains frozen predictions and avoids observed degradation from always-adapt, but makes no ADAPT decisions. These results support candidate-specific commitment control on evaluation cells; they do not establish interval coverage or population-risk protection on unseen natural shifts.

**Keywords:** test-time adaptation; abstention; distribution shift; label-free risk estimation; conformal calibration; selective decisions.

## Contents

|          |                                                          |          |
|----------|----------------------------------------------------------|----------|
| <b>1</b> | <b>Introduction</b>                                      | <b>5</b> |
| 1.1      | Three Scientific Contributions . . . . .                 | 5        |
| <b>2</b> | <b>Related Work</b>                                      | <b>6</b> |
| 2.1      | Candidate Adaptation and Update Selection . . . . .      | 6        |
| 2.2      | Label-Free Performance and Strategy Estimation . . . . . | 6        |
| 2.3      | Monitoring and Adaptation Controllers . . . . .          | 6        |
| 2.4      | Partial Identification and Robust Decisions . . . . .    | 7        |
| 2.5      | Marginal and Conditional Coverage . . . . .              | 7        |

|           |                                                                                  |           |
|-----------|----------------------------------------------------------------------------------|-----------|
| <b>3</b>  | <b>Setup, Notation, and Claim Levels</b>                                         | <b>8</b>  |
| 3.1       | Model Pair, Evidence, and Evaluation Unit . . . . .                              | 8         |
| 3.2       | Actions and Their Operational Meaning . . . . .                                  | 8         |
| 3.3       | Commitment, Interval Coverage, and Error . . . . .                               | 9         |
| <b>4</b>  | <b>Population Partial-Identification Frontier</b>                                | <b>9</b>  |
| 4.1       | Disagreement-Region Model . . . . .                                              | 9         |
| 4.2       | Label-Kernel Freedom on the Disagreement Region . . . . .                        | 10        |
| 4.3       | Benefit Reduction and Matched-Evidence Impossibility . . . . .                   | 10        |
| 4.4       | Boundary Case . . . . .                                                          | 11        |
| 4.5       | Exact Strict-Commitment Frontier . . . . .                                       | 12        |
| 4.6       | Interpretation and Limits of the Frontier . . . . .                              | 12        |
| <b>5</b>  | <b>What Label-Free Evidence Cannot Learn</b>                                     | <b>13</b> |
| 5.1       | Audit Setup and Evidence Fibres . . . . .                                        | 13        |
| 5.2       | Exact Fibre-Wise Audit Floor . . . . .                                           | 13        |
| 5.3       | Declared Budget and Abstention Radius . . . . .                                  | 14        |
| 5.4       | Why Labels at the Wrong Domain Buy Nothing . . . . .                             | 14        |
| 5.5       | What Additional Structure Can Buy . . . . .                                      | 14        |
| 5.6       | Implications for K-Bound . . . . .                                               | 15        |
| <b>6</b>  | <b>KGA: An Empirical Benefit-Interval Gate</b>                                   | <b>15</b> |
| 6.1       | What Transfers from the Population Analysis . . . . .                            | 15        |
| 6.2       | The Coverage-to-Action Proposition . . . . .                                     | 16        |
| 6.3       | Evidence, Calibration, and the Action Rule . . . . .                             | 19        |
| <b>7</b>  | <b>Primary Experimental Protocols</b>                                            | <b>21</b> |
| 7.1       | Two Primary Questions and the Evidence Hierarchy . . . . .                       | 21        |
| 7.2       | CCT-20 Design and Safe Utility . . . . .                                         | 21        |
| 7.3       | Scoring, Dependence, and Statistical Inference . . . . .                         | 22        |
| <b>8</b>  | <b>Primary Results</b>                                                           | <b>22</b> |
| 8.1       | Controlled Routing Depends on the Candidate . . . . .                            | 22        |
| 8.2       | Observed Interval Inclusion Is a Separate Diagnostic . . . . .                   | 23        |
| 8.3       | Synchronized Development Baseline Comparison . . . . .                           | 24        |
| 8.4       | CCT-20: Conservative Retention, Not Selective Routing . . . . .                  | 25        |
| 8.5       | Development Stops and Adverse Diagnostics . . . . .                              | 25        |
| <b>9</b>  | <b>Limitations and Broader Impact</b>                                            | <b>26</b> |
| 9.1       | Assumptions, Not an Unconditional Safety Guarantee . . . . .                     | 26        |
| 9.2       | Empirical Delimitations and Future Prospective Validation . . . . .              | 26        |
| 9.3       | Operational Scope and Potential Harm . . . . .                                   | 27        |
| <b>10</b> | <b>Conclusion</b>                                                                | <b>27</b> |
| <b>A</b>  | <b>A Rigorous Bridge from Cell Coverage to Population Coverage</b>               | <b>31</b> |
| A.1       | Unit Scaling and the Category Error of $\varepsilon = \beta$ . . . . .           | 31        |
| A.2       | Counterexample: Perfect Cell Prediction Does Not Imply Population Coverage . . . | 31        |
| A.3       | Formal Protocol for Cell-to-Population Transfer . . . . .                        | 33        |

|          |                                                                                   |           |
|----------|-----------------------------------------------------------------------------------|-----------|
| A.4      | Theorem and 4-Step Proof of Population Risk Control . . . . .                     | 33        |
| A.5      | The Price of Rigor: Sample Size vs. Abstention Analysis . . . . .                 | 35        |
| A.6      | Alternative Route: Direct Empirical Sampling of the Population Frontier . . . . . | 35        |
| A.7      | Explicit Implementation Boundaries and Operational Scope . . . . .                | 37        |
| <b>B</b> | <b>Supplementary Results and Non-Promoted Studies</b>                             | <b>39</b> |
| B.1      | Accuracy-Based Diagnostic Replays . . . . .                                       | 40        |
| B.2      | Balanced-Accuracy Diagnostic Replays . . . . .                                    | 40        |
| B.3      | Large-Scale Corruption: ImageNet-C . . . . .                                      | 40        |
| B.4      | Office-Home . . . . .                                                             | 45        |
| B.5      | Camelyon17 . . . . .                                                              | 46        |
| B.6      | RxRx1 . . . . .                                                                   | 46        |
| B.7      | iWildCam . . . . .                                                                | 47        |
| B.8      | Boundary and Transfer Diagnostics . . . . .                                       | 47        |
| B.9      | Native-Component Benchmark Comparison and Port Audits . . . . .                   | 47        |
| B.10     | So2Sat Development and Action-Unit Audit . . . . .                                | 49        |
| B.11     | Withdrawn Development Proxy for the Population Budget . . . . .                   | 49        |
| <b>C</b> | <b>Compact Assumption and Claim Ledger</b>                                        | <b>50</b> |
| <b>D</b> | <b>Calibration and Evaluation Contracts</b>                                       | <b>51</b> |
| <b>E</b> | <b>Decision Exposure and Repetition Audit</b>                                     | <b>51</b> |
| <b>F</b> | <b>Inference Discipline</b>                                                       | <b>52</b> |
| F.1      | Point Estimates, Paired Gaps, and Clusters . . . . .                              | 52        |
| F.2      | Current-Policy Corruption-Family Sensitivity . . . . .                            | 53        |
| F.3      | Retrospective Interval Inclusion by Family . . . . .                              | 53        |
| F.4      | False-Adapt Accounting . . . . .                                                  | 54        |
| F.5      | Multiplicity and Search History . . . . .                                         | 54        |
| <b>G</b> | <b>Track-Level Provenance and Remaining Evidence</b>                              | <b>55</b> |
| G.1      | CCT-20 Inference, Location, and Release Ledger . . . . .                          | 56        |
| <b>H</b> | <b>Extended KGA Implementation Contract</b>                                       | <b>57</b> |
| H.1      | State Isolation . . . . .                                                         | 57        |
| H.2      | Evidence Determinism . . . . .                                                    | 58        |
| H.3      | Calibration Identity . . . . .                                                    | 58        |
| H.4      | Runtime Accounting . . . . .                                                      | 58        |
| <b>I</b> | <b>Evidence Inventory</b>                                                         | <b>58</b> |
| I.1      | Protocol-Specific Feature and Model Schemas . . . . .                             | 58        |
| <b>J</b> | <b>Failure-Safe Deployment Checklist</b>                                          | <b>59</b> |
| <b>K</b> | <b>Prospective Natural-Shift Confirmation Protocol</b>                            | <b>59</b> |
| K.1      | Lock Before Target Labels . . . . .                                               | 60        |
| K.2      | Three-Way Data Separation . . . . .                                               | 60        |
| K.3      | Endpoints and Decision Rule . . . . .                                             | 60        |

|                                                                |           |
|----------------------------------------------------------------|-----------|
| K.4 Interpreting Outcomes . . . . .                            | 61        |
| <b>L Theorem Dependencies and Delimited Scope</b>              | <b>61</b> |
| <b>M Formalization and Artifact Scope</b>                      | <b>63</b> |
| <b>N Reproducibility and Release Procedures</b>                | <b>64</b> |
| N.1 Canonical Numerical Source of Truth . . . . .              | 64        |
| N.2 Reproduction Pipeline . . . . .                            | 65        |
| N.3 Calibration Identity and Implementation Contract . . . . . | 65        |
| N.4 Formalization and Lean Scope . . . . .                     | 66        |
| N.5 Release Checklist and Audit Logs . . . . .                 | 66        |
| <b>O Controlled Population-to-KGA Bridge</b>                   | <b>67</b> |
| <b>P Controlled Two-View MNIST Diagnostic</b>                  | <b>67</b> |

# 1 Introduction

Test-time adaptation changes a model using unlabeled deployment data. Entropy minimization, filtering, regularization, and related methods can recover accuracy under shift, but the same update can also make predictions worse [1, 4, 5, 6]. An adaptation method proposes a candidate; a deployment system must still decide whether to use it.

Our primary research question is: *When can evidence available without new target labels justify replacing a frozen predictor with a fixed adapted candidate?* Let  $f_0$  be the frozen predictor and  $f_a$  the candidate. Their population benefit is

$$\Delta = R_T(f_0) - R_T(f_a). \quad (1)$$

A positive benefit favors the update; a negative benefit favors the frozen model. We distinguish three conclusions: ADAPT supports positive benefit, FREEZE supports negative benefit, and ABSTAIN records insufficient evidence for either strict direction. Both non-adapt actions retain  $f_0$ , but only freeze asserts that the candidate is worse.

K-Bound studies this decision as a partial-identification problem. In a binary model, we hold the input law, predictors, and label-free evidence fixed and vary the target labels within a declared class. The resulting feasible benefit set gives an exact boundary for a uniformly justified strict decision. The bound on the hidden calibration residual is an assumption about that class, not a quantity that an unlabeled batch supplies for free. A second result makes this information limit explicit: an audit using the same evidence must still cover the unresolved residual range.

Knowability-Guided Adaptation (KGA) is an *empirical companion* to this analysis. It predicts the observed benefit of a model pair on an evaluation cell and acts only when a residual-based interval excludes zero. KGA is label-free at deployment, not during development and calibration: historical labeled outcomes train the benefit predictor and calibrate its errors. It neither estimates the population residual bound nor implements the population frontier numerically. Section 6.1 explains the common decision logic and the different guarantees.

## 1.1 Three Scientific Contributions

1. **An exact benefit frontier on a declared class.** We construct target label laws with identical label-free evidence, derive their feasible benefit interval, and characterize its strict positive, strict negative, and unresolved regions, including the zero-benefit boundary.
2. **An evidence-fibre audit limit.** We identify the least residual bound that the same observable law can support uniformly, explaining why historical labels help only through a stated link to the deployment domain.
3. **An empirical paired-benefit gate.** We give a reproducible three-way interval procedure and evaluate its commitment behavior under controlled and natural shifts. The standard coverage-to-action implication states exactly what follows if the interval covers its declared target.

The controlled CIFAR-10-C results show useful routing for some candidates, not a universal improvement: Tent and EATA have lower point regret than either fixed policy, whereas SAR favors always-adapt. The CCT-20 study shows a different behavior: KGA retains frozen predictions and avoids the measured loss from always adapting, but never selects ADAPT. These two studies separate selective routing from conservative retention. The supplementary results retain adverse, incomplete, and withheld studies, so this distinction does not depend on discarding unfavorable evidence.




Figure 1: KGA’s empirical workflow. A shadow candidate proposes an update; a learned benefit prediction and residual interval determine whether to accept it. Historical labels support development and calibration, while the new decision uses no evaluation labels.

## 2 Related Work

### 2.1 Candidate Adaptation and Update Selection

Tent, EATA, SAR, CoTTA, MEMO, SHOT, and test-time training improve candidate construction through entropy objectives, sample filtering, regularization, temporal stabilization, augmentation, or auxiliary tasks [1, 4, 5, 6, 7, 2, 3]. Selective methods also choose where or when to update. TTSA selects modality-specific adapters [10]; GALA selects layers and filters samples using gradient alignment [11]; Informed Adaptation learns topological features associated with useful updates [9]. KGA evaluates a fixed proposed candidate against its frozen baseline. Learned routing and rollback are established ideas; their existence is not our novelty claim.

### 2.2 Label-Free Performance and Strategy Estimation

ATC, agreement-on-the-line, AETTA, and prediction-matrix statistics estimate performance from unlabeled outputs and suitable development information [15, 17, 16, 18, 35, 34]. The closest empirical comparison is Kim et al.’s NeurIPS 2024 work, *Test-Time Adaptation Induces Stronger Accuracy and Agreement-on-the-Line* [51]. It uses agreement-based ALine estimators after adaptation, supports hyperparameter selection, and compares adaptation strategies without labeled OOD data.

KGA instead fits the paired observed benefit of one frozen/candidate pair and uses an interval to make a three-way decision. This is a difference in the fitted target and decision procedure, not a claim that accuracy estimates cannot be differenced to compare models. In Section 8.3 (Table 9), we compare current KGA with fixed policies, scalar evidence rules and a point gate sharing its exact predictions. Historical TTA/ALine, AETTA and POEM port labels do not establish faithful implementations of the published estimators or controllers. Authenticated native comparisons under a shared protocol remain distinct from this saved-feature diagnostic.

### 2.3 Monitoring and Adaptation Controllers

POEM protects adaptation through entropy matching and online betting [14]. Schirmer et al. monitor the running risk of an adapting model relative to a source-risk reference, under a proxy-informativeness condition [8]; the supervised precursor uses a labeled test stream [13]. Drift-to-Action

Table 1: Closest roles, targets, and remaining distinctions. Literature comparisons are conceptual; the paper does not report a synchronized official head-to-head experiment.

| Approach                              | Main target                                          | Relation to this paper                                                   |
|---------------------------------------|------------------------------------------------------|--------------------------------------------------------------------------|
| TTA/ALine; AETTA                      | adapted-model performance and selection              | closest empirical estimators; their scores can inform paired comparisons |
| POEM; risk monitors                   | protected adaptation or running risk                 | different monitored target and sequential assumptions                    |
| Partial-identification decisions      | uncertain payoffs across compatible laws             | context for the population benefit set and robust sign decision          |
| Selection/shift-conditional conformal | coverage after selection or within a shift class     | stronger, specialized targets not supplied by KGA’s marginal interval    |
| K-Bound / KGA                         | feasible population benefit / predicted cell benefit | shared update question, distinct uncertainty objects and guarantees      |

Controllers use delayed labels and online certificates to gate adaptation, abstention, rollback, and retraining [50]. These methods differ in the risk target, label access, and sequential guarantee. KGA concerns one paired benefit and does not establish an anytime-valid risk guarantee.

## 2.4 Partial Identification and Robust Decisions

An identified set describes the values compatible with the observations and assumptions. Decision theory under partial identification studies how to act when that set leaves payoffs ambiguous; Christensen et al. combine worst-case payoff uncertainty with statistical learning of identified parameters [54]. K-Bound addresses the narrower question of a uniformly correct strict benefit sign. It does not solve their general optimal-decision problem or claim that interval thresholding itself is new.

Indistinguishable adaptation problems are already central to domain-adaptation impossibility [20, 49]. K-Bound’s specific contribution is a disagreement-conditional label-kernel construction, its clipped benefit set, and the associated residual-audit floor. The proof does not introduce label-free impossibility. Risk-estimation assumptions based on conditional independence [19], label shift [44, 45], or disagreement discrepancy [21] restrict the compatible laws. Likewise, MMD-credal risk intervals impose RKHS and covariate-shift structure [12]. Such restrictions can change the identified set; they are not contradicted by a converse theorem for a richer class.

## 2.5 Marginal and Conditional Coverage

Selective classification and reject-option learning abstain on individual predictions [22, 46, 47]. KGA abstains on an update. Its coverage-to-action proposition applies standard interval logic [36, 37, 40, 41]. The main statistical work is to justify coverage, especially after selection or under shift.

Jin and Ren’s JOMI procedure targets coverage conditional on a unit being selected, under exchangeability and a compatible selection rule [52]. Gibbs et al. provide conditional guarantees over specified classes of shifts using a different calibration construction [53]. KGA’s ordinary residual interval implements neither procedure. Its marginal false-adapt bound therefore must be distinguished from error conditional on ADAPT and from coverage within each target environment. Learn Then Test separately addresses configuration selection and multiplicity [41].

Table 2: Notation and probability targets. Population and measured-cell quantities are kept separate.

| Quantity                    | Level      | Meaning                                                                |
|-----------------------------|------------|------------------------------------------------------------------------|
| $f_0, f_a$                  | both       | frozen predictor and realized shadow candidate                         |
| $\Delta$                    | population | expected risk reduction $R_T(f_0) - R_T(f_a)$                          |
| $D, \mu_T(D)$               | population | disagreement region and its target probability                         |
| $M$                         | population | label-free mean correctness-score margin on $D$                        |
| $\gamma, \beta$             | population | hidden calibration residual and its externally declared absolute bound |
| $\Delta^{\text{cell}}$      | empirical  | observed score difference on one predeclared evaluation unit           |
| $\hat{\Delta}, \varepsilon$ | empirical  | predicted cell benefit and residual-based interval radius              |
| $\text{FA}_u$               | empirical  | probability of ADAPT and a nonpositive cell benefit                    |
| $\text{FA}_c$               | empirical  | probability of nonpositive cell benefit conditional on ADAPT           |
| Commitment rate             | empirical  | fraction of ADAPT or FREEZE decisions; not interval coverage           |

### 3 Setup, Notation, and Claim Levels

#### 3.1 Model Pair, Evidence, and Evaluation Unit

Let  $P_S$  and  $P_T$  be the source and target laws and  $R_T(f) = \mathbb{E}_{(X,Y) \sim P_T}[\ell(f(X), Y)]$ . Population risks integrate a fresh draw from  $P_T$ , holding the realized model pair fixed. An adapter  $\mathcal{A}$  proposes  $f_a$  from  $f_0$  and unlabeled inputs  $X_{1:m}$ . The decision can read

$$Z = \phi(X_{1:m}, f_0, f_a),$$

including predictions, logits, and update statistics, but not the new evaluation labels.

For evaluation unit  $j$ , with records  $\mathcal{E}_j$  and a fixed higher-is-better score  $S$ , define

$$\Delta_j^{\text{cell}} = S(f_a; \mathcal{E}_j) - S(f_0; \mathcal{E}_j).$$

The two models are scored on the same records. This measured quantity is not population risk. For accuracy it is a sample average of paired correctness differences; for macro-F1 it is a difference of dataset-level scores, not an average of per-example losses. Metric scales are reported separately.

#### 3.2 Actions and Their Operational Meaning

**Definition 1** (Decision semantics). ADAPT commits  $f_a$  when the interval supports positive benefit. FREEZE rejects the candidate when it supports negative benefit. ABSTAIN retains  $f_0$  when neither strict direction is supported or the assessment is unavailable.

Freeze is a supported rejection. Abstention is an unresolved assessment. An operator can attach monitoring, label acquisition, or human review to the abstention record; the experiments here evaluate only the frozen fallback and uncertainty log, not those additional interventions. Retaining  $f_0$  after missing evidence must be logged as ABSTAIN, not as a certified FREEZE.

The observed-cell oracle chooses the larger score. Regret is the policy’s score shortfall from that oracle, averaged over the named units. Always-freeze regret is  $\max(\Delta_j^{\text{cell}}, 0)$ ; always-adapt regret is  $\max(-\Delta_j^{\text{cell}}, 0)$ . KGA uses the candidate score only after ADAPT. Lower regret is better, and a zero-benefit tie carries zero regret for either model even though it supports neither strict sign.

Table 3: Three claim levels and their validity obligations. This table collects the distinctions used throughout the paper.

| Claim level             | Required premise                                                                           | Conclusion and boundary                                                                                     |
|-------------------------|--------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| Population frontier     | binary 0/1 loss, fixed augmented evidence, declared residual class; richness for necessity | exact strict-sign frontier within that class, not a learned residual bound                                  |
| Operational proposition | marginal interval coverage for a named scalar target and unit                              | marginal false commitment for that target, not automatic conditional or anytime control                     |
| Empirical evaluation    | fixed recorded candidates, splits, metric and action rule                                  | measured cell regret, interval inclusion and action exposure; population transfer needs a separate argument |

### 3.3 Commitment, Interval Coverage, and Error

These are three different measurements:

$$\text{Commit} = \mathbb{P}(g \in \{\text{ADAPT}, \text{FREEZE}\}), \quad (2)$$

$$\text{Cov}_{\text{cell}} = \mathbb{P}(|\hat{\Delta} - \Delta^{\text{cell}}| \leq \varepsilon), \quad (3)$$

$$\text{FA}_u = \mathbb{P}(g = \text{ADAPT}, \Delta^{\text{cell}} \leq 0). \quad (4)$$

Conditional false adaptation is  $\text{FA}_c = \mathbb{P}(\Delta^{\text{cell}} \leq 0 \mid g = \text{ADAPT})$ , defined only when  $\mathbb{P}(g = \text{ADAPT}) > 0$ . The operational proposition below bounds the marginal error under interval coverage; it does not bound  $\text{FA}_c$ . A gate that never adapts has zero false adaptations mechanically, so every error report includes action exposure. Analogously, a false FREEZE is a FREEZE decision with  $\Delta^{\text{cell}} \geq 0$ ; its conditional frequency divides by FREEZE decisions and is undefined when that action is absent.

The population theorem uses binary labels and 0/1 loss. The empirical gate can instead target any predeclared scalar score difference, including multiclass metrics. This interface does not extend the binary theorem. In general multiclass classification,  $\Delta = \mathbb{P}_T(D)(p_a - p_0)$ , with  $p_a = \mathbb{P}_T(f_a = Y \mid D)$  and  $p_0 = \mathbb{P}_T(f_0 = Y \mid D)$ ; only the binary setting forces  $p_0 = 1 - p_a$ .

## 4 Population Partial-Identification Frontier

The result describes which benefit signs are compatible with a fixed label-free evidence law and a declared bound on the hidden residual.

### 4.1 Disagreement-Region Model

**Assumption 1** (Deployment setup). Let  $Y \in \{0, 1\}$ , use 0/1 loss, fix measurable binary predictors  $f_0, f_a$ , and assume  $\mu_T(D) > 0$  for  $D = \{x : f_0(x) \neq f_a(x)\}$ . Let  $s : \mathcal{X} \rightarrow [0, 1]$  be a fixed measurable candidate-correctness score, fitted using source information, and let the target law admit the measurable pointwise correctness kernel  $\eta_a(x) = \mathbb{P}_T(f_a(X) = Y \mid X = x)$ . On  $D$ , set

$$M = \mathbb{E}[s(X) \mid D] - \frac{1}{2}, \quad \gamma = \mathbb{E}[\eta_a(X) - s(X) \mid D]. \quad (5)$$

The declared population class is  $\mathcal{C}_\beta = \{P_T : |\gamma(P_T)| \leq \beta\}$  for an externally supplied  $\beta \geq 0$ . Because  $s \in [0, 1]$ , the feasible margin satisfies  $M \in [-\frac{1}{2}, \frac{1}{2}]$ . The identities and frontier below do not require a source-calibration property of  $s$ .

*Remark 1* (Residual identity). The definition gives

$$\begin{aligned} M + \gamma &= (\mathbb{E}[s(X) \mid D] - \tfrac{1}{2}) + \mathbb{E}[\eta_a(X) - s(X) \mid D] \\ &= \mathbb{E}[\eta_a(X) \mid D] - \tfrac{1}{2}. \end{aligned}$$

This is an identity, not empirical evidence that calibration transfers. We call  $\gamma$  a target disagreement-conditional calibration residual, not automatically source-to-target drift. Ordinary source calibration need not hold after conditioning on  $D$ , so  $\gamma$  can be nonzero even when  $P_T = P_S$ . A literal label-kernel shift interpretation would require an additional reference, for example the pointwise source correctness  $s(x) = \mathbb{P}_S(f_a(X) = Y \mid X = x)$ . The no-shift counterexample is given in Appendix L. Calibration notions also depend on the conditioning information: top-label calibration, for example, additionally conditions on the predicted class [43]. Neither score nor top-label calibration automatically supplies calibration conditional on  $D$ .

**Definition 2** (Risk alignment). Evidence  $Z$  is risk-aligned on a class  $\mathcal{P}$  if no two laws in  $\mathcal{P}$  with opposite nonzero benefit signs induce the same law of  $Z$ . This is weaker than uniform strict-sign identification: a fibre containing zero and positive benefit is risk-aligned by this definition, but does not support a uniformly positive claim. Population frontier claims take  $M$  to be fixed on each augmented evidence fibre, equivalently by using  $\tilde{Z} = (Z, M)$ .

**Definition 3** (Strict directional soundness). At a fixed augmented-evidence law, ADAPT is uniformly directionally sound when  $\Delta(P) > 0$  for every compatible  $P \in \mathcal{C}_\beta$ ; FREEZE is sound when  $\Delta(P) < 0$  for every such law. A pointwise maximal three-way rule commits whenever one strict action is sound and abstains otherwise. The strict convention treats  $\Delta = 0$  as insufficient for either directional claim, although the actions then have equal task risk.

## 4.2 Label-Kernel Freedom on the Disagreement Region

**Lemma 1** (Label-kernel freedom on disagreement). *Fix  $(\mu_T, P_S, f_0, f_a, s)$ , hence  $D$  and  $M$ , and fix the conditional label law outside  $D$ . For every measurable  $\eta : D \rightarrow [0, 1]$  there is a target law with input marginal  $\mu_T$  and  $\mathbb{P}(f_a(X) = Y \mid X = x) = \eta(x)$  on  $D$ . Every label-free observable formed from the input batch, the fixed models, and their outputs has the same law for all such  $\eta$ , while*

$$\gamma = \mathbb{E}[\eta(X) \mid D] - \tfrac{1}{2} - M.$$

*Proof.* On  $D$ , binary predictors disagree, so their outputs are 0 and 1. Assign conditional mass  $\eta(x)$  to  $Y = f_a(x)$  and  $1 - \eta(x)$  to  $Y = f_0(x)$ , retaining the fixed kernel outside  $D$ . This changes no input marginal or model output and therefore changes no label-free evidence law. The expression for  $\gamma$  follows from Eq. (5).  $\square$

## 4.3 Benefit Reduction and Matched-Evidence Impossibility

**Lemma 2** (Disagreement-region reduction). *Under Assumption 1, with  $\bar{a} := \mathbb{P}_T(f_a=Y \mid D)$ ,*

$$\begin{aligned} \Delta &= 2\mu_T(D)(\bar{a} - \tfrac{1}{2}), \\ \text{sign } \Delta &= \text{sign}(\bar{a} - \tfrac{1}{2}) = \text{sign}(M + \gamma). \end{aligned}$$

*The identity  $\text{sign } \Delta = \text{sign}(M + \gamma)$  holds for every target law; it does not assume  $|\gamma| \leq \beta$ .*

*Proof.* On  $D$  exactly one of  $f_0, f_a$  is correct under 0/1 loss, so the benefit equals  $2\mu_T(D)(\bar{a} - \frac{1}{2})$ . By the definitions of  $M$  and  $\gamma$ ,

$$\begin{aligned} M + \gamma &= \left( \mathbb{E}_{\mu_T|D}[s] - \frac{1}{2} \right) + \mathbb{E}_{\mu_T|D}[\eta_a - s] \\ &= \mathbb{E}_{\mu_T|D}[\eta_a] - \frac{1}{2} = \bar{a} - \frac{1}{2}, \end{aligned}$$

so  $\text{sign } \Delta = \text{sign}(\bar{a} - \frac{1}{2}) = \text{sign}(M + \gamma)$ .  $\square$

**Theorem 1** (Interior matched-evidence impossibility). *Fix  $\beta > 0$ . For any feasible margin  $M \in [-\frac{1}{2}, \frac{1}{2}]$  with  $|M| < \beta$  there exist  $P_T^1, P_T^2 \in \mathcal{C}_\beta$  such that  $\text{Law}(\tilde{Z} \mid P_T^1) = \text{Law}(\tilde{Z} \mid P_T^2)$  for  $\tilde{Z} = (Z, M)$ ,  $M(P_T^1) = M(P_T^2) = M$ , and  $\text{sign } \Delta_1 = -\text{sign } \Delta_2 \neq 0$ . (For  $\beta = 0$  the class  $\mathcal{C}_0$  forces  $\gamma = 0$ , so  $\text{sign } \Delta = \text{sign } M$  is determined and no such pair exists: matched-evidence ambiguity requires a positive declared residual bound.)*

*Proof.* This is Lemma 1 read at one radius. Choose  $0 < \delta < \min\{\beta - |M|, \frac{1}{2}\}$  and apply the lemma with the two constant kernels  $\eta \equiv \frac{1}{2} + \theta\delta$ ,  $\theta \in \{+1, -1\}$ , holding the off- $D$  kernel  $K^\circ$  common to both worlds. The lemma supplies the two laws  $P_T^\theta$ , and: *class membership*,  $\bar{a}_\theta = \frac{1}{2} + \theta\delta$  so  $\gamma_\theta = \theta\delta - M$  with  $|\gamma_\theta| \leq \delta + |M| < \beta$ , hence both worlds lie in  $\mathcal{C}_\beta$ ; *matched evidence*,  $\text{Law}(W \mid P_T^\eta)$  does not depend on  $\eta$  by the lemma and  $\tilde{Z} = (Z, M)$  is a function of  $W$ , so  $\text{Law}(\tilde{Z} \mid P_T^+) = \text{Law}(\tilde{Z} \mid P_T^-)$ ; *opposite signs*, Lemma 2 gives  $\Delta_\theta = 2\mu_T(D)\theta\delta \neq 0$  with  $\text{sign } \Delta_\theta = \theta$ .  $\square$

The two evidence laws here coincide *exactly*, not merely approximately: the construction changes only the label kernel on  $D$ , and  $Z$  is a function of unlabeled inputs and model outputs. We therefore do not present this as a Le Cam two-point argument. A Le Cam argument trades off a positive total-variation distance against the achievable error; here the distance is zero, which is the degenerate and strongest case, and the corresponding statement is exact non-identifiability rather than a rate.

**Corollary 1** (Matched-evidence abstention lower bound). *Fix  $\alpha \in (0, \frac{1}{2})$ . On the shared evidence law constructed in Theorem 1, any possibly randomized label-free rule  $g$  satisfying  $\mathbb{P}[g = \text{ADAPT}, \Delta \leq 0] \leq \alpha$  in every admissible world and  $\mathbb{P}[g = \text{FREEZE}, \Delta \geq 0] \leq \alpha$  in every admissible world must satisfy  $\mathbb{P}[g = \text{ABSTAIN}] \geq 1 - 2\alpha$ .*

*Proof.* The rule has the same action probabilities in the two evidence-identical worlds. Its adapt probability is at most  $\alpha$  in the negative world and its freeze probability is at most  $\alpha$  in the positive world. The three action probabilities sum to one.  $\square$

## 4.4 Boundary Case

**Proposition 1** (Boundary case and closed-band abstention). *Fix  $\alpha \in (0, \frac{1}{2})$ . Under Definition 3, abstention is the pointwise maximal sound three-way action throughout the closed band  $|M| \leq \beta$ , and any label-free rule controlling both marginal errors at level  $\alpha$  abstains there with probability at least  $1 - 2\alpha$ . Opposite nonzero signs are asserted only in the interior  $|M| < \beta$ .*

*Proof.* For  $|M| < \beta$  this is Theorem 1 together with Corollary 1. For  $|M| = \beta > 0$ , the declared class contains an augmented-evidence-identical zero-benefit world (take  $\gamma = -M$ , admissible since  $|\gamma| = \beta$ ) and a strict-benefit world of sign  $M$  (take  $\gamma = 0$ ). In the zero-benefit world  $\Delta = 0$ , so *both* directional error events  $\{\text{ADAPT}, \Delta \leq 0\}$  and  $\{\text{FREEZE}, \Delta \geq 0\}$  occur whenever the rule commits; since the rule's action probabilities are identical across the two evidence-identical worlds, dual control at level  $\alpha$  gives  $\mathbb{P}[\text{ADAPT}] \leq \alpha$  and  $\mathbb{P}[\text{FREEZE}] \leq \alpha$ , hence  $\mathbb{P}[\text{ABSTAIN}] \geq 1 - 2\alpha$ , and neither strict action is uniformly directionally sound. For  $M = \beta = 0$  the class  $\mathcal{C}_0$  forces  $\gamma = 0$ , so  $\Delta \equiv 0$  and the same conclusion holds directly.  $\square$

This directional convention is stricter than task regret: at  $\Delta = 0$ , adapt and freeze are risk-equivalent but neither certifies a strict sign.

#### 4.5 Exact Strict-Commitment Frontier

**Theorem 2** (Exact strict-commitment frontier). *Under Assumption 1, fix the observables  $(\mu_T, P_S, f_0, f_a, s)$  so  $M$  is determined, and let  $\beta \geq 0$ . The necessity and maximality clauses use the full class  $\mathcal{C}_\beta$  and therefore its richness: it contains the evidence-identical label kernels constructed in Theorem 1 and the zero-benefit boundary law. For a narrower declared subclass not closed under these constructions, only the sufficiency clause follows without an additional richness assumption.*

- (i) **Reduction.** *For every  $P_T$ ,  $\text{sign } \Delta = \text{sign}(M + \gamma)$  (Lemma 2).*
- (ii) **Sufficiency.** *If  $|M| > \beta$  and  $P_T \in \mathcal{C}_\beta$ , then  $\text{sign } \Delta = \text{sign } M$ .*
- (iii) **Necessity.** *If  $|M| < \beta$  (which requires  $\beta > 0$ ), there exist  $P_T^1, P_T^2 \in \mathcal{C}_\beta$  with the same augmented-evidence law and opposite nonzero benefit signs; at  $|M| = \beta > 0$  the admissible residual  $\gamma = -\text{sign}(M)\beta$  yields  $\Delta = 0$ , and in the degenerate case  $M = \beta = 0$  the class  $\mathcal{C}_0$  forces  $\Delta \equiv 0$ . Hence no strict benefit direction is uniformly supported by  $\text{Law}(\tilde{Z})$  over  $\mathcal{C}_\beta$  whenever  $|M| \leq \beta$ .*
- (iv) **Strict-commitment iff.** *At the fixed augmented-evidence law, a strict action in  $\{\text{ADAPT}, \text{FREEZE}\}$  is uniformly directionally sound in the sense of Definition 3 if and only if  $|M| > \beta$ ; on the closed band  $|M| \leq \beta$  (including  $M = \beta = 0$ , where  $\Delta \equiv 0$ ) abstention is the pointwise maximal sound action.*

*The population rule  $M > \beta \Rightarrow \text{ADAPT}$ ,  $M < -\beta \Rightarrow \text{FREEZE}$ ,  $|M| \leq \beta \Rightarrow \text{ABSTAIN}$  is the pointwise maximal sound certificate on  $\mathcal{C}_\beta$ .*

*Proof.* Part (i) is Lemma 2. For (ii), if  $M > \beta$  and  $|\gamma| \leq \beta$ , then  $M + \gamma > 0$ , so  $\text{sign } \Delta = \text{sign}(M + \gamma) = +1$ ; the negative case is symmetric. For (iii), Theorem 1 supplies two laws in the same augmented-evidence fibre with opposite nonzero signs whenever  $|M| < \beta$ . At  $|M| = \beta$  the boundary law with  $\gamma = -\text{sign}(M)\beta$  gives  $\Delta = 0$ , still precluding a determined nonzero sign. Part (iv) combines (ii), (iii), and Proposition 1. Pointwise maximality follows because (ii) certifies the unique strict direction outside the band, while Theorem 1 and Proposition 1 rule out both strict directions inside and on the boundary.  $\square$

The feasible range can also be read directly from the identified benefit set. Write  $d = \mu_T(D) > 0$ . Over the full declared class with the observables fixed,

$$\Delta(\mathcal{C}_\beta) = 2d \left[ \max\{-\tfrac{1}{2}, M - \beta\}, \min\{\tfrac{1}{2}, M + \beta\} \right].$$

Indeed,  $\bar{a} - \frac{1}{2} \in [-\frac{1}{2}, \frac{1}{2}]$  and  $|\bar{a} - \frac{1}{2} - M| \leq \beta$ ; every value in their intersection is attained by a constant correctness kernel on  $D$ . A strict direction is available exactly when this whole interval lies above or below zero.

| Region            | Meaning                        | Action  |
|-------------------|--------------------------------|---------|
| $M > \beta$       | strictly positive benefit      | ADAPT   |
| $M < -\beta$      | strictly negative benefit      | FREEZE  |
| $ M  < \beta$     | both strict signs are feasible | ABSTAIN |
| $ M  = \beta > 0$ | zero benefit is feasible       | ABSTAIN |
| $M = \beta = 0$   | benefit is identically zero    | ABSTAIN |

#### 4.6 Interpretation and Limits of the Frontier

The theorem depends on a declared class. Choosing  $\beta = 0$  assumes a zero target disagreement-conditional residual; it is not a conservative default. Without a credible upper bound on  $|\gamma|$ , the




Figure 2: Population strict-commitment frontier.  $M > \beta$  supports adapt;  $M < -\beta$  supports freeze; and  $|M| \leq \beta$  supports no uniform strict commitment. Interior ambiguity uses opposite nonzero benefit signs, whereas the boundary uses zero-benefit ambiguity. This population rule is not the numerical KGA rule used on the real-data panels.

population frontier cannot support a robust strict action. More unlabeled samples do not solve this problem because the missing difference lies in the target labels, not in the input batch.

Empirical abstention is not automatically evidence of structural unknowability. A wide finite-sample interval may instead reflect limited calibration data, estimator error, or transfer failure. The frontier therefore leaves one central question: can its required residual bound be learned from the same label-free evidence?

## 5 What Label-Free Evidence Cannot Learn

### 5.1 Audit Setup and Evidence Fibres

Let  $W$  collect everything a deployment-time audit may read: the unlabeled target batch, the fixed models and score function, label-free statistics derived from them, and any source or calibration sample whose law is specified independently of the target label kernel. An audit is a nonnegative random quantity  $\hat{\beta} = g(W, U)$ , where  $U$  is independent randomization.

For a declared class  $\mathcal{C}$  and an observable law  $z = \text{Law}(W)$ , define the evidence fibre

$$\mathcal{C}(z) = \{P \in \mathcal{C} : \text{Law}(W \mid P) = z\}$$

and its unresolved residual radius

$$\Gamma_z(\mathcal{C}) = \sup_{P \in \mathcal{C}(z)} |\gamma(P)|. \quad (6)$$

The fibre collects target laws that the audit cannot distinguish. Its radius is therefore the smallest uniform residual bound that those observations could support.

### 5.2 Exact Fibre-Wise Audit Floor

**Theorem 3** (Exact label-free audit floor). *Fix a nonempty fibre  $\mathcal{C}(z)$  and  $\delta \in [0, 1)$ . If  $\hat{\beta}$  satisfies  $\mathbb{P}_P(|\gamma(P)| \leq \hat{\beta}) \geq 1 - \delta$  for every  $P \in \mathcal{C}(z)$ , then*

$$\mathbb{P}(\hat{\beta} \geq \Gamma_z(\mathcal{C})) \geq 1 - \delta.$$

The constant audit  $\hat{\beta} \equiv \Gamma_z(\mathcal{C})$  is valid with  $\delta = 0$ . Consequently, on the high-probability event above, an audited frontier commits only where the frontier using the declared radius  $\Gamma_z(\mathcal{C})$  already commits.

*Proof.* Every law in the fibre induces the same law for  $W$ , and hence the same law for  $\hat{\beta}$ . Validity for each fixed  $P$  requires that common law to exceed  $|\gamma(P)|$  with probability at least  $1 - \delta$ . Taking an increasing sequence of attainable absolute residuals approaching the supremum gives the lower bound. The constant audit is valid by the definition of the supremum. If  $\hat{\beta} \geq \Gamma_z$ , then  $|M| > \hat{\beta}$  implies  $|M| > \Gamma_z$ .  $\square$

This result does not say that useful budgets never exist. It says that the same label-free fibre cannot create information about the target label kernel that the fibre definition removes. A valid audit must cover the unresolved radius; reading more functions of the same  $W$  does not shrink it. The constant  $\Gamma_z(\mathcal{C})$  is an oracle benchmark for the fixed evidence-law fibre. The theorem does not provide an algorithm that learns the unknown law  $z$  or its radius from a finite observed batch. We therefore describe this as a fibre-wise audit floor rather than an unspecified minimax estimation result.

### 5.3 Declared Budget and Abstention Radius

For the full declared class  $\mathcal{C}_\beta = \{P : |\gamma(P)| \leq \beta\}$  with  $0 \leq \beta \leq \frac{1}{2}$ , the label-kernel construction in Lemma 1 gives

$$\Gamma_z(\mathcal{C}_\beta) = \beta. \quad (7)$$

Thus the least fibre-wide budget a label-free audit can certify is the budget already declared by the class. The same number is also the population abstention radius:

|                                                                                   |
|-----------------------------------------------------------------------------------|
| least valid fibre-wide residual bound<br>= population abstention radius = $\beta$ |
|-----------------------------------------------------------------------------------|

This equality connects the audit problem to the strict-commitment frontier. It does not turn the declaration into an estimate.

### 5.4 Why Labels at the Wrong Domain Buy Nothing

Suppose the audit also observes a fully labeled calibration domain, but the declared class contains no coupling between that domain's label kernel and the deployment label kernel. The construction in Lemma 1 changes only deployment labels on  $D$  and leaves the calibration sample unchanged. The augmented evidence law is therefore identical across the same target fibre, so Theorem 3 applies without a smaller radius. Labels help only when they come from the deployment domain or when a stated assumption links the labeled and deployment domains.

### 5.5 What Additional Structure Can Buy

The missing information can be supplied, but each route changes the problem. Deployment labels can estimate target correctness directly. Historical labeled deployments can support an episode-marginal budget under an exchangeability model. A structural coupling can bound the change between calibration and target domains. Domain knowledge can declare a physically meaningful limit. These are purchases made with labels or assumptions; they are not consequences of the current unlabeled batch. Detailed rate results and extended constructions belong to the long-form theory record and are not required by the empirical claims here.

Table 4: Correspondence between population theory (K-Bound) and the empirical decision rule (KGA). A shared target functional does not make the quantities equal.

| Quantity       | Role in the Population Frontier                                       | Status in Empirical KGA                           |
|----------------|-----------------------------------------------------------------------|---------------------------------------------------|
| $M$            | Observable margin on disagreement region                              | Not explicitly estimated                          |
| $\gamma$       | Hidden calibration residual ( $M + \gamma = \bar{a} - \frac{1}{2}$ )  | Unobserved / unestimated                          |
| $\beta$        | Externally declared uniform bound on $ \gamma $                       | Not a numerical KGA input                         |
| $\Delta$       | Target risk reduction $2\mu_T(D)(M + \gamma)$                         | Estimand (cell or population)                     |
| $\hat{\Delta}$ | Theoretical regression function $\mathbb{E}[\Delta   Z]$              | Fitted prediction $h_\theta(Z)$ (empirical proxy) |
| $\varepsilon$  | Cell residual radius covering $ \hat{\Delta} - \Delta^{\text{cell}} $ | Calibrated conformal quantile                     |
| $b(m, \delta)$ | Hoeffding concentration radius for fresh $m$ draws                    | Population-transfer correction                    |

## 5.6 Implications for K-Bound

KGA does not estimate  $\beta$  and does not numerically apply  $|M| > \beta$  on the reported real-data benchmarks. Its practical certificate instead places a finite-sample uncertainty radius  $\varepsilon$  around the predicted cell benefit  $\hat{\Delta}$ . Population benefit  $\Delta$  and observed cell benefit  $\Delta^{\text{cell}}$  concern the same model pair at different levels. A population application would instead require coverage of the target population through a separate sampling-error bound. No such sampling-error bound is established for the reported panels.

## 6 KGA: An Empirical Benefit-Interval Gate

### 6.1 What Transfers from the Population Analysis

KGA is an empirical companion to K-Bound, not an implementation of its population frontier. Both ask whether to accept the same proposed update and both leave unresolved evidence as abstention. Their uncertainty objects differ. The binary population model gives

$$\Delta = 2\mu_T(D)(M + \gamma). \quad (8)$$

The observed cell outcome is

$$\Delta^{\text{cell}} = S(f_a; \mathcal{E}) - S(f_0; \mathcal{E}),$$

and the fitted empirical benefit prediction is

$$\hat{\Delta} = h_\theta(Z).$$

KGA does not estimate  $M$ ,  $\gamma$ , or  $\beta$ , and does not numerically apply  $|M| > \beta$ . Its radius  $\varepsilon$  describes prediction error for a measured cell outcome, not the range of population benefits compatible with a label-free evidence law.

Connecting the two formulations proceeds not through an automatic deductive chain, but through an explicit conditional concentration bridge:

$$\begin{aligned} \mathbb{P}(|\hat{\Delta} - \Delta^{\text{cell}}| \leq \varepsilon) &\geq 1 - \alpha_{\text{cell}}, \\ \mathbb{P}(|\Delta^{\text{cell}} - \Delta| \leq b(m, \delta)) &\geq 1 - \delta. \end{aligned}$$

The cell-coverage and fresh-sample concentration premises imply, by a union bound,

$$\begin{aligned} \mathbb{P}(|\hat{\Delta} - \Delta| \leq \varepsilon + b(m, \delta)) \\ \geq 1 - \alpha_{\text{cell}} - \delta. \end{aligned}$$

K-Bound establishes what an entire declared family of target worlds permits: if every allowed drift in  $\mathcal{C}_\beta$  preserves one strict direction ( $|M| > \beta$ ), that direction is population sound; otherwise the maximal uniform rule must abstain. KGA addresses the realized deployment unit under a separately declared operational transfer model: it predicts the empirical risk difference  $h_\theta(Z)$  and commits only when a calibrated interval excludes zero. The common object is  $\Delta$ ; the population theorem brackets it through  $(M, \gamma, \beta)$ , whereas KGA brackets it through  $(\hat{\Delta}, \varepsilon)$  at the cell level, and through  $(\hat{\Delta}, \varepsilon + b)$  at the population level via Proposition 3 (Theorem 4 in Appendix A.4). Table 4 makes this correspondence explicit.

KGA is label-free at deployment, not during development and calibration. Historical labeled cell outcomes fit  $h_\theta$  and supply residuals. At the new decision, KGA reads only  $Z$  and the frozen fitted artifacts. The interval premise for the empirical target is

$$\mathbb{P}(|\hat{\Delta} - \Delta^{\text{cell}}| \leq \varepsilon) \geq 1 - \alpha. \quad (9)$$

The proposition below concerns this target when this premise holds.

## 6.2 The Coverage-to-Action Proposition

The following standard implication makes the decision semantics explicit. It is not a new calibration theorem: coverage is the premise, and establishing it is the statistical burden.

**Proposition 2** (Coverage-to-action implication). *Fix the scalar benefit target and its evaluation unit before calibration. Let  $B$  and  $\hat{B}$  be finite real random variables, and let  $\varepsilon \in [0, \infty]$  be a measurable random radius on a common probability space containing all calibration, benefit-model fitting, candidate-construction, and evaluation-unit randomness. The prediction  $\hat{B}$  and radius  $\varepsilon$  must be available without reading the new unit’s evaluation labels. For  $\alpha \in (0, 1)$ , suppose the interval satisfies marginal coverage*

$$\mathbb{P}[|\hat{B} - B| \leq \varepsilon] \geq 1 - \alpha.$$

Define

$$g = \begin{cases} \text{ADAPT} & \hat{B} - \varepsilon > 0, \\ \text{FREEZE} & \hat{B} + \varepsilon < 0, \\ \text{ABSTAIN} & \text{otherwise.} \end{cases}$$

Then, marginally over that common probability space,

$$\begin{aligned} \mathbb{P}[g = \text{ADAPT}, B \leq 0] &\leq \alpha, \\ \mathbb{P}[g = \text{FREEZE}, B \geq 0] &\leq \alpha. \end{aligned}$$

*Proof.* On the event  $\{|\hat{B} - B| \leq \varepsilon\}$ , the condition  $\hat{B} - \varepsilon > 0$  implies  $B > 0$ , so a false adapt requires the complement of this event, which has probability at most  $\alpha$ . The freeze case is symmetric.  $\square$

*Remark 2* (Unit scaling: empirical radius versus theoretical budget). The empirical radius  $\varepsilon$  and theoretical residual bound  $\beta$  live on fundamentally different scales. The population identity is  $\Delta = 2d(M + \gamma)$  where  $d = P_T(D) > 0$  is the disagreement probability and  $2d$  is twice that probability. Thus,  $|\Delta - 2dM| = 2d|\gamma| \leq 2d\beta$ . The unclipped population envelope centered at  $2dM$  has radius  $2d\beta$  on the benefit scale, not  $\beta$ ; the exact population benefit set  $\Delta(\mathcal{C}_\beta)$  in Theorem 2 includes clipping to  $[-d, d]$ :  $\Delta(\mathcal{C}_\beta) = 2d[\max(-\frac{1}{2}, M - \beta), \min(\frac{1}{2}, M + \beta)]$ . Equating  $\varepsilon = \beta$  is therefore a conceptual category error: it conflates a conditional score residual on the disagreement region with an empirical cell-benefit prediction error. While numerical equality could occur by coincidence, their mathematical definitions remain distinct.

**Proposition 3** (Cell-to-population transfer via compound bound). *Let  $\Delta^{\text{cell}} = \frac{1}{m} \sum_{j=1}^m W_j$  be the measured cell benefit on  $m$  fresh, conditionally i.i.d. evaluation draws, where  $W_j = \mathbf{1}\{f_a(X_j) = Y_j\} - \mathbf{1}\{f_0(X_j) = Y_j\} \in [-1, 1]$ , with conditional mean  $\Delta$ . Suppose an independently fitted benefit predictor  $h$  and split-conformal radius  $\varepsilon$  achieve marginal cell coverage  $\mathbb{P}(|\hat{\Delta} - \Delta^{\text{cell}}| \leq \varepsilon) \geq 1 - \alpha_{\text{cell}}$  over exchangeable deployment episodes. For a sampling budget  $\delta \in (0, 1)$ , define  $b(m, \delta) = \min\{2, \sqrt{2 \log(2/\delta)/m}\}$ . Then the compound interval  $I_{\text{population}} = [\hat{\Delta} - (\varepsilon + b), \hat{\Delta} + (\varepsilon + b)]$  satisfies*

$$\mathbb{P}(\Delta \in I_{\text{population}}) \geq 1 - \alpha_{\text{cell}} - \delta.$$

The corresponding population decision rule  $g_{\text{pop}}$  using radius  $\varepsilon + b$  satisfies

$$\begin{aligned} &\mathbb{P}(\{g_{\text{pop}} = \text{ADAPT}, \Delta \leq 0\} \cup \\ &\{g_{\text{pop}} = \text{FREEZE}, \Delta \geq 0\}) \leq \alpha_{\text{cell}} + \delta. \end{aligned}$$

*Proof.* By Hoeffding’s inequality for bounded variables in  $[-1, 1]$  of range 2,  $\mathbb{P}(|\Delta^{\text{cell}} - \Delta| \leq b) \geq 1 - \delta$ . On the intersection of cell coverage and concentration, the triangle inequality yields  $|\hat{\Delta} - \Delta| \leq |\hat{\Delta} - \Delta^{\text{cell}}| + |\Delta^{\text{cell}} - \Delta| \leq \varepsilon + b$ . A union bound proves coverage at level  $1 - \alpha_{\text{cell}} - \delta$ . Event containment (as in Proposition 2) immediately bounds the probability of an incorrect strict direction by  $\alpha_{\text{cell}} + \delta$ .  $\square$

*Remark 3* (Direct empirical sampling of the population frontier). An alternative route directly operationalizes the population frontier  $|M| > \beta$ . Given fresh unlabeled target inputs, let  $n_D > 0$  fall in  $D$ . The empirical disagreement margin  $\widehat{M} = \frac{1}{n_D} \sum_{X_i \in D} s(X_i) - \frac{1}{2}$  concentrates with radius  $r_M = \sqrt{\log(2/\alpha)/(2n_D)}$ . Under the externally declared bound  $|\gamma| \leq \beta$ , the rule  $\widehat{M} - r_M > \beta \Rightarrow \text{ADAPT}$ ,  $\widehat{M} + r_M < -\beta \Rightarrow \text{FREEZE}$ , and  $\text{ABSTAIN}$  otherwise, guarantees directional error at most  $\alpha$ . This couples sampling uncertainty  $r_M$  additively with structural uncertainty  $\beta$ .

*Remark 4* (Marginal vs. conditional error). Proposition 2 bounds the *marginal, unconditional* false-commit probabilities for the declared target  $B$ . Its false-adapt rate is  $\text{FA}_u = \mathbb{P}[g = \text{ADAPT}, B \leq 0]$ . It does *not* bound  $\text{FA}_c = \mathbb{P}[B \leq 0 \mid g = \text{ADAPT}]$ , defined when  $\mathbb{P}[g = \text{ADAPT}] > 0$  and reported descriptively for  $B = \Delta^{\text{cell}}$ . Selection-conditional coverage requires a separate construction, such as that of Jin and Ren [52]; conditional guarantees over specified shift classes also require their own calibration procedure [53]. Neither is supplied by the marginal premise above. Simultaneous protection across candidates and time-uniform protection likewise require additional guarantees. Risk alignment (Definition 2) concerns population sign identifiability; coverage for  $B$  alone yields the *interval* implication above.

*Remark 5* (Fallback when assumptions are unsupported). If interval coverage for the declared target and evaluation unit—or the assumptions supporting it—cannot be justified, Proposition 2 supplies no level- $\alpha$  protection for either strict action. A separately declared operational fallback may retain the frozen model and record abstention. Serving the frozen model without a valid negative interval is not a certified FREEZE decision. This policy recommendation does not establish that the retained model is safe or optimal. If *risk alignment* (Definition 2) cannot be justified for an intended population sign claim, that interpretation remains unsupported. The interval implication of Proposition 2 still holds when coverage for the declared  $B$  is established (cf. Theorem 1).

Proposition 3 formalizes the conditional bridge to population risk  $\Delta$ : under independent fresh evaluation draws, adding the Hoeffding concentration radius  $b(m, \delta)$  to the empirical radius  $\varepsilon$  yields a compound population interval with radius  $\varepsilon + b$  and error budget  $\alpha_{\text{cell}} + \delta$ . The full protocol, counterexample, and proof are detailed in Appendix A. No such sampling-error bound is established

for the reported benchmark panels; in particular, when a candidate adapts to evaluation inputs, transductive normalization couples the candidate across batches.

The deterministic action layer for this bridge is now kernel-checked in `docs/research/kbound/formal/KBound/Population.lean`: an outcome-blind estimate, cell radius, and independently justified sampling radius define the compound interval, and strict ADAPT/FREEZE actions are sound whenever the target benefit is inside that interval. The companion replay interface is `docs/research/kbound/scripts/run_population_panel.py`; it rejects outcome-dependent inputs and records the input hash. This procedure is available for a benchmark panel. The historical canonical benchmark records do not contain a sealed per-cell tuple of fitted prediction, cell radius, and fresh-sample size, so no population-transfer number is promoted into those tables.

**Task 2 theorem–code reconciliation.** The formal map distinguishes the population and cell-level objects rather than treating the implementation as a proof of the transfer premises. In particular, `formal/KBound/Probability/HistoricalExtension.lean` proves `historical_one_bit_decoder_iff_fibre_consistent`: a label-free one-bit decoder exists only when the label is constant on each observable evidence fibre. The stronger historical claim that one representative per swap orbit is sufficient is refuted by the checked counterexample in `formal/KBound/Probability/ChannelCounterexample.lean`, not silently promoted to a theorem. The former module, `HistoricalExtension.lean`, formalizes `h_ratio_rate_transfer` only under an explicit episode-level contract ( $\text{miss}_T \leq H \text{miss}_S$ ,  $H \geq 0$ , and  $\text{miss}_S \leq \alpha$ ); it does not estimate  $H$ , prove benchmark exchangeability, or establish prospective population provenance. Thus K-Bound’s population interval  $(M, \gamma, \beta)$  and KGA’s fitted cell interval  $(\hat{\Delta}, \varepsilon)$  remain distinct, with the population procedure adding the separately justified sampling radius  $b(m, \delta)$ .

**ImageNet-C 108-cell population replay.** To audit this procedure on existing benchmark data without changing historical estimands, we replayed the outcome-blind compound interval  $[\hat{\Delta} \pm (\varepsilon + b)]$  across 108 ImageNet-C development cells (36 cells each for Tent, EATA, and SAR; five-fold cell-outcome-disjoint cross-fitting; sensitivity parameter  $m = 2,000$ , yielding uniform concentration radius  $b = 0.0607$ ; Table 5). Full fold breakdowns and radius distributions appear in Table 14 (Appendix A.7). Finite radii were obtained on 108/108 cells. Across the panel, the cell-level gate  $(\hat{\Delta} \pm \varepsilon)$  issued 40 ADAPT, 11 FREEZE, and 57 ABSTAIN decisions with zero false adaptations ( $\text{FA}_u = 0.0000$ ) and mean oracle regret 0.0167. Incorporating the sampling radius ( $b = 0.0607$ , with mean compound half-width  $\bar{\varepsilon} + b = 0.1272$ ) shifted compound decisions to 1 ADAPT, 10 FREEZE, and 97 ABSTAIN, preserving strictly zero false adaptations ( $\text{FA}_u = 0.0000$ ) while shifting regret to 0.0374. This concretely demonstrates the sensitivity to the added concentration radius  $b = 0.0607$ : because most empirical gains on this benchmark are smaller than the universal Hoeffding penalty for  $m = 2,000$ , the compound rule withholds action on nearly all adaptations, illustrating the numerical Price of Rigor rather than establishing prospective population truth. As clarified in Appendix A.7,  $m = 2,000$  is an evaluation sensitivity parameter, not fresh sampling; population truth remains unestablished on historical development records.

**Non-oracle population diagnostic.** While the synthetic diagnostic in Table 13 tested exact Bayes prediction ( $\hat{\Delta}(Z) = 0.6Z$ ), evaluating an imperfect predictor subject to estimation noise requires an extended non-oracle study. We conducted a diagnostic across 1,000 independent trials per configuration over 16  $(N_{\text{cal}}, m)$  pairs, fitting an imperfect linear predictor on finite development episodes ( $m = 32$ ; Table 12, Appendix A.7). An unhedged point rule ( $\text{sign}(\hat{\Delta})$ ) incurs a 30.9%–35.6% wrong-direction error rate ( $\text{FA}_u$ ) with 0% abstention; crucially, 100% of these errors occur on null/tie

Table 5: 108-cell ImageNet-C population replay: cell gate ( $\hat{\Delta} \pm \varepsilon$ ) vs. compound population interval ( $\hat{\Delta} \pm (\varepsilon + b)$ ) under 5-fold cross-fitting ( $\alpha_{\text{cell}} = \delta = 0.05, \alpha_{\text{pop}} = 0.10, m = 2,000, b = 0.0607$ ). Full 12-column breakdown with fold details and radii ranges in Table 14.

| Candidate                | $N$ | Regret (Cell / Pop)    | Cell Gate (A / F / U) | Compound Pop (A / F / U) | $\text{FA}_u$ |
|--------------------------|-----|------------------------|-----------------------|--------------------------|---------------|
| <b>Tent</b>              | 36  | 0.0162 / 0.0468        | 18 / 0 / 18           | 1 / 0 / 35               | 0.000         |
| <b>EATA</b>              | 36  | 0.0095 / 0.0386        | 20 / 0 / 16           | 0 / 0 / 36               | 0.000         |
| <b>SAR</b>               | 36  | 0.0245 / 0.0267        | 2 / 11 / 23           | 0 / 10 / 26              | 0.000         |
| <b>Total / Panel 108</b> |     | <b>0.0167 / 0.0374</b> | <b>40 / 11 / 57</b>   | <b>1 / 10 / 97</b>       | <b>0.000</b>  |

states ( $Z = 0, \Delta = 0$ ) due to estimation noise around zero, with zero errors on strictly harmful states ( $\Delta = -0.6$ ). Conformal cell calibration ( $\hat{\Delta} \pm \varepsilon_{\text{cell}}$ ) eliminates wrong-direction error (0.000) with empirical cell coverage matching nominal  $1 - \alpha_{\text{cell}} = 0.95$  (93.9%–96.5%), and already achieves  $\approx 100\%$  population coverage of  $\Delta$  across all configurations in this model. The compound population interval  $[\hat{\Delta} \pm (\varepsilon_{\text{cell}} + b)]$  likewise achieves 100% empirical population coverage and strictly zero wrong-direction error ( $0.000 \leq \alpha_{\text{pop}} = 0.10$ ), while enforcing conservative abstention (99.9%–100% at  $m = 32$ , contracting safely to  $\approx 33\%$  matching the null frequency for  $m \geq 128$ ) as a direct consequence of the universal Hoeffding penalty ( $b = 0.4802$ ), rather than because cell-level calibration failed in this simulated environment.

### 6.3 Evidence, Calibration, and the Action Rule

The controlled implementation stores 11 label-free features, with effective rank 9 because two features are deterministic differences. They summarize confidence, class balance, model disagreement, and update size. For each dataset and adapter, a separate `GradientBoostingRegressor` [33] uses squared-error loss, 250 trees, depth 2, learning rate 0.05, subsample 0.8, and seed 0. Natural protocols retain their own fitted model and feature schema; for example, CCT-20 uses a ridge model. Details are in Appendix I. A common interface does not imply that an estimator or radius transfers across adapters.

Given residuals  $r_i = |\hat{\Delta}_i - \Delta_i^{\text{cell}}|$ , ordered as  $r_{(1)} \leq \dots \leq r_{(n)}$ , the order-statistic rule is

$$k = \lceil (n+1)(1-\alpha) \rceil, \quad \varepsilon = \begin{cases} r_{(k)}, & k \leq n, \\ +\infty, & k > n. \end{cases} \quad (10)$$

Split-conformal coverage requires joint exchangeability of the calibration and next-unit scores under the fitting and candidate-construction protocol [36, 37]. An independently fitted predictor and exchangeable calibration/deployment units are one sufficient route. The rank counts evaluation units, not the images inside them.

We distinguish two calibration protocols across the paper: *Protocol B: Three-way cell-outcome-disjoint cross-fitting* (`kga.crossfit.controlled_grid_crossfit`, used in the current primary CIFAR table, interval diagnostics, the 108-cell ImageNet-C population replay, and the Appendix B.3 27-cell audit). For each score fold, its entire complement is partitioned into strictly disjoint estimator-fit and residual-calibration subsets, ensuring that scored cell outcomes enter neither predictor training nor conformal calibration pools. Historical saved LOO predictions and radii do not determine these current CIFAR results. *Protocol A: Leave-one-cell-out empirical order-statistic calibration* (used in earlier historical controlled grids). Benefit predictions  $\hat{\Delta}_i$  were cross-fitted on the remaining  $n-1$  cells, and each scored residual  $r_i = |\hat{\Delta}_i - \Delta_i^{\text{cell}}|$  was excluded from the pool used to compute its

---

**Algorithm 1** KGA for one proposed update

---

**Require:** Frozen  $f_0$ , adapter  $\mathcal{A}$ , unlabeled  $X$ , frozen  $h_\theta$  and calibration artifacts

- 1: Create  $f_a \leftarrow \mathcal{A}(f_0, X)$  in isolated shadow state
  - 2: Extract  $Z = \phi(X, f_0, f_a)$  and validate the assessment inputs
  - 3: **if** the assessment is unavailable **then**
  - 4:     retain  $f_0$ , log the reason, and return ABSTAIN
  - 5: **end if**
  - 6: Obtain  $\varepsilon$  from the calibrated residual rule in Eq. (10)
  - 7:  $\hat{\Delta} \leftarrow h_\theta(Z)$ ;  $L \leftarrow \hat{\Delta} - \varepsilon$ ;  $U \leftarrow \hat{\Delta} + \varepsilon$
  - 8: **if**  $L > 0$  **then**
  - 9:     commit  $f_a$  and return ADAPT
  - 10: **else if**  $U < 0$  **then**
  - 11:     discard  $f_a$  and return FREEZE
  - 12: **else**
  - 13:     retain  $f_0$ , log uncertainty, and return ABSTAIN
  - 14: **end if**
- 




Figure 3: Illustration of the three interval decisions for measured cell benefit. Only an interval strictly above or below zero supports a commitment. The figure illustrates the rule; coverage must be established for the intended evaluation unit.

own radius  $\varepsilon_i$ . However, this historical leave-one-out construction is peer-dependent: for  $j \neq i$ , the model  $h_{-j}$  was fitted on training data containing  $\Delta_i^{\text{cell}}$ , so the scored cell’s outcome could indirectly influence other residuals  $r_j$  in its calibration pool through model fitting. Where historical stress grids retain pooled LOO residuals (Protocol A), they are reported with this peer-dependence caveat rather than claiming full outcome disjointness. Natural protocols use separated development, calibration, and target partitions where the archived design records that separation.

For example, prediction 0.04 and radius 0.01 give  $[0.03, 0.05]$  and ADAPT; prediction  $-0.04$  gives  $[-0.05, -0.03]$  and FREEZE; prediction 0.005 gives  $[-0.005, 0.015]$  and ABSTAIN. These illustrate the rule, not measured results.

Table 6: Experiment status and role. Primary presentation does not imply prospective confirmation. All recorded unfavorable outcomes remain disclosed.

| Evidence                                                    | Design and status                                                                                                                            | Question it can answer                                                                                                   |
|-------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------|
| CIFAR-10-C: Tent,<br>EATA, SAR<br>CCT-20: Tent              | opened controlled grids; cross-fitted empirical calibration<br>internally sealed protocol; five checkpoints and nine target camera locations | candidate-specific routing on the evaluated cells<br>conservative retention under the stated score and safe-utility rule |
| So2Sat-LCZ42                                                | development feasibility stopped; no target access                                                                                            | whether a candidate met the pre-target eligibility checks                                                                |
| Other replayable tracks                                     | opened natural/semantic and candidate-configuration diagnostics                                                                              | transfer limits and one-sided retention; full results in the supplement                                                  |
| PACS; iWildCam; historical ports; focused Office-Home route | partial replay, withheld metric, superseded policy, or invalid route                                                                         | audit history only where the corresponding contract is incomplete or invalid                                             |

Candidate adaptation, frozen and candidate inference, evidence extraction, and shadow-state storage all contribute to runtime. The tree prediction and interval comparison are inexpensive, but the whole procedure is not one extra forward pass. Comparable end-to-end latency measurements are not available across the recorded tracks.

## 7 Primary Experimental Protocols

### 7.1 Two Primary Questions and the Evidence Hierarchy

The controlled study asks whether the gate distinguishes useful from harmful updates. The CCT-20 study asks whether a frozen gate avoids measured candidate harm under a natural camera-location shift. These are different claims. We predefine the interpretation of each study before presenting its results in Table 6.

CIFAR-10-C [23] supplies 2,160 recorded cells per candidate: five archived ResNet-18 run-seed groups, each evaluated across 432 corruption/composition/update conditions. The current gate is refitted from the saved label-free features  $Z$  and measured cell outcomes under Protocol B, rather than reusing historical stored predictions or radii. For this canonical evaluation, the analysis uses three-way cell-outcome-disjoint cross-fitting: each scored fold has separate estimator-fit and residual-calibration subsets, so its outcome enters neither its fitted prediction nor its radius. This is still a retrospective opened panel, not an untouched environment test; historical stored fields remain audit-only and do not determine the current decisions.

### 7.2 CCT-20 Design and Safe Utility

The CCT-20 camera-location study [48] crosses 5 independent source checkpoints with 9 target locations, giving 45 checkpoint–location cells. It uses 6950 probe images and 16325 evaluation images (23275 total). The primary metric is set-membership top-1 accuracy: a prediction is correct if it matches any distinct archived category for that image. It is not an official single-label leaderboard reproduction. The candidate uses the official Tent implementation, one non-episodic Adam step at learning rate  $10^{-3}$ , and state reset per cell.

The protocol was internally sealed before source training and outcome access, not publicly registered. Aggregate target label metadata had been inspected during candidate ranking; cell-level

outcomes remained unopened until one-shot scoring. This is outcome-blind execution, not a label-unopened study. The source ridge fit used 10 rows from two camera locations; calibration used 45 rows aggregated to nine cis-test locations. Transfer to trans-test locations remains an assumption: the protocol treats camera-location clusters as the population unit and assumes the nine trans-test clusters are exchangeable with the declared target population under the locked sampling rule. This is a design assumption, not an empirically proven fact about future deployments.

For either fixed policy  $b$ , define  $G_b$  as mean baseline regret minus KGA regret. Positive  $G_b$  means lower regret for KGA. The *safe-utility* rule is a predeclared, limited harm-avoidance check. Using nominal pointwise 95% bootstrap lower endpoints, it passes when

$$L_{\text{adapt}} > 0, \quad L_{\text{freeze}} > -0.005. \quad (11)$$

The second threshold is the frozen-policy noninferiority margin, on the accuracy-proportion scale. This rule can pass when KGA always retains  $f_0$ . The stronger routing criterion additionally requires improvement over both fixed policies at the declared inference level and nontrivial ADAPT/FREEZE exposure. The word *safe* names this protocol endpoint, not a population guarantee.

### 7.3 Scoring, Dependence, and Statistical Inference

All policies use the same evaluation records, and evaluation labels are joined only after actions are fixed. The candidate TTA step can be transductive: episodic candidates update on the unlabeled evaluation inputs, and online candidates may use evaluation-batch BatchNorm statistics. Development/calibration separation therefore does not imply independent fresh-sample inference.

We report oracle regret, action counts, commitment rate, and false-commit frequencies. Empirical interval inclusion and interval width are reported separately from commitment. A paired regret contrast always means baseline regret minus KGA regret, so positive values favor KGA. Accuracy and macro-F1 regrets are separated rather than pooled.

Bootstrap levels are nominal throughout. CIFAR sensitivity resamples six corruption families, conditional on an archived checkpoint. CCT-20 uses a paired two-way product bootstrap over checkpoint rows and location columns; sign-flip calculations instead use nine location means. For a vector  $G$  of cluster gaps, the sign-flip reference model is

$$G \stackrel{d}{=} s \odot G \quad \text{for every } s \in \{-1, +1\}^m. \quad (12)$$

Independent zero-symmetric cluster gaps suffice. Exchangeability, marginal symmetry, or a single global sign symmetry alone is not sufficient [38]. Enumeration removes Monte Carlo error but does not verify this null model. Holm adjustment requires valid input  $p$ -values; nominal family intervals also require valid marginal coverage of their constituent intervals. The full inference specification and retrospective search status are collected in Appendix F.

## 8 Primary Results

### 8.1 Controlled Routing Depends on the Candidate

Table 7 reports all three CIFAR-10-C candidates under the same current rule. Tent and EATA have lower measured point regret than always-adapt and always-freeze. SAR is the counterexample: always-adapt has lower regret than KGA. The mixed panel therefore supports candidate-specific routing, not a claim that gating improves every adapter.

Table 7: Primary controlled CIFAR-10-C results. Regret is the mean accuracy shortfall from the per-cell oracle, on the proportion scale; lower is better. A/F/U denotes ADAPT/FREEZE/ABSTAIN. Commitment rate is  $(A + F)/n$ , not interval coverage. Cells share checkpoints and residual pools.

| Candidate | $n$  | KGA    | Adapt  | Freeze | A/F/U        | FA <sub>u</sub> | Commitment rate |
|-----------|------|--------|--------|--------|--------------|-----------------|-----------------|
| Tent      | 2160 | 0.0018 | 0.0080 | 0.1239 | 1094/334/732 | 0.0000          | 0.6611          |
| EATA      | 2160 | 0.0016 | 0.0033 | 0.1313 | 1207/118/835 | 0.0000          | 0.6134          |
| SAR       | 2160 | 0.0017 | 0.0003 | 0.1405 | 1430/0/730   | 0.0005          | 0.6620          |

Table 8: Retrospective interval diagnostics for the controlled cells under Protocol B (disjoint fit/calibrate/score cross-fitting), at nominal target  $1 - \alpha = 0.90$ . Observed inclusion concerns  $\Delta^{\text{cell}}$ . The same collection supplies overlapping leave-one-out residual pools, so pooled inclusion is rank-constrained and is not independent validation of next-environment coverage. Candidate/family breakdowns and directional-error denominators are provided in the supplement.

| Candidate | $n$  | Nominal | Inclusion | Mean width | Median width | Commitment | FA <sub>u</sub> /FA <sub>c</sub> | FF <sub>u</sub> /FF <sub>c</sub> |
|-----------|------|---------|-----------|------------|--------------|------------|----------------------------------|----------------------------------|
| TENT      | 2160 | 0.9000  | 0.9046    | 0.0469     | 0.0483       | 0.6611     | 0.0000/0.0000                    | 0.0000/0.0000                    |
| EATA      | 2160 | 0.9000  | 0.9060    | 0.0387     | 0.0378       | 0.6134     | 0.0000/0.0000                    | 0.0005/0.0085                    |
| SAR       | 2160 | 0.9000  | 0.8949    | 0.0279     | 0.0277       | 0.6620     | 0.0005/0.0007                    | 0.0000/-                         |

Tent regret is 0.0018, compared with 0.0080 for always-adapt and 0.1239 for always-freeze. EATA’s values are 0.0016, 0.0033, and 0.1313. Their current actions are 1,094/334/732 and 1,207/118/835, respectively; both have zero observed false adaptations. SAR makes 1,430/0/730 decisions and has regret 0.0017, versus 0.0003 for always-adapt. A gate can avoid harmful updates yet lose opportunities when it abstains.

A historical cluster-bootstrap sensitivity analysis treats each of the five run-seed groups as a resampling unit ( $m=5$  groups, 5,000 replicates). Distinct source-training checkpoint identities are not established by the inspected replay manifest; this grouping must not be interpreted as independent-model inference. Across seeds 0–4, KGA Tent achieves regret [0.00194, 0.00164, 0.00142, 0.00150, 0.00151], consistently outperforming Always Adapt (regret [0.01098, 0.01105, 0.00729, 0.00783, 0.00744], with pairwise regret reductions of +0.00903, +0.00941, +0.00587, +0.00633, +0.00593) and Always Freeze (regret [0.12164, 0.12059, 0.12365, 0.12398, 0.12457], with pairwise reductions of +0.11970, +0.11894, +0.12223, +0.12248, +0.12306). Over all seeds, KGA Tent’s regret advantage over always-adapt has a cluster-bootstrap 95% interval of [+0.006, +0.009], and over always-freeze of [+0.120, +0.123]; both exclude zero in the favorable direction, with all 5,000 bootstrap replicates exhibiting positive advantage ( $\mathbb{P}^*(\text{gap} > 0) = 1.0$ , one-sided bootstrap replication probability, empirical  $p^* < 0.0002$ ). These reported historical group-resampling intervals do not establish independent-checkpoint or confirmatory superiority. Zero false adaptations on the recorded cells (FA<sub>u</sub>=0.000) is an empirical count, not a population guarantee. Corruption-family sensitivity breakdowns appear in Table 26.

## 8.2 Observed Interval Inclusion Is a Separate Diagnostic

Commitment rate measures how often the interval excludes zero. Interval inclusion measures how often it contains the observed cell outcome. Table 8 reports these separately, together with full interval width. The calculation replays the current order-statistic rule from saved predictions and already-opened cell outcomes; it does not reuse historical stored actions or radii.

Table 9: Current outcome-disjoint saved-feature comparison on 2,160 CIFAR-10-C Tent rows, generated from one authenticated paired-score authority. All rows have the same oracle and weights; both fitted gates use identical predictions. Lower regret is better.  $FA_u$  is the fraction of all rows with ADAPT and  $B \leq 0$ ; commitment is  $(A + F)/n$ , not interval coverage. No native external method or prospective significance claim is included.

| Policy                          | Regret | Mean acc. | $FA_u$ | Commitment | ADAPT | FREEZE | ABSTAIN |
|---------------------------------|--------|-----------|--------|------------|-------|--------|---------|
| Always Adapt                    | 0.0080 | 0.7869    | 0.334  | 1.000      | 2160  | 0      | 0       |
| Always Freeze                   | 0.1239 | 0.6709    | 0.000  | 1.000      | 0     | 2160   | 0       |
| Confidence increase             | 0.0080 | 0.7869    | 0.250  | 1.000      | 1834  | 326    | 0       |
| Entropy decrease                | 0.0081 | 0.7868    | 0.249  | 1.000      | 1812  | 348    | 0       |
| Point benefit (same prediction) | 0.0004 | 0.7945    | 0.041  | 1.000      | 1471  | 689    | 0       |
| Current KGA                     | 0.0018 | 0.7931    | 0.000  | 0.661      | 1094  | 334    | 732     |

The nominal target is not a verified coverage guarantee for these dependent cells. Inclusion near 90% can follow from the rank procedure itself, even when transfer to a new family is poor. Family differences are useful diagnostics, but independent held-out-family, checkpoint, and natural-environment coverage remains a separate test.

Observed inclusion rates closely track the 90% nominal level: Tent includes 1,954 of 2,160 outcomes (90.46%), EATA includes 1,957 (90.60%), and SAR includes 1,933 (89.49%). Zero observed false adaptations do not imply zero directional errors overall: EATA commits one false FREEZE among 118 FREEZE decisions ( $FF_c = 0.0085$ ), occurring in the Gaussian-noise family (1 error among 1 FREEZE decision). Tent commits zero false FREEZE decisions among 334 FREEZE decisions ( $FF_c = 0.0000$ ), while SAR commits zero FREEZE decisions ( $FF_c$  undefined) and one false ADAPT among 1,430 ADAPT decisions ( $FA_c = 0.0007$ , occurring in pixelate). These denominators matter more than a pooled near-nominal inclusion fraction alone.

### 8.3 Synchronized Development Baseline Comparison

Table 9 is regenerated from the authenticated current CIFAR-10-C Tent authority: 432 conditions in each of five run-seed groups, with equal weight on all 2,160 rows. These are stream/adaptation repetitions, not authenticated independently trained source models. Every rule uses the same saved frozen/candidate scores and candidate-specific oracle. Confidence increase and entropy decrease are fixed scalar rules, not renamed AETTA or POEM implementations. The point-benefit rule and KGA share exactly the same fitted predictions; only the decision rule differs.

The current gate uses Protocol B: five deterministic score folds per run-seed group, with 241 or 242 estimator-fit records and 104 separate residual-calibration records per fold. At  $\alpha = 0.10$ , the calibration rank is 95. The original literal candidate/seed/condition identifiers and GBR settings are retained. A pinned-runtime replay reproduced every prediction, radius and action for all five groups. In five additional full-pipeline checks, changing the first scored cell’s candidate accuracy to zero (one intervention per group) left its prediction, radius, zero decision threshold and action unchanged. These checks begin with saved features; they are not upstream label-independence or environment-held-out tests.

KGA has lower point regret than either fixed policy on these records, but the same-prediction point gate has lower regret than KGA. The latter commits 88 non-positive adaptations out of 2,160 rows, whereas KGA commits none while abstaining on 732 rows. This is an observed accuracy–commitment trade-off, not proof of future directional safety or superiority to a separately tuned abstaining threshold rule. No bootstrap intervals from the superseded table are transferred to this

Table 10: CCT-20 safe-utility endpoint defined in Eq. (11). Contrasts are baseline regret minus KGA regret on the set-membership accuracy-proportion scale. These are nominal pointwise 95% two-way bootstrap intervals, not interval coverage for individual cell-benefit predictions. The stronger comparisons and location-level records remain in the supplement.

| Comparator    | Mean contrast | Nominal 95% CI   | Required lower bound |
|---------------|---------------|------------------|----------------------|
| Always adapt  | 0.1819        | [0.0937, 0.2914] | $L > 0.0000$         |
| Always freeze | 0.0000        | [0.0000, 0.0000] | $L > -0.0050$        |

regenerated comparison.

The earlier synchronized and P4 threshold tables are retained in the repository as unpromoted historical text, not as current comparative evidence. Their printed fixed-policy regrets do not match the inspected paired arrays, and the previously cited aggressive archive contains 270 rather than 432 conditions per seed. The streaming-versus-episodic explanation is therefore withdrawn for these authorities. The source audit, exact-replay receipt and generated per-cell policy records are indexed in Appendix N. Authenticated native TTA/ALine [51], AETTA [18] and POEM [14] comparisons remain separate unfinished experiments.

#### 8.4 CCT-20: Conservative Retention, Not Selective Routing

Across 45 cells, KGA issued 0 ADAPT, 44 FREEZE, and 1 ABSTAIN decisions. It therefore served the frozen predictions in every cell. The candidate was helpful in 1 cell, tied in 0, and harmed in 44. The gate rejected most harmful updates but did not select the helpful one.

KGA matched always-freeze exactly and avoided the measured degradation of always-adapt. The unchanged locked rule therefore returns pass on the stored cells, with generated flag yes. The recorded verdict remains SAFE UTILITY ONLY: a scoped safe-utility result, not a finite-sample population-safety guarantee. The stronger routing criterion failed because the gate never selected ADAPT and matched freeze. The zero false-adapt count (0/45) follows mechanically from zero adapt exposure.

All development cells were harmful for Tent, so the fit supplied no helpful example. Only two locations contributed fitting data, and calibration-to-target transfer is unverified. The complete target-location ledger and primary inference are retained in Appendix G.1; no new natural-shift success is inferred from this result.

#### 8.5 Development Stops and Adverse Diagnostics

So2Sat-LCZ42 [28] tested 2 candidates on 9 development cities and 5 independent checkpoints. Tent had both helpful and harmful cities, but neither candidate passed all predeclared eligibility checks. The study stopped before gate calibration: no target inputs or labels were accessed, and there is no target natural-shift score. The development outcomes are now open, so retuning them cannot produce fresh confirmation.

The auxiliary studies also show limitations. ImageNet-R and PACS favor always-adapt; ImageNet-C is sensitive to the candidate and its operating point. Office-Home and RxRx1 mainly retain the frozen model, whereas the opened Camelyon17 OOD panel adapts throughout. The opened Camelyon17 panel shows that KGA can issue ADAPT decisions on natural data; it does not establish why CCT-20 produced no such decisions. On the observed CCT-20 cells, KGA retains the frozen model, whereas the one-sided Camelyon17 panel retains adaptation exposure. These candidate-specific outcomes do not demonstrate prospective selective routing. Transferring a fixed

benefit predictor across completely unseen wild domains without target calibration remains an open empirical challenge. Appendix B reports the valid diagnostic rows on separate metric scales, together with the reasons why iWildCam, historical POEM/AETTA ports, and the focused Office-Home route cannot support current headline comparisons.

## 9 Limitations and Broader Impact

### 9.1 Assumptions, Not an Unconditional Safety Guarantee

The population frontier is exact within its declared rich binary class. A credible residual bound must come from domain knowledge, target labels, or a stated transfer assumption. A narrower class may permit stronger decisions, while an unsupported bound cannot justify a strict population claim. The useful object is the class-dependent feasible set, not the algebraic inequality alone.

**The development-based surrogate failed.** We tested a source-development benefit-scale proxy as an operational surrogate for  $\beta$  rather than declaring the population budget directly; this proxy does not estimate the theorem’s disagreement-conditional calibration-residual bound. On CIFAR-10-C, the surrogate is 1.4–50 $\times$  too small for operational soundness: 24–73% of evaluation cells produced measured outcomes outside the proxy-declared interval, and commit error is 0.4–16.6%. On ImageNet-C, 5 of 10 configurations return zero commitments on all 405 evaluation cells. These historical numbers summarize the failure of an empirical proxy heuristic, not a verified numerical measurement of theoretical-class violations (the latent residual  $\gamma$  is unobservable without labels, and the historical raw sweep artifact is unavailable). We therefore withdraw this surrogate estimation route and declare  $\beta$  from domain knowledge or an explicit transfer assumption instead. KGA’s empirical prediction radius  $\varepsilon$  is not an estimator of the population residual bound  $\beta$ . A population-benefit guarantee can instead be obtained through an explicit statistical protocol: valid conformal calibration for measured cell benefit, a justified cell-to-population sampling bound  $b(m, \delta)$ , and an enlarged decision interval  $\varepsilon + b$  (Proposition 3). These conditions are not established for the currently reported benchmark panels. KGA neither estimates nor numerically uses  $\beta$ .

For KGA, empirical interval coverage requires separate support at the intended decision unit. A larger radius reduces commitment but can increase regret when the frozen model is worse. A smaller radius increases exposure and makes calibration failures more consequential. Changing the adapter, checkpoint, feature schema, or update settings changes the prediction problem and requires recalibration or a justified transfer argument.

**Repeated deployment needs a separate guarantee.** A predeclared policy does not by itself supply exchangeability or time-uniform safety. Reusing calibration data, changing the frozen baseline after accepted updates, and selecting later candidates from earlier outcomes change the joint process. Adaptive conformal inference studies online coverage frequency with feedback [39]; that is not time-uniform control of KGA’s false commitments. Selection-conditional and group-conditional guarantees likewise require additional methods, as discussed in Section 2.

### 9.2 Empirical Delimitations and Future Prospective Validation

The current controlled results share checkpoints and related corruption cells. The new interval-inclusion table is a retrospective diagnostic, not an independent calibration test. The natural results primarily evaluate retention. CCT-20 is a protocol-scoped prospective, outcome-unopened real-population study at held-out camera locations, but it supplied no ADAPT exposure; So2Sat

stopped in development. Empirical abstention may reflect weak prediction, limited data, wide intervals, or transfer failure. It does not establish population non-identifiability.

The current same-input saved-feature comparison in Section 8.3 (Table 9) evaluates fixed policies, scalar evidence rules, and internal benefit-estimation heuristics. The scoped native-component comparison in Appendix B.9 (Table 23) is available for the declared 45-cell subset, but does not close the broader official-comparison requirement because Kim et al.’s TTA/ALine source is unavailable and the Baek ALine controller is not estimable on every cell. Broader evidence would strengthen the empirical claim: a natural-shift study with helpful and harmful updates and nonzero ADAPT/FREEZE exposure; interval inclusion and error evaluated on untouched groups; and inference across independently trained checkpoints and held-out environments. For tracks other than CCT-20, these remain future prospective tests, not results obtained by rewriting or reclassifying an existing panel. The future protocol is specified in Appendix K.

As a release-audit check, we replayed the outcome-blind population procedure on a three-record synthetic panel and on 108 cells of the authenticated local ImageNet-C development panel (five-fold cross-fitting; 36 cells each for Tent, EATA, and SAR). The resulting records are outcome-blind at decision time. These checks verify that the population action construction and the common metric implementation execute end to end; they are not independent prospective population evidence, and they do not convert the protocol-matched TTA/ALine, AETTA, or POEM rows into official native reproductions.

### 9.3 Operational Scope and Potential Harm

A rejected update can prevent adaptation harm, but retaining a frozen model is not proof that the model is accurate or suitable for a high-stakes setting. Misinterpreting an abstention as a certificate would be particularly risky. Deployment should preserve the assessment reason and connect unresolved cases to a domain-appropriate review or evidence-acquisition process. The reported experiments do not validate clinical, safety-critical, or autonomous operation. No held-out physical camera deployment or common end-to-end latency result is claimed.

The maintained software distinguishes an unavailable assessment from a supported negative interval. Its tests and proof artifacts document a scoped validity, commitment-control, and rollback layer; they do not certify every data pipeline or production system. Reproduction instructions, formal scope, excluded historical extensions, and release checks are provided in the supplement.

## 10 Conclusion

K-Bound separates information that is absent from uncertainty that can be estimated. For a fixed model pair in the declared binary class, the feasible population benefit set determines when a strict direction is supported. The evidence-fibre audit result shows why the same unlabeled observations cannot remove the unresolved residual bound.

KGA is an empirical companion: it predicts a measured cell benefit and uses an interval to accept, reject, or defer an update. The controlled results show candidate-specific routing value. The CCT-20 result shows conservative retention, while the So2Sat development stop exposes a transfer limit. Together they motivate evidence-aware commitment, with the validity of each conclusion tied to its stated target and assumptions. Stronger natural-shift routing and independent coverage validation remain open empirical questions.

## References

- [1] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, T. Darrell. Tent: Fully test-time adaptation by entropy minimization. *ICLR*, 2021.
- [2] J. Liang, D. Hu, J. Feng. Do we really need to access the source data? Source hypothesis transfer for unsupervised domain adaptation. *ICML*, 2020.
- [3] Y. Sun, X. Wang, Z. Liu, J. Miller, A. Efros, M. Hardt. Test-time training with self-supervision for generalization under distribution shifts. *ICML*, 2020.
- [4] S. Niu, J. Wu, Y. Zhang, Y. Chen, S. Zheng, P. Zhao, M. Tan. Efficient test-time model adaptation without forgetting. *ICML*, 2022.
- [5] S. Niu, J. Wu, Y. Zhang, Z. Wen, Y. Chen, P. Zhao, M. Tan. Towards stable test-time adaptation in dynamic wild world. *ICLR*, 2023.
- [6] Q. Wang, O. Fink, L. Van Gool, D. Dai. Continual test-time domain adaptation. *CVPR*, 2022.
- [7] M. Zhang, S. Levine, C. Finn. MEMO: Test time robustness via adaptation and augmentation. *NeurIPS*, 2022.
- [8] M. Schirmer, M. Jazbec, C. A. Naesseth, E. Nalisnick. Monitoring risks in test-time adaptation. *NeurIPS*, 2025. arXiv:2507.08721.
- [9] O. C. Mutlu, M. Honarmand, P. Azizian, S. Surabhi, D. P. Wall. Topology-guided test-time adaptation via persistent homology: From affective behavior analysis to autonomous driving. *CVPR Workshops*, pages 5331–5340, 2026.
- [10] M. Chen, B. Zhang, Z. Han, W. Jiang, Y. Wang, S. Feng, Y. Du, B. Bao. Test-time selective adaptation for uni-modal distribution shift in multi-modal data. *Proceedings of the 42nd International Conference on Machine Learning*, PMLR 267:9711–9727, 2025. <https://proceedings.mlr.press/v267/chen25ch.html>.
- [11] S. Sahoo, M. ElAraby, J. Ngnawe, Y. B. Pequignot, F. Precioso, C. Gagné. A layer selection approach to test time adaptation. *Proceedings of the AAAI Conference on Artificial Intelligence* 39(19):20237–20245, 2025. doi:10.1609/aaai.v39i19.34229.
- [12] A. H. Ariq. MMD-balls as credal sets: A PAC-Bayesian framework for epistemic uncertainty in test-time adaptation. arXiv:2605.21783, 2026.
- [13] A. Podkopaev, A. Ramdas. Tracking the risk of a deployed model and detecting harmful distribution shifts. *ICLR*, 2022. arXiv:2110.06177.
- [14] Y. Bar, S. Shaer, Y. Romano. Protected test-time adaptation via online entropy matching: A betting approach. *NeurIPS*, 2024. arXiv:2408.07511.
- [15] S. Garg, S. Balakrishnan, Z. Lipton, B. Neyshabur, H. Sedghi. Leveraging unlabeled data to predict out-of-distribution performance. *ICLR*, 2022. arXiv:2201.04234.
- [16] J. Miller, R. Taori, A. Raghunathan, S. Sagawa, P. W. Koh, V. Shankar, P. Liang, Y. Carmon, L. Schmidt. Accuracy on the line: On the strong correlation between out-of-distribution and in-distribution generalization. *ICML*, 2021. arXiv:2107.04649.
- [17] C. Baek, Y. Jiang, A. Raghunathan, J. Z. Kolter. Agreement-on-the-line: Predicting the performance of neural networks under distribution shift. *NeurIPS*, 2022. arXiv:2206.13089.
- [18] T. Lee, S. Chottananurak, T. Gong, S.-J. Lee. AETTA: Label-free accuracy estimation for test-time adaptation. *CVPR*, 2024. arXiv:2404.01351.
- [19] J. Steinhardt, P. Liang. Unsupervised risk estimation using only conditional independence structure. *NeurIPS*, 2016. arXiv:1606.05313.
- [20] S. Ben-David, J. Blitzer, K. Crammer, A. Kulesza, F. Pereira, J. Wortman Vaughan. A theory of learning from different domains. *Machine Learning* 79(1–2):151–175, 2010a.
- [21] E. Rosenfeld, S. Garg. (Almost) provable error bounds under distribution shift via disagreement discrepancy. *NeurIPS*, 2023. arXiv:2306.00312.
- [22] Y. Geifman, R. El-Yaniv. Selective classification for deep neural networks. *NeurIPS*, 2017.

- [23] D. Hendrycks, T. Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. *ICLR*, 2019.
- [24] D. Hendrycks, S. Basart, N. Mu, S. Kadavath, F. Wang, E. Dorundo, R. Desai, T. Zhu, S. Parajuli, M. Guo, D. Song, J. Steinhardt, J. Gilmer. The many faces of robustness: A critical analysis of out-of-distribution generalization. *ICCV*, pages 8340–8349, 2021.
- [25] P. W. Koh, S. Sagawa, H. Marklund, S. M. Xie, M. Zhang, A. Balsubramani, W. Hu, M. Yasunaga, et al. WILDS: A benchmark of in-the-wild distribution shifts. *ICML*, 2021.
- [26] P. Bándi, O. Geessink, Q. Manson, M. van Dijk, M. Balkenhol, M. Hermesen, et al. From detection of individual metastases to classification of lymph node status at the patient level: The CAMELYON17 challenge. *IEEE Trans. Med. Imaging* 38(2):550–560, 2019.
- [27] M. Sypetkowski, M. Rezanejad, S. Saberian, O. Kraus, J. Urbanik, J. Taylor, B. Mabey, M. Victors, J. Yosinski, A. Rezazadeh Sereshkeh, I. Haque, B. Earnshaw. RxRx1: A dataset for evaluating experimental batch correction methods. *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, pages 4285–4294, 2023. arXiv:2301.05768.
- [28] X. X. Zhu, J. Hu, C. Qiu, Y. Shi, J. Kang, L. Mou, H. Bagheri, M. Häberle, Y. Hua, R. Huang, L. Hughes, H. Li, Y. Sun, G. Zhang, S. Han, M. Schmitt, Y. Wang. So2Sat LCZ42: A benchmark data set for the classification of global local climate zones [Software and Data Sets]. *IEEE Geoscience and Remote Sensing Magazine* 8(3):76–89, 2020. doi:10.1109/MGRS.2020.2964708.
- [29] H. Venkateswara, J. Eusebio, S. Chakraborty, S. Panchanathan. Deep hashing network for unsupervised domain adaptation. *CVPR*, 2017. (Office-Home dataset release).
- [30] D. Li, Y. Yang, Y.-Z. Song, T. M. Hospedales. Deeper, broader and artier domain generalization. *ICCV*, pages 5543–5551, 2017. doi:10.1109/ICCV.2017.591.
- [31] B. Recht, R. Roelofs, L. Schmidt, V. Shankar. Do CIFAR-10 classifiers generalize to CIFAR-10? *arXiv preprint*, 2018. arXiv:1806.00451.
- [32] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, et al. PyTorch: An imperative style, high-performance deep learning library. *NeurIPS*, 2019.
- [33] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in Python. *JMLR* 12:2825–2830, 2011.
- [34] R. Xie, A. Odonnat, V. Feofanov, W. Deng, J. Zhang, B. An. MaNo: exploiting matrix norm for unsupervised accuracy estimation under distribution shifts. *NeurIPS*, 2024. arXiv:2405.18979.
- [35] W. Deng, Y. Suh, S. Gould, L. Zheng. Confidence and dispersity speak: characterizing prediction matrix for unsupervised accuracy estimation. *ICML*, 2023. arXiv:2302.01094.
- [36] R. J. Tibshirani, R. F. Barber, E. Candès, A. Ramdas. Conformal prediction under covariate shift. *NeurIPS*, 2019. arXiv:1904.06019.
- [37] R. F. Barber, E. J. Candès, A. Ramdas, R. J. Tibshirani. Conformal prediction beyond exchangeability. *Annals of Statistics* 51(2):816–845, 2023. arXiv:2202.13415.
- [38] J. Hemerik, J. Goeman. Exact testing with random permutations. *TEST* 27:811–825, 2018. doi:10.1007/s11749-017-0571-1.
- [39] I. Gibbs, E. Candès. Adaptive conformal inference under distribution shift. *NeurIPS*, 2021. arXiv:2106.00170.
- [40] S. Bates, A. Angelopoulos, L. Lei, J. Malik, M. I. Jordan. Distribution-free, risk-controlling prediction sets. *JACM* 68(6), 2021. arXiv:2101.02703.
- [41] A. N. Angelopoulos, S. Bates, E. J. Candès, M. I. Jordan, L. Lei. Learn then test: calibrating predictive algorithms to achieve risk control. *Annals of Applied Statistics*, 2025. arXiv:2110.01052.
- [42] S. Sonoda, K. Kasaura, Y. Mizuno, K. Tsukamoto, N. Onda. Lean formalization of generalization error bound by Rademacher complexity. *arXiv preprint*, 2025. arXiv:2503.19605v3.
- [43] C. Gupta, A. Ramdas. Top-label calibration and multiclass-to-binary reductions. *ICLR*, 2022. arXiv:2107.08353.

- [44] Z. C. Lipton, Y.-X. Wang, A. Smola. Detecting and correcting for label shift with black box predictors. *ICML*, 2018.
- [45] S. Garg, Y. Wu, S. Balakrishnan, Z. C. Lipton. A unified view of label shift estimation. *NeurIPS*, 2020.
- [46] C. K. Chow. On optimum recognition error and reject tradeoff. *IEEE Trans. Information Theory*, 16(1):41–46, 1970.
- [47] C. Cortes, G. DeSalvo, M. Mohri. Learning with rejection. *ALT*, 2016.
- [48] S. Beery, G. Van Horn, and P. Perona, “Recognition in Terra Incognita,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, 2018, pp. 472–489, doi: 10.1007/978-3-030-01270-0\_28.
- [49] S. Ben-David, T. Lu, T. Luu, D. Pál. Impossibility theorems for domain adaptation. *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, PMLR 9:129–136, 2010b. <https://proceedings.mlr.press/v9/david10a.html>.
- [50] I. Lamaakal, C. Yahyati, K. El Makkaoui, I. Ouahbi, Y. Maleh. Drift-to-Action Controllers: Budgeted interventions with online risk certificates. *arXiv preprint*, 2026. arXiv:2603.08578v1.
- [51] E. Kim, M. Sun, C. Baek, A. Raghunathan, J. Z. Kolter. Test-Time Adaptation Induces Stronger Accuracy and Agreement-on-the-Line. *Advances in Neural Information Processing Systems*, 37, 2024. doi:10.52202/079017-3820.
- [52] Y. Jin, Z. Ren. Confidence on the focal: conformal prediction with selection-conditional coverage. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 87(4):1239–1259, 2025. doi:10.1093/jrsssb/qkaf016.
- [53] I. Gibbs, J. J. Cherian, E. J. Candès. Conformal prediction with conditional guarantees. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 87(4):1100–1126, 2025. doi:10.1093/jrsssb/qkaf008.
- [54] T. Christensen, H. R. Moon, F. Schorfheide. Optimal Decision Rules When Payoffs are Partially Identified. *The Review of Economic Studies*, rdag017, 2026. Advance article, corrected proof. doi:10.1093/restud/rdag017.

## A A Rigorous Bridge from Cell Coverage to Population Coverage

The formal theory of K-Bound (Section 4 through Section 5) characterizes the exact population partial-identification frontier under a declared bound  $\beta$  on the hidden disagreement-conditional residual  $\gamma$ . In contrast, the empirical gate KGA (Section 6) predicts measured evaluation-cell benefit  $\Delta^{\text{cell}}$  and constructs a prediction interval  $[\hat{\Delta} - \varepsilon, \hat{\Delta} + \varepsilon]$ . This appendix provides the rigorous mathematical bridge between these two formulations, corrects a common dimensional misconception, provides a formal counterexample to naive transfer, and establishes sufficient conditions under which an enlarged conformal interval provably covers population benefit  $\Delta$ .

### A.1 Unit Scaling and the Category Error of $\varepsilon = \beta$

It is tempting to conjecture that an empirical gate could directly enforce the population frontier by setting its empirical conformal radius equal to the population residual bound,  $\varepsilon = \beta$ . This conjecture is fundamentally false:  $\varepsilon$  and  $\beta$  represent different mathematical objects living on different scales.

Recall the exact disagreement-region reduction under binary 0/1 loss (Lemma 2):

$$\Delta = 2d(M + \gamma), \quad d = P_T(D) > 0, \quad (13)$$

where  $D = \{x : f_0(x) \neq f_a(x)\}$  is the disagreement region,  $M = \mathbb{E}[s(X) \mid D] - \frac{1}{2} \in [-\frac{1}{2}, \frac{1}{2}]$  is the observable disagreement margin, and  $\gamma = \mathbb{E}[\eta_a(X) - s(X) \mid D]$  is the hidden calibration residual. Under the externally declared residual class  $\mathcal{C}_\beta = \{P_T : |\gamma(P_T)| \leq \beta\}$ , the feasible population benefit set satisfies:

$$|\Delta - 2dM| = 2d|\gamma| \leq 2d\beta. \quad (14)$$

Consequently, the unclipped population-benefit envelope centered at  $2dM$  has radius

$$r_{\text{pop}} = 2d\beta, \quad (15)$$

not  $\beta$ ; the exact population benefit set  $\Delta(\mathcal{C}_\beta)$  in Theorem 2 includes clipping to  $[-d, d]$ :  $\Delta(\mathcal{C}_\beta) = 2d[\max(-\frac{1}{2}, M - \beta), \min(\frac{1}{2}, M + \beta)]$ . Here  $d = P_T(D) \in (0, 1]$  is the disagreement probability, so the scaling factor  $2d \in (0, 2]$  is twice that probability. For example, with typical values  $d = 0.20$ ,  $M = 0.15$ , and  $\beta = 0.05$ , the population-benefit interval is:

$$\begin{aligned} 2(0.20)[0.15 - 0.05, 0.15 + 0.05] &= 0.40[0.10, 0.20] \\ &= [0.04, 0.08]. \end{aligned}$$

The interval center is 0.06 and its half-width on the risk-reduction scale is 0.02. The declared residual budget is  $\beta = 0.05$ , whereas the benefit-scale radius is 0.02. Equating  $\varepsilon = \beta$  is therefore a dimensional and conceptual category error:  $\beta$  bounds a conditional score bias on  $D$ , whereas  $\varepsilon$  measures prediction error for an empirical cell outcome on the entire input domain. While numerical equality could occur by accidental coincidence, identifying their mathematical definitions is fundamentally incorrect.

### A.2 Counterexample: Perfect Cell Prediction Does Not Imply Population Coverage

We now prove that even an ideal empirical predictor with zero cell-prediction error ( $\varepsilon = 0$ ) can fail to cover population benefit and can commit to false adaptation with substantial probability.

**Proposition 4** (Failure of automatic cell-to-population transfer). *There exists a data-generating process, a fixed model pair  $(f_0, f_a)$ , and a benefit predictor  $\hat{\Delta}$  such that  $\mathbb{P}(|\hat{\Delta} - \Delta^{\text{cell}}| = 0) = 1$  (perfect cell coverage with  $\varepsilon = 0$ ), yet the cell gate selects ADAPT with probability 0.45 despite the true population benefit being strictly negative ( $\Delta = -0.10$ ).*

*Proof.* Fix binary predictors  $f_0(x) \equiv 0$  and  $f_a(x) \equiv 1$ . Let the target input space consist of two points,  $\mathcal{X} = \{a, b\}$ , with input distribution

$$\mathbb{P}(X = a) = 0.45, \quad \mathbb{P}(X = b) = 0.55,$$

and deterministic target labels  $Y(a) = 1$  and  $Y(b) = 0$ . The zero-one population risks are:

$$R_T(f_0) = \mathbb{P}(f_0(X) \neq Y) = \mathbb{P}(X = a) = 0.45,$$

$$R_T(f_a) = \mathbb{P}(f_a(X) \neq Y) = \mathbb{P}(X = b) = 0.55.$$

The true population benefit of candidate adaptation is strictly negative:

$$\Delta = R_T(f_0) - R_T(f_a) = 0.45 - 0.55 = -0.10 < 0.$$

Now consider an evaluation cell consisting of a single independent target sample ( $m = 1$ ). The measured cell benefit is

$$\begin{aligned} \Delta^{\text{cell}} &= \mathbf{1}\{f_a(X) = Y\} - \mathbf{1}\{f_0(X) = Y\} \\ &= \begin{cases} +1, & X = a, \\ -1, & X = b. \end{cases} \end{aligned}$$

Define the fixed label-free cell predictor  $h(a) = +1$  and  $h(b) = -1$ . This predictor observes only the unlabeled input  $X$  and achieves perfect prediction of the cell outcome:

$$|\hat{\Delta} - \Delta^{\text{cell}}| = |h(X) - \Delta^{\text{cell}}| = 0 \quad \text{almost surely.}$$

A conformal interval with radius  $\varepsilon = 0$  achieves 100% marginal coverage of  $\Delta^{\text{cell}}$ . However, whenever  $X = a$  (which occurs with probability 0.45),  $h(a) - \varepsilon = +1 > 0$ , triggering an ADAPT decision. The probability of false adaptation is therefore:

$$\mathbb{P}(g = \text{ADAPT}, \Delta \leq 0) = \mathbb{P}(X = a) = 0.45.$$

Specializing this to residual notation with constant score  $s(x) \equiv 1/2$ , the disagreement region is  $D = \{a, b\}$ ,  $d = 1$ ,  $M = 0$ , and  $\gamma = \mathbb{E}[\eta_a(X) - s(X) \mid D] = 0.45 - 0.50 = -0.05$ . The tightest valid residual bound is  $\beta = 0.05$ , whereas the empirical cell-prediction radius is  $\varepsilon = 0$ . Thus  $\varepsilon = 0 \neq 0.05 = \beta$ .  $\square$

Proposition 4 demonstrates why empirical cell-level validity cannot be equated with population-level safety: a gate that accurately predicts small-batch or single-cell advantages can repeatedly commit to globally harmful adaptations.

### A.3 Formal Protocol for Cell-to-Population Transfer

To rigorously bridge cell prediction to population benefit, we specify a formal statistical protocol:

- (P1) Independent Estimator Fitting:** The benefit predictor  $h$  and feature extractor  $\phi$  are fitted on historical development episodes  $\mathcal{D}_{\text{dev}}$  and frozen prior to calibration and deployment.
- (P2) Representative Episode Exchangeability:** Conditional on the frozen development data  $\mathcal{D}_{\text{dev}}$  (which determines the fitted predictor  $h$  and feature representation  $\phi$ ), the  $N$  calibration episodes  $(Z_1, B_1), \dots, (Z_N, B_N)$  and the deployment episode  $(Z_*, B_*)$  are exchangeable draws from a deployment meta-distribution  $\mathcal{Q}$ . Each episode  $e$  defines an environment  $P_e$ .
- (P3) Pre-Evaluation Candidate Isolation:** For each episode  $e$ , candidate  $f_{a,e} = \mathcal{A}(f_0, X_{e,\text{adapt}})$  is constructed from its unlabeled adaptation window, and its parameter and normalization states are strictly frozen prior to scoring.
- (P4) Fresh Evaluation Sample:** Evaluation cell  $e$  comprises  $m$  fresh, conditionally i.i.d. draws  $(X_{e,j}, Y_{e,j})_{j=1}^m \sim P_e$  that are independent of the adaptation data.
- (P5) Matched Risk Metric:** The cell outcome is the sample paired accuracy difference  $B_e = \Delta_e^{\text{cell}} = \frac{1}{m} \sum_{j=1}^m W_{e,j}$ , where  $W_{e,j} = \mathbf{1}\{f_{a,e}(X_{e,j}) = Y_{e,j}\} - \mathbf{1}\{f_0(X_{e,j}) = Y_{e,j}\} \in [-1, 1]$ , matching the population risk reduction  $\Delta_e = R_{P_e}(f_0) - R_{P_e}(f_{a,e})$ .

### A.4 Theorem and 4-Step Proof of Population Risk Control

Under protocol **(P1)–(P5)**, we prove that an enlarged decision interval achieves finite-sample population-benefit coverage and strict directional error control.

**Theorem 4** (Finite-sample population benefit certificate). *Under protocol **(P1)–(P5)**, let  $\alpha_{\text{cell}} \in (0, 1)$  and  $\delta \in (0, 1)$  be predeclared error budgets such that  $\alpha_{\text{cell}} + \delta < 1$ . Let  $\varepsilon$  be the order-statistic split-conformal radius over  $N$  calibration residuals:*

$$k = \lceil (N+1)(1 - \alpha_{\text{cell}}) \rceil, \quad \varepsilon = \begin{cases} r(k), & k \leq N, \\ +\infty, & k > N. \end{cases} \quad (16)$$

Define the finite-sample concentration radius

$$b(m, \delta) = \min \left\{ 2, \sqrt{\frac{2 \log(2/\delta)}{m}} \right\}. \quad (17)$$

Then for a new deployment episode  $* \sim \mathcal{Q}$ , the compound interval  $I_{\text{population}} = [\hat{\Delta}_* - (\varepsilon + b), \hat{\Delta}_* + (\varepsilon + b)]$  satisfies:

$$\mathbb{P}(\Delta_* \in I_{\text{population}}) \geq 1 - \alpha_{\text{cell}} - \delta. \quad (18)$$

Furthermore, the three-way decision rule

$$g_* = \begin{cases} \text{ADAPT}, & \hat{\Delta}_* - (\varepsilon + b) > 0, \\ \text{FREEZE}, & \hat{\Delta}_* + (\varepsilon + b) < 0, \\ \text{ABSTAIN}, & \text{otherwise}, \end{cases} \quad (19)$$

controls the marginal probability of committing to an incorrect strict direction:

$$\begin{aligned} \mathbb{P}(\{g_* = \text{ADAPT}, \Delta_* \leq 0\} \cup \{g_* = \text{FREEZE}, \Delta_* \geq 0\}) \\ \leq \alpha_{\text{cell}} + \delta. \end{aligned} \quad (20)$$

*Proof.* The proof proceeds in four discrete steps:

**Step 1: Conformal coverage of measured cell benefit.** By (P1) and (P2), the residuals  $r_e = |h(Z_e) - B_e|$  for  $e \in \{1, \dots, N, *\}$  are exchangeable random variables conditional on the frozen development data  $\mathcal{D}_{\text{dev}}$  (and hence conditional on  $h$  and  $\phi$ ). By standard split-conformal quantile properties [36], the  $k$ -th order statistic guarantees marginal coverage:

$$\mathbb{P}(\mathcal{E}_{\text{conf}}) \geq 1 - \alpha_{\text{cell}},$$

where  $\mathcal{E}_{\text{conf}} = \{|h(Z_*) - B_*| \leq \varepsilon\}$ .

**Step 2: Hoeffding concentration of sample benefit to population benefit.** Conditioned on environment  $P_*$  and the frozen model pair  $(f_0, f_{a,*})$ , the variables  $W_{*,j} = \mathbf{1}\{f_{a,*}(X_{*,j}) = Y_{*,j}\} - \mathbf{1}\{f_0(X_{*,j}) = Y_{*,j}\}$  are independent and bounded in  $[-1, 1]$  with range  $c = 1 - (-1) = 2$ . Their conditional expectation is the true population benefit:

$$\mathbb{E}[W_{*,j} \mid P_*, f_0, f_{a,*}] = R_{P_*}(f_0) - R_{P_*}(f_{a,*}) = \Delta_*.$$

By Hoeffding's inequality for the sample average  $B_* = \frac{1}{m} \sum_{j=1}^m W_{*,j}$ :

$$\begin{aligned} \mathbb{P}(|B_* - \Delta_*| > t \mid P_*, f_0, f_{a,*}) &\leq 2 \exp\left(-\frac{2m^2 t^2}{m(2)^2}\right) \\ &= 2 \exp\left(-\frac{mt^2}{2}\right). \end{aligned}$$

Setting the right-hand side equal to  $\delta$  yields  $t = \sqrt{2 \log(2/\delta)/m}$ . Since  $B_*$  and  $\Delta_*$  both lie in  $[-1, 1]$ , their difference cannot exceed 2. Hence, defining  $b(m, \delta) = \min\{2, \sqrt{2 \log(2/\delta)/m}\}$ , we have unconditionally:

$$\mathbb{P}(\mathcal{E}_{\text{conc}}) \geq 1 - \delta,$$

where  $\mathcal{E}_{\text{conc}} = \{|B_* - \Delta_*| \leq b(m, \delta)\}$ .

**Step 3: Compound interval coverage.** On the intersection event  $\mathcal{E}_{\text{conf}} \cap \mathcal{E}_{\text{conc}}$ , the triangle inequality gives:

$$|h(Z_*) - \Delta_*| \leq |h(Z_*) - B_*| + |B_* - \Delta_*| \leq \varepsilon + b(m, \delta).$$

Applying the union bound to the complements:

$$\begin{aligned} \mathbb{P}(\Delta_* \in I_{\text{population}}) &\geq \mathbb{P}(\mathcal{E}_{\text{conf}} \cap \mathcal{E}_{\text{conc}}) \\ &\geq 1 - \mathbb{P}(\mathcal{E}_{\text{conf}}^c) - \mathbb{P}(\mathcal{E}_{\text{conc}}^c) \\ &\geq 1 - \alpha_{\text{cell}} - \delta. \end{aligned}$$

Crucially, this step requires no independence assumption between  $\mathcal{E}_{\text{conf}}$  and  $\mathcal{E}_{\text{conc}}$ .

**Step 4: Strict directional error bound.** Suppose  $g_* = \text{ADAPT}$ , which requires  $h(Z_*) - (\varepsilon + b) > 0$ . On the coverage event  $\{\Delta_* \in I_{\text{population}}\}$ , this implies  $\Delta_* \geq h(Z_*) - (\varepsilon + b) > 0$ . Therefore, a false adaptation ( $g_* = \text{ADAPT}$  and  $\Delta_* \leq 0$ ) can occur only on the complement event  $\Delta_* \notin I_{\text{population}}$ . Symmetrically, a false freeze ( $g_* = \text{FREEZE}$  and  $\Delta_* \geq 0$ ) also requires  $\Delta_* \notin I_{\text{population}}$ . Since the actions ADAPT and FREEZE are mutually exclusive, the union of false commitments satisfies:

$$\begin{aligned} \{g_* = \text{ADAPT}, \Delta_* \leq 0\} \cup \{g_* = \text{FREEZE}, \Delta_* \geq 0\} \\ \subseteq \{\Delta_* \notin I_{\text{population}}\}. \end{aligned}$$

Taking probabilities yields the claimed bound  $\leq \alpha_{\text{cell}} + \delta$ .  $\square$

## A.5 The Price of Rigor: Sample Size vs. Abstention Analysis

Theorem 4 explains mathematically why empirical test-time adaptation gates frequently abstain when rigorous population validity is demanded. The total interval half-width is  $\varepsilon + b(m, \delta)$ . Table 11 computes the Hoeffding sampling radius  $b(m, \delta)$  at confidence level  $1 - \delta = 0.95$  across representative evaluation sample sizes  $m$ .

Table 11: Sampling error radius  $b(m, \delta)$  for paired zero-one accuracy differences at  $\delta = 0.05$  (95% concentration confidence) across evaluation sample sizes  $m$ .

| Sample Size $m$ | Hoeffding Radius $b(m, \delta = 0.05)$ | Total Radius ( $\varepsilon = 0.015$ ) |
|-----------------|----------------------------------------|----------------------------------------|
| 64              | 0.3396                                 | 0.3546                                 |
| 200             | 0.1921                                 | 0.2071                                 |
| 512             | 0.1200                                 | 0.1350                                 |
| 1,000           | 0.0859                                 | 0.1009                                 |
| 5,000           | 0.0384                                 | 0.0534                                 |
| 10,000          | 0.0272                                 | 0.0422                                 |
| 20,000          | 0.0192                                 | 0.0342                                 |

As an illustrative scenario, consider an evaluation cell with  $m = 512$  samples (a common benchmark batch size), a forecast benefit of  $\hat{\Delta} = +0.04$  (a 4% point gain), conformal prediction error  $\varepsilon = 0.015$ , and concentration risk budget  $\delta = 0.05$ . Here the Hoeffding radius is  $b = 0.1200$ , yielding a total half-width  $\varepsilon + b = 0.1350$ . Even with a positive point prediction, the population lower bound is:

$$\hat{\Delta} - (\varepsilon + b) = 0.04 - 0.1350 = -0.0950 < 0.$$

The interval spans zero, forcing the population decision rule to retain the frozen predictor and record **ABSTAIN**. Holding the prediction  $\hat{\Delta} = 0.04$  and cell radius  $\varepsilon = 0.015$  fixed, certifying a positive population benefit requires  $\hat{\Delta} - \varepsilon - b(m, \delta) > 0$ , or:

$$0.04 - 0.015 - \sqrt{\frac{2 \log(2/\delta)}{m}} > 0 \iff m > \frac{2 \log(40)}{(0.025)^2} \approx 11,804.41.$$

Thus  $m = 11,805$  is the smallest integer satisfying this specific condition. We emphasize that this is an illustrative calculation from this particular Hoeffding bound holding the prediction and cell radius fixed, not an experimental finding or a universal sample requirement for certifying a 4% gain under other procedures or assumptions.

This calculation illustrates the potential commitment cost of the proposed population interval. It does not identify the cause of abstention or freezing in the reported benchmarks, whose gates did not use this sampling correction. Establishing a necessary sample-size requirement would require a separate lower-bound argument. Furthermore, the empirical outcome on CCT-20 was not universal abstention, but conservative retention: KGA issued 44 FREEZE decisions, 1 ABSTAIN, and 0 ADAPT. When evidence of positive benefit is insufficient, a prudent gate must retain the frozen baseline and record ABSTAIN; FREEZE specifically requires a supported negative interval, and uncertainty alone does not justify a directional freeze claim.

## A.6 Alternative Route: Direct Empirical Sampling of the Population Frontier

Finally, for completeness, we formalize an alternative operationalization that directly samples the theoretical population frontier  $|M| > \beta$  rather than training a learned benefit regressor  $\hat{\Delta} = h(Z)$ .

**Proposition 5** (Sampled population frontier). *Let  $(f_0, f_a)$  be fixed binary predictors and let  $s : \mathcal{X} \rightarrow [0, 1]$  be a source-fitted score. Suppose target inputs  $X_1, \dots, X_m \sim \mu_T$  are drawn independently, and let  $D_m = \{i \in \{1, \dots, m\} : f_0(X_i) \neq f_a(X_i)\}$  be the observed disagreement indices with count  $n_D = |D_m|$ . If  $n_D > 0$ , define the empirical disagreement margin:*

$$\widehat{M} = \frac{1}{n_D} \sum_{i \in D_m} s(X_i) - \frac{1}{2}. \quad (21)$$

For confidence level  $1 - \alpha \in (0, 1)$ , define the margin concentration radius:

$$r_M(n_D, \alpha) = \sqrt{\frac{\log(2/\alpha)}{2n_D}}. \quad (22)$$

Under the declared residual class  $\mathcal{C}_\beta = \{P_T : |\gamma(P_T)| \leq \beta\}$ , the decision rule

$$g_{\text{margin}} = \begin{cases} \text{ADAPT}, & \widehat{M} - r_M > \beta, \\ \text{FREEZE}, & \widehat{M} + r_M < -\beta, \\ \text{ABSTAIN}, & \text{otherwise (or if } n_D = 0), \end{cases} \quad (23)$$

satisfies  $\mathbb{P}(g_{\text{margin}} = \text{ADAPT}, \Delta \leq 0) \leq \alpha$  and  $\mathbb{P}(g_{\text{margin}} = \text{FREEZE}, \Delta \geq 0) \leq \alpha$ .

*Proof.* Conditional on  $n_D > 0$ , the disagreement samples  $\{X_i : i \in D_m\}$  are i.i.d. draws from the conditional marginal  $\mu_T \mid D$ . Since  $s(X) \in [0, 1]$ , Hoeffding's inequality for variables of range 1 gives:

$$\begin{aligned} \mathbb{P}\left(|\widehat{M} - M| > r_M \mid n_D\right) &\leq 2 \exp\left(-\frac{2n_D^2 r_M^2}{n_D(1)^2}\right) \\ &= 2 \exp\left(-2n_D r_M^2\right) = \alpha. \end{aligned}$$

On the coverage event  $|\widehat{M} - M| \leq r_M$ , the condition  $\widehat{M} - r_M > \beta$  implies:

$$M \geq \widehat{M} - r_M > \beta \geq |\gamma| \geq -\gamma \implies M + \gamma > 0.$$

By the disagreement reduction (Lemma 2),  $\Delta = 2d(M + \gamma) > 0$ . Thus, a false adaptation requires  $|\widehat{M} - M| > r_M$ , which has probability at most  $\alpha$ . The freeze case is identical by symmetry.  $\square$

This direct route demonstrates the additive relationship between uncertainties:

$$\begin{aligned} \text{decision threshold} = & \underbrace{r_M(n_D, \alpha)}_{\text{finite-sample margin uncertainty}} \\ & + \underbrace{\beta}_{\text{structural calibration uncertainty}}. \end{aligned}$$

While this rule directly embodies Theorem 2, it still requires an externally declared budget  $\beta$ , proving once again that label-free data alone cannot resolve structural non-identifiability.

## A.7 Explicit Implementation Boundaries and Operational Scope

To maintain absolute transparency regarding theoretical constructs versus empirical evaluations, we explicitly delineate three distinct objects formulated in this paper:

1. **Empirical KGA ( $\hat{\Delta} \pm \varepsilon$ ):** The cell-benefit prediction interval uses the residual construction specified for each study (Section 6). Current primary CIFAR results, the 108-cell ImageNet-C replay, and the corrected 27-cell SAR audit use Protocol B’s separate fitting, calibration, and scoring roles. Historical pooled leave-one-out analyses retain Protocol A’s outcome-dependence limitation; natural studies retain their separately documented data partitions. These are cell-benefit decisions, not population certificates.
2. **Population Extension ( $\hat{\Delta} \pm (\varepsilon + b)$ ):** The compound decision interval incorporating finite-sample Hoeffding concentration  $b(m, \delta)$  (Theorem 4). This procedure is exercised under controlled synthetic diagnostic simulations with exact ground truth (an observed 100% empirical population-inclusion rate and zero wrong-direction commitments in the finite trials); it is not a population guarantee for the historical empirical benchmark tables, which evaluate transductive evaluation-cell benefit rather than independent population draws. The separate 108-cell benchmark replay below exercises the enlarged interval as a sensitivity analysis, not as fresh-sample validation.
3. **Sampled Frontier Gate ( $\hat{M} \pm r_M$  under external  $\beta$ ):** The direct finite-sample operationalization of the knowability frontier (Proposition 5). This provides a formal alternative requiring an externally supplied residual bound  $\beta$ ; it has no real-data benchmark instantiation in this paper.

Neither the population transfer theorem nor the sampled frontier retroactively converts the empirical benchmark results into population-safety claims.

**Synthetic population simulation protocol and evidence.** To empirically evaluate the compound interval  $I_{\text{population}} = [\hat{\Delta} - (\varepsilon + b), \hat{\Delta} + (\varepsilon + b)]$  under exact ground truth, we conducted a diagnostic simulation (`experiments/kbound/synthetic_population_diagnostic.py`). In each episode, an environment covariate state  $Z \in \{-1, 0, 1\}$  is drawn with equal probability (1/3 each) and held fixed while  $m$  test examples are sampled with conditional label distributions  $\mathbb{P}(Y = 1 \mid Z = -1) = 0.20$  ( $\Delta = -0.60$ , harmful),  $\mathbb{P}(Y = 1 \mid Z = 0) = 0.50$  ( $\Delta = 0.00$ , tie/null), and  $\mathbb{P}(Y = 1 \mid Z = 1) = 0.80$  ( $\Delta = +0.60$ , helpful). The benefit estimator  $\hat{\Delta}(Z) = 0.6Z$  is the exact Bayes-optimal predictor. Calibration draws  $N_{\text{cal}}$  independent episodes to form residual pool  $\{r_c = |\hat{\Delta}(Z_c) - b_c|\}$ . Risk parameters are  $\alpha_{\text{cell}} = 0.05$  and  $\delta_{\text{sampling}} = 0.05$  (total population error budget  $\alpha_{\text{pop}} = 0.10$ ). Across calibration sizes  $N_{\text{cal}} \in \{18, 19, 39, 99\}$  and evaluation batch sizes  $m \in \{32, 128, 512, 2048\}$  over 300 independent trials per configuration (Table 13), empirical coverage and action distributions confirm: (1) at  $N_{\text{cal}} = 18$ , exact rank  $k = \lceil 19 \times 0.95 \rceil = 19 > 18$  is infeasible, setting  $\varepsilon = \infty$  and 100% abstention (fail-closed rank protection); (2) at  $m = 32$ , the sampling radius  $b(32, 0.05) = 0.4802$  combined with calibration radius  $\varepsilon_{\text{cell}} \approx 0.28$  yields compound width  $\varepsilon + b \approx 0.76 > 0.60$ , bracketing zero across all three states and enforcing 100% abstention (the Price of Rigor); (3) at  $m \geq 128$ , the gate safely commits on non-zero states while abstaining on ties ( $\approx 33.3\%$ , matching the null frequency); (4) wrong-direction error is 0.0000 across all 16 configurations ( $\leq \alpha_{\text{pop}} = 0.10$ ). An observed 100% population coverage rate on finite simulation trials reflects consistency under the simulated model, not a mathematical proof of universal coverage. Full summary: `experiments/kbound/results/synthetic_diagnostic/diagnostic_summary.json`.

**Non-oracle population diagnostic.** To evaluate the setting where the benefit predictor  $\hat{\Delta}$  is not an oracle Bayes rule but an imperfect empirical model fitted on finite development episodes, we conducted an extended non-oracle diagnostic (1,000 independent trials per configuration; `experiments/`

kbound/results/synthetic\_diagnostic/nonoracle\_diagnostic\_summary.json). In each trial, an imperfect linear predictor is fitted on 15 independent development episodes ( $m = 32$ ) subject to estimation noise. We then evaluate three decision mechanisms across the  $(N_{\text{cal}}, m)$  grid (Table 12): (1) **Point Rule** ( $\text{sign}(\hat{\Delta})$ ): Blindly committing to adaptation whenever  $\hat{\Delta} > 0$  achieves 0% abstention, but incurs a 31%–36% wrong-direction commitment error rate; crucially, 100% of these errors occur on null/tie states ( $Z = 0, \Delta = 0$ ) due to random estimation noise around zero, with strictly zero errors on harmful states ( $\Delta = -0.6$ ). (2) **Cell Interval** ( $\hat{\Delta} \pm \varepsilon_{\text{cell}}$ ): Conformal residual calibration over the imperfect predictions guarantees valid empirical cell coverage (93.9%–96.5%, tightly tracking  $1 - \alpha_{\text{cell}} = 0.95$ ) and eliminates wrong-direction error (0.0000). Furthermore, in this simulated model, the cell interval already achieves  $\approx 100\%$  empirical population coverage of  $\Delta$  across all configurations. (3) **Population Interval** ( $\hat{\Delta} \pm (\varepsilon_{\text{cell}} + b)$ ): By incorporating the Hoeffding sampling concentration  $b(m, \delta)$ , the compound interval achieved an observed 100% empirical population-inclusion rate and zero wrong-direction error in these finite trials ( $0.0000 \leq \alpha_{\text{pop}}$ ). At  $m = 32$ , it enforces the Price of Rigor through 99.9%–100% abstention as a consequence of the universal Hoeffding penalty ( $b(32) = 0.4802$ ), while at  $m \geq 128$  it safely commits on non-zero states while withholding action on ties ( $\approx 32\%$ – $34\%$  abstention, matching the 33.3% null state frequency).

Table 12: Non-oracle population diagnostic (1,000 trials/cell;  $\alpha_{\text{cell}} = \delta = 0.05$ ,  $\alpha_{\text{pop}} = 0.10$ ). Evaluates an imperfect benefit predictor fitted on 15 independent development episodes ( $m = 32$ ) with estimation error. Entries report **Cell Coverage / Wrong-Direction Error / Abstention Rate**. The point rule commits on ties (31%–36% wrong-direction error entirely on null states  $Z = 0$ , 0% abstention). The cell interval eliminates wrong-direction error (0.000) with empirical cell coverage matching nominal  $1 - \alpha = 0.95$ , and achieves  $\approx 100\%$  population coverage of  $\Delta$  in this model. The population compound interval  $[\hat{\Delta} \pm (\varepsilon + b)]$  achieved an observed 100% empirical population-inclusion rate in these finite trials and zero wrong-direction error ( $0.000 \leq \alpha_{\text{pop}}$ ); it is not a universal coverage guarantee. It enforces the Price of Rigor (99.9%–100% abstention at  $m = 32$  via Hoeffding radius  $b = 0.480$ , dropping to  $\approx 33\%$  matching the null frequency for  $m \geq 128$ ). Artifact: experiments/kbound/results/synthetic\_diagnostic/nonoracle\_diagnostic\_summary.json.

| $N_{\text{cal}}$                                                                                                 | $m = 32$ ( $b = 0.480$ ) | $m = 128$ ( $b = 0.240$ ) | $m = 512$ ( $b = 0.120$ ) | $m = 2048$ ( $b = 0.061$ ) |
|------------------------------------------------------------------------------------------------------------------|--------------------------|---------------------------|---------------------------|----------------------------|
| <i>Point Rule</i> ( $\text{sign}(\hat{\Delta})$ ): <i>Err / Abstain</i>                                          |                          |                           |                           |                            |
| All $N$                                                                                                          | 0.309–0.356 / 0.000      | 0.313–0.332 / 0.000       | 0.314–0.353 / 0.000       | 0.326–0.340 / 0.000        |
| <i>Cell Interval</i> ( $\hat{\Delta} \pm \varepsilon_{\text{cell}}$ ): <i>CellCov / Err / Abstain</i>            |                          |                           |                           |                            |
| 18                                                                                                               | 1.000 / 0.000 / 1.000    | 1.000 / 0.000 / 1.000     | 1.000 / 0.000 / 1.000     | 1.000 / 0.000 / 1.000      |
| 19                                                                                                               | 0.961 / 0.000 / 0.327    | 0.948 / 0.000 / 0.313     | 0.948 / 0.000 / 0.314     | 0.940 / 0.000 / 0.336      |
| 39                                                                                                               | 0.960 / 0.000 / 0.348    | 0.946 / 0.000 / 0.319     | 0.956 / 0.000 / 0.345     | 0.952 / 0.000 / 0.326      |
| 99                                                                                                               | 0.961 / 0.000 / 0.357    | 0.947 / 0.000 / 0.325     | 0.939 / 0.000 / 0.321     | 0.965 / 0.000 / 0.326      |
| <i>Population Interval</i> ( $\hat{\Delta} \pm (\varepsilon_{\text{cell}} + b)$ ): <i>PopCov / Err / Abstain</i> |                          |                           |                           |                            |
| 18                                                                                                               | 1.000 / 0.000 / 1.000    | 1.000 / 0.000 / 1.000     | 1.000 / 0.000 / 1.000     | 1.000 / 0.000 / 1.000      |
| 19                                                                                                               | 1.000 / 0.000 / 0.999    | 1.000 / 0.000 / 0.342     | 1.000 / 0.000 / 0.314     | 1.000 / 0.000 / 0.336      |
| 39                                                                                                               | 1.000 / 0.000 / 1.000    | 1.000 / 0.000 / 0.338     | 1.000 / 0.000 / 0.345     | 1.000 / 0.000 / 0.326      |
| 99                                                                                                               | 1.000 / 0.000 / 0.999    | 1.000 / 0.000 / 0.340     | 1.000 / 0.000 / 0.321     | 1.000 / 0.000 / 0.327      |

**108-Cell ImageNet-C development replay.** As a separate release replay, we instantiated the same outcome-blind action construction on 108 existing ImageNet-C development cells (five-fold cross-fitting, 36 cells each for Tent, EATA, and SAR; sample-size parameter  $m = 2,000$ ;

Table 13: Synthetic population diagnostic under exact Bayes-optimal prediction (300 trials/cell;  $\alpha_{\text{cell}} = \delta = 0.05$ ,  $\alpha_{\text{pop}} = 0.10$ ). Entries report empirical Cell Coverage / Abstention rate. Across all 16 cells, the finite simulation observed empirical population inclusion of 1.000 and wrong-direction error of 0.000; this is not a universal coverage guarantee. Sampling radii:  $b(32) = 0.4802$ ,  $b(128) = 0.2401$ ,  $b(512) = 0.1200$ ,  $b(2048) = 0.0600$ . Artifact: `experiments/kbound/results/synthetic_diagnostic/diagnostic_summary.json`.

| $N_{\text{cal}}$ | $m = 32$      | $m = 128$     | $m = 512$     | $m = 2048$    |
|------------------|---------------|---------------|---------------|---------------|
| 18               | 1.000 / 1.000 | 1.000 / 1.000 | 1.000 / 1.000 | 1.000 / 1.000 |
| 19               | 0.953 / 1.000 | 0.963 / 0.330 | 0.950 / 0.333 | 0.973 / 0.333 |
| 39               | 0.963 / 1.000 | 0.977 / 0.370 | 0.953 / 0.340 | 0.973 / 0.353 |
| 99               | 0.967 / 1.000 | 0.967 / 0.397 | 0.950 / 0.397 | 0.960 / 0.347 |

Table 14: Outcome-blind ImageNet-C 108-cell population-transfer replay across three candidates (36 cells each, five-fold cell-outcome-disjoint cross-fitting;  $\alpha_{\text{cell}} = 0.05$ ,  $\delta_{\text{samp}} = 0.05$ ,  $\alpha_{\text{pop}} = 0.10$ , sensitivity parameter  $m = 2,000$ , giving uniform sampling radius  $b = 0.0607$ ). For each fold, 7–8 cells are scored while the remaining 28–29 complement cells are split into 19 residual-calibration cells and 9–10 estimator-fit cells. Finite  $\varepsilon$  indicates non-degenerate calibration (108/108 cells). Decisions report counts of ADAPT/FREEZE/ABSTAIN before (cell gate,  $\hat{\Delta} \pm \varepsilon$ ) versus after (compound population interval,  $\hat{\Delta} \pm (\varepsilon + b)$ ). Across all 108 cells,  $\text{FA}_{\text{u}} = 0.000$  (zero false adaptations). Regret reports mean oracle shortfall on cell outcomes for Cell Gate vs. Compound Population Interval ( $0.0167 \rightarrow 0.0374$ , illustrating the Price of Rigor). Population truth is not established on historical development replays. Receipts: `experiments/kbound/results/task2_imagenetc_population_replay_20260911/population_panel.json`.

| Candidate            | $N$        | Cell radius $\varepsilon$ : mean [min, max] | Finite $\varepsilon$ |
|----------------------|------------|---------------------------------------------|----------------------|
| <b>Tent</b>          | 36         | 0.0519 [0.0331, 0.0659]                     | 36 / 36              |
| <b>EATA</b>          | 36         | 0.0292 [0.0196, 0.0393]                     | 36 / 36              |
| <b>SAR</b>           | 36         | 0.1184 [0.0365, 0.2106]                     | 36 / 36              |
| <b>Total / Panel</b> | <b>108</b> | 0.0665 [0.0196, 0.2106]                     | <b>108 / 108</b>     |

| Candidate            | Regret<br>cell / compound | Cell gate<br>A / F / U | Compound<br>A / F / U | $\text{FA}_{\text{u}}$ | Population truth       |
|----------------------|---------------------------|------------------------|-----------------------|------------------------|------------------------|
| <b>Tent</b>          | 0.0162 / 0.0468           | 18 / 0 / 18            | 1 / 0 / 35            | 0.000                  | Not established        |
| <b>EATA</b>          | 0.0095 / 0.0386           | 20 / 0 / 16            | 0 / 0 / 36            | 0.000                  | Not established        |
| <b>SAR</b>           | 0.0245 / 0.0267           | 2 / 11 / 23            | 0 / 10 / 26           | 0.000                  | Not established        |
| <b>Total / Panel</b> | <b>0.0167 / 0.0374</b>    | <b>40 / 11 / 57</b>    | <b>1 / 10 / 97</b>    | <b>0.000</b>           | <b>Not established</b> |

Table 14). The parameter  $m = 2,000$  is a sensitivity input to the enlarged-radius formula, not a count of independently drawn evaluation samples. All 108 cells are existing development records; this replay supplies no new independent sampling. The authenticated input and output receipts are `experiments/kbound/results/task2_imagenetc_population_replay_20260911/TASK2_IMAGENETC_POPULATION_REPLAY_RECEIPT.json` and `experiments/kbound/results/task2_population_panel_rerun_20260911_final.json`. This is a benchmark replay, not independent sampling or prospective population evidence.

## B Supplementary Results and Non-Promoted Studies

This section preserves the evidence outside the two primary protocols. The presentation separates ordinary accuracy from balanced accuracy; no value is pooled across metrics, datasets, candidates,

or protocols. The source-manifest identity is 3ac2c7d14d44803e4ff8a19258d036dc86327adb02fdb4cd34536e5eeee9cb24. Historical, incomplete, and invalid studies keep their original scope instead of being silently removed.

## B.1 Accuracy-Based Diagnostic Replays

Table 15 collects the remaining ordinary-accuracy replays; the dataset-specific interpretations follow below.

Table 15: Supplementary accuracy-regret rows. Lower regret is better. Commitment rate counts ADAPT or FREEZE and is not interval coverage. These are protocol-specific point summaries: Office-Home’s primary and stream-seed rows are different from its invalidated five-checkpoint audit; ImageNet-C is configuration-sensitive; PACS has aggregate-only replay. For PACS,  $n = 12$  counts domain-seed inference units, whereas action frequencies use 216 repeated decisions.

| Protocol                 | $n$ | KGA    | Adapt  | Freeze | FA <sub>u</sub> | Commitment rate |
|--------------------------|-----|--------|--------|--------|-----------------|-----------------|
| Office-Home primary      | 35  | 0.0158 | 0.0468 | 0.0158 | 0.0000          | 0.3143          |
| Office-Home stream seeds | 54  | 0.0215 | 0.0458 | 0.0217 | 0.0000          | 0.2778          |
| ImageNet-C Tent          | 135 | 0.0139 | 0.0191 | 0.0145 | 0.0000          | 0.0074          |
| ImageNet-C EATA          | 135 | 0.0031 | 0.0001 | 0.0342 | 0.0000          | 0.8000          |
| ImageNet-C SAR           | 135 | 0.0348 | 0.0529 | 0.0319 | 0.0148          | 0.0741          |
| PACS (aggregate)         | 12  | 0.0431 | 0.0176 | 0.0446 | 0.0093          | 0.8148          |
| CIFAR-10.1               | 48  | 0.0017 | 0.0190 | 0.0017 | 0.0000          | 0.4583          |

## B.2 Balanced-Accuracy Diagnostic Replays

Table 16 reports the balanced-accuracy replays on their own metric scale, with each recorded candidate and evaluation unit kept distinct.

Table 16: Supplementary balanced-accuracy-regret rows, kept separate from ordinary accuracy. The archived runner contracts name balanced accuracy for Camelyon17, ImageNet-R, and RxRx1; ordinary accuracy can coincide only under the recorded class-balancing conditions. ImageNet-R aggregates candidate backbones rather than one deployed gate. Camelyon17 OOD was already opened; its B-v2 rows are separate within-seed diagnostics. RxRx1’s displayed row is model seed 0.

| Protocol             | $n$ | KGA    | Adapt  | Freeze | FA <sub>u</sub> | Commitment rate |
|----------------------|-----|--------|--------|--------|-----------------|-----------------|
| ImageNet-R backbones | 480 | 0.0136 | 0.0064 | 0.0325 | 0.0063          | 0.4417          |
| Camelyon17 OOD       | 18  | 0.0000 | 0.0000 | 0.1381 | 0.0000          | 1.0000          |
| Camelyon17 B-v2 Tent | 108 | 0.0399 | 0.0097 | 0.0820 | 0.0185          | 0.3241          |
| Camelyon17 B-v2 EATA | 108 | 0.0235 | 0.0040 | 0.0911 | 0.0741          | 0.5648          |
| Camelyon17 B-v2 SAR  | 108 | 0.0240 | 0.0016 | 0.1001 | 0.0093          | 0.4630          |
| RxRx1 model seed 0   | 60  | 0.0000 | 0.2531 | 0.0000 | 0.0000          | 1.0000          |

## B.3 Large-Scale Corruption: ImageNet-C

ImageNet-C is candidate-dependent. SAR has KGA/adapt/freeze regrets 0.0348/0.0529/0.0319; on this 135-cell historical panel, KGA improves upon always-adapt (0.0348 vs. 0.0529) but trails always-freeze (0.0348 vs. 0.0319), with one false adaptation among 135 cells and without seed-robust inferential support. Tent ties the always-freeze point estimate and makes no strict decisions in the reconciled row. EATA nearly tracks always-adapt but has a larger KGA regret (0.0007 versus

Table 17: Completed corrected-BN SAR development comparison: five stream seeds over 27 aligned ImageNet-C conditions, not five independently trained models. Accuracy and paired-oracle regret average the 135 seed-condition records. A/F/U denote ADAPT/FREEZE/ABSTAIN; false A means a non-positive measured benefit on an ADAPT decision. FREEZE and ABSTAIN both serve the frozen model. This new comparison is separate from the historical 135-cell panel and 60-cell arm study.

| Policy             | Accuracy | Oracle regret | A / F / U    | False A |
|--------------------|----------|---------------|--------------|---------|
| Always freeze      | 0.4062   | 0.0320        | 0 / 135 / 0  | 0       |
| Always adapt       | 0.4107   | 0.0275        | 135 / 0 / 0  | 21      |
| Point-benefit gate | 0.4334   | 0.0047        | 108 / 27 / 0 | 6       |
| KGA                | 0.4120   | 0.0262        | 16 / 6 / 113 | 1       |

0.0001). These results show why ImageNet-C must be interpreted at the candidate-policy level rather than through one dataset-wide success label.

**Completed Corrected-BN Five-Stream-Seed Comparison (135 Cells).** The corrected series completed all 135 cells: the same 27 declared conditions for each of stream seeds 0–4, around one shared source checkpoint (ResNet-50). Saved-record digests, adjacent execution manifests, condition order, and raw-score projections were authenticated before aggregation. These checks do not authenticate the model and code bytes originally loaded by the historical producer processes. The maintained head-to-head CLI was run on the final seed, repeated byte-identically, and rerun after changing one scored candidate accuracy within its valid range. That cell’s prediction, radius, interval endpoints, threshold, and both gate actions remained unchanged. Seeds 0–3 retain their separately verified CLI replays. This is saved-feature, cell-outcome-disjoint verification, not environment-held-out or prospective validation. All point-policy decisions use KGA’s identical fitted predictions and the same zero threshold.

Table 17 shows a lower KGA point regret than either fixed policy (0.0262 vs. 0.0275 for always-adapt and 0.0320 for always-freeze), but substantially better point accuracy and regret for uncalibrated benefit regression (0.4334 acc, 0.0047 regret). KGA’s one false ADAPT occurs among 16 ADAPT decisions; the point gate’s six occur among 108. Thus fewer false events do not establish a lower conditional error rate or overall superiority. The declared bootstrap first averages five seeds within each of 27 aligned conditions. Its two-comparison Holm check supports lower regret than always-freeze ( $p = 0.0496$ ), but not always-adapt ( $p = 0.9397$ ); the latter regret-difference interval crosses zero. These are conditional, exploratory uncertainty summaries: they do not resample source models or establish independence among related corruption, severity, and composition conditions. The two-comparator success criterion is not met. Full input hashes, replay/intervention outputs, loss-less normalization, and aggregation are preserved in `experiments/kbound/results/sar_release_gate_20260913/five_seed_completion_20260914T042454Z/`. The aggregator’s legacy Protocol-A schema label does not describe the bound current three-way gate replay; the replay receipts specify that construction.

**Historical 60-Cell Matched-Control Matrix (Seed 0).** To resolve whether KGA’s decision value on SAR was confounded by an aggressive setting, an earlier development-only matched control matrix was reported on the same ImageNet-C source checkpoint (ResNet-50; shot and impulse noise, all five severities, three batch regimes with batch sizes {16, 32, 64}, 64 images per cell, two sample windows per condition, 60 cells per arm, seed 0). The raw 60-cell records and reconciliation manifest

Table 18: Historical configuration matrix for the five ImageNet-C SAR matched-control arms (ResNet-50, seed 0). The raw 60-cell records are not resident in the current checkout, so this matrix is not release-promotable. All reported arms perform test-time Batch Normalization statistics updates with per-batch model resets. Learning rate  $\eta$ , adapted affine parameters, and layer 4 status define each arm.

| Arm | Description              | Learning Rate $\eta$ | Adapted Layers    | Layer 4 State |
|-----|--------------------------|----------------------|-------------------|---------------|
| 1   | BN-statistics only       | 0.0 (no grad)        | None              | Frozen        |
| 2   | Reference SAR (official) | $2.5 \times 10^{-4}$ | Layers 1–3 affine | Frozen        |
| 3   | Rate only                | $4.0 \times 10^{-3}$ | Layers 1–3 affine | Frozen        |
| 4   | Mask only                | $2.5 \times 10^{-4}$ | Layers 1–4 affine | Adapted       |
| 5   | Historical combined      | $4.0 \times 10^{-3}$ | Layers 1–4 affine | Adapted       |

Table 19: Historical, non-promotable ImageNet-C SAR control summary on 60 development cells per arm (ResNet-50, shot/impulse noise [2 types], all five severities [5 levels], three batch sizes {16, 32, 64} [3 sizes], evaluated across 2 distinct sample windows per condition:  $2 \times 5 \times 3 \times 2 = 60$  cells; seed 0). The raw cell records and reconciliation manifest are not resident in the current checkout. The 5 arms and Protocol A definitions are retained for historical context only; current authenticated evidence is the separate 27-cell audit in Table 20.

| Arm                                    | Always Adapt | Always Freeze | KGA Acc | Regret A | Regret F | Regret KGA |
|----------------------------------------|--------------|---------------|---------|----------|----------|------------|
| BN-statistics only                     | 0.3984       | 0.3734        | 0.4125  | 0.0141   | 0.0391   | 0.0000     |
| Reference SAR ( $2.5 \times 10^{-4}$ ) | 0.3953       | 0.3734        | 0.4086  | 0.0156   | 0.0375   | 0.0023     |
| Rate only ( $4 \times 10^{-3}$ )       | 0.3594       | 0.3734        | 0.4031  | 0.0500   | 0.0359   | 0.0063     |
| Mask only (incl. layer4)               | 0.3945       | 0.3734        | 0.4055  | 0.0161   | 0.0372   | 0.0052     |
| Historical combined                    | 0.3534       | 0.3734        | 0.4031  | 0.0560   | 0.0359   | 0.0063     |

for that report are not resident in the current checkout, so it is retained as historical context and is not release-promoted. Table 18 defines the reported five-arm configuration matrix isolating learning rate and parameter mask factors. Table 19 reproduces that historical summary; it is not a current authenticated benchmark result.

The historical matrix was intended to isolate the rate-versus-mask question. Its reported reference SAR follows the official upstream implementation [5] on ResNet-50: SGD with learning rate  $\eta = 2.5 \times 10^{-4}$ , momentum 0.9, entropy threshold  $E_0 = 0.4 \times \ln(1000)$ , and gradient updates restricted to affine parameters in layers 1–3 (freezing layer 4), with per-batch model resets. The two aggressiveness settings isolate: (1) learning rate ( $\eta = 2.5 \times 10^{-4}$  reference vs.  $4.0 \times 10^{-3}$  aggressive), and (2) parameter mask (layers 1–3 adapted vs. all layers 1–4 adapted). The historical report stated that across 60 cells per arm, action exposure was: Reference SAR issued 46 ADAPT, 10 FREEZE, 4 ABSTAIN (zero false ADAPT; all 10 non-positive cells were frozen); Rate only issued 42 ADAPT, 14 FREEZE, 4 ABSTAIN; Mask only issued 45 ADAPT, 11 FREEZE, 4 ABSTAIN; Historical combined issued 42 ADAPT, 14 FREEZE, 4 ABSTAIN; and BN-statistics only issued 50 ADAPT, 6 FREEZE, 4 ABSTAIN. In that historical report, relative to reference SAR, the rate-only arm drops always-adapt accuracy by 0.0359 and increases its regret by 0.0344, whereas the mask-only arm changes adaptation accuracy by only  $-0.0008$ . The historical combination lowers always-adapt accuracy by 0.0419. The historical report further stated that under the published reference SAR parameters, KGA achieves mean accuracy 0.4086 (vs. 0.3953 for Always Adapt and 0.3734 for Always Freeze) with regret 0.0023, strictly outperforming both fixed policies on that development seed-0

Table 20: Saved-feature outcome-disjoint calibration audit on 27 historical, pre-repair seed-0 local SAR records (ResNet-50, reference-matched learning rate  $\eta = 2.5 \times 10^{-4}$  and layers 1–3 affine; 3 corruptions  $\times$  3 severities  $\times$  3 compositions = 27 cells; small batch regime, 25 helpful  $B > 0$ , 0 harmful  $B < 0$ , 2 neutral  $B = 0$ ). This replay does not repair upstream BatchNorm behavior. Clean Cross-Fit KGA uses Protocol B (5-fold outcome-disjoint cross-fitting,  $\alpha = 0.10$ ); Historical LOO KGA uses Protocol A (leave-one-cell-out exact rank,  $\alpha = 0.10$ ).  $\text{FA}_u = \mathbb{P}[\text{ADAPT}, B \leq 0]$ . Helpful Captured is the fraction of the 25 helpful cells ( $B > 0$ ) receiving ADAPT. Baseline accuracy  $a_0 = 0.3976$ .

| Policy                                  | Accuracy      | Regret        | $\text{FA}_u$       | Helpful Captured       | Actions (A / F / U) |
|-----------------------------------------|---------------|---------------|---------------------|------------------------|---------------------|
| Always Adapt                            | 0.4342        | 0.0000        | 0.074 (2/27)        | 25 / 25 (100.0%)       | 27 / 0 / 0          |
| Always Freeze                           | 0.3976        | 0.0366        | 0.000 (0/27)        | 0 / 25 (0.0%)          | 0 / 27 / 0          |
| Historical LOO KGA (Protocol A)         | 0.4318        | 0.0024        | 0.000 (0/27)        | 20 / 25 (80.0%)        | 20 / 0 / 7          |
| <b>Clean Cross-Fit KGA (Protocol B)</b> | <b>0.4301</b> | <b>0.0041</b> | <b>0.000 (0/27)</b> | <b>19 / 25 (76.0%)</b> | <b>19 / 0 / 8</b>   |

ResNet-50 checkpoint. Importantly, Table 19 also reported that the BN-statistics-only arm + KGA achieves mean accuracy 0.4125 (zero oracle regret), outperforming reference SAR + KGA (0.4086). If its underlying records are recovered and re-authenticated, this would make normalization-only adaptation an important alternative in this setting: test-time normalization statistics account for the bulk of adaptation gains on these noise families, indicating that SAR’s gradient updates do not add value beyond normalization adaptation alone on this panel. Cross-candidate regret must not be pooled as though candidates share a single oracle; task accuracy on common records provides the valid basis for comparison. Historical source pointer (raw 60-cell manifest not resident): `experiments/kbound/results/imagenetc_sar_official_control_v1/`.

**Historical Pre-Repair 27-Cell Calibration Audit (Seed 0).** As a calibration-audit check, we replayed the completed seed-0 reference-SAR records through the canonical three-way cell-outcome-disjoint cross-fitting implementation (Protocol B, `kga.crossfit.controlled_grid_crossfit`), without altering historical records or the matched-control table above. This saved-feature audit evaluates the 27 completed conditions from the historical, pre-repair seed-0 grid (3 corruptions [gaussian, shot, impulse noise]  $\times$  3 severities [1, 3, 5]  $\times$  3 compositions [iid, imbalanced, single\_class]; small batch regime; baseline  $a_0 = 0.3976$ ). Table 20 reports the full performance metrics across all decision policies.

On these 27 audit cells, Clean Cross-Fit KGA (Protocol B) achieves accuracy 0.4301 (shortfall of only 0.0041 regret vs. Always Adapt), zero false adaptations ( $\text{FA}_u = 0.000$ ), and captures 19 of 25 beneficial updates (76.0%), issuing 19 ADAPT, 0 FREEZE, and 8 ABSTAIN decisions. Historical LOO KGA (Protocol A) achieves accuracy 0.4318 (0.0024 regret, 20/25 helpful captured, actions 20/0/7). A whole-pipeline perturbation test changed one scored outcome by a large amount while leaving that cell’s fitted prediction, radius, and action unchanged, confirming that each scored cell’s outcome is strictly excluded from its own fit and calibration path under Protocol B. We emphasize that the resident 27-cell audit in Table 20 and the historical 60-cell matched-control summary in Table 19 represent distinct experimental configurations: the historical summary described 2 noise types  $\times$  5 severities  $\times$  3 batch sizes  $\times$  2 windows (and is not currently raw-record authenticated), whereas the resident audit evaluates 3 noise types  $\times$  3 severities  $\times$  3 compositions at small batch regime. The 27-cell audit is the current formal pipeline verification; it does not validate or replace the historical 60-cell summary, and it does not establish that seeds 1–4 are complete. An additional September 12 replay retains the original audit’s literal condition-string sample identifiers. It exactly reproduces every seed-0 prediction, radius, and action. For seeds 0/1/2, Protocol-B accuracy is

43.01/43.20/41.07%, with ADAPT/FREEZE/ABSTAIN counts 19/0/8, 21/0/6, and 13/0/14. For each of the 81 cells individually, replacing its saved benefit by a value within  $[-1, 1]$  and rerunning fitting and calibration leaves its prediction, radius, interval endpoints, and action unchanged; the decision boundary remains fixed at zero. This verifies exclusion from the saved-feature replay, not upstream feature construction, environment-held-out validity, or prospective coverage. None of these gates beats always-adapt. The full replay is stored in the September 12 condition-ID replay receipt in the task-closure evidence directory.

**Development Replay of Corrected-BN Seed 0.** An execution-path audit also found that these historical local runs excluded the final block from both affine optimization and BatchNorm reconfiguration. The pinned SAR source instead configures all BatchNorm layers to use batch statistics and excludes final-block parameters separately from optimization. Consequently, historical train-mode scoring could update final-block running buffers before eval-mode evidence extraction. The local implementation now separates these operations, with a regression-tested checkpoint-version guard preventing continuation from old records. The corrected series’ seed 0 completes the same declared 27 condition identifiers, with 23 positive and four negative candidate benefits. Its saved features and outcomes were replayed through the maintained eight-fold three-way construction, with  $\alpha = 0.10$ , fixed zero decision threshold, and the existing 11-feature gradient-boosted regressor (250 trees, depth 2, learning rate 0.05, subsample 0.8, fit seed 0). Each fold has nine residual-calibration cells and rank nine; the other roles are 14 fit/four score cells or 15 fit/three score cells. The point-benefit policy uses exactly the same fitted predictions as KGA. KGA issues 5 ADAPT, 1 FREEZE, and 21 ABSTAIN decisions, with accuracy 0.4146 and paired-oracle regret 0.0271. Always-adapt accuracy is 0.4186 (regret 0.0232), always-freeze accuracy is 0.4095 (regret 0.0322), and the same-prediction point policy reaches 0.4290 (regret 0.0128). KGA therefore does not beat both fixed policies. One of its five ADAPT decisions is false ( $FA_u = 1/27$ ,  $FA_c = 1/5$ ); the point policy makes three false ADAPT decisions among 23 ADAPT actions. For every scored cell, replacing its candidate accuracy by a different endpoint in  $\{0, 1\}$  with the frozen accuracy and features fixed, then rerunning all fits and calibration, leaves that cell’s prediction, radius, interval, threshold, and action unchanged. These 27 synthetic validity probes do not create new observations. The original records remain unchanged. This is a dependent, opened-grid development replay, not a native-method comparison, environment-held-out result, or independent source-training repetition. Full-precision records, software identities, and the executable replay are in `experiments/kbound/results/sar_bncorrected_seed0_disjoint_20260912/`. The maintained head-to-head CLI now also accepts an explicitly hash-bound saved-row checkpoint and a fresh output directory. Its independent CLI replay in `experiments/kbound/results/sar_bncorrected_seed0_cli_20260912/` reproduces the same predictions, radii, actions, and paired-outcome metrics. Its fallback POEM/AETTA rows remain local ports and are not external native evidence.

### B.3.1 ImageNet-C Decision Baselines Head-to-Head Comparison

To directly evaluate whether 1D heuristic gates (ATC confidence, predictive entropy progress, EATA reliable-sample filtering, and drift monitors) can match KGA on large-scale vision corruptions, we replayed all decision rules across the 108 ImageNet-C noise conditions (36 conditions each for Tent, EATA, and SAR) using identical sample streams and feature vectors. Table 21 reports the head-to-head results. The Drift Monitor row is a protocol-matched port monitoring prediction-distribution drift under the synchronized streaming harness, not native upstream POEM execution.

On the Tent and EATA panels, the harmful-update base rates are 6% and 3%, respectively. On the SAR panel, the harmful-update base rate is 44%; ATC confidence, entropy progress,

Table 21: ImageNet-C head-to-head decision baseline comparison across 36 noise conditions per candidate ( $N = 108$  total evaluations). Regret is shortfall relative to per-condition oracle; lower is better. Adapt rate is  $\mathbb{P}[\text{ADAPT}]$ . Directional commitment rate is  $\mathbb{P}[\text{ADAPT or FREEZE}]$ .  $\text{FA}_u = \mathbb{P}[\text{ADAPT}, B \leq 0]$ ;  $\text{FA}_c = \mathbb{P}[B \leq 0 \mid \text{ADAPT}]$ . Harmful adapted is the proportion of all harmful conditions ( $B < 0$ ) that received adaptation updates.

| Candidate                                      | Policy                                   | Mean Acc | Regret | $\text{FA}_u$ | $\text{FA}_c$ | Harmful Adapted | Worst Case | Adapt Rate | Commitment Rate |
|------------------------------------------------|------------------------------------------|----------|--------|---------------|---------------|-----------------|------------|------------|-----------------|
| <b>Tent</b><br>(Oracle: 0.455,<br>Harmful: 6%) | Always Adapt                             | 0.454    | 0.0007 | 0.056         | 0.056         | 1.00            | 0.164      | 1.00       | 1.00            |
|                                                | Always Freeze                            | 0.405    | 0.0494 | 0.000         | —             | 0.00            | 0.088      | 0.00       | 1.00            |
|                                                | ATC (confidence)                         | 0.443    | 0.0114 | 0.056         | 0.067         | 1.00            | 0.088      | 0.83       | 1.00            |
|                                                | Entropy Progress                         | 0.446    | 0.0089 | 0.056         | 0.062         | 1.00            | 0.088      | 0.89       | 1.00            |
|                                                | EATA Filter                              | 0.454    | 0.0007 | 0.056         | 0.056         | 1.00            | 0.164      | 1.00       | 1.00            |
|                                                | Drift Monitor (POEM port)                | 0.452    | 0.0028 | 0.056         | 0.057         | 1.00            | 0.088      | 0.97       | 1.00            |
|                                                | Benefit Regression ( $\varepsilon = 0$ ) | 0.455    | 0.0000 | 0.000         | 0.000         | 0.00            | 0.164      | 0.94       | 1.00            |
|                                                | <b>KGA (Ours)</b>                        | 0.449    | 0.0060 | <b>0.000</b>  | <b>0.000</b>  | <b>0.00</b>     | 0.164      | 0.69       | 0.69            |
| <b>EATA</b><br>(Oracle: 0.444,<br>Harmful: 3%) | Always Adapt                             | 0.444    | 0.0000 | 0.028         | 0.028         | 1.00            | 0.169      | 1.00       | 1.00            |
|                                                | Always Freeze                            | 0.405    | 0.0386 | 0.000         | —             | 0.00            | 0.088      | 0.00       | 1.00            |
|                                                | ATC (confidence)                         | 0.411    | 0.0328 | 0.028         | 0.100         | 1.00            | 0.088      | 0.28       | 1.00            |
|                                                | Entropy Progress                         | 0.412    | 0.0321 | 0.028         | 0.091         | 1.00            | 0.088      | 0.31       | 1.00            |
|                                                | EATA Filter                              | 0.444    | 0.0000 | 0.028         | 0.028         | 1.00            | 0.169      | 1.00       | 1.00            |
|                                                | Drift Monitor (POEM port)                | 0.441    | 0.0023 | 0.028         | 0.029         | 1.00            | 0.088      | 0.97       | 1.00            |
|                                                | Benefit Regression ( $\varepsilon = 0$ ) | 0.444    | 0.0000 | 0.028         | 0.028         | 1.00            | 0.169      | 1.00       | 1.00            |
|                                                | <b>KGA (Ours)</b>                        | 0.439    | 0.0047 | <b>0.000</b>  | <b>0.000</b>  | <b>0.00</b>     | 0.169      | 0.75       | 0.75            |
| <b>SAR</b><br>(Oracle: 0.432,<br>Harmful: 44%) | Always Adapt                             | 0.319    | 0.1124 | 0.444         | 0.444         | 1.00            | 0.061      | 1.00       | 1.00            |
|                                                | Always Freeze                            | 0.405    | 0.0267 | 0.000         | —             | 0.00            | 0.088      | 0.00       | 1.00            |
|                                                | ATC (confidence)                         | 0.365    | 0.0673 | 0.111         | <b>1.000</b>  | 0.25            | 0.088      | 0.11       | 1.00            |
|                                                | Entropy Progress                         | 0.355    | 0.0765 | 0.139         | <b>1.000</b>  | 0.31            | 0.088      | 0.14       | 1.00            |
|                                                | EATA Filter                              | 0.395    | 0.0369 | 0.028         | <b>1.000</b>  | 0.06            | 0.088      | 0.03       | 1.00            |
|                                                | Drift Monitor (POEM port)                | 0.394    | 0.0382 | 0.056         | <b>1.000</b>  | 0.12            | 0.088      | 0.06       | 1.00            |
|                                                | Benefit Regression ( $\varepsilon = 0$ ) | 0.431    | 0.0010 | 0.028         | 0.050         | 0.06            | 0.103      | 0.56       | 1.00            |
|                                                | <b>KGA (Ours)</b>                        | 0.409    | 0.0229 | <b>0.000</b>  | <b>0.000</b>  | <b>0.00</b>     | 0.088      | 0.06       | 0.39            |

EATA filtering, and drift monitoring each have observed  $\text{FA}_c = 1.000$ , meaning all their ADAPT commitments on this panel decrease accuracy. KGA has no observed false ADAPT decisions across these three candidate panels, with  $\text{FA}_u = 0.000$  and  $\text{FA}_c = 0.000$ . On SAR its commitment rate is 0.39 (61% abstention). These are descriptive results on the evaluated cells, not proof of protection on future shifts.

We also examine the point-benefit regressor ( $\varepsilon = 0$ ). On the SAR panel, uncalibrated benefit regression achieves mean accuracy 0.431 and regret 0.0010 with  $\text{FA}_u = 0.028$  ( $\text{FA}_c = 0.050$ , deploying harmful updates on 1 of 16 harmful conditions), whereas KGA achieves accuracy 0.409 and regret 0.0229 with strictly zero false adaptation ( $\text{FA}_u = 0.000$ ). On this specific SAR diagnostic, the point-benefit predictor achieves lower regret by accepting a non-zero false adaptation risk. The interval gate adds conservatism (issuing 2 ADAPT, 12 FREEZE, and 22 ABSTAIN decisions out of 36 conditions, corresponding to an adapt rate of 0.06 and commitment rate of 0.39), with no observed false ADAPT on this panel and a higher abstention rate. Zero observed false ADAPT under two ADAPT decisions is descriptive rather than large-sample confirmation. The interval rule controls the declared marginal directional error when its coverage premise holds. This diagnostic does not establish that premise under arbitrary drift.

## B.4 Office-Home

The primary Office-Home row [29] is conservative. KGA regret is 0.0158, tying always-freeze (0.0158) and below always-adapt (0.0468), with 0 adapt decisions. This is a descriptive point-estimate match that evaluates conservative retention rather than selective routing. The archived protocol does not supply a predeclared uniform stability guarantee for the full-fit versus out-of-fold benefit predictor.

The 54-cell Office-Home stream-seed replication in Table 15 has KGA regret 0.0215 versus 0.0458 and 0.0217, a small point advantage relative to each fixed policy. The three run seeds do not record independent checkpoint identities, and the seed-level interval includes zero, so the comparison

Table 22: Office-Home protocols are different experiments, not conflicting versions of one row. Entries are measured accuracy regret to the per-condition oracle; lower is better. The first row is the canonical single-candidate KGA replay with fixed SAR. The focused route is invalid, so only the post-hoc fixed-SAR opportunity regret (.0027) and freeze regret (.0230) are retained; its actions have no routing or certificate interpretation. A/F/U are adapt/freeze/abstain actions.

| Protocol                        | $n$ | Router / Adapt / Freeze | A/F/U   |
|---------------------------------|-----|-------------------------|---------|
| Archived stream-seed replay     | 54  | .0215 / .0458 / .0217   | 1/14/39 |
| Five-checkpoint candidate audit | 15  | invalid / .0027 / .0230 | n/a     |

remains descriptive.

The post-hoc five-checkpoint Office-Home study is a candidate opportunity audit, not a valid KGA result. It uses five independently trained source checkpoints, the three target-test domains, one fixed IID/tiny condition, and Tent, EATA, and SAR. Checkpoint identity passed, but routing and serialization did not. The identifying agreement relation is binary-only and does not apply to the 65-class task; the old spectral estimate was sign-indeterminate and unbounded. Thirteen of 15 anchors were negative, 9 of 60 stored  $\hat{b}$  values fell outside  $[0, 1]$ , and EATA and SAR produced exactly the same predictions in 10 of 15 cells. Each checkpoint has only three LOO cells, whereas a 90% empirical order-statistic single-candidate radius requires at least 10 total cells. All five result files fail strict JSON and contain 57 literal `Infinity` values in total.

The archived exploratory implementation printed ABSTAIN in all 15 cells, but those outputs have no safety, routing, or certificate meaning. The valid descriptive quantities are checkpoint accuracies and candidate opportunity: freeze is 0.5711, SAR is 0.5914, their paired difference is 0.0203 (nominal post-hoc 95%  $t$ -interval  $[0.0123, 0.0283]$ ), and the per-cell oracle is 0.5941. Oracle minus SAR is 0.0027 (nominal post-hoc 95% interval  $[-0.0024, 0.0077]$ ). The best adapted candidate improves over freeze in 13 cells, ties in one, and harms in one. Because the target was already opened and SAR was chosen after comparing the audit means, these statistics are post-hoc opportunity evidence only.

## B.5 Camelyon17

The archived Camelyon17 OOD evaluation [26] yields zero KGA and always-adapt regret against 0.1381 for freeze. All 18 cells are helpful, so this already-opened panel is a descriptive always-adapt one-sided diagnostic rather than prospective evidence of selective routing. Within-seed candidate diagnostics remain separate from this OOD endpoint. In the B-v2 SAR diagnostic, KGA regret is 0.0006, compared with 0.0016 for always-adapt and 0.1001 for always-freeze. All three run-seed point estimates favor KGA relative to each fixed policy, but the archived paired interval against always-adapt includes zero and the Holm-adjusted comparison does not pass. The run seeds also do not record independent checkpoint identities, and the target was already open. We therefore retain this as within-seed stress evidence rather than untouched natural-domain transfer.

## B.6 RxRx1

The displayed RxRx1 row [27] is model seed 0: KGA and freeze regret are zero against 0.2531 for always-adapt, with no adapt decisions. A separate audit across three independently trained model checkpoints repeats the same all-freeze tie. It strengthens model-variation reproducibility

but adds no adapt exposure. RxRx1 therefore evaluates reproducible conservative retention rather than selective routing.

## B.7 iWildCam

iWildCam [25] is not assigned a canonical numerical row. The archived scorer used sklearn macro-F1, which averages prediction-only classes, rather than the official WILDS metric, which averages labels present in the target truth. The archived runtime, sample identities, and target population were not sealed under the current contract, so its diagnostic scores and routing actions remain audit-only. A pinned rerun with the official metric, sealed sample identifiers, and a population hash is required before promotion. The older binary dominance flag and the later replay under the same invalid metric contract remain hash-locked only as superseded audit history.

## B.8 Boundary and Transfer Diagnostics

### B.8.1 PACS

PACS [30] KGA regret is 0.0431, compared with 0.0176 for always-adapt. Its archived aggregate lacks the residual detail needed for full gate replay, so it remains an aggregate performance diagnostic with incomplete certificate replay.

### B.8.2 ImageNet-R Backbone Breakdown

ImageNet-R [24] KGA regret is 0.0150, compared with 0.0064 for always-adapt, across the ten-backbone architecture panel. No backbone has lower KGA point regret than each fixed comparator. KGA regret is higher than always-adapt regret on 8 of 10 backbones; four of those eight have a 0% harmful base rate, so routing has no harm to avoid. On `efficientnet_b0` (100% harmful) and `swin_t` (60.4% harmful), KGA improves over always-adapt but exactly matches always-freeze and makes no adapt decisions. The pattern is consistent with the regime map: routing can add value only where candidate harm is present, while these two rows evaluate conservative retention rather than selective routing.

### B.8.3 CIFAR-10.1

CIFAR-10.1 [31] matches always-freeze at 0.0017 and is retained as a low-margin cross-seed transfer diagnostic. This pattern alone does not prove structural non-identifiability; weak evidence, low margin, estimator inadequacy, or calibration transfer can produce the same behavior.

## B.9 Native-Component Benchmark Comparison and Port Audits

We report a hash-bound native-component comparison on a retrospective 45-cell ImageNet-C subset (three noise corruptions [gaussian, shot, impulse], severities 1, 3, and 5, and five stream seeds; 15 score cells; Table 23). The panel uses 256 target images and 64 clean reference images per cell, with separate fit, calibration, and score partitions. The hash-bound score receipt is `experiments/kbound/results/task3_native_common_45_20260914_score/receipt.json`. It records 45 final records and 15 score cells, with `benchmark_completed=false` and `scoped_saved_output_comparison_completed=true`.

On the 15 available score cells, KGA and always-freeze both achieve equal-cell mean accuracy 0.4107; KGA commits on 10 of 15 cells (all FREEZE, zero false ADAPT,  $FA_u = 0.0000$ ), whereas the point-benefit gate reaches accuracy 0.4073 with two ADAPT actions ( $FA_u = 0.1333$ ,  $FA_c = 1.0000$ )

Table 23: Hash-bound native-component comparison on the retrospective 45-cell ImageNet-C subset (three noise corruptions, severities 1/3/5, five stream seeds; 15 score cells). Accuracy is the equal-cell mean of the served candidate or frozen predictions. “Commit” is the commitment rate  $(A + F)/n$ ; FA-U and FA-C are unconditional and conditional false-ADAPT rates. A method marked UNAVAILABLE is not silently substituted or averaged over a subset. Values are displayed to four decimals; full-precision values remain in the hash-bound JSON/CSV records. The AETTA row is an explicit controller over native AETTA component estimates; Kim TTA/ALine had no authenticated available score cells. Native numerical warnings were retained in the per-cell receipts.

| Method                | Status      | Available | Total | Accuracy | Commit | FA-U   | FA-C   |
|-----------------------|-------------|-----------|-------|----------|--------|--------|--------|
| AETTA_controller      | AVAILABLE   | 15        | 15    | 0.2953   | 1.0000 | 0.8667 | 0.8667 |
| KGA                   | AVAILABLE   | 15        | 15    | 0.4107   | 0.6667 | 0.0000 | –      |
| always_adapt          | AVAILABLE   | 15        | 15    | 0.2953   | 1.0000 | 0.8667 | 0.8667 |
| always_freeze         | AVAILABLE   | 15        | 15    | 0.4107   | 1.0000 | 0.0000 | –      |
| confidence_gain       | AVAILABLE   | 15        | 15    | 0.4107   | 1.0000 | 0.0000 | –      |
| entropy_progress      | AVAILABLE   | 15        | 15    | 0.4107   | 1.0000 | 0.0000 | –      |
| point_benefit         | AVAILABLE   | 15        | 15    | 0.4073   | 1.0000 | 0.1333 | 1.0000 |
| Baek_ALine_controller | UNAVAILABLE | 11        | 15    | –        | –      | –      | –      |
| Kim_TTALine           | UNAVAILABLE | 0         | 15    | –        | –      | –      | –      |

and the explicit AETTA controller collapses to accuracy 0.2953 with 13 false ADAPT actions ( $FA_u = FA_c = 0.8667$ ). The AETTA row evaluates an explicit controller over native component estimates, not an assumption that an accuracy estimator is itself a deployment policy. Baek ALine controller is available on 11 of 15 cells, and Kim et al.’s TTA/ALine source has no authenticated available score cells. Methods marked UNAVAILABLE are retained with explicit availability counts rather than imputed. These descriptive values are regenerated from hash-bound score records and do not establish superiority beyond this scoped subset panel.

**Separation from Historical Scalar Proxies.** The repository also retains an earlier historical head-to-head artifact on the locked CIFAR-10-C stress stream, whose KGA actions came from an earlier policy and served as developmental history. Historical protocol-matched POEM and AETTA rows are retained as non-native diagnostics; they are not official upstream executions and are not promoted as external-method evidence. For local decision rules, Section 8.3 (Table 9) reports a regenerated Protocol-B comparison across 2,160 CIFAR-10-C rows using fixed policies, scalar evidence rules, and a same-prediction point gate. The historical port names refer to local scalar comparators inspired by TTA/ALine [51], AETTA [18], and POEM [14], not validated equivalents of those algorithms: `poem_port` tests whether entropy decreases and `aetta_port` tests whether post-adaptation confidence exceeds pre-adaptation confidence. These 1D scalar heuristics do not execute a native protector, accuracy-estimation procedure, or multi-model agreement fit, and cannot substitute for an authentic comparison.

**Component Engineering Smoke Tests.** Separate CPU harnesses execute authenticated POEM and AETTA components on a fixed 104-image Gaussian-noise severity-5 stream (batch size 4, separate fresh-process repetitions; receipts under `experiments/kbound/results/task_closure_20260911/`). POEM used its full pretrained ResNet-50-GN model, 32 clean reference images, and the native 100-sample warmup; a reduced reference set and disclosed CPU constructor change make this a smoke test, not unchanged upstream-driver execution. AETTA used the full torchvision ResNet-50 ImageNet V2 checkpoint and its native four-stage dropout estimator with ten draws. All

26 batches reached the native corrected-error cap. These checks removed a blanket CPU-execution objection, but do not establish CUDA/MPS parity, full ImageNet-C coverage, or external-method superiority.

**Execution Barriers and Audit Location Map.** Full native execution across the complete ImageNet-C benchmark remains prevented by concrete upstream software and hardware barriers: (1) *Dependency integration*: the official POEM release omits `models/Res.py`; (2) *Device integration*: hardcoded, unconditional CUDA assertions in upstream entrypoints prevent execution on Apple Silicon MPS; (3) *Data readiness*: the clean ImageNet validation archive (50,000 images, 1,000 classes) was extracted on external storage and verified against the official ILSVRC2012 development kit; (4) *Stream standardization*: native drivers iterate over fixed corruption sequences rather than accepting arbitrary streaming windows. To avoid cluttering the scientific presentation, the complete execution history, environment hashes, source-authentication logs, and independent witness receipts are preserved in Appendix N (*Reproducibility and Release Procedures*).

The standalone head-to-head analysis script uses the repository’s disjoint fit/calibration/score construction (Protocol B). A synthetic 40-cell regression reruns the full gate calculation after perturbing each scored outcome and confirms that predictions, radii, interval endpoints, and actions remain invariant. Saved-feature outcome exclusion confirms computational separation, but does not establish exchangeability of dependent corruption cells or population coverage under future shifts.

## B.10 So2Sat Development and Action-Unit Audit

The locked So2Sat-LCZ42 v1 study evaluated 2 candidates, Tent and SAR, on 9 development cities and 5 independent source checkpoints, giving 45 city-checkpoint cells per candidate. This was a candidate-feasibility stage, not target evaluation. Tent had 4 helpful and 4 harmful cities and a cell-oracle opportunity of 0.8848 percentage points. Its leave-one-city-out sign agreement was 51.11%, and its routed score was 0.1544 percentage points below always-adapt, its better fixed policy. SAR had 0 qualifying helpful cities, 6 harmful cities, and an oracle opportunity of 0.5640 percentage points. It selected ADAPT in only 5 of 45 cells, and its routed score was 0.0842 percentage points below always-freeze. Tent passed 7 of 9 feasibility checks and SAR passed 6 of 9, so neither candidate was eligible.

The locked status records no feasible candidate and a stop before gate calibration. No held-out gate-calibration rows or target inputs were accessed; the recorded counts were 0 gate-calibration rows, 0 target inputs, 0 target pixels, and 0 target labels. The study therefore documents failure to meet the v1 cross-city development-transfer feasibility criteria while also showing that a cell-oracle routing opportunity existed for Tent. It contains no target natural-shift score. All 90 development outcomes are now open, so retuning on them cannot produce fresh confirmatory evidence. In addition, v1 declares one target action per city while the target runner creates city-checkpoint actions. This mismatch did not affect the stop, because no target action was created, but target execution remains blocked until a new versioned protocol and implementation use the same action unit.

## B.11 Withdrawn Development Proxy for the Population Budget

K-Bound characterizes population benefit over a declared residual class  $\mathcal{C}_\beta$ . KGA instead predicts measured evaluation-cell benefit  $\Delta^{\text{cell}}$  and uses a separately calibrated residual interval. Its direct mathematical action implication is Proposition 2, conditional on coverage for that declared target. It neither estimates the population residual budget  $\beta$  nor numerically evaluates the population frontier.

Table 24: Claim-to-support ledger for the compact submission. Support pairs each evidence reference with the conditions required for its stated scope.

| Claim and type                                         | Support and requirement                                            | Boundary                                                                                                                          |
|--------------------------------------------------------|--------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| Interior evidence ambiguity (theorem)                  | Thm. 1; $\beta > 0$ , $ M  < \beta$ , and full fibre richness      | binary, 0/1 loss                                                                                                                  |
| Closed-band abstention (proposition)                   | Prop. 1; strict directional semantics and a rich class             | zero benefit blocks a strict claim                                                                                                |
| Strict-commitment frontier (theorem)                   | Thm. 2; fixed augmented evidence and declared $\mathcal{C}_\beta$  | $\beta$ is externally supplied                                                                                                    |
| Marginal false-adapt control (operational proposition) | Prop. 2 under Eq. (9) for the measured cell target                 | neither population-risk nor conditional $\text{FA}_c$ control follows                                                             |
| CIFAR-10-C Tent regret (empirical)                     | Table 7 and 26; frozen Tent controlled-grid protocol               | lower empirical order-statistic point regret; nominal family-bootstrap intervals favor KGA; retrospective Holm is nonconfirmatory |
| CIFAR-10-C EATA regret (empirical)                     | Table 7 and 26; frozen EATA controlled-grid protocol               | adapt-side family interval includes zero; retrospective Holm is nonconfirmatory                                                   |
| Natural-shift conservative rows (empirical)            | Tables 15, 16, and 10; each split and candidate remain fixed       | often zero ADAPT exposure                                                                                                         |
| So2Sat development stop (diagnostic)                   | Sec. 8.5; locked v1 feasibility rules and unopened target boundary | no eligible candidate, gate calibration, or target score                                                                          |
| ImageNet-C candidate scope (diagnostic)                | Table 15; candidate rows interpreted separately                    | candidate-specific evidence                                                                                                       |
| PACS/ImageNet-R/CIFAR-10.1 scope (diagnostic)          | Tables 15 and 16; canonical replay scope retained                  | observed behavior does not prove unknowability                                                                                    |
| Deployment generalization (future evidence)            | Sec. K; diverse unopened natural targets                           | outside current evidence scope                                                                                                    |
| Real-camera protocol (prospective)                     | locked sessions and fresh held-out recordings                      | protocol infrastructure only; no result in this paper                                                                             |

Appendix A gives a separate population-transfer implication requiring an additional sampling-error bound  $b(m, \delta)$  and an enlarged radius  $\varepsilon + b$ . The currently reported transductive benchmark panels do not establish these additional premises.

For the historical  $\beta$  proxy, Section 9 discloses the negative result. The historical summary (1.4–50× undersizing, 24–73% proxy interval violations, and 0.4–16.6% decision errors on CIFAR-10-C) represents the operational failure of a development-based proxy heuristic, not a verified numerical measurement of latent theoretical-class violations on evaluation cells. The raw sweep artifact is unavailable, and this estimation route is completely withdrawn. A credible budget  $\beta$  must come from domain knowledge, a declared coupling, or historical labeled deployments under an explicit transfer assumption. Setting  $\beta = 0$  assumes zero target disagreement-conditional residual; it is not a safe default.

## C Compact Assumption and Claim Ledger

Table 24 separates formal claims from empirical findings and records the conditions and scope limits for each.

## D Calibration and Evaluation Contracts

**Controlled-grid contract.** The fitting unit is the condition cell. The current primary CIFAR and synchronized internal-baseline tables use Protocol B: each score fold is excluded from both estimator fitting and residual calibration, and those two subsets are disjoint. Historical Protocol-A tables instead use leave-one-cell-out predictions and peer-dependent residual calibration; excluding a cell’s own residual does not remove its indirect influence through other fitted predictions. The execution mapping in N identifies the construction for each table. Neither construction alone establishes exchangeability across corruption families, checkpoints, or future environments. Saved-feature outcome-exclusion tests verify their stated computational boundary, not full upstream independence.

**Natural-shift contract.** Development conditions fit the benefit model; a disjoint calibration partition supplies residuals; choices are frozen; and the target partition is scored once where the archived protocol records all three stages. When the saved evidence does not support that complete chain, the row is marked descriptive or diagnostic. Replaying a table does not strengthen its original design.

The candidate TTA step itself is transductive. Episodic candidates update on each unlabeled evaluation batch before predicting that batch. Online candidates update on a separate auxiliary stream, but prediction still uses evaluation-batch BatchNorm statistics. Thus “disjoint” above refers to development, residual-calibration, and target partitions (and to the recorded auxiliary stream versus evaluation sample identities); it does not mean that candidate inference is inductive. No target labels enter adaptation, prediction, evidence, or the action.

**Action contract.** Every action is computed from  $Z$ ,  $\hat{\Delta}$ , and  $\varepsilon$  without target labels. Offline labels determine  $\Delta^{\text{cell}}$ , the measured oracle action, empirical regret, and false-commit indicators only after the action is fixed. Existing artifact fields named `delta` store this cell outcome; their schema and numerical values are unchanged. Abstention resolves to the frozen prediction for measured task-performance accounting, while its log retains the distinction from a certified freeze.

**Configuration contract.** The checkpoint, candidate adapter, optimizer, adapted parameter subset, update count, evidence schema, feature normalization, benefit estimator, calibration split,  $\alpha$ , and software version are part of one calibration identity. Changing any of them requires either a declared transfer argument or recalibration. The fail-closed default on identity mismatch is non-commitment.

## E Decision Exposure and Repetition Audit

Point regret alone can conceal whether a conservative match exercised the ADAPT branch. Table 25 therefore reconstructs action counts from the same canonical panel used in Table 7, 15, and 16. A/F/U denote adapt, freeze, and abstain. PACS is omitted from the action table because the released aggregate cross-checks regret but does not contain the per-cell benefit estimates and residuals required for a complete gate replay. iWildCam is omitted because its archived actions were scored under the wrong macro-F1 contract.

The largest adapt exposure occurs on the controlled CIFAR-10-C rows. This is why those rows are the main empirical test of the adapt-side certificate. Primary Office-Home and RxRx1 have zero adapt exposure, while the 54-cell Office-Home stream-seed replication has one adapt decision. Their zero false-adapt values barely test whether the gate can distinguish a harmful adaptation from a

Table 25: Action exposure in replayable canonical rows. Counts are not independent experiments: they are evaluation units within the protocol shown. A zero false-adapt count is informative only when adaptation was attempted and the residual pool did not make the inequality automatic. Harm and false adaptation are defined by measured cell outcomes, not population risks. Counts were regenerated from source manifest SHA-256 3ac2c7d14d44803e4ff8a19258d036dc86327adb02fdb4cd34536e5eeee9cb24.

| Row                           | $n$  | A    | F   | U   | false A | FA <sub>u</sub> | repetition / scope                                                  |
|-------------------------------|------|------|-----|-----|---------|-----------------|---------------------------------------------------------------------|
| Office-Home primary           | 35   | 0    | 11  | 24  | 0       | .0000           | 2 test-stream seeds; ties freeze                                    |
| Office-Home stream-seed repl. | 54   | 1    | 14  | 39  | 0       | .0000           | 3 test-stream seeds; point-only edge                                |
| ImageNet-C SAR                | 135  | 13   | 15  | 107 | 1       | .0074           | 5 seeds $\times$ 27 cells                                           |
| ImageNet-C Tent               | 135  | 0    | 0   | 135 | 0       | .0000           | 5 seeds $\times$ 27 cells                                           |
| ImageNet-C EATA               | 135  | 120  | 1   | 14  | 1       | .0074           | 5 seeds $\times$ 27 cells                                           |
| ImageNet-R                    | 480  | 165  | 29  | 286 | 0       | .0000           | 4 seeds $\times$ 10 backbones $\times$ 12 cells                     |
| CIFAR-10-C Tent               | 2160 | 1094 | 334 | 732 | 0       | .0000           | 5 seeds $\times$ 432 cells                                          |
| CIFAR-10-C EATA               | 2160 | 1207 | 118 | 835 | 0       | .0000           | 5 seeds $\times$ 432 cells                                          |
| CIFAR-10-C SAR                | 2160 | 1430 | 0   | 730 | 1       | .0005           | 5 seeds $\times$ 432 cells                                          |
| Camelyon17 OOD                | 18   | 18   | 0   | 0   | 0       | .0000           | archived opened helpful-only OOD subset                             |
| Camelyon17 B-v2 Tent          | 108  | 40   | 0   | 68  | 1       | .0093           | 3 within-seed grids; diagnostic                                     |
| Camelyon17 B-v2 EATA          | 108  | 61   | 0   | 47  | 0       | .0000           | 3 within-seed grids; diagnostic                                     |
| Camelyon17 B-v2 SAR           | 108  | 71   | 0   | 37  | 0       | .0000           | 3 within-seed grids; diagnostic                                     |
| RxRx1 J                       | 60   | 0    | 60  | 0   | 0       | .0000           | displayed model seed 0; separate 3-checkpoint audit also all-freeze |
| CIFAR-10.1 K                  | 48   | 0    | 22  | 26  | 0       | .0000           | locked cross-seed replay                                            |

helpful one; they show that the controller declines the candidate under those evidence/calibration contracts. The invalidated five-checkpoint Office-Home route and the withheld iWildCam actions are not counted. Camelyon17 OOD has the opposite limitation: every archived condition is helpful and every action adapts, so there is no harmful opportunity to reject. ImageNet-C EATA has substantial adapt exposure and one false adaptation, while Tent abstains everywhere. Those three patterns are operationally different even when their regrets are small.

The repetition column also prevents a common inflation of evidence. Five seeds of one checkpoint and grid protocol are not five independently designed studies. Ten ImageNet-R backbones are a panel of candidate policies, not ten target domains for one frozen gate. Three Camelyon17 within-seed grids permit a stability diagnostic but do not create an untouched hospital-domain replication. RxRx1 is stronger on model variation because the source checkpoint is independently trained three times, but its all-freeze behavior still provides no adapt exposure.

## F Inference Discipline

### F.1 Point Estimates, Paired Gaps, and Clusters

For a baseline  $b \in \{\text{always-adapt}, \text{always-freeze}\}$ , the paired routing gap at condition  $j$  is

$$G_{bj} = \text{regret}_{b,j} - \text{regret}_{\text{KGA},j}.$$

A positive mean gap means lower regret for KGA. All paired routing-gap intervals reported in this paper use this baseline-minus-KGA sign convention. The point-estimate two-comparator criterion requires both means to be positive. Interval support additionally requires the lower confidence limit for both paired means to exceed zero at the declared resampling unit. These definitions are intentionally separate: Table 15 can establish a point ordering, but not an inference claim without

Table 26: Retrospective current-policy sensitivity over six observed CIFAR-10-C corruption families. Each cell reports baseline regret minus KGA regret and a nominal, unadjusted 95% family-bootstrap interval; positive values favor KGA. Holm values adjust retrospective sign-flip reference values over six prospectively named candidate-by-baseline contrasts. The replay, corruption-family unit, sign-flip test, and Holm adjustment were not preregistered. The within-Tent two-contrast Holm value is 0.03125 but is post hoc. This table is conditional on one archived checkpoint and does not provide verified finite-sample coverage, simultaneous intervals, independent-checkpoint inference, prospective confirmation, natural-shift evidence, or official POEM/AETTA superiority.

| Candidate | Adapt gap [95% CI]            | Freeze gap [95% CI]        | Six-way Holm $p$ (A/F) |
|-----------|-------------------------------|----------------------------|------------------------|
| Tent      | 0.00618 [0.00231, 0.00954]    | 0.12214 [0.06304, 0.17767] | 0.09375/0.09375        |
| EATA      | 0.00167 [-0.00127, 0.00438]   | 0.12973 [0.06901, 0.18748] | 0.31250/0.09375        |
| SAR       | -0.00133 [-0.00274, -0.00016] | 0.13881 [0.08003, 0.19685] | 0.90625/0.09375        |

the associated unit, interval artifact, and justified coverage assumptions. A positive computed lower endpoint is numerical interval support under the resampling model, not verified coverage. Multiplicity adjustments do not repair invalid marginal intervals or tests.

Condition-level resampling treats cells as exchangeable and can understate uncertainty when severities and compositions share a corruption family. Corruption-family resampling preserves observed within-family dependence, but few clusters limit reliability; conservativeness is not guaranteed. Seed resampling measures variation across repeated training or test streams only when the seeds represent the intended deployment randomness. The compact paper therefore does not use one generic “bootstrap CI” label. An inference claim must state whether the unit is cell, corruption-severity pair, corruption family, seed, backbone, domain, or physical session.

## F.2 Current-Policy Corruption-Family Sensitivity

The six candidate-by-baseline contrasts were named before this analysis; the replay policy, corruption-family unit, sign-flip reference test, and Holm adjustment were selected retrospectively. The current-policy artifact averages run seeds and condition cells within each of the six families. Table 26 retains every contrast, including EATA’s adapt-side interval crossing zero and SAR’s disadvantage against always-adapt. The within-Tent two-contrast Holm calculation is post hoc and is not the declared six-contrast family.

## F.3 Retrospective Interval Inclusion by Family

The following audit replays the current policy from stored, cross-fitted predictions and measured cell outcomes. It uses each candidate’s original residual rule and excludes the scored residual from its own radius. No fitting, tuning, new target access, or training is performed. The input, policy, code, and output digests are recorded in the generated diagnostic JSON.

The pooled inclusion fraction is constrained by the overlapping leave-one-cell-out order statistics. It is not an independent estimate of coverage on new environments. Family rows reveal the distribution of those same misses, widths, and decisions; they do not supply group-conditional guarantees. A nominal 90% target is a calibration setting, not a promise validated by this replay.

Table 27: Current-policy interval and decision diagnostics by corruption family. Width is full interval width,  $2\epsilon$ . Inclusion is measured on already-open, dependent cross-fitted cells. Conditional errors use the number of the corresponding action as denominator and are undefined when that action is absent. These diagnostics are not held-out group-coverage validation.

| Candidate | Family           | $n$ | Inclusion | Mean width | Commitment | False A/A; FA <sub>c</sub> | False F/F; FF <sub>c</sub> |
|-----------|------------------|-----|-----------|------------|------------|----------------------------|----------------------------|
| TENT      | contrast         | 360 | 0.8111    | 0.0468     | 0.6917     | 0/180; 0.0000              | 0/69; 0.0000               |
| TENT      | defocus_blur     | 360 | 0.9389    | 0.0472     | 0.7111     | 0/180; 0.0000              | 0/76; 0.0000               |
| TENT      | fog              | 360 | 0.9278    | 0.0469     | 0.6944     | 0/180; 0.0000              | 0/70; 0.0000               |
| TENT      | gaussian_noise   | 360 | 0.8750    | 0.0471     | 0.8500     | 0/305; 0.0000              | 0/1; 0.0000                |
| TENT      | jpeg_compression | 360 | 0.9611    | 0.0466     | 0.3500     | 0/69; 0.0000               | 0/57; 0.0000               |
| TENT      | pixelate         | 360 | 0.9139    | 0.0466     | 0.6694     | 0/180; 0.0000              | 0/61; 0.0000               |
| EATA      | contrast         | 360 | 0.8417    | 0.0381     | 0.5778     | 0/180; 0.0000              | 0/28; 0.0000               |
| EATA      | defocus_blur     | 360 | 0.9194    | 0.0385     | 0.5944     | 0/180; 0.0000              | 0/34; 0.0000               |
| EATA      | fog              | 360 | 0.9000    | 0.0387     | 0.5778     | 0/180; 0.0000              | 0/28; 0.0000               |
| EATA      | gaussian_noise   | 360 | 0.8722    | 0.0386     | 0.9750     | 0/350; 0.0000              | 1/1; 1.0000                |
| EATA      | jpeg_compression | 360 | 0.9806    | 0.0390     | 0.4056     | 0/137; 0.0000              | 0/9; 0.0000                |
| EATA      | pixelate         | 360 | 0.9222    | 0.0392     | 0.5500     | 0/180; 0.0000              | 0/18; 0.0000               |
| SAR       | contrast         | 360 | 0.8389    | 0.0280     | 0.5000     | 0/180; 0.0000              | 0/0; –                     |
| SAR       | defocus_blur     | 360 | 0.9222    | 0.0278     | 0.5000     | 0/180; 0.0000              | 0/0; –                     |
| SAR       | fog              | 360 | 0.9167    | 0.0278     | 0.5000     | 0/180; 0.0000              | 0/0; –                     |
| SAR       | gaussian_noise   | 360 | 0.8028    | 0.0279     | 1.0000     | 0/360; 0.0000              | 0/0; –                     |
| SAR       | jpeg_compression | 360 | 0.9722    | 0.0279     | 0.9306     | 0/335; 0.0000              | 0/0; –                     |
| SAR       | pixelate         | 360 | 0.9167    | 0.0281     | 0.5417     | 1/195; 0.0051              | 0/0; –                     |

## F.4 False-Adapt Accounting

The event in Eq. (4) is unconditional and concerns measured cell benefit. Its reported frequency divides ADAPT decisions with  $\Delta_j^{\text{cell}} \leq 0$  by all evaluated units, including freezes and abstentions. This need not estimate the population-harm event  $\{g = \text{ADAPT}, \Delta \leq 0\}$ . Conditional false-adapt divides by adapt decisions and is undefined when no adaptation occurs. A zero observed count should be accompanied by the number of adapt decisions; otherwise a controller that never adapts appears perfect. Table 25 exposes that denominator.

If the scored residual is included in the same finite residual vector whose order statistic decides the action, the number of residuals above the order statistic is algebraically constrained. Such coverage is not an independent empirical validation. Excluding the scored residual removes direct self-inclusion, but pooled LOO inclusion remains rank-constrained. It is not independent validation and does not supply the exchangeability or stability needed for a finite-sample theorem. Clean split calibration and LOO empirical calibration therefore remain different protocol labels.

## F.5 Multiplicity and Search History

An earlier repository census counted many historical two-comparator flags, but it did not freeze the file list and commit needed to reproduce a current total. We therefore do not promote that count or its positive rate as release evidence. The panel in this paper is a post-hoc reconciled audit panel assembled from individually locked and historical protocols; it is not a globally preregistered confirmatory family. A result selected after a search cannot become confirmatory by correcting only the final positive rows. Future confirmatory families must name every candidate-track test before any test value is computed, and must retain failures and incomplete runs.

**Search history.** The repository contains many exploratory configurations. One superseded exploratory census reported 1,387 determinations and 326 positives (23.5%); a broader archived

Table 28: Provenance and availability status of experimental tracks in the compact release. “Replay” means the stored quantities suffice to recompute the displayed gate output under the released scorer.

| Track           | Released status         | Present evidence                                                                                   | Remaining limitation                                                                           |
|-----------------|-------------------------|----------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------|
| CIFAR-10-C      | replayable              | per-condition $\Delta^{\text{cell}}$ , $\hat{\Delta}$ , radius, action; 5 seeds; 3 candidates      | related grid cells; no untouched corruption-family test in canonical table                     |
| ImageNet-C      | replayable              | per-condition records for 5 seeds and Tent/EATA/SAR                                                | 27-cell controlled configuration; candidate-specific, not one policy                           |
| Office-Home     | replayable              | canonical primary and 54-cell stream-seed records; later five-checkpoint audit reported separately | full-fit-versus-OOF transfer-stability premise not predeclared                                 |
| iWildCam        | withheld                | archived opened test-split records; non-promoted official-metric diagnostic                        | wrong archived metric; population was not sealed and overlap is unauditale                     |
| Camelyon17      | replayable diagnostic   | archived opened OOD row plus separate B-v2 diagnostics                                             | OOD row is one-sided and non-prospective; B-v2 is not untouched-domain evidence                |
| RxRx1           | replayable              | displayed model-seed-0 row plus three model-checkpoint records and source hashes                   | all three checkpoints freeze throughout; no helpful test conditions                            |
| CCT-20          | sealed complete         | 5 independent checkpoints, 9 target locations, actions, predictions, score, and inference hashes   | only two fit locations; no ADAPT exposure; trans-target extrapolation                          |
| So2Sat-LCZ42    | sealed development stop | 2 candidate bundles, 9 development cities, 5 checkpoints, selection record, and 0 target inputs    | no feasible candidate; all 90 gate-fit outcomes now open; no gate calibration or target result |
| PACS            | partial                 | three-seed aggregate and source cross-checks                                                       | missing per-cell $\hat{\Delta}$ and residuals prevent full gate replay                         |
| ImageNet-R      | replayable              | 4 seeds, 10 candidate backbones, 12 conditions each                                                | architecture aggregate is diagnostic, not one deployable router                                |
| CIFAR-10.1      | replayable              | locked cross-seed condition records                                                                | low-margin conservative match; selective routing not exercised                                 |
| Physical camera | protocol only           | dashboard, logging schema, planned session split                                                   | no fresh held-out physical recordings or result table                                          |

scan reported 1,427 determinations and 345 positives. Because neither census froze one reproducible family definition, file list, and commit, neither total is promoted as the release multiplicity family. We instead report a current post-hoc reconciled panel and do not relabel its favorable rows as confirmatory discoveries. Future claims should declare the candidate, calibration unit, target unit, seed structure, inference unit, and two-comparator success criterion before test labels are opened.

## G Track-Level Provenance and Remaining Evidence

Separate archival records preserve two additional scope checks outside the canonical panel. FMoW Protocol L reached held-out scoring and matched always-freeze, while always-adapt had lower regret; its status remains not-cleared. PovertyMap stopped at its preregistered development screen because the harm-detection AUC was 0.637, below the 0.65 threshold, so held-out validation and test scoring did not run. These records remain in the claim inventory and are not pooled with the canonical rows.

The provenance table is part of the empirical record, not administrative metadata. An auditor or reader can trace a canonical row to compact records and their original archive hashes. If the original large artifact is not shipped, the manifest still records its byte count and digest. This guards against silent replacement of a result while allowing the release to omit raw datasets and oversized transient checkpoints.

Three upgrades would materially strengthen the empirical case. First, hold out entire shift

Table 29: Locked CCT-20 primary contrasts. The estimand is baseline regret minus KGA regret, so positive values mean lower regret for KGA. Intervals use a paired two-way product bootstrap that resamples checkpoint rows and location columns independently. Two nominal 97.5% per-comparison intervals target a nominal 95% Bonferroni family level; actual family coverage requires valid marginal coverage. Reference  $p$ -values enumerate one-sided sign flips of the nine location means under Eq. (12). Holm adjusts only these reference values; the threshold flags reproduce the locked calculation, not a verified population guarantee.

| Comparator    | Baseline regret – KGA regret | Nominal 97.5% CI         | Sign-flip $p$ | Holm $p$   | Holm flag (.05) |
|---------------|------------------------------|--------------------------|---------------|------------|-----------------|
| Always adapt  | 0.18190227                   | [0.08501071, 0.31066374] | 0.00195312    | 0.00390625 | yes             |
| Always freeze | 0                            | [0, 0]                   | 1             | 1          | no              |

families rather than individual controlled cells and fit all hyperparameters before opening the family. Second, add independent model and test-session seeds to the natural tracks that currently have one target replay. Third, run the physical protocol with days, devices, and objects assigned at the session level, then score every policy on the same locked streams. These upgrades address dependence and transfer; simply adding more rows from the same grid does not.

## G.1 CCT-20 Inference, Location, and Release Ledger

**Provenance ledger.** Cell-level outcomes remained unopened before target execution. Aggregate target label metadata had already been inspected during candidate ranking, so this is not a label-unopened study. Item-level target labels were deserialized once, after the one-shot scoring marker was spent. The release chain also records the numeric-v2 stabilization in a sealed addendum created before outcome access. That stabilization replaced an unstable divergence calculation with a stable Gaussian KL computation; it did not change the model, data split, or scoring rule. We treat it as protocol provenance, not result-driven tuning. All primary outcome values and action/effect counts come from the sealed, content-addressed, receipt-linked CCT-20 release manifest and its generated macros (canonicalized inference-document SHA-256 5d1ba74d9a8b16ef24746ca3d95969e635d34fd4112806f6bf91b4b45a6cf907).

Finite-sample sign-flip validity requires the null joint law of the nine location gaps to satisfy Eq. (12). Independent zero-symmetric location gaps conditional on the evaluated checkpoints would suffice, but that assumption is unverified here. Exchangeability alone does not suffice. The sign-flip analysis conditions on the evaluated checkpoints, whereas the two-way bootstrap also resamples checkpoint levels. These are different uncertainty models, and neither establishes distribution-free safety for unseen locations.

**Primary result.** KGA exactly matched always-freeze. Its contrast against always-freeze was 0 with nominal Bonferroni-adjusted 97.5% bootstrap interval  $[0, 0]$ , enumerated location sign-flip  $p = 1$ , and Holm  $p = 1$ . Against always-adapt, the contrast was 0.18190227 with interval  $[0.08501071, 0.31066374]$ , sign-flip reference  $p = 0.00195312$ , and Holm  $p = 0.00390625$ . Measured cell benefit is adapted accuracy minus frozen accuracy: 1 cell was helpful, 0 were exactly zero, and 44 were harmful. These counts explain why freezing was useful here. They do not show that the gate learned when to adapt.

**Limits and next step.** All 55 development cells showed lower performance after Tent than under the frozen model, so the gate had no helpful ADAPT example. The ridge fit used 10 checkpoint–location rows from only two camera units for 11 recorded features. Calibration used

Table 30: Complete locked CCT-20 location ledger. A/F/U denotes ADAPT/FREEZE/ABSTAIN, and  $+/0/-$  counts positive/zero/negative measured-benefit cells. Difference columns are equal-weight means over five checkpoint cells and are rounded to four decimals for display; the sealed manifest retains full precision. Cells are repeated-model observations, and the inferential cluster unit is the camera location.

| Location | $n/\text{checkpoint}$ | A/F/U | Benefit $+/0/-$ | Adapt $-$ freeze | KGA $-$ adapt | KGA $-$ freeze |
|----------|-----------------------|-------|-----------------|------------------|---------------|----------------|
| 0        | 1302                  | 0/5/0 | 0/0/5           | -0.5418          | 0.5418        | 0.0000         |
| 7        | 1283                  | 0/5/0 | 1/0/4           | -0.0293          | 0.0293        | 0.0000         |
| 28       | 911                   | 0/5/0 | 0/0/5           | -0.0562          | 0.0562        | 0.0000         |
| 40       | 800                   | 0/5/0 | 0/0/5           | -0.1232          | 0.1232        | 0.0000         |
| 46       | 4432                  | 0/4/1 | 0/0/5           | -0.0723          | 0.0723        | 0.0000         |
| 78       | 1633                  | 0/5/0 | 0/0/5           | -0.0839          | 0.0839        | 0.0000         |
| 100      | 3441                  | 0/5/0 | 0/0/5           | -0.2035          | 0.2035        | 0.0000         |
| 105      | 1540                  | 0/5/0 | 0/0/5           | -0.2892          | 0.2892        | 0.0000         |
| 130      | 983                   | 0/5/0 | 0/0/5           | -0.2376          | 0.2376        | 0.0000         |

45 rows collapsed to nine cis-test camera locations, while the locked targets came from trans-test locations. Coverage therefore needs an unverified cross-split exchangeability assumption. Target Mahalanobis diagnostics ranged from 7.12 to 485.67 (median 28.73), showing strong extrapolation; this diagnostic never changed an action and is not a formal coverage guarantee. One natural target study also cannot establish transfer of selective routing across natural shifts. A follow-on locked study on So2Sat-LCZ42 retained an outcome-blind boundary but stopped at development feasibility, as described next. A new versioned study still needs diverse development units, useful exposure to both ADAPT and FREEZE, and untouched confirmation environments. Cell-level 16-indicator macro-F1 is archived as descriptive secondary evidence; no post-hoc aggregate or inference claim is made.

Table 30 reports all nine location clusters. Evaluation  $n$  is the number of evaluation images per checkpoint at that location. Each A/F/U triplet and each helpful/zero/harmful triplet therefore has denominator five, one cell for each checkpoint. These labels describe measured cell outcomes. Accuracy differences use one sign throughout: adapted minus frozen for adaptation benefit, and KGA minus the named fixed policy for the two policy columns.

The release ledger contains all 45 cells: 0 ADAPT, 44 FREEZE, and 1 ABSTAIN actions; 1 helpful, 0 zero, and 44 harmful adaptation effects. The release verdict remains SAFE UTILITY ONLY. The location table is descriptive; sign-flip enumeration uses the nine location means under the reference model in Eq. (12), and no post-hoc aggregate is claimed for the secondary macro-F1 outcome.

## H Extended KGA Implementation Contract

### H.1 State Isolation

The frozen and candidate predictors must not share mutable normalization buffers, optimizer state, or parameter storage during preview. The candidate is created from a checkpoint, adapted on the unlabeled window, and evaluated without overwriting the active state. Only an adapt decision may replace the active checkpoint. Freeze and abstain discard the temporary state. A unit test should hash the active parameters before and after each non-adapt decision and require equality.

## H.2 Evidence Determinism

Evidence extraction must define preprocessing, batch order, numerical precision, class ordering, normalization constants, and missing-value behavior. Paired features require frozen and candidate outputs on the same samples in the same order. Source statistics and ATC thresholds are immutable calibration artifacts. Evidence schema versions should be explicit rather than inferred from vector length, because two schemas can share a dimension while changing semantics.

## H.3 Calibration Identity

The calibration identity should hash the source checkpoint, candidate adapter class, adapted parameter names, optimizer and learning rate, update steps, evidence schema, benefit-model hyperparameters, feature scaler, residual construction,  $\alpha$ , and inference-unit definition. The radius belongs to that identity. Reusing it after a candidate or evidence change is not a benign software optimization; it changes the random variable whose residual is being covered.

## H.4 Runtime Accounting

A deployment report should time capture/preprocessing, frozen inference, candidate adaptation, candidate inference, evidence extraction, benefit prediction, interval decision, copy creation, and rollback separately. Mean and tail latency should be reported at the decision-window level. Peak memory should include both model states and optimizer buffers. The compact manuscript leaves these cells unfilled because the canonical cross-dataset manifest does not provide one comparable runtime measurement. Omitting an unsupported number is preferable to quoting controller-only latency as end-to-end cost.

# I Evidence Inventory

## I.1 Protocol-Specific Feature and Model Schemas

The controlled-grid implementation records 11 label-free coordinates with effective rank 9. Two coordinates are deterministic differences: `entropy_drop=pre_entropy-post_entropy` and `pbal_drop=pre_pbal-post_pbal`. The panel therefore contains 11 stored values, not 11 independent information sources. Together they describe confidence, class balance, change from the frozen model, and update size.

Natural-shift artifacts do not all use the same schema. iWildCam, RxRx1, and most archived tracks use the 11-feature panel. Camelyon17 uses 17 features by adding disagreement, entropy gap, free-energy shift, batch-statistic divergence, a source-fixed ATC estimate, and confidence drop. Office-Home uses 18 features and also records source-prior KL and mean candidate disagreement. A model and its radius are valid only for the schema on which they were calibrated. Missing, reordered, or invalid features force non-commitment.

The controlled and archived GBRT protocols use scikit-learn’s `GradientBoostingRegressor` [33] with squared-error loss, 250 trees, depth 2, learning rate 0.05, subsample 0.8, and random seed 0. The fit is specific to each dataset and adapter. CCT-20 instead uses its sealed ridge model, as described in Section 7.2. Thus “adapter agnostic” describes the interface: any adapter can propose a candidate that is evaluated and either committed or rolled back. It does not mean that a benefit model fitted for Tent transfers validly to EATA, SAR, or another adapter.

The base panel has 11 coordinates; the canonical Camelyon17 rich-Z artifact has 17 and the Office-Home artifact has 18. Other natural tracks retain the 11-coordinate base panel. These

Table 31: Feature families used by the maintained implementation. All quantities are aggregate statistics of unlabeled inputs, frozen/candidate outputs, or the proposed update.

| Family                  | Representative coordinates                            | Required state                      | Failure handling                               |
|-------------------------|-------------------------------------------------------|-------------------------------------|------------------------------------------------|
| Predictive uncertainty  | mean entropy, confidence, quantiles                   | frozen/candidate soft-max           | reject NaN or malformed probabilities          |
| Marginal balance        | normalized class-marginal entropy and change          | predicted-class marginals           | fixed class count and normalization            |
| Model change            | entropy gap, confidence drop, prediction disagreement | paired frozen/candidate outputs     | require aligned sample order                   |
| Collapse signal         | fraction above adapted confidence 0.9                 | candidate probabilities             | threshold fixed before scoring                 |
| Distribution shift      | free-energy shift, batch/running-statistic divergence | frozen logits and source statistics | missing source statistic forces non-commitment |
| Update magnitude        | parameter-update $\ell_2$ norm                        | shadow candidate delta              | configuration hash must match calibration      |
| Source-calibrated proxy | ATC-style threshold estimate                          | source-fixed threshold              | target labels never used                       |

Table 32: Operational failures and conservative responses.

| Failure                       | Detection                     | Response                     |
|-------------------------------|-------------------------------|------------------------------|
| Checkpoint/adaptor mismatch   | stored hash mismatch          | rollback and recalibrate     |
| Evidence outside support      | distance or support audit     | abstain                      |
| Missing/invalid feature       | schema and finite-value check | retain $f_0$ ; abstain       |
| Insufficient calibration size | rank $k > n$                  | infinite radius              |
| Batch too small               | minimum-size contract         | wait or abstain              |
| Residual transfer concern     | group/residual diagnostic     | widen or withhold claim      |
| Continual-state drift         | state and stream monitor      | restore last certified state |

protocol-specific dimensions are part of the fitted estimator contract. They should not be described as universal sufficient statistics: a different shift may hide the benefit direction from all listed features. Feature selection can improve empirical prediction, but cannot defeat the label-kernel construction of Lemma 1 over an unrestricted target fibre.

## J Failure-Safe Deployment Checklist

The audit log should store a timestamp, dataset or stream identifier, checkpoint and adapter hashes, evidence schema version, evidence vector, benefit estimate, radius, interval endpoints, action, rollback outcome, latency components, and configuration hash. Labels, if later revealed, are joined in a separate offline audit table. This separation makes it possible to verify that the live decision was target-label-free.

## K Prospective Natural-Shift Confirmation Protocol

The sealed CCT-20 study supplies a completed natural-shift target evaluation. Its 45 checkpoint-location cells contain 1 helpful and 44 harmful candidate updates, but KGA makes 44 FREEZE, 1 ABSTAIN, and 0 ADAPT decisions. It therefore matches always-freeze and has lower regret than always-adapt, providing evidence of conservative harm avoidance under natural shift. The predeclared stronger routing criterion was not met because the policy never selected ADAPT and did not select the single helpful cell. A second locked attempt on So2Sat-LCZ42 stopped at its development-feasibility gate with no eligible Tent or SAR candidate. No held-out gate-calibration rows or target inputs were accessed. The following protocol therefore describes a new versioned

confirmation, not a continuation or reinterpretation of v1.

All So2Sat v1 gate-fit outcomes are open and cannot be used to tune a fresh confirmatory arm. Its target protocol also defines one action per city, whereas the current runner creates one action per city-checkpoint cell. That mismatch did not affect the development stop, but the target runner must remain disabled until a versioned lock resolves the action unit. The design below is future evidence and is not counted as a completed target experiment in this paper.

### K.1 Lock Before Target Labels

The protocol should name one source checkpoint, one candidate adapter, one evidence schema, one benefit-model class, one calibration rule, one  $\alpha$ , and one target-domain family. Adapter hyperparameters must be selected on development domains and then frozen. The lock should hash the checkpoint, code revision, configuration, calibration records, and planned scoring script. It must also name the test unit and the multiplicity family. A later change creates a new protocol version; it does not overwrite an earlier development-stopped or incomplete arm.

Natural heterogeneity should be built through independently defined environments rather than by pooling whichever benchmark conditions produce opposite signs. Examples include hospitals, camera locations, laboratory batches, acquisition days, physical devices, or visual domains. Environment membership may be observed at deployment, but the measured cell-benefit label remains hidden. The target test environments should be absent from estimator fitting, residual calibration, threshold tuning, and candidate selection.

### K.2 Three-Way Data Separation

Development data fit  $h_\theta$  and select the evidence pipeline. A separate calibration set produces residuals for Eq. (10). The held-out target set is opened exactly once after all software and configuration hashes match the lock. Splitting adjacent frames, patches, or samples from the same session across these stages is leakage even when filenames differ. The split must occur at the environment or session level that defines deployment exchangeability.

The target stream is replayed identically for always-freeze, always-adapt, simple gates, KGA without a radius, and the full certificate. Model randomness should either be frozen for paired replay or repeated through independently trained checkpoints declared in advance. Every policy receives the same ordered windows. Labels are stored behind a separate audit interface and are unavailable to the live decision process.

### K.3 Endpoints and Decision Rule

The primary endpoint is paired regret against each fixed policy. Success requires both baseline minus KGA regret gaps to be positive and both predeclared intervals to exclude zero, with coverage justified for the declared estimand and sampling model. If the intended claim concerns population risk rather than measured cells, the protocol must also justify that transfer. The false-adapt endpoint is  $FA_u$  with the adapt count reported beside it. Secondary endpoints are  $FA_c$ , adapt rate, freeze rate, abstain rate, commitment rate, task accuracy, and latency. Conditional false-adapt is descriptive unless a separate selective-risk result is proved.

The primary certificate rule remains unchanged: adapt if  $\hat{\Delta} - \varepsilon > 0$ , freeze if  $\hat{\Delta} + \varepsilon < 0$ , and abstain otherwise. The experiment must not lower  $\varepsilon$ , change  $\alpha$ , or select a different benefit regressor after observing target labels. Exploratory sensitivity analyses may be reported, but they remain outside the primary bar.

Table 33: Minimum lock for a confirmatory natural-shift KGA experiment.

| Component              | Must be fixed                                                                               | Audit evidence                                        |
|------------------------|---------------------------------------------------------------------------------------------|-------------------------------------------------------|
| Model and candidate    | checkpoint, adapter, optimizer, steps, adapted parameters                                   | hashes and immutable configuration                    |
| Evidence and estimator | schema, normalization, model class, hyperparameters, seed                                   | feature contract and fitted-artifact hash             |
| Calibration            | split unit, residual construction, exact rank, $\alpha$                                     | calibration IDs, residuals, radius, no target overlap |
| Target evaluation      | environments, windowing, order, exclusion rules, transductive/inductive candidate semantics | sealed manifest and one-time scoring log              |
| Primary inference      | two paired regret gaps, cluster unit, interval method                                       | code and random seed fixed before label reveal        |
| Safety                 | FA <sub>u</sub> and adapt exposure                                                          | action log joined to labels only offline              |
| Multiplicity           | all candidate-track arms in family                                                          | time-stamped inventory including failures             |
| Runtime                | end-to-end components and peak memory                                                       | raw timing rows on the target hardware                |

## K.4 Interpreting Outcomes

Four outcomes are scientifically useful. An interval-supported reduction in regret relative to each fixed comparator, with nontrivial ADAPT and FREEZE exposure, supports routing utility under the locked natural-shift class. Matching the better fixed policy with few errors supports conservative commitment but not selective routing. High abstention with low error indicates a conservative certificate and motivates better evidence or more calibration data. A comparator shortfall fails the predeclared routing-utility criterion; excess false adaptation challenges the corresponding safety or calibration criterion. Both outcomes must remain in the result family. None of these outcomes alone establishes the population band: that conclusion still requires the declared class and its structural assumptions.

## L Theorem Dependencies and Delimited Scope

The compact paper uses three related logical statements. First, the disagreement reduction is an identity under the binary setup. Second, the population frontier is a uniform statement over a rich, declared class. Third, the empirical certificate is a marginal implication from interval coverage. The population statements concern  $\Delta$ ; the reported empirical calibration concerns  $\Delta^{\text{cell}}$ . Connecting these targets requires an additional validity argument. Neither the assumptions nor the quantifiers are interchangeable.

**Why the boundary is separate.** For  $|M| < \beta$ , the construction chooses two latent residuals that yield opposite nonzero benefits. At  $|M| = \beta > 0$ , one admissible residual can make the benefit zero while another retains the sign of  $M$ ; opposite nonzero signs need not exist. Both cases block a uniform strict commitment, but for different reasons. At  $M = \beta = 0$ , every law in the zero-residual class has zero benefit under the model. Abstention remains the maximal action only under the paper’s strict-direction convention, not because either task-risk choice is worse.

**Why a calibration residual need not be drift.** Let  $P_T = P_S$ , with equiprobable inputs  $a, b$ , deterministic labels  $Y(a) = 1$ ,  $Y(b) = 0$ , and candidate  $f_a \equiv 1$ . Set  $f_0(a) = 0$ ,  $f_0(b) = 1$ , and  $s \equiv \frac{1}{2}$ . The candidate is correct on half the source inputs, so this constant score is source-calibrated in the ordinary sense. Yet disagreement occurs only at  $a$ , where the candidate is always correct.

Table 34: Dependency map. A check in one row does not validate premises in another row.

| Statement                      | Inputs                                                         | Conclusion                                             | Not concluded                                                     |
|--------------------------------|----------------------------------------------------------------|--------------------------------------------------------|-------------------------------------------------------------------|
| Disagreement reduction         | binary labels, $0/1$ loss, $\mu_T(D) > 0$                      | $\Delta = 2\mu_T(D)(M + \gamma)$                       | calibration transfer or observability of $\gamma$                 |
| Interior impossibility         | label-free fibre, $\beta > 0$ , $ M  < \beta$ , class richness | evidence-identical laws with opposite nonzero benefits | every empirical abstention is structural                          |
| Closed-band result             | strict directional semantics and boundary law                  | no uniformly sound strict action on $ M  \leq \beta$   | adapt/freeze have different task risk at $\Delta = 0$             |
| Frontier sufficiency           | declared $ \gamma  \leq \beta$ and fixed $M$                   | direction is fixed when $ M  > \beta$                  | label-free estimation of $\beta$                                  |
| Coverage-to-action proposition | marginal interval coverage for a declared target $B$           | marginal false-adapt/freeze control for $B$            | conditional error control or risk alignment                       |
| KGA benchmark                  | fitted $h_\theta$ , declared residual rule, frozen protocol    | measured regret and actions on saved units             | population-benefit coverage or universal transfer to a new domain |
| Lean build                     | encoded definitions and assumptions                            | kernel-checked formal statements                       | correctness of unencoded data and deployment assumptions          |

Thus  $M = 0$ ,  $\gamma = \frac{1}{2}$ , and  $\Delta = \frac{1}{2}$ , although the source and target laws are identical. Conditioning on disagreement, not a distribution change, explains the nonzero residual. The frontier requires a bound on this residual regardless of its cause.

**Why coverage is the premise.** The certificate proof is elementary once Eq. (9) holds: on the coverage event, an interval wholly above zero cannot contain a nonpositive  $\Delta^{\text{cell}}$ . The scientific burden is therefore to justify coverage for the declared outcome and calibration/deployment pair. Joint exchangeability of the residual scores is one route; known shift correction or another valid construction may be another. Risk alignment can explain why a useful predictor transfers, but it is not an additional logical premise after coverage is given.

**Why cell coverage is not population coverage.** Consider the fixed binary predictors  $f_0 \equiv 0$  and  $f_a \equiv 1$ . Let the target input be  $a$  with probability 0.45 and  $b$  with probability 0.55, with deterministic labels  $Y(a) = 1$  and  $Y(b) = 0$ . Population benefit is  $\Delta = -0.10$ . In a one-item evaluation cell, a fixed predictor that maps  $a$  to  $+1$  and  $b$  to  $-1$  predicts  $\Delta^{\text{cell}}$  perfectly without reading the new label. With radius zero, cell coverage is perfect, yet the gate selects ADAPT with probability 0.45 despite negative benefit under the fixed population law. This constructed example does not contradict Proposition 2; it shows why its covered target cannot be changed.

**Why exchangeability alone does not validate sign flips.** As a separate null example, let all nine cluster gaps equal the same fair random sign  $S$ . This vector is exchangeable, each gap is symmetric about zero, and a global sign change preserves its law. Nevertheless, when  $S = +1$ , independent-coordinate sign-flip enumeration for the mean gives a reference value of  $1/512$ . That occurs with probability  $1/2$ , not  $1/512$ . The vector fails Eq. (12). This example explains the required assumption; it is not a model fitted to the CCT-20 observations.

**Why a null is not an impossibility result.** A null benchmark row can arise because the true effects are small, the candidate is one-sided, the evidence map is weak, the benefit model underfits, the radius is wide, or calibration does not transfer. The population impossibility requires evidence-identical target laws inside the declared class with incompatible strict directions. The canonical benchmark table does not estimate that full fibre. It therefore reports weak or adverse empirical regimes without calling them proven unknowable worlds.

**What would falsify the practical claim.** The practical claim is conditional and testable. A clean target replay with the locked pipeline can show excessive  $\text{FA}_u$ , poor coverage, or regret above both fixed policies. A checkpoint or adapter change can show that calibration does not transfer. A family-held-out split can show that cell-level cross-fitting was optimistic. Such failures do not refute the population identity; they refute the operational assumptions used to obtain a useful empirical interval.

## M Formalization and Artifact Scope

The Lean project under `docs/research/kbound/formal/` provides a kernel-checked inventory for the theorem names it exports. The compact paper relies on the LaTeX statements above; it does not claim that every sentence has a one-to-one Lean theorem.

In particular, deployment-class richness, risk alignment of a learned evidence map, exchangeability of a benchmark split, and correctness of data preprocessing remain outside the proof assistant unless separately encoded.

The current strict-core audit kernel-checks 150 named declarations: 70 finite, algebraic, and deterministic population-transfer results, and 80 probability, construction, counterexample, and corrected conditional-extension results. The added modules derive one-shot residual coverage from exchangeable measurable scores, including ties; filtered Ville and bounded optional-stopping bounds; general KL/TV testing bounds for finite product experiments; and independent bounded and martingale-difference concentration inequalities. Their assumptions are explicit. In particular, these results do not establish exchangeability or conditional nulls for the benchmark data, nor do they turn batch-level coverage into population-risk coverage.

The measurable target-law construction also goes beyond the original two-atom example. It fixes the input law and off-disagreement label kernel, varies a measurable correctness field on disagreement, and integrates the actual zero-one loss difference. With labels supported on the two predictions on disagreement, feasible margins, and positive disagreement mass, it proves the clipped identified interval and exact strict ADAPT/FREEZE frontiers without assuming `RichAt`. It does not prove richness for an arbitrary restricted deployment subclass.

The finite-sample population-transfer procedure is also separated into a deterministic machine-checked layer. `KBound/Population.lean` defines the compound radius  $\varepsilon + b$ , its interval, and the strict action rule, and proves that a committed action has the correct benefit sign whenever the interval-containment premise holds. The outcome-blind replay script `scripts/run_population_panel.py` computes the same interval from fitted predictions, cell radii, and fresh-sample sizes while refusing target or outcome fields. This closes the procedure-to-action interface, not the statistical premises: the benchmark tables still lack a sealed per-cell sampling and exchangeability record, so they are not relabeled as population-certified results.

The separate synthetic episode runner `scripts/run_population_locked_panel.py` checks disjoint episode identities across fitting, calibration, and scoring. Its evidence writer reproduces the supplied decisions from the stated inputs and risk budgets before saving them, refuses an occupied output directory, and reruns the scored-outcome perturbation check with those same budgets. The perturbation stays within the valid benefit range  $[-1, 1]$ . With one calibration episode at  $\alpha_{\text{cell}} = 0.05$ , the requested rank is unavailable; the implemented interval is unbounded and the procedure abstains. Saved JSON represents these endpoints explicitly as strings `-inf` and `+inf`, rather than nonstandard JSON numbers. These boundary and evidence-integrity checks concern the specified synthetic procedure, not an additional benchmark or evidence of real-population exchangeability.

One historical extension remains outside the verified scope. A measurable label swap preserves

label-free evidence and reverses population benefit, but selecting one representative from each swap orbit need not identify the sign: distinct orbits can share the same evidence law and have opposite selected signs. Lean verifies this counterexample and the corresponding set-theoretic fibre-consistency criterion, not a general measurable decoder. A corrected conditional H/ratio-rate contract also proves that a source miss budget at most  $\alpha$  transfers to a target miss budget at most  $H\alpha$ , assuming  $H \geq 0$  and the supplied domination premise  $\text{targetMiss} \leq H \text{sourceMiss}$ . It does not derive that premise, estimate  $H$ , or establish benchmark exchangeability. The unrestricted historical claims are excluded from this compact submission. The **-full-foundations** gate therefore remains failing for the sixth layer; the scoped kernel audit checks the registered results and rejects nonstandard proof axioms.

The formal build and empirical build answer different questions. Lean checks that a proposition follows from formal assumptions. The reconciliation scripts check that displayed aggregates match saved records under implemented definitions. Neither checks the truth of an unencoded deployment assumption, and neither substitutes for independent experimental replication.

The authoritative compact artifacts are layered: (1) this source and its compiled PDF; (2) shared theorem inputs for exact proofs; (3) canonical JSON and generated LaTeX for numbers; (4) the source manifest for record-level provenance; and (5) package tests and Lean logs for implementation and formal checks.

Historical drafts, dashboards, and exploratory search folders can remain useful development records, but they are not evidence sources for this compact submission unless referenced by the canonical manifest.

## N Reproducibility and Release Procedures

### N.1 Canonical Numerical Source of Truth

The CIFAR primary rows and the supplementary reconciled point estimates have one numerical source: `experiments/kbound/results/reconciled_panels_v1/canonical_panel_results.json`. Its compact source manifest has SHA-256 `3ac2c7d14d44803e4ff8a19258d036dc86327adb02fdb4cd34536e5eeee9cb24`. Named auxiliary artifacts support only their stated roles. Current-policy family sensitivity is retrospective: Holm adjustment over six prospectively named contrasts gives reference values of 0.09375 for both Tent comparisons, but the corruption-family unit and test were not preregistered. The earlier Tent cluster artifact and archived POEM/AETTA comparison remain historical. The So2Sat-LCZ42 selection artifact is a separate locked development-gate record with no promoted target row. None of these overwrites the canonical point-estimate panel. The generated result manifest indexes every promoted claim and its scope.

**Source Traceability and Execution Mapping.** To resolve provenance and pipeline differences across benchmark tables: (1) **Table 7 (Primary CIFAR-10-C)**: ResNet-18, Tent/EATA/SAR. The canonical manifest binds compact per-condition records to their archived source files. The current generator, `scripts/reconcile_result_panels.py`, refits from saved  $Z, B$  under Protocol B and ignores historical prediction/radius/action fields. Per seed: 432 records, five score folds of 86 or 87, 104 residual-calibration records and 242 or 241 estimator-fit records. Calibration rank is  $\lceil (104 + 1)(1 - 0.10) \rceil = 95$ . Exact literal candidate/seed/condition identifiers determine disjoint folds; each pooled row has weight  $1/2160$ . The earlier Protocol-A description with pool 431 and rank 389 describes a historical procedure, not this current table. A pinned-runtime replay reproduced all Tent predictions, radii and actions for five seeds; changing one scored outcome per seed left that cell’s decision inputs and action unchanged. This is saved-feature cell-level verification, not upstream or environment-held-out independence. (2) **Table 9 (Current Synchronized Baselines)**: Tent, six internal decision rules on the same canonical paired scores and current predictions

as item (1), under Protocol B. `docs/research/kbound/scripts/replay_cifar_current_arithmetic.py` emits JSON, per-cell decisions and LaTeX into `paper/generated/current_cifar_baselines_20260912/`; it authenticates inputs against an explicitly hash-bound audit and does not refit or tune. Its generated policy arrays and metrics were independently recomputed. The separate `PINNED_CANONICAL_REPLAY_VERIFICATION.json` records exact five-seed refitting and five outcome interventions in Python 3.12.13, NumPy 2.4.4 and scikit-learn 1.8.0. Both fixed-policy regrets match item (1) exactly. The historical synchronized/P4 tables are excluded from current results: their cited aggressive archive has 270 cells per seed, and all inspected 432-cell synchronized paired arrays equal the primary arrays. Source hashes, comparisons and preserved historical text are under `experiments/kbound/results/cifar_paired_array_audit_20260912.Q4GKLg/`. This authenticates saved-record calculations, not image-level sample identity, independent source training or native baseline execution. (3) **Table 19 (Matched SAR Controls)**: Historical ResNet-50, 5 SAR arms (Table 18). The cited 60-cell raw records and reconciliation manifest are not resident in the current checkout, so this table is not release-promoted. Historical configuration: per-cell adaptation with  $b \in \{16, 32, 64\}$  on 64-image windows, evaluated under Protocol A (LOO residual calibration,  $n = 60$ , pool  $n_{\text{cal}} = 59$ , rank  $k = 54$  at  $\alpha = 0.10$ , equal weights  $1/60$ ). The current resident 27-cell reference-SAR audit is reported separately in Table 20. (4) **Table 21 (ImageNet-C Baselines)**: ResNet-50, Tent/EATA/SAR (8 decision rules). *Records*: `decision_baselines_sarfix/decision_baselines.json` and `checkpoint.json`. 36 noise conditions  $\times$  3 candidates ( $N = 108$ ). Evaluated under Protocol A (leave-one-condition-out GBR on 11-dim features,  $\alpha = 0.10$ , equal weights). (5) **Table 5 / 14 and Table 20 (Outcome-Disjoint Studies)**: Evaluated under Protocol B (three-way cell-outcome-disjoint cross-fitting, `kga.crossfit.controlled_grid_crossfit`), where each complement is partitioned into strictly disjoint estimator-fit and residual-calibration subsets, ensuring scored cell outcomes enter neither predictor fitting nor conformal calibration pools.

## N.2 Reproduction Pipeline

The macOS release profile pins Python 3.12.13 and package hashes in `requirements-release-macos-arm64.lock.txt`, with installed-content and publication-tool identity checks. The Linux API image has a separate Python 3.11 lock; an API or paper-only environment does not satisfy the repository-wide research dependencies, including PyTorch [32]. After providing the authenticated evidence and required toolchains, the intended full release command from the repository root is:

```
KBOUND_PYTHON=.venv/bin/python bash docs/research/kbound/runbooks/release_candidate.sh
all
```

It runs the read-only preflight, validates and regenerates result authorities, executes software and formal checks, builds the two PDFs and compact DOCX, renders every PDF page, and emits release checksums. It consumes frozen model outputs and never launches training. The complete integrated release gate has not yet passed; successful manuscript builds and focused software tests must not be interpreted as that result.

For presentation-only rebuilding from already validated authorities, the command is:

```
BUILD_LONG_TMLR=1 BUILD_DOCX=1 bash docs/research/kbound/scripts/build_pdfs.sh
```

This shorter command validates the frozen authorities, regenerates presentation tables and figures, and builds the maintained PDFs and DOCX. Before generation, the current CIFAR arithmetic checker authenticates the seven saved inputs and regenerates all three table artifacts in a fresh temporary directory; a mismatch stops the build. It does not replace the full release gate or create new empirical evidence.

## N.3 Calibration Identity and Implementation Contract

The maintained API binds a certificate to the checkpoint, adapter configuration, evidence schema, benefit estimator, residual construction, safety level, and scoring unit. The deployment path must

fail closed when any identity differs. `KGA.certify_evidence` requires a frozen estimator, the exact schema, a protocol hash, and a separate residual set. The lower-level `certify` function is an audit interface; it does not construct evidence or validate transfer by itself.

## N.4 Formalization and Lean Scope

Lean artifacts and validators check the encoded theorem cores and selected finite constructions. Formalization of learning-theory results has precedent, including a Rademacher-complexity generalization bound [42]; we cite its pinned 2025 version as context, not as a proof of K-Bound’s assumptions. These checks do not certify dataset correctness, preprocessing, empirical transfer, exchangeability, or every deployment assumption stated in prose. Machine checking strengthens the statements it encodes; it does not validate the benchmark pipeline outside those statements.

## N.5 Release Checklist and Audit Logs

The Phase-1 provenance audit matches all 106 canonical source/compact hash pairs, all 10 recoverable configuration hashes, and all 6 archived checkpoint hashes, alongside a 21-file current-code snapshot. Historical runs lacking full serialized identity remain explicitly partial rather than being backfilled. The release gate verifies reconciliation tests, claim-ledger and result-manifest validation, manuscript checks, LaTeX builds, checksum verification, and visual PDF inspection. Build logs, the canonical manifest, the provenance seal, and final PDFs are archived together.

## O Controlled Population-to-KGA Bridge

The population frontier and empirical KGA have related three-way action spaces but different inputs. The population rule acts on the identified set induced by  $(M, \beta)$ , where  $\Delta = M - \gamma$  and  $\beta$  is externally declared. KGA acts on an empirical interval built from  $(\hat{\Delta}, \varepsilon)$ . The seven constructed cells below make the distinction concrete. They are an algebraic interface diagnostic, not evaluation evidence and not a method for estimating  $\beta$ .

Table 35: Population side of the seven-example bridge. The strict population rule abstains on the closed band  $|M| \leq \beta$ .

| Cell      | $M$   | $\gamma$ | $\Delta$ | $\beta$ | Population action |
|-----------|-------|----------|----------|---------|-------------------|
| bridge-00 | -0.30 | 0.04     | -0.26    | 0.10    | FREEZE            |
| bridge-01 | -0.10 | 0.02     | -0.08    | 0.10    | ABSTAIN           |
| bridge-02 | -0.05 | -0.02    | -0.07    | 0.10    | ABSTAIN           |
| bridge-03 | 0.00  | 0.00     | 0.00     | 0.10    | ABSTAIN           |
| bridge-04 | 0.05  | 0.02     | 0.07     | 0.10    | ABSTAIN           |
| bridge-05 | 0.10  | -0.02    | 0.08     | 0.10    | ABSTAIN           |
| bridge-06 | 0.30  | -0.04    | 0.26     | 0.10    | ADAPT             |

Table 36: Empirical side of the same bridge. Five actions agree with the population rule and two differ because  $\varepsilon$  is not a substitute for  $\beta$ .

| Cell      | $\hat{\Delta}$ | $\varepsilon$ | Empirical action | Same action? |
|-----------|----------------|---------------|------------------|--------------|
| bridge-00 | -0.24          | 0.05          | FREEZE           | yes          |
| bridge-01 | -0.07          | 0.08          | ABSTAIN          | yes          |
| bridge-02 | -0.09          | 0.04          | FREEZE           | no           |
| bridge-03 | 0.01           | 0.04          | ABSTAIN          | yes          |
| bridge-04 | 0.09           | 0.04          | ADAPT            | no           |
| bridge-05 | 0.07           | 0.08          | ABSTAIN          | yes          |
| bridge-06 | 0.24           | 0.05          | ADAPT            | yes          |

The two disagreements are informative rather than erroneous. In bridge-02 and bridge-04, the declared population class still permits both benefit signs, while the constructed empirical interval excludes zero. That comparison shows why population non-identifiability does not imply empirical abstention on every realized cell, and why empirical interval commitment does not remove the population assumption. The authoritative artifact is `experiments/kbound/results/frontier_kga_bridge_v1/bridge_results.json` (claim KB-CLAIM-043).

## P Controlled Two-View MNIST Diagnostic

The D33 experiment studies a deliberately controlled two-view setting with injected corruption. Across 130 conditions, KGA makes 9 ADAPT decisions, 119 FREEZE decisions, and 2 ABSTAIN decisions. Table 37 reports the recorded aggregate outcomes.

The recorded KGA-minus-single-A difference is 0.3231 percentage points, whereas the KGA-minus-always-fuse difference is 27.3554 percentage points. These values describe this finite, injected-corruption diagnostic only. In particular, zero observed false ADAPT decisions among only nine ADAPT decisions does not establish a small conditional error rate, and this experiment does not

Table 37: Controlled two-view MNIST (D33). Accuracy is averaged over the 130 recorded conditions.

| Quantity                    | Recorded value         |
|-----------------------------|------------------------|
| Conditions                  | 130                    |
| ADAPT/ FREEZE/ ABSTAIN      | 9 / 119 / 2            |
| KGA mean accuracy           | 85.6785%               |
| Single-view A mean accuracy | 85.3554%               |
| Always-fuse mean accuracy   | 58.3231%               |
| Oracle mean accuracy        | 85.7308%               |
| Observed false ADAPT        | 0 of 9 ADAPT decisions |

establish performance under a natural distribution shift. The authoritative artifact is `experiments/kbound/results/controlled_multimodal_d33/results.json` (claim KB-CLAIM-027).