

When Is Label-Free Adaptation Knowable? Helpful, Harmful, and Unknowable Regimes Under Distribution Shift

Anonymous Author(s)

Affiliation withheld for double-blind review

Abstract—Test-time adaptation (TTA) can improve a model on unlabeled target data, but the same update that helps under one distribution shift can silently harm under another—and with no target labels, the system cannot tell which case it faces. The decision that matters therefore comes *before* adaptation: not how to adapt, but whether to commit the update at all.

We cast this as a label-free *adapt/freeze/abstain* decision problem. The adaptation benefit is the drop in target risk from replacing the frozen baseline f_0 with the adapted candidate f_a ($\Delta = R_T(f_0) - R_T(f_a)$), and the system must recover its sign from unlabeled evidence Z alone. That deciding adaptation label-free requires an assumption on the shift is already known: identifying OOD accuracy is as hard as identifying the optimal predictor [13], and the optimal abstain-versus-predict rule under a shift budget has a known decision-theoretic form [14]. We *specialize and sharpen* this to the benefit sign. Our central result is an *exact* benefit-sign frontier: over a declared target class, sign Δ is identifiable from label-free evidence *if and only if* an observable disagreement-region margin exceeds the admitted calibration-drift budget ($|M| > \beta$)—a clean, computable geometric criterion. Below this frontier two target worlds can induce identical evidence yet demand opposite decisions, so no label-free rule is correct in both and abstention is information-theoretically necessary. Standard label-free accuracy estimators (ATC, DoC, GDE, COT, AETTA, agreement-on-the-line) are the zero-budget face of this frontier, sound precisely when calibration does not drift.

On the identifiable side of the frontier, we give an adapt/freeze/abstain certificate that controls the unconditional false-adapt and false-freeze probabilities under stated risk-alignment, coverage, and calibration assumptions. We instantiate it as Knowability-Guided Adaptation (KGA), a mechanism-agnostic wrapper around Tent, EATA, and SAR.

Headline empirical result: on five held-out natural shifts (Office-Home, iWildCam, Camelyon17, RxRx1, PACS), KGA is *uniformly no-harm* under a valid out-of-fold radius—it beats the worse fixed policy and ties the better at $FA_u=0$. This is the deployment claim the theory targets when real shifts are one-sided or near-oracle: conditional insurance against detectable harm, not a universal accuracy gain. In *identifiable* mixed harmful+helpful regimes (pre-registered CIFAR-10-C stress grid, ImageNet-C SAR), KGA further beats both trivial policies and the no-harm SOTA baselines POEM and AETTA on regret-to-oracle (Holm-corrected CIs; Table IV), confirming the certificate exactly where




Fig. 1. **K-Bound previews the decision.** From an unlabeled batch under shift, the certificate estimates the adaptation benefit $\hat{\Delta}$ with a finite-sample radius ϵ and returns ADAPT (knowably helpful), FREEZE (knowably harmful), or ABSTAIN (sign unidentifiable), with false-adapt/false-freeze $\leq \alpha$.

the frontier says the benefit sign is knowable. Where evidence is weak, it abstains or freezes conservatively. K-Bound thus reframes TTA as an identifiability-limited deployment decision: adapt when benefit is certifiable, freeze when harm is certifiable, abstain when label-free evidence cannot support a valid commitment.

Index Terms—test-time adaptation, distribution shift, selective prediction, identifiability, label-free risk estimation

I. INTRODUCTION

Should a model adapt to unlabeled target data, or would adapting make it worse? Test-time adaptation (TTA) updates a model on the test stream—through entropy minimization or stability-aware updates—and in favorable shifts recovers accuracy; in unfavorable shifts the same objective silently degrades it. With no target labels, the system usually cannot tell which case it faces. The decision that matters therefore comes *before* the update: not how to adapt, but whether to commit it at all. Given a frozen baseline f_0 , an adapted candidate f_a , and an unlabeled batch, the system must choose one of three actions—**adapt** when adaptation is knowably helpful, **freeze** when knowably harmful, or **abstain** when label-free evidence cannot identify the sign of the benefit.

This is a fundamental limit, not an engineering gap—and one already charted in the literature that we specialize. That label-free deployment quantities cannot in general decide adaptation without an assumption on the shift is established: identifying OOD accuracy is as hard as identifying

the optimal predictor [13], and the optimal predict-versus-abstain rule under a shift budget has a known form [14]. We do not re-derive this general impossibility; we *sharpen* it for the benefit sign into a computable criterion. Whether adaptation helps or hurts is label-free identifiable *if and only if* an observable margin exceeds the calibration-drift budget on the disagreement region (our spine theorem, $|M| > \beta$). On the unidentifiable side of this frontier, two target worlds carry identical label-free evidence but opposite benefit, so no label-free rule—however the adapter is built—can be correct in both, and abstention is information-theoretically forced. K-Bound supplies the constructive companion: when evidence is risk-aligned and the benefit estimate carries a valid uncertainty radius, a finite-sample certificate adapts on the identifiable-helpful side, freezes on the identifiable-harmful side, and abstains exactly where the frontier makes knowing impossible—holding false-adapt and false-freeze at level α under exchangeable calibration.

We instantiate **one** method, KGA (Algorithm 1): a gradient-boosted benefit estimator $\hat{\Delta}(Z)$ plus a conformal radius ε fit on a labeled calibration split, then adapt/freeze/abstain. The wrapper is mechanism-agnostic (Tent, EATA, SAR), and the headline experiments use no separate multicandidate or alternate certificate. K-Bound does *not* claim universal improvement. When adaptation is already safe a certificate can only tie always-adapt; its value is highest precisely when harmful adaptation is frequent, detectable, and costly. KGA is, in short, a safety layer for TTA—not an accuracy booster—and the experiments are organized to show exactly that.

a) Contributions.:

- 1) **A new problem framing.** We recast test-time adaptation as a label-free *adapt/freeze/abstain* decision—whether to adapt at all, not how to adapt. The underlying impossibility (deciding label-free needs an assumption on the shift) is inherited from prior work [13], [14], which we *specialize* to the benefit sign.
- 2) **One exact result (the spine).** We sharpen that known landscape into the *benefit-sign frontier*: the sign of the adaptation benefit is label-free identifiable *if and only if* an observable margin exceeds the calibration-drift budget on the disagreement region, i.e. $|M| > \beta$ (Theorem 1)—a clean, computable geometric criterion. The “only if” direction is an information-theoretic impossibility for the benefit sign—matched evidence, opposite benefit (Lemma 1); the “if” direction yields a finite-sample certificate with false-adapt and false-freeze each $\leq \alpha$. The same proofs hold for any *split-observable* functional (sign $\Delta = \text{sign}(M + \gamma)$, M observable, $|\gamma| \leq \beta$), with TTA the demonstrated instance.
- 3) **Unification.** ATC, DoC, GDE, COT, AETTA, and agreement-on-the-line are the zero-budget ($\beta=0$) face of the frontier—sound precisely when calibration does not drift.
- 4) **A drop-in method.** We instantiate the certificate

TABLE I

Dataset/regime evidence summary. KGA’s STATUS PER DATASET AND THE STRENGTH OF EVIDENCE BEHIND IT. **Headline** ROWS ARE HELD-OUT NATURAL SHIFTS (UNIFORM NO-HARM). **Confirmatory** ROWS ARE IDENTIFIABLE SYNTHETIC/MIXED REGIMES WHERE THE FRONTIER PREDICTS KNOWABILITY. “CI-ROBUST” = BOTH REGRET-GAP BOOTSTRAP CIs EXCLUDE ZERO; “NO-HARM” = MATCHES THE BETTER FIXED POLICY WITHOUT BEATING BOTH. PER-DATASET NUMBERS ARE IN TABLES II–IX; DIRECT COMMITTAL-GATE BASELINES ON CONTROLLED TRACKS ARE IN APPENDIX M (TABLE XXIX).

Dataset / regime	KGA result	Evidence
<i>Headline: natural / held-out shifts</i>		
Office-Home M v2	no-harm	beats adapt (CI), ties freeze
iWildCam H v2	no-harm	beats adapt (CI), ties freeze
Camelyon17 genuine OOD	ties adapt	no-harm
RxRx1	ties freeze	no-harm
PACS (domain-gen)	no-harm (3/4 domains)	beats adapt, ties freeze
<i>Confirmatory: identifiable synthetic/mixed</i>		
CIFAR-10-C stress	beats both	CI-robust, Holm corrected
ImageNet-C SAR	beats both	mechanism-faithful operating point
<i>Unknowable / diagnostic (not headline)</i>		
ImageNet-R	weak / split-specific	not headline
CIFAR-10.1	transfer failure	negative evidence

as KGA, a mechanism-agnostic wrapper around Tent/EATA/SAR (Algorithm 1).

- 5) **Headline empirical evidence (natural shifts).** On five held-out natural benchmarks (Office-Home, iWildCam, Camelyon17, RxRx1, PACS), KGA is *uniformly no-harm* under a valid out-of-fold radius: it beats the worse fixed policy and ties the better at $FA_u=0$ (Table I). This is the primary deployment claim—conditional safety when adaptation is risky, not universal improvement.
- 6) **Confirmatory identifiable regimes.** On pre-registered mixed harmful+helpful stress grids (CIFAR-10-C, ImageNet-C SAR), KGA beats both trivial policies and **beats POEM and AETTA** on regret-to-oracle (Table IV), where the theory predicts the benefit sign is knowable.

Empirically, KGA follows the regime map the theory draws. **Natural shifts** (Office-Home, iWildCam, Camelyon17, RxRx1, PACS) are uniformly no-harm: KGA matches the better fixed policy and prevents damage from always-adapt where it is harmful (Table I). **Identifiable synthetic/mixed regimes** (CIFAR-10-C stress, ImageNet-C SAR) supply the policy wins where harmful adaptation is frequent and detectable, including the pre-registered head-to-head over POEM/AETTA. **Unknowable regimes** (e.g. ImageNet-R, CIFAR-10.1) correctly force abstention or conservative freeze (Appendix).

Table note. Per-dataset quantitative results (regret, false-

adapt, coverage) and the direct committal-gate baselines appear in the per-dataset tables and Appendix M.

II. RELATED WORK

Across every adjacent line, one distinction recurs—the one the theory formalizes: *improving* an unlabeled objective versus *certifying the sign* of its benefit before acting.

a) *TTA methods.*: Tent [1], SHOT [2], TTT [3], EATA [4], SAR [5], CoTTA [6], and MEMO [7] make adaptation more robust but commit to adapting; none certifies the sign of benefit *before* updating, and none abstains on the adapt/freeze decision itself.

b) *Guarded and monitored TTA.*: POEM [11] and sequential risk monitors [9], [10] act during or after adaptation and control false alarms in one failure direction. K-Bound decides *before* committing, poses a three-way decision with an explicit unknowable region, and controls false-adapt on the benefit sign.

c) *Label-free accuracy and error estimation.*: ATC [13], Agreement-on-the-Line [16], [15], and AETTA [18] estimate target performance without labels. Garg et al. [13] also establish the limit governing this whole line—identifying OOD accuracy is as hard as identifying the optimal predictor—so every such estimate is sound only under an assumption on the shift. We *organize* these estimators rather than compete with them: ATC, DoC, GDE, COT, AETTA, and agreement-on-the-line are all the zero-budget ($\beta=0$) face of our frontier, sound precisely when calibration does not drift. Most directly, [17] applies agreement-on-the-line to decide, at the *population* level, when TTA helps or hurts and to select an adaptation strategy—a plug-in use on that same $\beta=0$ face; it issues no per-deployment certificate, no ABSTAIN action, and no identifiability limit, and its linear trend is known to break on shifts (e.g. Gaussian-noise corruption)—precisely the regime where K-BOUND abstains. Estimating $R_T(f)$ differs from certifying $\text{sign}(R_T(f_0) - R_T(f_a))$ before committing the candidate update, with controlled false-adapt rate. Thus, K-Bound is not an accuracy estimator used as a gate; it changes the target estimand from target accuracy to the sign of the frozen-versus-adapted risk difference, with abstention when that sign is not identifiable.

d) *Impossibility, abstention under shift, and identifiability.*: The general impossibility we specialize is owned by prior work: Garg et al. [13] prove identifying OOD accuracy is as hard as identifying the optimal predictor, and Kalai & Kanade [14] give the matching decision theory—the optimal abstain-versus-predict rule under an observable-quantity-versus-shift-budget threshold, with a minimax certificate. We cite these as the sources and do not re-derive them. Unsupervised risk estimation [21], [22], [25] and selective prediction [37] likewise show performance is identifiable only under structure. K-Bound *specializes* this landscape to the *sign* of adaptation benefit, certifying it before committing with an explicit abstain region and false-adapt control at level α . What none of these lines




Fig. 2. **The benefit-sign frontier.** The sign of the adaptation benefit is identifiable from label-free evidence iff the observable margin $|\hat{M}|$ exceeds the calibration-drift budget β . Outside the band KGA ADAPTS (knowably helpful) or FREEZES (knowably harmful); inside the band—the unknowable region—abstention is information-theoretically forced.

supplies—and what the next section makes exact—is the computable boundary ($|\hat{M}| > \beta$) between the adaptation decisions one can certify label-free and those no rule ever can.

III. PROBLEM SETUP AND THEORY

Source $P_S(X, Y)$ has labels at train/validation; target $P_T(X, Y)$ exposes only $P_T(X)$ at test time. Fix loss ℓ , frozen f_0 , and candidate f_a . Target risk $R_T(f) = \mathbb{E}_{P_T}[\ell(f(X), Y)]$. The *adaptation benefit* is

$$\Delta = R_T(f_0) - R_T(f_a), \quad \Delta > 0 \iff \text{adapting helps.} \quad (1)$$

Observable evidence $Z = \phi(X_{1:m}, f_0, f_a)$ is any statistic computable without target labels (entropy, confidence, drift, disagreement, update norms). The central question is when sign Δ is recoverable from Z alone. The certificate needs two *separate* requirements. (i) **Risk-alignment** is a property of Z and the class: sign Δ is a measurable function of the population law of Z over the declared class—equivalently, no two admissible worlds with opposite benefit sign induce the *same* Z -law. (ii) **Coverage** is a property of the calibrator: the benefit interval $[\hat{\Delta} - \varepsilon, \hat{\Delta} + \varepsilon]$ has marginal coverage $\geq 1 - \alpha$. Risk-alignment makes sign Δ recoverable in principle; coverage makes the finite-sample certificate valid. When Z is *not* risk-aligned, no label-free rule can recover sign Δ (Lemma 1) and the certificate must abstain.

What “label-free” means. No target/test labels at deploy; labeled source and held-out validation may calibrate $\hat{\Delta}$ and ε . Evidence Z uses only unlabeled batches, predictions, entropy, drift, and update statistics. The method is **target-label-free**, not label-free of all supervision: labels may be used only in the source/validation calibration split, never in the deployment batch or held-out test decision.

Definition 1 (Regimes). *Knowably helpful*: evidence certifies $\Delta > 0$ (adapt). *Knowably harmful*: evidence certifies $\Delta < 0$ (freeze). *Unknowable*: the same Z is compatible with both signs (abstain).

Lemma 1 (Non-identifiability forces abstention). *There exist worlds P_T^1, P_T^2 with $\text{Law}(Z | P_T^1) = \text{Law}(Z | P_T^2)$ but*

sign $\Delta_1 = -\text{sign } \Delta_2 \neq 0$. No label-free rule can be correct in both; any valid rule must abstain on matched evidence with probability at least $1 - 2\alpha$ when false-adapt and false-freeze are each controlled at α .

Theorem 1 (K-Bound certificate and benefit-sign frontier). Suppose $\hat{\Delta}(z)$ satisfies $\mathbb{P}[|\hat{\Delta} - \Delta| \leq \varepsilon] \geq 1 - \alpha$. The rule

$$g(Z) = \begin{cases} \text{ADAPT} & \hat{\Delta} - \varepsilon > 0, \\ \text{FREEZE} & \hat{\Delta} + \varepsilon < 0, \\ \text{ABSTAIN} & \text{otherwise,} \end{cases}$$

controls the marginal, unconditional false-adapt and false-freeze probabilities at level α — $\text{FA}_u = \mathbb{P}[\text{adapt}, \Delta \leq 0] \leq \alpha$ under the conformal coverage assumption (§IV; the conditional rate FA_c is not α -bounded and is reported empirically). On the disagreement region $D = \{x : f_0(x) \neq f_a(x)\}$, write $M = \mathbb{E}[s \mid D] - \frac{1}{2}$ for a source-calibrated score s of f_a and $\gamma = \mathbb{E}[\eta_a - s \mid D]$ for the (unobserved) calibration drift, so that $\text{sign } \Delta = \text{sign}(M + \gamma)$. Over a declared target class $\mathcal{C}_\beta = \{P_T : |\gamma| \leq \beta\}$ with a supplied drift budget $\beta \geq 0$, the benefit sign is label-free identifiable if and only if $|M| > \beta$ (exact frontier). Here β is the externally justified worst-case drift the deployment class admits, γ is the unknown population-drift realization it bounds, and the empirical radius ε that KGA uses is an operational uncertainty radius calibrated from held-out residuals—under the stated exchangeability and risk-alignment assumptions it provides finite-sample coverage for the deployment class, but it is not a direct estimate of the worst-case drift β . Intuitively, the decision is safe when the observable margin clears the admitted drift and impossible when drift can dominate it—the obstruction no checkable assumption removes in full generality (Lemma 1; full dichotomy: Appendix A).

Corollary 1 (Unjustified assumptions $\Rightarrow$ abstain). The certificate’s false-adapt guarantee holds only under the stated coverage/exchangeability and risk-alignment assumptions. When those assumptions cannot be justified for the deployment batch (e.g. the evidence is not risk-aligned, or the calibration sample is not exchangeable with deployment), the guarantee does not apply, and the conservative operational action is to abstain or freeze rather than commit to adaptation: committing adapt without a valid positive certified lower bound is not protected by the target false-adapt level α . (At deployment Δ is unobservable, so this is a statement about assumptions, not about observing $|\hat{\Delta} - \Delta|$.)

Full proofs, rate theorems, one-bit dichotomy, and validators are in Appendix A. One previously open piece is now settled: the weakest label-free class on which one declared bit certifies the benefit sign is characterized *unconditionally* as an explicit finite family of dominance polytopes (Appendix P, Theorem 8), with General Position recovered as the collapsing face. The certificate is, moreover, *minimax order-optimal* on the identifiable side—its sample complexity matches a two-point information-theoretic lower bound up to a universal constant—and its false-adapt

guarantee extends to anytime-valid streaming and to family-wise selection over a candidate panel (extended version). Two pieces of the knowability program remain open: tight (constant-sharp) evidence-channel rates, and a computable general-capacity characterization—though the latter is now partly delimited, since the binary observable margin is provably insufficient once $K \geq 3$, so regularity conditions on the contrast are necessary.

IV. METHOD: KGA

KGA turns the frontier certificate into a deployable rule. It is the only headline method—a decision wrapper around a fixed adapter, not a new adaptation objective. We fit the benefit estimator and its radius *out of fold*, so the data used to fit $\hat{\Delta}$ and the data used to calibrate its residual are disjoint for every point. On the per-cell stress grids (CIFAR-10-C and the synthetic grids): for each cell i a gradient-boosted $\hat{\Delta}_{-i}$ is trained on the other $N-1$ cells and evaluated on cell i , and ε is the $(1-\alpha)$ empirical quantile of the leave-one-out residuals $\{|\hat{\Delta}_{-i}(Z_i) - \Delta_i|\}_{i=1}^N$ —leave-one-out (jackknife) cross-conformal, so no cell’s residual uses an estimator that saw it. On the natural-shift protocols (Camelyon17, iWildCam, Office-Home) the split is across domains: $\hat{\Delta}$ and ε are fit once on a dev split and the held-out test domain is scored a single time—ordinary split-conformal. At deploy we extract Z from an unlabeled batch and return adapt, freeze, or abstain (Algorithm 1). Labels enter only through the calibration split; the deployment batch and held-out test scorer never use target labels for the adapter, ε , evidence, or rule selection. The coverage guarantee is split-dependent. The natural-shift protocols use a clean dev/test split: *split-conformal* coverage $\mathbb{P}[|\hat{\Delta} - \Delta| \leq \varepsilon] \geq 1 - \alpha$, exact in finite samples under exchangeability of the calibration and test cells. The per-cell grids set ε to the $(1 - \alpha)$ empirical quantile of the leave-one-out residuals, so the *calibration-set* coverage is $\geq 1 - \alpha$ by construction (realized 0.898 at nominal 0.90, $\alpha=0.10$)—the standard jackknife radius, which is distribution-free only asymptotically under estimator stability; its formal finite-sample counterpart is the jackknife+ of Barber et al. (2021) at level $1 - 2\alpha$, which we note but do not implement. The decision rule’s false-adapt guarantee is stated under the coverage hypothesis (§IV), which the empirical coverage meets.

On the good event $|\hat{\Delta} - \Delta| \leq \varepsilon$, committing adapt only when $L > 0$ implies $\Delta > 0$ (Theorem 1). Assumptions: exchangeable calibration pairs, bounded benefits, and risk-aligned Z ; no target labels at deploy. Overhead: one forward pass of f_0 and f_a plus lightweight $\hat{\Delta}$ evaluation.

a) *The threshold is derived, not tuned.*: KGA’s commit boundary is not a swept hyperparameter, but the population frontier and the empirical radius must be kept distinct. The *population* frontier of Theorem 1—over the declared class $\mathcal{C}_\beta$, $\text{sign } \Delta$ is identifiable iff the disagreement margin $|M|$ exceeds the supplied budget β —fixes the *required* budget at population level; the drift realization $\gamma \leq \beta$ is unknown and not observed at deploy. KGA does not

Algorithm 1 KGA: Knowability-Guided Adaptation

Require: Frozen f_0 ; adapter $\mathcal{A}$; batch $X_{1:m}$; $\hat{\Delta}$, ε , α
Ensure: Decision $g \in \{\text{ADAPT}, \text{FREEZE}, \text{ABSTAIN}\}$

```

1:  $f_a \leftarrow \mathcal{A}(f_0, X_{1:m})$ 
2:  $Z \leftarrow \phi(X_{1:m}, f_0, f_a)$ 
3:  $\hat{\Delta} \leftarrow \hat{\Delta}(Z)$ ;  $L \leftarrow \hat{\Delta} - \varepsilon$ ;  $U \leftarrow \hat{\Delta} + \varepsilon$ 
4: if  $L > 0$  then
5:    $g \leftarrow \text{ADAPT}$ 
6: else if  $U < 0$  then
7:    $g \leftarrow \text{FREEZE}$ 
8: else
9:    $g \leftarrow \text{ABSTAIN}$ 
10: end if
11: return  $g$ 

```

observe γ ; it commits under a *conservative empirical radius* ε , the $(1-\alpha)$ residual quantile on the calibration split, fixed there and *never* selected on target-test data. The two are related but not identical: ε is a finite-sample, assumption-dependent surrogate that controls false-adapt under the stated exchangeability and risk-alignment assumptions—not the frontier itself. Confidence-thresholding (ATC) and agreement-trend selection (AGL) fit a scalar against a labeled or held-out signal and predict *accuracy*; KGA certifies the benefit *sign* with a boundary *forced* by the identifiability budget, and—unlike any thresholded estimator—ABSTAINS on the region where Lemma 1 *proves* no label-free rule can decide. The impossibility construction, not a tuned cutoff, is what separates KGA from “estimate accuracy, then threshold.”

V. EXPERIMENTS

The experiments test one claim: KGA is a safety layer, beating both fixed policies only where harmful adaptation is frequent and detectable, and doing no harm everywhere else. We evaluate it on two controlled corruption/stress tracks (CIFAR-10-C stress, ImageNet-C SAR) and three held-out natural-shift replications (Protocols G, H v2, and M v2; Table I). Calibration and held-out scoring follow Algorithm 2 (Appendix A). Metrics: mean accuracy; regret-to-oracle against the per-condition oracle $g^* = \mathbf{1}[\Delta > 0]$; *adapt rate* $\mathbb{P}[\text{adapt}]$ and *coverage* $\mathbb{P}[\text{adapt or freeze}]$ (one minus the abstain rate); and two false-adapt measures, reported separately—the *unconditional* false-adapt probability $\text{FA}_u = \mathbb{P}[\text{adapt and } \Delta \leq 0]$ and the *conditional* false-adapt rate $\text{FA}_c = \mathbb{P}[\Delta \leq 0 \mid \text{adapt}]$. Theorem 1 bounds the *unconditional* probability, $\text{FA}_u \leq \alpha$, under its coverage and risk-alignment assumptions; the per-dataset tables additionally list the empirical FA_c among committed adapt decisions. The two false-adapt measures coincide whenever no harmful adapt is committed ($\text{FA}_u = \text{FA}_c = 0$ on every headline dataset, since KGA commits none) and diverge only for rules that commit some; the full four-metric breakdown— FA_u , FA_c , adapt-rate, and coverage—where the leaky gates separate from the certificate is therefore

candidate	policy	mean acc	regret $\downarrow$	FA_c	coverage	beats both?
Tent	always-adapt	0.786	0.0086	—	1.00	
	always-freeze	0.671	0.1232	—	0.00	
	K-Bound	0.793	0.0016	0.00	0.67	yes
EATA	always-adapt	0.798	0.0037	—	1.00	
	always-freeze	0.671	0.1311	—	0.00	
	K-Bound	0.801	0.0015	0.00	0.63	yes
SAR	always-adapt	0.806	0.0018	—	1.00	
	always-freeze	0.671	0.1372	—	0.00	
	K-Bound	0.806	0.0018	0.00	—	ties adapt

TABLE II

CIFAR-10-C STRESS GRID (432 CONDITIONS/METHOD, FIVE SEEDS, $\alpha=0.1$, PROTOCOL A).

reported side by side in the decision-baseline comparison (Table III), so the reader sees both what the theorem bounds and what happens operationally among adapted decisions.

A. CIFAR-10-C stress grid (5 seeds)

The per-corruption suite is helpful-dominated; the *stress* grid crosses severity, batch size, composition, and update aggressiveness (432 conditions/method, ResNet-18, pre-registered Protocol A, seeds 0–4). Harmful base rates are 34% (Tent), 27% (EATA), 16% (SAR). KGA **beats both** trivial policies for Tent and EATA at 0% false-adapt; it ties collapse-resistant SAR (Table II). Among harmful cells KGA adapts 0 and freezes them; among helpful cells it adapts almost all; on marginal cells it abstains. The regret-gap 95% design-based bootstrap CIs (described below) exclude zero for both candidates—Tent (vs always-adapt +0.019 [0.011, 0.026], vs always-freeze +0.080 [0.052, 0.107]) and EATA (vs always-adapt +0.015 [0.009, 0.021], vs always-freeze +0.075 [0.049, 0.100])—with Holm correction over the comparison family and false-adapt 0/2160. (These are *design-based* intervals: they resample the pre-registered mixing-ratio distribution of Protocol A—the 11-point Pareto curve over harmful fraction, $N_{\text{boot}}=5000$ —so they quantify uncertainty over the deployment mixture. A conventional paired *per-condition* bootstrap that resamples the 432 conditions at their realized composition (`scripts/percondition_bootstrap.py`) returns the same verdict: the regret gap to both fixed policies excludes zero for Tent and EATA (beats-both), while helpful-dominated SAR does not improve on always-adapt and beats always-freeze—all at $\text{FA}_u=0$.) The mixing-ratio Pareto (Fig. 3) shows KGA on the regret-vs-harmful-fraction front, strictly beating both once the harmful fraction exceeds $p^* \approx 0.1$ (Tent/SAR) or 0 (EATA).

B. Decision-baseline comparison: a finite-sample false-adapt guarantee

KGA’s claim is not merely “better than always-adapt/always-freeze”—it is better than *simple label-free decision rules* under the same locked protocol. On the CIFAR-10-C stress grid (Tent candidate, identical cells), we score six rules from the same evidence (Table III): a




Fig. 3. CIFAR-10-C stress grid: regret-to-oracle vs harmful fraction p per candidate. KGA ties always-adapt at $p=0$ and beats both for $p \gtrsim p^*$.




Fig. 4. **The safety knob.** Sweeping the level α trades decision coverage against the false-adapt rate: the certificate holds false-adapt at or below target while coverage grows with α , so a deployer chooses the operating point rather than accepting a fixed one.

confidence gate (adapt iff adaptation raised mean confidence), an entropy gate (adapt iff it lowered entropy—Tent’s own objective), a drift/KL threshold, an ATC-style accuracy-prediction gate, KGA *without* its conformal radius (adapt iff $\hat{\Delta} > 0$), and the full certificate. Only the certificate provides a finite-sample false-adapt *guarantee*, and it attains zero observed FA_u on this grid while staying near-oracle on regret (the radius-free variant keeps $\text{FA}_u=0.049 < \alpha$ empirically, but with no such guarantee). The confidence/entropy gates false-adapt on $\sim 74\%$ of harmful cells; the drift threshold, tuned to the budget, degenerates to always-freeze. The KGA-without-radius row isolates the radius: removing it lowers regret slightly ($0.0017 \rightarrow 0.0004$) but raises FA_u from 0 to 0.049 (to 0.141 on harmful cells)—the radius is what turns a good benefit *estimate* into a false-adapt *guarantee*. Reproduced by `scripts/gate_baseline_comparison.py` on the locked per-cell dump.

TABLE III

Decision-baseline comparison, CIFAR-10-C stress grid (TENT CANDIDATE, $n=432$ CELLS, 149 HARMFUL, $\alpha=0.1$). LOWER REGRET AND LOWER FA_u ARE BETTER. ONLY THE CONFORMAL CERTIFICATE KEEPS $\text{FA}_u=0$ ON THE FULL GRID AND ON HARMFUL CELLS WHILE STAYING NEAR-ORACLE.

Decision rule	regret $\downarrow$	FA_u	FA_c	adapt	cov.	$\text{FA}_u(\text{harm})$
confidence gate	0.0084	0.257	0.301	0.85	1.00	0.745
entropy gate	0.0086	0.255	0.304	0.84	1.00	0.738
drift/KL gate	0.1232	0.000	0.000	0.00	1.00	0.000 [†]
ATC-style gate	0.0045	0.116	0.172	0.67	1.00	0.336
KGA (no radius)	0.0004	0.049	0.071	0.68	1.00	0.141
KGA (certificate)	0.0017	0.000	0.000	0.51	0.68	0.000

[†]Degenerate: the budget-tuned drift gate never adapts (adapt-rate 0), so its “0 false-adapt” is just always-freeze (regret 0.123).

TABLE IV

Pre-registered mixed head-to-head (CIFAR-10-C STRESS, TENT, 432 COND./SEED, 5 SEEDS). REGRET-TO-ORACLE $\downarrow$; FA_u = UNCONDITIONAL FALSE-ADAPT. KGA VS. X : $\text{MEAN}(\text{regret}_{\text{KGA}} - \text{regret}_X)$ WITH 95% PAIRED-BOOTSTRAP CI; WIN IFF BOTH CIs EXCLUDE 0 AND SURVIVE HOLM.

policy	mean regret $\downarrow$	FA_u	decisive rate	KGA gap	Holm beats?
always-adapt	0.0079	0.329	1.00	-0.0064	yes
always-freeze	0.1241	0.000	1.00	-0.1225	yes
POEM	0.0088	0.245	1.00	-0.0072 [-0.0088, -0.0057]	yes
AETTA	0.0073	0.240	1.00	-0.0058 [-0.0071, -0.0044]	yes
KGA	0.0016	0.000	0.68	—	WIN
oracle	0.0000	0.000	1.00	—	ceiling

C. Mixed head-to-head vs. POEM and AETTA (pre-registered)

The stress-grid wins above beat the *trivial* policies always-adapt and always-freeze. A harder question is whether KGA dominates the *no-harm SOTA*—POEM [11] (betting martingale entropy protection) and AETTA [18] (label-free accuracy gate)—on a *mixed* harmful+helpful stream. We pre-registered the benchmark, metrics, and WIN/TIE/LOSE criterion *before* computing any number (MIXED_BENCHMARK_PROTOCOL.md). All six policies consume the *same* logged per-condition signals Z and benefit B from the locked CIFAR-10-C stress grid (Tent, 432 conditions/seed, five seeds; harmful fraction $\approx 16\text{--}18\%$); only the decision rule differs. POEM and AETTA are faithful ports of their published protectors (documented simplifications: batch-summary entropy rather than per-sample streams).

Verdict: WIN. KGA lowers mean regret-to-oracle vs. both POEM and AETTA with paired-bootstrap 95% CIs entirely below zero and Holm correction ($\alpha=0.05$), at $\text{FA}_u=0 \leq \alpha$ (Table IV). POEM and AETTA false-adapt at $\sim 24\text{--}25\%$ on the same stream—they lack the certificate’s marginal false-adapt control. Secondary pooled sets (Tent+EATA, EATA alone) yield the same WIN verdict (experiments/kbound/results/mixed_headtohead_v1/).

D. ImageNet-C SAR (mechanism-faithful)

On ImageNet-C noise (ResNet-50, 36 cells), SAR is harmful on 44% of cells under a mechanism-faithful re-implementation (SAM, recovery reset, EMA) at the learning rate shared across candidates. Always-adapt collapses

candidate	policy	mean acc	regret $\downarrow$	harmful cells	beats both?
Tent	always-adapt	0.454	0.0007	6%	no (ties adapt)
	always-freeze	0.405	0.0494		
	K-Bound	0.446	0.0082		
EATA	always-adapt	0.444	0.0000	3%	no (ties adapt)
	always-freeze	0.405	0.0386		
	K-Bound	0.440	0.0038		
SAR	always-adapt	0.319	0.1124	44%	
	always-freeze	0.405	0.0267		
	K-Bound	0.409	0.0229		yes

TABLE V

IMAGENET-C, MECHANISM-FAITHFUL SAR (36 CELLS, RESNET-50, $\alpha=0.1$).




Fig. 5. **Natural-shift safety at a glance.** Regret reduction of KGA over each fixed policy (positive $\rightarrow$ KGA better) with 95% bootstrap CIs under a valid out-of-fold radius. On both Office-Home and iWildCam, KGA beats always-adapt (CI excludes zero) and ties always-freeze (CI includes zero): no-harm / damage-prevention, not a beats-both.

to 0.319 accuracy while KGA reaches 0.409, **beating both** trivial policies (regret 0.0229 vs 0.1124/ 0.0267). Tent and EATA remain benign; KGA ties always-adapt (Table V). Under the same logged conditions, an ATC-style confidence gate *false-adapts on all harmful SAR cells* (false-adapt 1.00) while KGA abstains on them—so the certificate beats not only the two fixed policies but a simple label-free decision gate under the same protocol; direct committal-gate baselines are in Appendix M (Table XXIX). The ImageNet-C result uses our mechanism-faithful SAR reimplement at a shared learning rate; SAR’s official gentler schedule is left as a robustness check, so the claim is about this harmful-SAR operating point rather than all SAR configurations (Appendix).

E. Natural-shift results: uniformly no-harm under a valid out-of-fold radius

This is the deployment claim the theory targets. On every held-out natural shift, under a valid out-of-fold radius, KGA matches the better fixed policy and beats the worse—uniform no-harm, with no natural-shift beats-both. To keep that claim clean we separate two layers that must not be conflated: **Headline** = canonical KGA only (Algorithm 1; dev-locked adapter from a pre-registered panel, or fixed adapter for G; GBR $\hat{\Delta}$ + global in-domain ε ; held-out test once). **Diagnostic** = earlier Camelyon/iWildCam runs (sparse Z , multicandidate τ^* , Protocol B/F route audits) that explain failed operating points—Appendix A, §F—and

TABLE VI

Natural-shift protocols at a glance. THE ADAPTER IS LOCKED ON A DEV SPLIT, NEVER ON HELD-OUT TEST; KGA’S $\hat{\Delta}$ AND ε ARE CALIBRATED ON THE DEV SPLIT, AND THE HELD-OUT TEST IS SCORED ONCE. STATUS: CI-ROBUST = BOTH REGRET-GAP BOOTSTRAP CIs EXCLUDE ZERO; POINT = POINT-ESTIMATE BEATS-BOTH THAT TIES THE BETTER FIXED POLICY UNDER RESAMPLING; NO-HARM = MATCHES THE BETTER FIXED POLICY WITHOUT BEATING BOTH.

Dataset	Adapter lock	Calibration	Held-out test	Candidate	Status
Office-Home M v2	target-val panel	dev seeds {0, 1}	target-test {0, 1}	SAR aggressive	no-harm
iWildCam H v2	id_val panel	cal seed {0}	test seed {1}	Tent episodic	no-harm
Camelyon17 G	dev hospitals	dev seeds {0, 1}	genuine OOD {2, 3, 4}	EATA online	no-harm
RxRx1 (Prot. J)	locked protocol	protocol cal	OOD test	fixed online adapter	no-harm

TABLE VII

Camelyon17 Protocol G: withdrawn beats-both. THE LOCKED $n=54$ “HELD-OUT” SET POOLS IN-DISTRIBUTION `id_val` WITH THE OOD DOMAINS. ON GENUINE OOD TEST ALONE ($n=18$) ONLINE EATA NEVER HARMS, SO KGA TIES ALWAYS-ADAPT—A NO-HARM RESULT, NOT A WIN.

policy	regret $\downarrow$	FA $_c$	beats both?
<i>pooled, incl. in-distribution id_val (n=54)</i>			
always-adapt	0.0013	—	
always-freeze	0.0749	—	
K-Bound	3.6×10^{-5}	2.6%	artifact
<i>genuine OOD test only (n=18)</i>			
always-adapt	0.0000	—	
always-freeze	0.1381	—	
K-Bound	0.0000	0%	no (ties adapt)

motivated G/H/M v2. Diagnostics are not competing wins. No held-out test labels are used to select the adapter, tune ε , choose evidence features, or revise the decision rule.

F. Camelyon17 Protocol G (no-harm; withdrawn beats-both)

WILDS Camelyon17 is a hospital-to-hospital histopathology shift; the deployed adapter is EATA-online (fixed) with a rich label-free panel (17 statistics). An earlier Protocol-G analysis reported a beats-both at $n=54$, but that “held-out” set splits records *by random seed* and **pools in-distribution id_val cells** into the OOD test/val domains. On the *genuine* OOD held-out test alone ($n=18$), online EATA never harms—always-adapt is near-oracle (regret 0.0000)—so KGA **ties always-adapt** and does *not* beat both. We therefore **withdraw** the Protocol-G beats-both and reclassify Camelyon17 as an honest *no-harm* result: KGA matches the better fixed policy at false-adapt $\leq \alpha$ (Table VII). Audit and re-run: audits/integrity_2026-06-20/camelyon_reconciliation/.

G. iWildCam Protocol H v2

WILDS iWildCam is a geographic camera-trap shift (182 classes). Protocol H v2 dev-locks `tent_episodic` from a six-arm Tent/EATA/SAR panel on `id_val` (cal seed {0}, eval seed {1}), then scores held-out test once (72 cells). Under a valid *out-of-fold* conformal radius (§IV), KGA is **no-harm**: regret 0.0041 *equals* always-freeze 0.0041 and beats always-adapt 0.103 (gap +0.099 [0.080, 0.118],

TABLE VIII
iWildCam PROTOCOL H v2: DEV-LOCKED TENT_EPISODIC, HELD-OUT
TEST ($n=72$).

policy	regret $\downarrow$	FA $_c$	beats both?
always-adapt	0.1028	—	
always-freeze	0.0041	—	
K-Bound	0.0041	0%	no (ties freeze)

TABLE IX
OFFICE-HOME PROTOCOL M v2: DEV-LOCKED
SAR_ONLINE_AGGRESSIVE, TARGET-TEST ($n=35$).

policy	regret $\downarrow$	FA $_c$	beats both?
always-adapt	0.0468	—	
always-freeze	0.0158	—	
K-Bound	0.0157	0%	no (ties freeze)

bootstrap CI excludes zero), at false-adapt $0\% \leq \alpha$ (Table VIII). KGA correctly freezes the harmful camera-trap adaptation; this is damage-prevention no-harm, *not* a beats-both. (An earlier “point-estimate beats-both” used an in-sample radius; the valid out-of-fold radius abstains slightly more and lands exactly on always-freeze.)

H. Office-Home Protocol M v2

Office-Home is a domain shift across art/clipart/product/real. Protocol M v2 dev-locks **sar_online_aggressive** from a six-arm gradient-TTA panel on target-val, then scores held-out target-test (cal seeds $\{0,1\}$, test seeds $\{0,1\}$; 35 cells). Under a valid out-of-fold radius, KGA is **no-harm**: regret 0.0157 ties always-freeze 0.0158 (gap $+0.0001$ [0, 0.0003], CI includes zero) and beats always-adapt 0.047 (gap $+0.031$ [0.004, 0.062], CI excludes zero), at false-adapt $0\% \leq \alpha$ (Table IX). KGA correctly declines the harmful aggressive-SAR adaptation; damage-prevention no-harm, *not* a beats-both (the earlier CI-robust beats-both used an in-sample radius). Distinct from the deployed **eata_online_mild** audit (helpful-dominated null; Appendix).

a) Statistical robustness (bootstrap CIs).: Resampling the held-out test conditions ($B=2000$) with the dev calibration—estimator and *out-of-fold* radius—held *fixed*, no natural shift is a CI-robust beats-both; all are *dominate-adapt/tie-the-better-policy* no-harm. Office-Home beats always-adapt ($+0.031$ [0.004, 0.062]) but ties always-freeze ($+0.0001$ [0, 0.0003], CI includes zero); iWildCam beats always-adapt ($+0.099$ [0.080, 0.118]) and equals always-freeze *exactly* ($+0.0000$ [0, 0]); Camelyon17 beats always-freeze ($+0.072$ [0.053, 0.093]) and ties always-adapt (-0.0016 [-0.005 , 0.001]). On every natural shift KGA matches the better fixed policy and strictly beats the worse one at false-adapt 0—uniformly no-

harm, with *no* natural-shift beats-both. The earlier Office-Home/iWildCam beats-both used an in-sample conformal radius; this is the corrected out-of-fold re-run (**research_lock/KBOUND_WIN_BOOTSTRAP_CIS_oof.json**).

b) Cross-protocol aggregate (not a universal mixed deployment).: Each single dataset is one-sided, so KGA’s simultaneous margin over *both* policies is necessarily small there. To evaluate the policy benefit of per-condition adapt/freeze/abstain routing across heterogeneous shift regimes, we *aggregate three independently dev-calibrated held-out deployment protocols*: OfficeHome (35), iWildCam (72), and Camelyon17 (36 genuine OOD **test/val** cells; in-distribution **id_val** excluded to avoid the Protocol-G pooling artifact), for $n=143$ conditions. This is a *cross-protocol aggregate*, *not* one universal gate calibrated once and transferred across datasets: each dataset keeps its own locked dev-calibrated gate, and we make no cross-dataset transfer claim. (Camelyon17 contributes its 36 OOD **test+val** cells here, whereas the single-dataset headline of Table VII uses the $n=18$ OOD **test** split alone; both exclude in-distribution **id_val**.) Because this aggregate genuinely mixes helpful conditions (Camelyon) with harmful ones (Office-Home, iWildCam), it is the one setting where per-condition routing can in principle beat both fixed policies. The earlier figures (regret 0.0026; $\sim 24\times$ vs always-adapt, $\sim 13\times$ vs always-freeze) used an in-sample conformal radius and are **withdrawn**. The corrected out-of-fold re-run (**mixed_protocol_oof_v2**) yields $\text{regret}_{\text{KGA}}=0.0059$ vs. 0.0632 (always-adapt) and 0.0342 (always-freeze), with regret gaps vs. always-freeze $+0.0283$ [0.019, 0.038] and vs. always-adapt $+0.0573$ [0.043, 0.072] (both CIs exclude zero; pooled false-adapt 0). This is a **beats-both** on a *constructed* cross-protocol aggregate—not a universal mixed-deployment win and not a natural-shift headline.

c) Supporting and negative regimes.: These tracks support Table I but are not headline policy wins (details: Appendix A). **Controlled multimodal D33** (130 conditions): injected two-view corruption; KGA beats both trivial fusion policies at $\text{FA}_u=0$ (mechanism confirmation; Appendix). **CIFAR-10-C per-corruption** (65 cells) is helpful-dominated; KGA ties always-adapt. **RxRx1** is harmful-dominated at 1,139-class scale; always-adapt is destructive and KGA matches freeze-oracle safety. **ImageNet-R** remains weak and split-specific; KGA is not yet a stable headline win there. The deployed Office-Home adapter audit remains a helpful-dominated null (Appendix E), distinct from Protocol M v2. **CIFAR-10.1** is low-margin: within-seed calibration ties freeze with 0% false-adapts on Tent/EATA, but cross-seed certificate transfer does not clear the headline bar.

I. Real-camera deployment: a pre-registered protocol (future work)

To show the certificate extends beyond benchmark data, we *pre-register* (but do not yet report) a two-phone physical package-inspection study; it is deliberately *not* part of

this paper’s empirical claims and is included only to fix the evaluation in advance. The protocol tables are intentionally outcome-free until the source, calibration-fit, conformal, held-out, and external-device replication sessions are recorded and pass the audit in Table XXIV. Table X fixes the deployment protocol; Table XI is the only primary result table. The physical-shift breakdown, recording inventory, audit trail, runtime profile, and ablations remain in the supplement (Tables XXII–XXVII). No placeholder is used as evidence.

J. What exactly is guaranteed?

Statement	Status	Requirement
Abstention necessary in matched-evidence opposite-benefit worlds	theorem	stated model class
$FA_u \leq \alpha$ (unconditional false-adapt)	theorem	valid coverage + exchangeability + risk alignment
$FA_c \leq \alpha$ (conditional false-adapt)	not claimed	report FA_c descriptively only
Natural-shift no-harm (5 benchmarks)	headline empirical	dev lock + held-out OOF radius
Stress-grid (Tent/EATA) beats-both	confirmatory empirical	locked Protocol A v1 stress grid
Mixed vs POEM/AETTA (mixed harmful+helpful)	confirmatory empirical	pre-registered (SIV) WIN
Universal improvement	not claimed	impossible in general
ε estimates population drift budget β	false	β is declared; ε is empirical

VI. LIMITATIONS AND HONEST SCOPE

Scope of empirical wins. On the stress grid KGA beats both for Tent/EATA and ties SAR (Table II); on ImageNet-C the pattern flips by candidate—Tent/EATA are robust while SAR collapses and KGA beats both under a mechanism-faithful re-run (Table V). Per-corruption CIFAR-10-C and helpful natural shifts are insurance regimes, not policy wins.

Real-shift scope. Under a valid out-of-fold conformal radius, no natural shift is a beats-both: Office-Home and iWildCam are *no-harm* (they beat always-adapt and tie always-freeze; Tables IX, VIII), so the beats-both wins are on the synthetic stress grids (CIFAR-10-C, ImageNet-C SAR). Camelyon17 Protocol G is reclassified to an honest *no-harm* result—its earlier $n=54$ beats-both pooled in-distribution `id_val` into the held-out set and is withdrawn (Table VII). Earlier Camelyon Protocol B/F runs and multicandidate routes are **diagnostics**, not alternate headline wins.

Theory and calibration. The conformal radius assumes exchangeable calibration pairs; sign bracketing without structure is impossible (Theorem 1 and Appendix A). KGA’s value scales with how often catastrophic, detectable harm occurs—not with universal accuracy gains. On

helpful-dominated shifts the resulting near-*tie* with always-adapt reflects the certificate abstaining rather than committing harmful updates—KGA approximately matches the better policy, at most paying a small abstention cost when it declines a helpful but uncertifiable update—so the value proposition is conditional, like insurance against detectable harm.

Detectability and transfer. ImageNet-R still sits in a weak, split-specific evidence regime (Appendix E)—an instance of the *unknowable* region Theorem 1 predicts, where the label-free margin falls below the calibration-drift budget; KGA correctly declines to commit, so this is the theory operating as designed, not a tuning failure, and richer label-free evidence is the natural future lever. iWildCam H v2 clears the certificate bar despite modest margin over always-freeze; intermediate-domain multicandidate routes do not (pooled false-adapt $\gg \alpha$).

The limitations are the frontier, not flaws. The regime dependencies above are not three independent weaknesses but one phenomenon—the benefit-sign frontier (Theorem 1) seen from the operating side. Tying always-adapt on a helpful-dominated shift is the no-harm guarantee *holding*: one cannot beat an oracle that is already always-adapt. Declining ImageNet-R is the *unknowable* branch firing, where the observable margin falls below the drift budget and *no* label-free rule—however good the adapter—can decide the benefit sign (Lemma 1). The coverage drop is mandatory abstention, not a tuning loss: committing in the unknowable region would be a false-adapt. None can be removed without violating the impossibility result—a method that “won everywhere” would be claiming to decide the sign of an unidentifiable quantity. The one axis that genuinely improves all three is a *richer label-free evidence map*, which raises the observable margin $|M|$ so more shifts clear the budget (more knowable shifts, fewer abstentions, higher coverage), together with a *tighter radius* (jackknife+ or more calibration data) that reclaims coverage on knowable cells abstained only out of conservatism—both bounded by the frontier they cannot cross.

Missing baselines and independent validation. ImageNet-C committal gates (ATC-style confidence; Appendix M, Table XXIX) false-adapt at 100% on harmful SAR cells, while KGA abstains there. The pre-registered mixed head-to-head vs. POEM/AETTA (Table IV) is complete; an optional official-repo arm (per-sample POEM, dropout-AETTA) remains a camera-ready faithfulness check. Independent replication of ImageNet-R and Office-Home M v2 across more seeds would strengthen stability claims. Abstention lowers *decision* coverage by construction but never withholds a prediction—an ABSTAIN keeps the safe default f_0 (the source model), so every input still receives an output, and the coverage/false-adapt trade-off is tunable via α .

TABLE X

R1 – Pre-registered real-world deployment protocol. THE SPLIT AND ONLINE-LABEL CONSTRAINTS ARE FIXED BEFORE PHYSICAL HELD-OUT CAPTURE. SESSION IDENTIFIERS ARE POPULATED FROM THE SIGNED RECORDING MANIFEST.

Item	Locked design
Task	Four-class package/label inspection: correct, missing label, misaligned label, damaged label
Camera	Two phone cameras and real physical objects; the second phone is reserved for external-device replication
Base model	MobileNetV3-Small
Candidate adaptation	Episodic Tent; one update step per decision window
Decision window	32 frames
Evidence	14 label-free frozen/candidate disagreement and shift features
Calibration-fit sessions	S03, S04 (distinct recorded session; IDs locked in manifest)
Conformal calibration sessions	S05, S06 (different day/session from calibration-fit)
Held-out physical test sessions	S07, S08 (different day/session; inaccessible during tuning)
Physical shifts	Lighting, shadow, blur, background, viewpoint, distance, and batch composition
Policies compared	Always-freeze, always-adapt, confidence gate, entropy gate, KGA without radius, KGA certificate
Deployment labels used live?	No; labels are joined only during offline evaluation
Miscoverage level	$\alpha = 0.10$
Official live model	Frozen fallback unless the candidate is certified for adaptation

TABLE XI

R2 – Primary held-out real-world result (RESULT PENDING). ALL POLICIES WILL BE SCORED ON THE IDENTICAL LOCKED PHYSICAL STREAM AFTER REAL CAPTURES PASS THE SOURCE-MODEL GATE (≥ 0.80 BALANCED ACCURACY AND MACRO-F1 ON S02). UNTIL THEN EVERY METRIC CELL READS *RESULT PENDING* — *NO MEASURED DATA AVAILABLE*; DEVELOPMENT REPLAYS ON MOCK CLIPS ARE AUDIT-ONLY AND MUST NOT BE CITED.

Policy	Balanced acc. $\uparrow$	Macro-F1 $\uparrow$	Regret $\downarrow$	$FA_u \downarrow$	$FA_c \downarrow$	Adapt rate	Abstain rate	Mean latency (ms) $\downarrow$
Always-freeze	—	—	—	—	—	—	—	—
Always-adapt	—	—	—	—	—	—	—	—
Confidence gate	—	—	—	—	—	—	—	—
Entropy gate	—	—	—	—	—	—	—	—
KGA, no radius	—	—	—	—	—	—	—	—
KGA certificate	—	—	—	—	—	—	—	—

$FA_u = \mathbb{P}(\text{adapt} \wedge \Delta \leq 0)$ is the certificate-aligned unconditional false-adapt rate; $FA_c = \mathbb{P}(\Delta \leq 0 \mid \text{adapt})$ is descriptive. Ground-truth labels are unavailable during live decisions and are revealed only for offline evaluation. Latency reporting includes candidate adaptation, the KGA decision, and full end-to-end window latency.

What we do not claim. Universal accuracy gains; adapter selection by test regret; multicandidate routes with violated false-adapt bounds; or wins on helpful-dominated shifts where always-adapt is already oracle-optimal.

VII. REPRODUCIBILITY

All code, locked protocol yamls (`research_lock/`), and result JSONs backing every reported number are in the anonymized submission repository. Controlled and natural-shift headline runs are reproduced via `bash docs/research/kbound/scripts/kbtrain.sh` (protocol-h-v2, protocol-m-v2, stress-grid targets); theory validators and CPU experiments run from the bundled reproduction driver. The core certificate module is packaged as an installable module in the anonymized repository (`docs/research/kbound/kbound_pkg/`).

VIII. CONCLUSION

K-Bound reframes test-time adaptation as an identifiability-limited decision problem: adapt when helpful, freeze when harmful, and abstain when label-free evidence cannot support a valid decision. The theory shows that this abstention region is not merely

conservative; under matched label-free evidence with opposite adaptation benefit, no valid label-free rule can safely commit. **Our headline empirical result is uniform no-harm on five held-out natural shifts** under a valid out-of-fold radius—KGA prevents harmful adaptation without sacrificing the near-oracle policy on helpful-dominated shifts. Wins over trivial policies and POEM/AETTA appear only in identifiable mixed harmful+helpful regimes, exactly where the frontier predicts knowability. These results suggest that future TTA systems should expose not only an adaptation objective, but also a validity layer that decides whether adaptation is knowable before deployment.

APPENDIX

Algorithm 2 is the shared evaluation protocol behind Tables II, VII, VIII, IX, and the supplementary tracks below.

This appendix is not required for the main empirical claim; it preserves the full theorem stack and validator mapping behind the consolidated Theorem 1. The main paper states a consolidated headline theorem (Theorem 1).

Algorithm 2 Calibration and held-out evaluation for K-Bound

Require: Conditions $\mathcal{C}$; frozen f_0 ; adapters $\mathcal{A}_1, \dots, \mathcal{A}_K$; calibration split $\mathcal{C}_{\text{cal}}$; test split $\mathcal{C}_{\text{test}}$; level α

Ensure: Regret-to-oracle, false-adapt rate, coverage, beats-both verdict

```

1: for all  $c \in \mathcal{C}_{\text{cal}}$  do
2:   for all adapters  $\mathcal{A}_k$  do
3:      $f_a^{(k,c)} \leftarrow \mathcal{A}_k(f_0, X_c)$ 
4:      $Z_{k,c} \leftarrow \phi(X_c, f_0, f_a^{(k,c)})$ 
5:      $\Delta_{k,c} \leftarrow R_c(f_0) - R_c(f_a^{(k,c)})$ 
6:   end for
7: end for
8: Out-of-fold radius (no leakage): for each calibration
   pair  $i$ , fit  $\hat{\Delta}_{-i}$  on all calibration pairs except  $i$  and set
    $\rho_i \leftarrow |\hat{\Delta}_{-i}(Z_i) - \Delta_i| \triangleright$  estimator/residual data disjoint
9:  $\varepsilon \leftarrow (1 - \alpha)$  empirical quantile of  $\{\rho_i\} \triangleright$  leave-one-out
   (jackknife) radius
10: Fit deployment estimator  $\hat{\Delta}(Z)$  on all calibration pairs
     $\triangleright$  scores the held-out test split
11: for all  $c \in \mathcal{C}_{\text{test}}$  do
12:   for all adapters  $\mathcal{A}_k$  do
13:     Run Algorithm 1 with  $(f_0, \mathcal{A}_k, X_c, \hat{\Delta}, \varepsilon)$ 
14:     Record decision, regret-to-oracle, coverage, false-
       adapt indicator
15:   end for
16: end for
17: Aggregate metrics vs. always-adapt, always-freeze, per-
    condition oracle

```

Below we give the full five-theorem stack, lemmas, and proof sketches with validator references.

The body is organized as four pillars, with five core theorems serving as the load-bearing statements inside that scaffold. **Pillar I (impossibility floor):** Theorem 2 plus the Le Cam lower bound and regret identity establish when abstention is information-theoretically forced. **Pillar II (certificate + audit):** Theorem 3 gives finite-sample false-adapt control, and the one-bit audit machinery of Theorem 5(iv) supplies a common falsifier for the legacy budget premises (formerly presented as separate propositions, now treated as corollaries of one audit framework). **Pillar III (identifiable regimes):** Theorem 4 and Theorem 5 characterize exactly what label-free evidence identifies and what one-bit supplement remains irreducible. **Pillar IV (rates):** Theorem 6 gives matching evidence-channel rates. The open piece—a weakest falsifiable class on which one declared bit suffices—is settled in Appendix P: conditionally for the single class $\mathcal{C}_{\text{mono}}$, and unconditionally as an explicit finite family of dominance polytopes (Theorem 8). Supporting variants (multiclass/regression reductions, label-shift and drift boundaries, ambiguity-reach unification, router capacity, anytime certificates, the γ -meter, and the agreement-accuracy law) are presented as corollaries or companions in the companion manuscript,

each with a validator; every numeric claim here traces to a script in the public repository.

Lemma 2 (Reduction to the disagreement region). *Let $D = \{x : f_0(x) \neq f_a(x)\}$ (observable). For binary Y and 0/1 loss, with $\bar{a} = \mathbb{P}_T(f_a=Y \mid D)$, $\Delta = 2\mu(D)(\bar{a} - \frac{1}{2})$, so for $\mu(D) > 0$: $\text{sign } \Delta = \text{sign}(\bar{a} - \frac{1}{2})$. Writing s for any source-calibrated score of f_a , $M := \mathbb{E}_{\mu|D}[s] - \frac{1}{2}$ (observable) and $\gamma := \mathbb{E}_{\mu|D}[\eta_a - s]$ (unobservable), $\boxed{\text{sign } \Delta = \text{sign}(M + \gamma)}$. For K -class 0/1 loss $\Delta = \mu(D)(p_a - p_0)$, and for squared loss the sign reduces to one moment on D ; both extend the identity verbatim (manuscript, App. C).*

Proof sketch. Off D the losses coincide; on D exactly one of f_0, f_a is correct (binary), giving $\mathbb{E}[\delta \mid D] = 2\bar{a} - 1$; linearity splits $\bar{a} - \frac{1}{2} = M + \gamma$. Validated on synthetic draws (identity exact to machine precision; multiclass to 10^{-16} over 4000 trials). $\square$

Lemma 3 (Plug-in regret identity). *For any estimator $\hat{\Delta}$, the gate $\hat{g} = \mathbf{1}[\hat{\Delta} > 0]$ satisfies $R(\hat{g}) - R(g^*) = \mathbb{E}[\Delta(Z) \mathbf{1}\{\text{sign } \hat{\Delta} \neq \text{sign } \Delta\}]$ with $g^* = \mathbf{1}[\Delta > 0]$ Bayes; on $\{|\Delta| < \varepsilon\}$ committal regret is $\leq \varepsilon$, justifying abstention in the low-margin band. Validated to machine precision on synthetic regret identities.*

Proposition 1 (Finite-sample minimax regret floor for label-free gating). *Intuition. The identity of Lemma 3 is an equality, not a hardness statement; pairing it with Le Cam’s two-point method turns it into a finite- n lower bound on the regret of any label-free gate, which does not vanish as $n \rightarrow \infty$ when the two worlds are rescaled to stay equally hard. Let $\{P_\theta\}_{\theta \in \{-1, +1\}}$ be two target worlds sharing the same disagreement structure with constant benefit magnitude $|\Delta(P_\theta)| \equiv \Lambda > 0$ and opposite signs $\text{sign } \Delta(P_\theta) = \theta$, and let $Q_\theta^{(n)}$ denote the law of any label-free evidence Z computed from n unlabeled observations under P_θ . Then for every label-free committal gate $\hat{g} : Z \mapsto \{\text{ADAPT}, \text{FREEZE}\}$ and every finite n ,*

$$\begin{aligned}
& \inf_{\hat{g}} \max_{\theta \in \{\pm 1\}} [R_{P_\theta}(\hat{g}) - R_{P_\theta}(g^*)] \\
& \geq \frac{\Lambda}{2} \left(1 - \text{TV}(Q_{-1}^{(n)}, Q_{+1}^{(n)})\right) \\
& \geq \frac{\Lambda}{4} e^{-\text{KL}(Q_{-1}^{(n)} \| Q_{+1}^{(n)})}.
\end{aligned}$$

In particular there is an explicit Gaussian two-point family ($X \sim \mathcal{N}(\theta d_n, 1)$ on D , $f_0 = \mathbf{1}[x > 0]$, $f_a = \mathbf{1}[x < 0]$, with the target label rule giving $\Delta = \theta$, so $\Lambda = 1$) for which, taking $d_n = c/\sqrt{n}$, both bounds are constant in n : $\text{TV}(Q_{-1}^{(n)}, Q_{+1}^{(n)}) = 2\Phi(c) - 1$ and the floor equals $\Lambda \Phi(-c)$, with the Bretagnolle–Huber certificate $\frac{\Lambda}{4} e^{-2c^2} > 0$. Theorem 2 is the $c \rightarrow 0$ ($\text{TV} \rightarrow 0$) face, where the floor saturates at $\Lambda/2$.

Proof. By Lemma 3, since $|\Delta| \equiv \Lambda$ and g^* is the constant correct action $\text{sign } \theta$ in world θ , $R_{P_\theta}(\hat{g}) - R_{P_\theta}(g^*) = \Lambda Q_\theta^{(n)}(\hat{g} \text{ commits to the wrong sign})$. Averaging over the

uniform prior on θ gives $\frac{1}{2} \sum_{\theta} [R_{P_{\theta}}(\hat{g}) - R_{P_{\theta}}(g^*)] = \frac{\Delta}{2} M(\hat{g})$, where $M(\hat{g}) = Q_{+1}^{(n)}(\hat{g}=\text{FREEZE}) + Q_{-1}^{(n)}(\hat{g}=\text{ADAPT})$ is the mixed committal error of Theorem 2(ii). Le Cam’s testing-affinity identity yields $\inf_{\hat{g}} M(\hat{g}) = 1 - \text{TV}(Q_{-1}^{(n)}, Q_{+1}^{(n)})$, and $\max \geq \text{avg}$ gives the minimax form; Bretagnolle–Huber ($1 - \text{TV} \geq \frac{1}{2} e^{-\text{KL}}$) gives the closed-form certificate. For the Gaussian family the sufficient statistic $\bar{X} \sim \mathcal{N}(\theta d_n, 1/n)$ has $\text{TV} = 2\Phi(\sqrt{n} d_n) - 1$ and per-observation $\text{KL} = 2d_n^2$, so $d_n = c/\sqrt{n}$ makes both n -independent; $|\Delta| = 1$ exactly because each label rule makes one of f_0, f_a correct on a probability-one event. Validated (`val_thm2_lecam_finite_n.py`, seeds fixed): across $c \in \{0.25, 0.5, 1, 1.5\}$ and $n \in \{10, 50, 200, 1000\}$ the empirical floor is constant in n , the Bayes (sample-mean threshold) gate *meets* it (so it is tight), and a brute-force minimax over a panel of label-free rules (all thresholds/polarities on five statistics) never beats it within Monte-Carlo error. $\square$

Theorem 2 (The unknowable regime is real and tight). *Intuition. Identical label-free evidence can support opposite adaptation benefits; any valid rule must then abstain rather than guess.* (i) *There exist target worlds P_T^1, P_T^2 with identical evidence laws, $\text{Law}(Z \mid P_T^1) = \text{Law}(Z \mid P_T^2)$, and sign $\Delta_1 = -\text{sign } \Delta_2 \neq 0$ (explicit witness: $X \sim \mathcal{N}(0, 1)$, $f_a = 1 - f_0$, flipped labels). No label-free rule distinguishes them; the worst-case-optimal committal-regret action is ABSTAIN.* (ii) *Quantitatively, for any committal rule g the mixed error obeys the Le Cam identity $\inf_g M(g) = 1 - \text{TV}(P_{T,Z}^{1,(n)}, P_{T,Z}^{2,(n)})$, equal to 1 in the witness for every n and decaying only as the evidence laws separate.* (iii) *Any rule with false-adapt and false-freeze rates $\leq \alpha$ must abstain on matched evidence with probability $\geq 1 - 2\alpha$: abstention is mandatory, not conservative.*

Proof sketch. (i) The witness makes Z a function of X alone with $P^1(X) = P^2(X)$. (ii) A committal rule is a test between the two evidence laws; Le Cam’s testing affinity plus data processing. (iii) Matched evidence forces common action probabilities q_a, q_f ; validity caps each at α . Validated: empirical $\inf_g M$ tracks $1 - \text{TV}$ across all n , 100% abstention on the witness. $\square$

Theorem 3 (The adapt/freeze/abstain certificate). *Intuition. When a benefit estimate carries a valid uncertainty radius, committing only outside the abstain band controls false-adapt and false-freeze at level α . Suppose $\hat{\Delta}(z)$ admits a radius $\varepsilon(z)$ with $\mathbb{P}[|\hat{\Delta} - \Delta| \leq \varepsilon] \geq 1 - \alpha$. The rule ADAPT if $\hat{\Delta} - \varepsilon > 0$, FREEZE if $\hat{\Delta} + \varepsilon < 0$, else ABSTAIN, has marginal, unconditional false-adapt and false-freeze rates $\leq \alpha$ each (over the calibration draw under exchangeability). With $\hat{\Delta}$ an empirical-Bernstein mean of n bounded paired benefits (variance proxy σ^2 , range $b - a$), the certificate begins committing (one-sided, radius $\varepsilon < |\Delta|$) once $n \gtrsim (2\sigma^2/\Delta^2) \log(2/\alpha) + 3(b - a)|\Delta|^{-1} \log(2/\alpha)$, and certifies the sign two-sided — both the miss and the wrong-commit probability $\leq \alpha$, which requires the*

whole $(1 - \alpha)$ ball to clear 0, i.e. $\varepsilon < |\Delta|/2$ — once $n \gtrsim (8\sigma^2/\Delta^2) \log(2/\alpha) + O((b - a)|\Delta|^{-1} \log(1/\alpha))$: sample complexity scales as Δ^{-2} , logarithmically in $1/\alpha$, and is minimax order-optimal in the extended theory appendix. A sequential, anytime-valid e-process version (valid at any stopping time; β -drift-robust variant included) is Theorem B.1 of the manuscript.

Proof sketch. On the good event, ADAPT implies $\Delta \geq \hat{\Delta} - \varepsilon > 0$; the complement has mass $\leq \alpha$; the one-sided commit threshold $\varepsilon(n) < |\Delta|$ inverts the Bernstein radius, while two-sided certification of the sign (miss $\leq \alpha$) requires $\varepsilon(n) < |\Delta|/2$, hence the $8\sigma^2/\Delta^2$ rate proved in the minimax extension. Note that this marginal conformal guarantee applies to the radius and decision rather than the accuracy of the benefit model itself, and relies on calibration-deployment exchangeability at the locked-protocol granularity, degrading under task shift. The residual burden — a calibration sample exchangeable with deployment — is exactly what Theorem 5(iv) makes auditable. Validated on synthetic e-process tests (false-adapt $\leq \alpha$ at nominal level). $\square$

Theorem 4 (Exact benefit-sign frontier). *Intuition. On the disagreement region, the benefit sign is knowable iff an observable margin $|M|$ exceeds a calibration-drift budget β ; otherwise abstention is forced. Fix the observables (μ, P_S, f_0, f_a) with $\mu(D) > 0$, so M is determined, and let $\mathcal{C}_{\beta} = \{P_T : |\gamma| \leq \beta\}$. Noting that γ is a population quantity not observed at deploy, the statements below hold unconditionally for every drift budget $\beta \geq 0$ and every target with $\mu(D) > 0$ — the sign reduction of Lemma 2 is an algebraic identity requiring no regularity — with β a free parameter of the frontier rather than an assumption (and (iii) shows every identifying class is of this form). The location/alignment regularity invoked elsewhere is required only for the separate graded capacity statement in the appendix, not for the sign frontier here.* (i) γ is a one-dimensional sufficient statistic for sign Δ given the observables. (ii) sign Δ is identifiable over $\mathcal{C}_{\beta}$ iff $|M| > \beta$, and then equals sign M ; the minimax committal error is 0 if $|M| > \beta$ and $\frac{1}{2}$ otherwise, so the three-way rule (ADAPT if $M > \beta$, FREEZE if $M < -\beta$, else ABSTAIN) is the unique maximal sound certificate. (iii) Necessity: any class on which sign Δ is identifiable lies on one side of the hyperplane $\gamma = -M$; every identifying assumption is a one-sided constraint on γ . (iv) The label-free heuristics ATC [13], DoC, GDE [19], COT [20], AETTA [18], and agreement-on-the-line [16], [17] are exactly the $\beta=0$ face: sound iff $\gamma = 0$, sharing the single failure mode $\gamma \neq 0$.

Proof sketch. All from Lemma 2: (ii) a $\pm\beta$ perturbation crosses 0 iff $|M| \leq \beta$, plus the two-point argument of Theorem 2; (iii) constancy of $\text{sign}(M + \gamma)$ is one-sidedness; (iv) each heuristic thresholds an observable surrogate, i.e. commits to sign M . Validated: frontier law exact over 5×10^5 draws; the unconditional family characterization

(identifiable over $\mathcal{C}_\beta$ iff $|M| > \beta$) holds for every tested β ; ATC error matches the predicted event to machine precision at drift 0.04–0.28 (synthetic validators for Theorems T1–T5 and AGL). $\square$

Corollary 2 (Knowable regimes). (a) *Under covariate shift ($P_S(Y|X) = P_T(Y|X)$, bounded density ratio), importance weighting identifies Δ and its sign from labeled source plus unlabeled target (γ -budget discharged with $\beta = 0$ on the importance-weighted score); the premise itself is not label-free testable, re-entering Theorem 2 when it fails.* (b) *Conversely the budget is irreducible: no label-free certificate identifies sign Δ without a supplied bound on γ — the assumption-free certificate sought by prior heuristics does not exist. (Label shift, bounded drift, smooth drift, and the unified ambiguity-reach hierarchy: manuscript, App. C.)*

Theorem 5 (One-bit dichotomy; Conjecture 1 resolved). *Intuition. Label-free observables fix every adaptation statistic except one global orientation bit; no testable assumption class removes that bit in full generality. Consider a panel $f_0, f_1, \dots$ on D with correctness indicators conditionally independent given Y (CEI) and per-class symmetric accuracies (jointly, hypothesis H ; advantages $b_j = 2a_j - 1$, agreements $c_{ij} = 2A_{ij} - 1$; the rank-one algebra and anchor are classical [26], [27], [29], [30] — new here are the deficit, the dichotomy, and the audit). (i) H is minimal: under CEI alone, $c_{ij} - b_i b_j = \pi(1 - \pi)\delta_i \delta_j$ exactly (δ_j the per-class accuracy gap), so the rank-one law holds for a $K \geq 3$ panel iff every $\delta_j = 0$. (ii) One bit: under H the full evidence law determines every $|b_j|$, every product $b_i b_j$, and every benefit magnitude $|b_j - b_0|$ — everything except a single global flip $b \mapsto -b$, which is free at every moment order: the label complement preserves the entire evidence law ($\text{TV} = 0$) while negating every benefit sign. The drift $\gamma = b_a/2 - M$ is thereby point-identified up to that bit. (iii) No testable bracketing (the dichotomy): in any output space (binary, multiclass, regression) there is a swap involution $P \mapsto P^*$ — conditional mass exchange between the two disagreeing predictions on D , midpoint reflection for regression — with $\text{TV}(L(P), L(P^*)) = 0$ for the law L of all label-free observables (labeled source included) and $\Delta(P^*) = -\Delta(P)$. Hence no evidence-definable (label-free testable) assumption class identifies sign Δ ; every identifying assumption is a bit-selector (γ -budgets select sign M ; majority-above-chance selects $\text{sign}(\text{median } b)$; the agreement-line slope selects $\text{sign}(1 - 2\bar{w})$); and the minimal untestable supplement is exactly one bit, which suffices under H . (iv) Budgets are falsifiable, never verifiable: the bit-robust test rejecting “ $|\gamma| \leq \beta$ ” when $\min_{\text{flip}} |\hat{\gamma}| > \beta + r_n$ is level- α with power $\geq 1 - \alpha$ beyond $2r_n$; verification fails exactly on the blind set $\{ ||M| - |b_a|/2| \leq \beta < |M| + |b_a|/2 \}$, and the $M \geq 4$ rank-one residual τ is sound but not complete (strictly interior stealth laws with $\tau = 0$ exist that bias b and can flip the recovered decision).*

Proof sketch. (i) Expand c_{ij} over Y and cancel cross

terms. (ii) Triple products $c_{ik}c_{il}/c_{kl} = b_i^2$; the complement $Y' = 1 - Y$ fixes predictions (maps of X) hence the whole evidence law, while $b \mapsto -b$ and H is preserved. (iii) The swap preserves the X -marginal and source, so L is fixed at every order; on D it exchanges $p_a \leftrightarrow p_0$ (or negates $f_0 + f_a - 2Y$), so Δ negates; identifiability on a swap-closed class is impossible, and selecting one member per swap orbit is by definition one bit. Algebraically: observables generate the even correctness subalgebra ($s_i s_j = u_i u_j$), the decision is odd ($s_j = u_j v$), and the swap is $v \mapsto -v$. (iv) Empirical-Bernstein for $\hat{M}$, a product-ratio perturbation bound for $\hat{b}_a$, and the flip-min for bit-robust soundness; the stealth construction is a strictly interior LP on the 2^4 correctness patterns matching all pairwise moments to a rank-one target with a shifted first moment. Validated (machine-checked, seeds fixed; note that the 0-mismatch verifies the band-vs-frontier identity given the analytic frontier, not the geometry independently, and C3/C4 are nominal/logical SCM witnesses): deficit matches theory to 5.7×10^{-4} ; flip $\text{TV} = 0$ with $\max_j |b'_j + b_j| = 1.1 \times 10^{-16}$; swap evidence change exactly 0 with $\Delta^* = -\Delta$ (multiclass $K=5$ and regression); audit level $0.000 \leq \alpha=0.05$, power 1.000; stealth $\tau = 2 \times 10^{-17}$ with recovery bias 0.229 (Figs. 6–9). Full proofs appear in the supplementary theory appendix. $\square$

Conjecture 1 (Label-free bracketing — resolved). *The unconditional label-free bracketing of the ordinal comparison on D (posed as the open piece in earlier versions) does not exist: by Theorem 5(iii) no evidence-definable assumption identifies sign Δ , and the minimal untestable supplement is one bit. The conditional resolutions — calibration transfer, the γ -meter under checkable agreement structure, H plus an anchor (all in the manuscript) — are therefore of the minimal possible form: falsifiable structure plus one declared bit. On the margin-monotone relative-calibration class $\mathcal{C}_{\text{mono}}$ (Definition 2) one declared bit certifies sign Δ unconditionally (Theorem 7(i)–(ii)); and $\mathcal{C}_{\text{mono}}$ is a weakest such class under a general-position assumption (Theorem 7(iii)). The unconditional weakest classes (general position removed) are characterized as an explicit finite family of dominance polytopes (Theorem 8); there is no unique weakest class, consistent with the bracketing non-existence above.*

Theorem 6 (The evidence-channel rate: labels buy constants and the bit, not the rate). *Intuition. Unlabeled agreement statistics and labeled paired benefits share the same $1/\sqrt{m}$ rate; labels change constants and supply the missing bit, not the exponent. Within audited H with margins ($\min_{i \neq j} |c_{ij}| \geq c_{\min}$, $\min_j |b_j| \geq b_{\min}$) and a declared bit, m unlabeled samples on D give $|\hat{b}_j - b_j| \leq C \sqrt{\log(K/\delta)/m}$ with $C = \Theta(1/(b_{\min} c_{\min}^2))$, so $\text{sign}(b_j - b_0)$ is certified at $m \asymp C^2 (b_j - b_0)^{-2} \log(K/\delta)$; a Le Cam two-point bound inside H shows $m^{-1/2}$ is minimax for the evidence channel. The labeled-paired-benefit channel has the same rate: the efficiency ratio is constant in m but scales as $1/c_{\min}$, diverg-*




Fig. 6. One-bit identification (Synthetic validator; Thm. 5(ii)).




Fig. 7. Budget audit (Synthetic validator; Thm. 5(iv)).

ing as agreement approaches chance. Relatedly, accuracy is an exact affine function of agreement-with-reference iff the disagreement win rate is constant (slope $1 - 2\bar{w}$), the anchored form of agreement-on-the-line; its slope sign is the same bit as in Theorem 5(iii) (manuscript, §AGL).

Proof sketch. Hoeffding on agreements, the product-ratio perturbation lemma, and a χ^2 /KL bound between K -bit pattern laws differing by ε in one coordinate ($\text{KL} \leq C'\varepsilon^2$), tensorized; Bretagnolle–Huber finishes. Validated: recovery slope -0.51 ; minimax slope -0.525 tracking $c_0/\sqrt{m}$; $\text{KL}/\varepsilon^2 \in [0.164, 0.171]$; efficiency $2.15 \rightarrow 4.49 \rightarrow \infty$ as $c_{\min} 0.21 \rightarrow 0.067 \rightarrow \leq 0.022$ (Fig. 8; AGL: $R^2 = 1.000$ under constant w , slope $1 - 2w$ exact on synthetic checks). $\square$

The weakest one-bit class. Theorem 5 fixes the supplement at one bit; the *weakest falsifiable class* on which that bit suffices is settled conditionally—a margin-monotone relative-calibration class on which one declared bit certifies sign Δ unconditionally, and a weakest such class under a general-position assumption (Appendix P, Thm 7). The *unconditional* weakest classes form an explicit finite family of dominance polytopes (Thm 8).

A. Validation figures

This appendix holds **non-headline** tracks: helpful-dominated controls, honest nulls, and the Camelyon/iWildCam **diagnostic ladder** (Protocols B/F, multicandidate routes, sparse-Z failures) that motivated the headline Protocols G / H v2 / M v2 in the main text.

B. ELARA-U as a KGA instantiation (retrospective audit)

K-Bound supplies the framework, KGA supplies the adapt/freeze/abstain certificate, and ELARA-U supplies a validation-fitted multimodal routing candidate. We execute this composition through one canonical runner on five previously opened multimodal score caches (85




Fig. 8. Evidence-channel rate (Synthetic validator; Thm. 6).



Fig. 9. Conjecture 1 closure validator.

valid categories). Because full target labels construct the per-category Brier-benefit certificates in this audit, it is **retrospective—not a label-free or headline claim**. At $\alpha=0.10$, KGA returns 7 adapt, 2 freeze, and 76 abstain decisions, with 0 observed Brier-defined false-adapts and 10.6% coverage. Mean AUROC is 0.7941 (freeze), 0.8220 (ELARA-U candidate), and 0.7939 (KGA); thus this opened-data audit does not beat both fixed policies and cannot support promotion. The table is generated directly from the versioned result artifact.

C. Online non-stationary CIFAR-10

D. CIFAR-10-C per-corruption (helpful-dominated control)

E. Excluded wins and regime boundaries

Several benchmarks flag beats-both under alternate routes or adapters but do *not* qualify as headline natural-shift wins under the Protocol G/H/M v2 bar (canonical KGA, dev-locked adapter from pre-registered panel, in-domain calibration, held-out test, false-adapt $\leq \alpha$).

a) *Superseded protocols.*: iWildCam Protocol H v1 (fixed **sar_online**) and OfficeHome Protocol M v1 (14-candidate win-finder) are audit-only; headline claims use H/M v2.

b) *Office-Home (deployed adapter).*: The source-selected deployed adapter (mild online EATA) is *helpful-dominated* on target test (0% harmful cells; always-adapt is oracle). KGA has 0% false-adapts but does not beat both (regret 0.015 vs always-adapt 0). An aggressive online SAR arm can beat both on the same test split, but that arm was not the deployed adapter; reporting it as a win would be adapter shopping. The multicandidate deploy route abstains on 100% of target-test cells (certificate transfer-AUC ≈ 0.49). Protocol M v2 (main text) uses a separate dev-lock panel and held-out target-test scorer.

TABLE XII
KGA OVER ELARA-U ON OPENED MULTIMODAL CACHES. THIS IS A RETROSPECTIVE AUDIT, NOT A LABEL-FREE OR HEADLINE CLAIM. AUROC IS AVERAGED OVER VALID CATEGORIES.

Track	n	freeze	ELARA	KGA	A/F/U	status
3D-ADAM	23	0.796	0.866	0.796	0/0/23	audit only
Real-IAD-D3	20	0.873	0.888	0.873	7/0/13	audit only
MulSen-AD	15	0.922	0.919	0.922	0/0/15	audit only
MVTec-3D	8	0.846	0.853	0.846	0/0/8	audit only
Real-IAD-NatDeg	19	0.585	0.609	0.585	0/2/17	audit only

Mode: retrospective_audit; promotion eligible: no.

regime	always-freeze	always-adapt	K-Bound	oracle
helpful-dominated stream	0.472	0.548	0.553	0.704
harm-dominated stream	0.440	0.339	0.426	0.630
mean across regimes	0.456	0.444	0.490	0.667

TABLE XIII

ONLINE CIFAR-10 TTA (3 SEEDS/REGIME). KGA IS NEAR-BEST IN BOTH REGIMES.

method	mean acc (65 cells)	false-adapt	regret-to-oracle
frozen	0.412	0.000	0.217
Tent (online)	0.626	0.015	0.0030
K-Bound (KGA)	0.626	0.015	0.0032
oracle (per-cell)	0.629	—	—

TABLE XIV

PER-CORRUPTION CIFAR-10-C: KGA TIES ALWAYS-ADAPT (SAFETY INSURANCE).

c) ImageNet-R.: The diverse-backbone panel is helpful-leaning (mean $\bar{B}$ =+0.017, harmful rate 13.9%, best harm-AUC 0.662). The multicandidate route abstains on all 48 conditions; single-candidate KGA often ties always-adapt (oracle) on helpful backbones. This is the weak-detectability / unknowable regime predicted by Corollary 2.

d) iWildCam intermediate domains.: On id_val, a multicandidate router can flag beats-both on subsets, but pooled false-adapt on committed adapts is $\approx 78\% \gg \alpha$ —disqualified. Headline claim: Protocol H v2 on held-out test only (Table VIII).

e) Camelyon17 four-seed hospital slice.: The original four held-out hospital seeds are helpful-dominated (0% harmful); KGA beats always-freeze but not always-adapt (Table XV). Protocol G uses a composition-stress grid with rich evidence where harmful cells are mixed (30%) and detectable—a different regime on the same benchmark family.

f) CIFAR-10.1.: Within each seed, Tent/EATA abstain heavily with 0% false-adapts and tie always-freeze. Cross-seed calibration (fit on seeds 0–2, test on 3–4) does *not* clear the headline bar (false-adapt 44%); low-margin benefits do not transfer across independent stream seeds. Report CIFAR-10.1 as a within-seed safety probe, not a fourth natural headline win.

F. Camelyon17 diagnostic ladder

a) Protocol F route audit (diagnostic only—not headline).: Protocol G (main-text headline) uses canonical GBR

candidate	policy	mean balanced acc	regret↓	beats both?
Tent	always-adapt	0.817	0.000	
	always-freeze	0.786	0.031	
	K-Bound	0.797	0.020	no
EATA	always-adapt	0.836	0.000	
	always-freeze	0.786	0.050	
	K-Bound	0.818	0.018	no

TABLE XV

WILDS CAMELYON17 (4 SEEDS): HELPFUL-DOMINATED; KGA BEATS FREEZE, NOT ADAPT.

+ global ε on EATA-online only. The table below audits *alternate routes* on the same GPU records: multicandidate routing **fails**; pooled GBR+global and PPI+Mondrian *can* beat both but are not the headline method. Sparse-evidence Protocol B also fails (false-adapt 0.33 $\gg \alpha$). These negatives motivated G.

TABLE XVI

Diagnostic only. PROTOCOL F ROUTE AUDIT (SAME RECORDS AS PROTOCOL G).

variant	regret↓	false-adapt	beats both?
multicandidate router	0.093	71%	no
GBR + global ε	0.0019	3.3%	yes
PPI + Mondrian	0.0018	3.3%	yes

G. RxRx1 (harmful, freeze-oracle)

RxRx1 9+ grid (360 conditions)	value
harmful base rate	81.5%
mean benefit $\bar{B}$	−0.063
harm detectability (best AUC)	0.78–0.84
SAR-online always-adapt (collapse)	0.024
KGA false-adapts	0

TABLE XVII

RxRx1: KGA MATCHES FREEZE-ORACLE; DAMAGE PREVENTION, NOT BEATS-BOTH.

H. ImageNet-R (unknowable regime)

I. Office-Home and iWildCam

J. Controlled multimodal corruption (Protocol D33)

Pre-registered MNIST two-view fusion (130 conditions; one modality corrupted across 13 severities). When fusion

ImageNet-R (48 condition–seed pairs)	value
harmful base rate	13.9%
mean benefit $\bar{B}$	+0.017
harm detectability (best AUC)	0.662
KGA decisions	48/48 abstain

TABLE XVIII
IMAGENET-R: MANDATORY ABSTENTION (COROLLARY 2).

Office-Home (gradient-TTA)	val	test
Harmful base rate	0.17	0.12
Certificate transfer-AUC	0.33	0.49
KGA false-adapts / beats both	0 / no	0 / no

TABLE XIX
OFFICE-HOME: ZERO FALSE-ADAPTS; NO BEATS-BOTH (INVALID TRANSFER EVIDENCE).

flips from helpful to harmful, KGA routes correctly: mean accuracy 0.857 vs. always-fuse 0.583 and always-single-A 0.854 (both gaps $P=1.0$), false-adapt 0/9 commits ($\leq \alpha$). This is a *controlled* mechanism confirmation—not a natural benchmark—but shows the certificate dominates both trivial policies when detectable multimodal harm is present (CONTROLLED_MULTIMODAL_PROTOCOL_D33_v1).

K. CIFAR-10.1 (low-margin natural shift, Protocol K audit)

Over five seeds, KGA abstains heavily on Tent/EATA (59–86%), makes zero false adapts there under *within-seed* calibration, and ties always-freeze while always-adapt overpays; SAR is mostly helpful (5% harmful). Cross-seed transfer of the certificate (Protocol K) does not meet the headline bar; see §E.

L. ImageNet-C diagnostics

An earlier ImageNet-C SAR port (31% harmful) motivated the mechanism-faithful re-run (Table V in the main text). Decision baselines and α /evidence ablations appear in the frontier appendix (App. M).

All experiments are reproducible from the anonymized submission repository: theory validators, pre-registered Protocols A/G/H v2/M v2, and locked result JSONs (experiments/kbound/results/, research_lock/KBOUND_HEADLINE_FINDINGS_v3.json). Environment: Python 3.12 and PyTorch 2.5.1 for GPU runs; kbtrain.sh wraps venv activation and natural-shift protocol scorers. Core certificate code: `pip install -e docs/research/kbound/kbound_pkg/`.

M. Real-camera deployment supplement (pre-registered templates)

Tables XXII–XXVII define the physical-shift breakdown, recording inventory, anti-leakage audit, per-condition release, runtime profile, and ablations. Measured cells remain blank until populated directly from locked real-camera artifacts; the synthetic dashboard smoke test is not admissible evidence for these tables.

iWildCam	val (multicand)	test
Harmful base rate	59%	83%
Best harm-AUC	≈ 0.62	0.69
Multicandidate beats both?	no	no
Protocol H v2 canonical KGA (tent_episodic)	—	yes (0% FA)

TABLE XX
IWILDCAM: **diagnostic** MULTICANDIDATE ROUTE ON VAL/TEST VS **headline** PROTOCOL H V2 CANONICAL KGA ON TEST (TABLE VIII IN MAIN TEXT).

Track	Repro command / artifact	Status
CIFAR-10-C stress (5 seeds)	kbtrain.sh stress targets	locked JSON
POEM/AETTA head-to-head	poem_aetta/run_all_headtohead.sh	WIN
Mixed OOF aggregate	mixed_stream_kbound.py	locked
Assumption audit v1	run_assumption_audit_v1.py	cached arm
ImageNet-C SAR faithful	result JSON + Table V	locked
Camelyon17 Protocol G	analyze_F on protocol G bundle	locked
iWildCam Protocol H v2	kbtrain.sh protocol-h-v2	locked
Office-Home Protocol M v2	kbtrain.sh protocol-m-v2	locked
Theory validators	theory_validation/ suite	pass

TABLE XXI
REPRODUCIBILITY STATUS FOR HEADLINE TABLES (JUNE 2026 LOCK).

Every theorem has a numerical sanity-check in the theory-validation suite; multi-seed experiments report pooled means with paired bootstrap where pre-registered.

The benefit-sign frontier (Thm. 4) predicts that on the low-margin band any *committal* label-free rule must commit sign errors, whereas a valid certificate must *abstain* there. This appendix demonstrates that prediction on the two *synthetic-corruption* streams whose operating point the certificate is calibrated on (CIFAR-10-C, ImageNet-C); it makes *no* external-validation claim. The honest natural-shift (Camelyon17) finding—a limitation, not a win—closes the appendix (App. P).

N. CIFAR-10-C: K-Bound beats both trivial policies at zero false-adapt

On the pre-registered CIFAR-10-C stress grid (432 conditions/candidate), KGA adapts the helpful conditions and freezes the harmful ones with a *zero* false-adapt rate, so it strictly beats both always-adapt and always-freeze for the two benign candidates (Tent, EATA) and ties the collapse-resistant SAR (Table XXVIII).




Fig. 10. **CIFAR-10-C frontier (regret-to-oracle, log scale).** KGA is below *both* always-adapt and always-freeze for Tent and EATA at 0% false-adapt, and ties always-adapt for SAR (Table XXVIII).

TABLE XXII

R3 – Held-out physical-shift breakdown (results pending). THIS TABLE IS PROMOTED TO THE MAIN PAPER ONLY IF THE LOCKED RESULT SUPPORTS A CLEAR, NON-CHERRY-PICKED REGIME ANALYSIS.

Held-out shift family	Windows	Always-adapt regret	KGA regret	KGA decision pattern	Interpretation
Mild lighting shift	—	—	—	—	—
Strong side shadow	—	—	—	—	—
Blur / vibration	—	—	—	—	—
Background / viewpoint shift	—	—	—	—	—
Batch-composition shift	—	—	—	—	—

TABLE XXIII

S1 – Dataset and recording inventory (results pending). NO FRAME-LEVEL RANDOM SPLIT IS PERMITTED ACROSS RECORDING SESSIONS.

Split	Recording day/session	Object IDs	Conditions	Windows	Frames	Label distribution
Source train	—	—	Stable bright conditions	—	—	—
Source validation	—	—	Stable source conditions	—	—	—
Calibration-fit	—	—	Moderate physical shifts	—	—	—
Calibration-conformal	—	—	Disjoint moderate shifts	—	—	—
Held-out test	—	—	Hard physical shifts	—	—	—
External-device replication	—	—	Second phone; hard shifts	—	—	—

TABLE XXIV

S2 – Calibration and anti-leakage audit (results pending).

Check	Audited result
Frozen checkpoint unchanged after candidate adaptation?	—
Held-out labels inaccessible to the live runtime?	—
Feature schema unchanged after protocol lock?	—
ϵ calibrated only from the conformal split?	—
Held-out sessions excluded from adapter, feature, and threshold tuning?	—
Identical held-out stream replayed for every policy?	—
Config hash stored in every log row?	—
Model hash stored in every log row?	—

TABLE XXV

S3 – Per-condition held-out audit trail (results pending). THE RELEASED CSV/JSON CONTAINS ONE ROW FOR EVERY LOCKED HELD-OUT CONDITION.

Condition ID	Shift	True Δ	Oracle action	Freeze result	Adapt result	KGA action	Correct decision?	Regret
H01	Dim light	—	—	—	—	—	—	—
H02	Strong shadow	—	—	—	—	—	—	—
H03	Lens blur	—	—	—	—	—	—	—
H04	Background shift	—	—	—	—	—	—	—

TABLE XXVI

S4 – Resource and live-runtime profile (results pending).

Component	Mean (ms)	p95 (ms)	Memory (MB)	Notes
Frozen inference	—	—	—	Official prediction
Candidate Tent update	—	—	—	One step per window
Candidate inference	—	—	—	Shadow candidate
Evidence extraction	—	—	—	14 label-free features
Benefit estimator and KGA gate	—	—	—	Certificate decision
Full decision window	—	—	—	End-to-end
Video capture / preprocessing	—	—	—	Phone-to-Mac stream

TABLE XXVII
S5 – Locked physical evidence and conformal gate ablations (results pending).

Ablation variant	Regret ↓	Uncond. false adapt (FA_u) ↓	Adapt rate	Abstain rate	Conformal threshold ε
Full KGA	—	—	—	—	—
No radius ($\varepsilon = 0$)	—	—	—	—	—
No blur/brightness features	—	—	—	—	—
No disagreement features	—	—	—	—	—
Confidence features only	—	—	—	—	—
Entropy features only	—	—	—	—	—

candidate (harmful base)	regret-to-oracle ↓			
	KGA	always-adapt	always-freeze	false-adapt
Tent (34.5%)	0.0016	0.0086	0.1232	0.00
EATA (27.1%)	0.0015	0.0037	0.1311	0.00
SAR (15.7%)	0.0018	0.0018	0.1372	0.00

TABLE XXVIII

CIFAR-10-C stress grid (432 conditions/candidate, $\alpha=0.1$).
KGA BEATS BOTH FOR TENT AND EATA AT A 0% FALSE-ADAPT RATE, AND TIES ALWAYS-ADAPT FOR SAR.

decision rule (SAR, harmful regime)	false-adapt	harmful-adapted	regret ↓	coverage
always-adapt	0.44	1.00	0.1124	1.00
always-freeze	—	0.00	0.0267	1.00
ATC-style confidence (plug-in)	1.00	0.25	0.0673	1.00
ATC-style confidence (LOO)	1.00	0.06	0.0375	1.00
K-Bound (KGA)	0.00	0.00	0.0229	0.39

TABLE XXIX

Committal label-free heuristics false-adapt exactly where KGA abstains (ImageNet-C, REFERENCE-PARITY SAR, 36 CELLS, $\alpha=0.1$). ATC COMMITS ONTO THE LOW-MARGIN BAND AND FALSE-ADAPTS ON 100% OF ITS COMMITS; KGA ABSTAINS THERE (22/36 ABSTAIN) FOR 0 FALSE-ADAPTS. ON THE BENIGN CANDIDATES KGA ALSO KEEPS 0 FALSE-ADAPTS AT HIGHER COVERAGE (TENT: REGRET 0.0060, COVERAGE 0.69; EATA: REGRET 0.0047, COVERAGE 0.75).

O. ImageNet-C/SAR: the frontier on real corruptions—committal rules fail, K-Bound does not

On the logged ImageNet-C conditions (36 cells/candidate; SAR is harmful on 44.4% of cells), every *committal* label-free heuristic must place its commits somewhere on the low-margin band and therefore false-adapts: a plug-in ATC [13] confidence rule adapts into harmful cells on 100% of its commits (false-adapt 1.00, harmful-cells-adapted 0.25), and even its leave-one-out-tuned variant keeps false-adapt = 1.00. Among the rules compared, KGA reaches zero false-adapt and zero harmful-cells-adapted by *abstaining* on the band (coverage 0.39), consistent with the frontier prediction (Table XXIX).

P. Historical note (Camelyon17, frozen synthetic threshold)

Under a **frozen, Camelyon-free** multicandidate threshold $\tau^*=0.52$ (calibrated only on synthetic grids), the certificate does *not* beat both trivial policies on an early Camelyon17 composition grid (regret KGA 0.076 vs always-freeze 0.071). That failure motivated in-domain calibration and rich evidence; the headline result is Protocol G (main




Fig. 11. **ImageNet-C/SAR frontier.** Committal label-free rules false-adapt on the harmful regime; KGA reaches zero false-adapt by abstaining the low-margin band (coverage 0.39; Table XXIX).

text, Table VII), which uses dev-hospital calibration and does not transplant τ^* from synthetic runs.

The weakest one-bit class (resolving the open piece of Conjecture 1). Theorem 5 fixes the supplement at *one bit*; what remained open is the *weakest falsifiable class* on which that bit suffices. We settle it conditionally. On D write $a(x) := 2\eta_a(x) - 1$ with $\eta_a(x) = \mathbb{P}_T(f_a=Y | x)$, so $\Delta = \int_D a d\mu$ (Lemma 2). Let $m(x) := 2s(x) - 1$ be the observable relative (calibrated-score) margin, so $M = \frac{1}{2}\mathbb{E}_{\mu|D}[m]$, and let the observable *relative-calibration field* be $c(x) := w(x)(m(x) - m^*)$ with observable nonnegative weight $w(x)$ and observable neutral anchor m^* (the margin at which $s = \frac{1}{2}$). Put $T_E := \int_D c d\mu$ (observable), and let E denote the law of all label-free observables; two targets are *evidence-identical* when they share E (equivalently (μ, m, s) , since predictions and s are fixed maps of x).

Definition 2 (Margin-monotone relative-calibration class). $\mathcal{C}_{\text{mono}}$ is the class of targets for which $a(x) = \sigma c(x)$ μ -a.e. on D for a single global orientation $\sigma \in \{\pm 1\}$; equivalently the benefit is single-crossing in the observable margin at m^* ($\text{sign } a(x) = \sigma \text{sign}(m(x) - m^*)$) with calibrated magnitude $|a(x)| = |c(x)|$. The class is *falsifiable but not verifiable*: a finite labelled probe can reject “ $a = \pm c$ ”, yet σ is never label-free identified.

Assumption 1 (General position). Within an evidence fibre the regions of D whose benefit sign is free carry comparable, not-fully- E -pinned benefit mass: $T_E \neq 0$ and no proper subcollection of free regions has an E -certifiable dominant total $\int |a| d\mu$.

Lemma 4 (Structural sign-freedom = bit complexity). *Fix E and a falsifiable class $\mathcal{C}$, and let $K(\mathcal{C}, E)$ be the number*

of disjoint μ -positive regions of D whose benefit signs are jointly free in $\mathcal{C}$ on the fibre E . Under Assumption 1, a falsifiable structural supplement of b bits pins sign Δ across the fibre iff $b \geq K(\mathcal{C}, E)$.

Proof. If $b \geq K$, declaring the orientation of each free region fixes the sign field; the rest of a being E -pinned, $\Delta = \int_D a d\mu$ then has E -computable sign. If $b < K$, some free region R is undeclared; the two targets agreeing off R and differing in sign a on R are evidence-identical and share all b declared bits, while by comparability of $\int_R |a|$ with the complement (Assumption 1) their benefits have opposite sign—no decoder of the b declared bits separates them. $\square$

Theorem 7 (Conditional resolution of Conjecture 1: a weakest one-bit class). (i) Sufficiency (*unconditional*): every target in $\mathcal{C}_{\text{mono}}$ has sign $\Delta = \sigma \text{sign}(T_E)$ with T_E observable, so E and the single orientation bit σ determine sign Δ . (ii) Necessity (*unconditional*): σ and $-\sigma$ are evidence-identical ($\text{Law}(E \mid \sigma) = \text{Law}(E \mid -\sigma)$) yet give Δ and $-\Delta$, so by *Le Cam* at least one bit is required. Thus exactly one bit certifies sign Δ on $\mathcal{C}_{\text{mono}}$, with no further assumption. (iii) Minimality (*under Assumption 1*): every falsifiable $\mathcal{C} \supsetneq \mathcal{C}_{\text{mono}}$ contains a non-margin-monotone target with $K(\mathcal{C}, E) \geq 2$, so by Lemma 4 at least two bits are required. Hence $\mathcal{C}_{\text{mono}}$ is a weakest falsifiable one-bit class—minimal in the single-crossing, calibrated-magnitude direction; it is not unique (Remark 1).

Proof. (i) $\Delta = \int_D a d\mu = \sigma \int_D c d\mu = \sigma T_E$. (ii) Sending $\sigma \mapsto -\sigma$ is the D -local label complement of Theorem 5: it fixes (μ, m, s) , hence E , while negating a and thus Δ ; the two-point bound of Theorem 2 gives minimax committal error $\frac{1}{2}$. (iii) A target outside $\mathcal{C}_{\text{mono}}$ violates sign $a = \sigma \text{sign}(m - m^*)$ on a μ -positive region, so its sign field is not one global orientation; in the swap-closed (falsifiable) fibre at least two regions are independently signable, i.e. $K \geq 2$, and Lemma 4 applies. $\square$

Proposition 2 (Explicit non-identifiability witness). For each candidate single structural bit there are two targets with identical evidence law ($\text{TV} = 0$) sharing that bit's value but with opposite sign Δ . Take $D = R_1 \cup R_2$, $\mu(R_1) = \mu(R_2) = \frac{1}{2}$, with R_1 calibrated ($c_1=1$) and R_2 at the anchor ($c_2=0$, no observable benefit signal). For $\beta = \text{sign } a_1$: $a = (+1, -2)$ vs. $(+1, +2)$ give $\Delta = -\frac{1}{2}$ vs. $+\frac{3}{2}$. For $\beta = \text{sign } a_2$: $a = (+1, +\frac{1}{5})$ vs. $(-1, +\frac{1}{5})$ give $\Delta = +\frac{3}{5}$ vs. $-\frac{2}{5}$. No single bit certifies sign Δ , so two bits are necessary off $\mathcal{C}_{\text{mono}}$.

Remark 1 (What is unconditional, and what rests on general position). Parts (i)–(ii)—that one declared bit certifies sign Δ on $\mathcal{C}_{\text{mono}}$ —hold *unconditionally*. Only minimality (iii) uses Assumption 1: if it fails (one region's benefit mass is E -certifiably dominant) that region's orientation alone certifies sign Δ even outside $\mathcal{C}_{\text{mono}}$, moving the boundary. Equivalently, Assumption 1 is the requirement that the certificate be *domination-free* (minimax-sound over the

unobserved magnitudes); so the General-Position statement is only *conditional*; the unconditional weakest classes are characterized in Theorem 8 as an explicit finite family of dominance polytopes. Even under Assumption 1, $\mathcal{C}_{\text{mono}}$ is only a weakest one-bit class: any calibrated sign-aligned field $\{a = \sigma h : h \text{ observable}\}$ is equally one-bit and is incomparable to $\mathcal{C}_{\text{mono}}$. All six checks are machine-checked by a deterministic release validator: the unconditional core (sufficiency/necessity 20000/20000, $|\Delta - \sigma T_E| = 0$ exactly), the explicit minimality witnesses of (iii), the general-position ablation, and the non-uniqueness check.

Why general position cannot simply be dropped (the precise obstruction, and exactly what is open). The minimal counterexample is two equal-mass regions $D = R_0 \cup R_1$, $\mu(R_0) = \mu(R_1) = \frac{1}{2}$, with the observable field $c = +1$ on the calibrated region R_0 and $c = 0$ on the anchor region R_1 (no observable benefit signal there, so R_1 is sign-free in any evidence fibre). Enlarge $\mathcal{C}_{\text{mono}}$ to the falsifiable class $\mathcal{C}_{\text{dom}} = \{a_{R_0} = \sigma, |a_{R_1}| \leq \rho\}$ for a declared constant $\rho < 1$. Since $\Delta = \frac{1}{2}(a_{R_0} + a_{R_1})$ and $|a_{R_1}| \leq \rho < 1 = |a_{R_0}|$, the calibrated region dominates and sign $\Delta = \sigma$ for every member, so one declared bit certifies sign Δ on all of $\mathcal{C}_{\text{dom}}$. Yet $\mathcal{C}_{\text{dom}} \supsetneq \mathcal{C}_{\text{mono}}$ (it contains non-margin-monotone targets with $a_{R_1} \neq 0$, e.g. $a = (+1, \pm\frac{1}{2})$ giving $\Delta = \frac{3}{4}, \frac{1}{4}$, both positive) and is falsifiable (a labelled probe rejects targets that are uncalibrated on R_0 or violate $|a_{R_1}| \leq \rho$). This contradicts minimality once an E -certifiable dominant region is permitted: dropping $a = (+1, \pm 2) \Rightarrow \Delta = +\frac{3}{2}, -\frac{1}{2}$ shows that the moment $\rho \geq 1$ the anchor region can flip sign Δ and two bits become necessary ($K=2$). The obstruction to an *unconditional* weakest class is therefore exactly this: characterising minimality across the *entire* lattice of dominant-region thresholds ρ (where the weakest one-bit class is not $\mathcal{C}_{\text{mono}}$ but a ρ -dependent family), which Assumption 1 sidesteps by fiat. Theorem 8 (below) resolves it: the ρ -lattice is the slice $\{r_0=1, r_1=\rho\}$ of the single dominance polytope $W^* = \{r_0 \geq r_1\}$, with $\rho=1$ its boundary. The construction is machine-checked (`val_conj1_genpos.py`, seeds fixed): $\mathcal{C}_{\text{dom}}$ strictly contains $\mathcal{C}_{\text{mono}}$, is falsifiable (reject rate 1.0), is one-bit-certified (error rate 0 over 2×10^5 targets), and loses one-bit certifiability without the domination bound (error rate $\approx \frac{1}{4}$).

Unconditional resolution (General Position removed). Dropping Assumption 1 reveals, for each orientation pattern, a single maximal one-bit class—a *dominance polytope*—of which the ρ -family of Remark 1 is one slice.

Definition 3 (Orientation pattern and canonical class). On a fibre E an *orientation pattern* is a tied set $P \subseteq \{i : c_i \neq 0\}$ with a tied sign-pattern $s \in \{\pm\}^P$; the free set is $Z = D \setminus P$. For an observable magnitude region $W \subseteq [0, 1]^{|D|}$ the *canonical class* is $\mathcal{C}(P, s, W) = \{a : \text{sign } a_i = \sigma s_i \text{sign } c_i \ (i \in P), (|a_j|)_j \in W\}$ with one declared bit σ . Write $T(r) = \sum_{i \in P} s_i \text{sign}(c_i) \mu_i r_i$ and $G(r) = \sum_{i \in Z} \mu_i r_i$.

Lemma 5 (Canonical form of a falsifiable class). *A class is falsifiable iff it equals some $\mathcal{C}(P, s, W)$; anchors ($c_i = 0$) lie necessarily in Z . (Swap-closure, Thm. 5(iii)), forces the constraint on Z to depend only on $|a|$; forcing an anchor sign is not evidence-definable. Both aligned ($s_i = +$) and anti-aligned ($s_i = -$) ties are falsifiable since σ is unidentified.)*

Theorem 8 (Unconditional one-bit criterion; weakest classes). *Let $\mathcal{C}(P, s, W)$ be falsifiable. (i) (Dominance margin.) One declared bit certifies sign Δ on $\mathcal{C}(P, s, W)$ iff $\inf_{r \in W} (T(r) - G(r)) \geq 0$ (with some member $\Delta > 0$) or $\sup_{r \in W} (T(r) + G(r)) \leq 0$. (ii) (Weakest class.) For a fixed pattern (P, s) the unique maximal one-bit class is the dominance polytope $\mathcal{C}^* = \mathcal{C}(P, s, W^*)$, $W^* = \{r \in [0, 1]^{|D|} : T(r) \geq G(r)\}$; $\mathcal{C}_{\text{mono}}$, every $\mathcal{C}_{\text{dom}}(\rho \leq 1)$, and every box one-bit class are proper subclasses of one $\mathcal{C}^*$. (iii) (Finite family.) The set of weakest falsifiable one-bit classes is the explicit finite family $\{\mathcal{C}^*\}$ indexed by $(P \subseteq \{c \neq 0\}, s \in \{\pm\}^P)$ modulo the global flip $\sigma \mapsto -\sigma$ (at most $3^{\lfloor c \neq 0 \rfloor}$ patterns); distinct patterns give pairwise incomparable maxima, so there is no unique weakest class. $\mathcal{C}_{\text{mono}}$ is the all-tied all-aligned member, and Assumption 1 is exactly the face on which the family collapses to $\mathcal{C}_{\text{mono}}$, recovering Thm. 7(iii).*

Proof. Δ is linear, so over $\mathcal{C}(P, s, W)$ at bit $\sigma = +1$ it ranges over $[\inf_W (T - G), \sup_W (T + G)]$ (free signs independent by swap-closure; the $\sigma = -1$ fibre is its negation); a single bit certifies iff each fibre’s strict sign set is a singleton, giving (i). For W^* , $\inf_{W^*} (T - G) \geq 0$ by construction, and any r_0 with $T(r_0) < G(r_0)$ admits, with opposing free signs, a member with $\Delta < 0$ while W^* admits $\Delta > 0$, so W^* is maximal; $\mathcal{C}_{\text{dom}}(\rho)$ sits in W^* iff $\rho \leq 1$, with $\rho = 1$ the boundary, giving (ii). Incomparability (iii) is witnessed on $c = (1, 1, 0)$ by $a = (0, \frac{1}{3}, 0)$ and $a = (1, -\frac{1}{3}, \frac{1}{3})$. On the general-position face the free block can overpower the tied block unless Z -magnitudes vanish, leaving only $\mathcal{C}_{\text{mono}}$. $\square$

Verification. The criterion (i) was checked against exact brute-force vertex truth on 2.8×10^5 box fibres and 3.2×10^3 non-box integer polytopes (exact-LP inf/sup) with 0 mismatches; sufficiency on W^* over 1.2×10^5 members (0 errors), necessity just outside W^* (error rate $\approx \frac{1}{2}$), and the $\mathcal{C}_{\text{dom}}(\rho)$ collapse ($\rho \leq 1$ one-bit, $\rho > 1$ lost) are machine-checked (`val_unconditional_weakest.py`, seeds fixed; fails loudly). The single non-tautological input is Lemma 5, which rests on the swap involution (Thm. 5(iii)).

REFERENCES

- [1] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, T. Darrell. Tent: Fully test-time adaptation by entropy minimization. *ICLR*, 2021.
- [2] J. Liang, D. Hu, J. Feng. Do we really need to access the source data? Source hypothesis transfer for unsupervised domain adaptation. *ICML*, 2020.
- [3] Y. Sun, X. Wang, Z. Liu, J. Miller, A. Efros, M. Hardt. Test-time training with self-supervision for generalization under distribution shifts. *ICML*, 2020.
- [4] S. Niu, J. Wu, Y. Zhang, Y. Chen, S. Zheng, P. Zhao, M. Tan. Efficient test-time model adaptation without forgetting. *ICML*, 2022.
- [5] S. Niu, J. Wu, Y. Zhang, Z. Wen, Y. Chen, P. Zhao, M. Tan. Towards stable test-time adaptation in dynamic wild world. *ICLR*, 2023.
- [6] Q. Wang, O. Fink, L. Van Gool, D. Dai. Continual test-time domain adaptation. *CVPR*, 2022.
- [7] M. Zhang, S. Levine, C. Finn. MEMO: Test time robustness via adaptation and augmentation. *NeurIPS*, 2022.
- [8] S. Jeon, Z. Mai, H. Tian, V. Bakshi, L. Wang, J. Hou, P. Zhang, A. Chowdhury, W.-L. Chao. Continually adapt or not (CAN)? A continual learning benchmark of camera trap species classification over time. *NeurIPS Workshop (Imageomics)*, 2025. OpenReview: [jEItIzAgHs](https://openreview.net/forum?id=jEItIzAgHs).
- [9] M. Schirmer, M. Jazbec, C. A. Naesseth, E. Nalisnick. Monitoring risks in test-time adaptation. *NeurIPS*, 2025. arXiv:2507.08721.
- [10] A. Podkopaev, A. Ramdas. Tracking the risk of a deployed model and detecting harmful distribution shifts. *ICLR*, 2022. arXiv:2110.06177.
- [11] Y. Bar, S. Shaer, Y. Romano. Protected test-time adaptation via online entropy matching: A betting approach. *NeurIPS*, 2024. arXiv:2408.07511.
- [12] A. Sreeram, M. N. Huh, J. K. Wong, et al. Tempora: Characterising the time-contingent utility of online test-time adaptation. arXiv:2602.06136, 2026.
- [13] S. Garg, S. Balakrishnan, Z. Lipton, B. Neyshabur, H. Sedghi. Leveraging unlabeled data to predict out-of-distribution performance. *ICLR*, 2022. arXiv:2201.04234.
- [14] A. T. Kalai, V. Kanade. Towards optimally abstaining from prediction with OOD test examples. *NeurIPS*, 2021. arXiv:2105.14119.
- [15] J. Miller, R. Taori, A. Raghunathan, S. Sagawa, P. W. Koh, V. Shankar, P. Liang, Y. Carmon, L. Schmidt. Accuracy on the line: On the strong correlation between out-of-distribution and in-distribution generalization. *ICML*, 2021. arXiv:2107.04649.
- [16] C. Baek, Y. Jiang, A. Raghunathan, J. Z. Kolter. Agreement-on-the-line: Predicting the performance of neural networks under distribution shift. *NeurIPS*, 2022. arXiv:2206.13089.
- [17] E. Kim, M. Sun, C. Baek, A. Raghunathan, J. Z. Kolter. Test-time adaptation induces stronger accuracy and agreement-on-the-line. *NeurIPS*, 2024. arXiv:2310.04941.
- [18] T. Lee, S. Chottananurak, T. Gong, S.-J. Lee. AETTA: Label-free accuracy estimation for test-time adaptation. *CVPR*, 2024. arXiv:2404.01351.
- [19] Y. Jiang, V. Nagarajan, C. Baek, J. Z. Kolter. Assessing generalization of SGD via disagreement. *ICLR*, 2022.
- [20] Y. Lu, Z. Wang, R. Zhai, S. Kolouri, J. Campbell, K. Sycara. Characterizing out-of-distribution error via optimal transport. *NeurIPS*, 2023.
- [21] J. Steinhardt, P. Liang. Unsupervised risk estimation using only conditional independence structure. *NeurIPS*, 2016. arXiv:1606.05313.
- [22] S. Ben-David, J. Blitzer, K. Crammer, A. Kulesza, F. Pereira, J. Wortman Vaughan. A theory of learning from different domains. *Machine Learning* 79(1-2):151–175, 2010.
- [23] S. Ben-David, T. Lu, T. Luu, D. Pál. Impossibility theorems for domain adaptation. *AISTATS*, 2010.
- [24] Z. Lipton, Y.-X. Wang, A. Smola. Detecting and correcting for label shift with black box predictors. *ICML*, 2018.
- [25] E. Rosenfeld, S. Garg. (Almost) provable error bounds under distribution shift via disagreement discrepancy. *NeurIPS*, 2023. arXiv:2306.00312.
- [26] F. Parisi, F. Strino, B. Nadler, Y. Kluger. Ranking and combining multiple predictors without labeled data. *PNAS* 111(4):1253–1258, 2014. arXiv:1303.3257.
- [27] A. Jaffe, B. Nadler, Y. Kluger. Estimating the accuracies of multiple classifiers without labeled data. *AISTATS*, 2015. arXiv:1407.7644.

- [28] A. Jaffe, E. Fetaya, B. Nadler, T. Jiang, Y. Kluger. Unsupervised ensemble learning with dependent classifiers. *AISTATS*, 2016. arXiv:1510.05830.
- [29] E. A. Platanios, A. Blum, T. Mitchell. Estimating accuracy from unlabeled data. *UAI*, 2014.
- [30] S. Ibrahim, X. Fu, N. Sidiropoulos. Crowdsourcing via pairwise co-occurrences: Identifiability and algorithms. *NeurIPS*, 2019. arXiv:1909.12325.
- [31] A. Corrada-Emmanuel. The logic of NTQR evaluations of noisy AI agents: Complete postulates and logically consistent error correlations. *NeurIPS*, 2024. arXiv:2312.05392; see also arXiv:2409.11052, arXiv:2412.16238.
- [32] A. P. Dawid, A. M. Skene. Maximum likelihood estimation of observer error-rates using the EM algorithm. *JRSS-C* 28(1):20–28, 1979.
- [33] L. Le Cam. *Asymptotic Methods in Statistical Decision Theory*. Springer, 1986.
- [34] A. B. Tsybakov. *Introduction to Nonparametric Estimation*. Springer Series in Statistics, 2009.
- [35] C. Gupta, A. Ramdas. Top-label calibration and multiclass-to-binary reductions. *ICLR*, 2022 (arXiv:2107.08353).
- [36] C. Manski. Nonparametric bounds on treatment effects. *Amer. Econ. Rev. P&P* 80(2):319–323, 1990.
- [37] Y. Geifman, R. El-Yaniv. Selective classification for deep neural networks. *NeurIPS*, 2017.
- [38] V. Vovk, A. Gammerman, G. Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- [39] A. Angelopoulos, S. Bates. Conformal prediction: A gentle introduction. *Found. Trends Mach. Learn.*, 2023.
- [40] G. Shafer. Testing by betting: A strategy for statistical and scientific communication. *JRSS-A*, 2021.
- [41] A. Ramdas, P. Grünwald, V. Vovk, G. Shafer. Game-theoretic statistics and safe anytime-valid inference. *Statist. Sci.*, 2023.
- [42] V. Vovk, I. Petej, I. Nouretdinov, E. Ahlberg, L. Carlsson, A. Gammerman. Retrain or not retrain: Conformal test martingales for change-point detection. *COPA*, 2021. arXiv:2102.10439.
- [43] S. R. Howard, A. Ramdas, J. McAuliffe, J. Sekhon. Time-uniform, nonparametric, nonasymptotic confidence sequences. *Annals of Statistics* 49(2):1055–1080, 2021.
- [44] J. Ville. *Étude critique de la notion de collectif*. Gauthier-Villars, Paris, 1939.
- [45] I. Waudby-Smith, A. Ramdas. Estimating means of bounded random variables by betting. *J. R. Stat. Soc. B* 86(1):1–27, 2024.
- [46] V. Vovk, A. Gammerman, G. Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.
- [47] A. Maurer, M. Pontil. Empirical Bernstein bounds and sample variance penalization. *COLT*, 2009.
- [48] W. Hoeffding. Probability inequalities for sums of bounded random variables. *JASA* 58(301):13–30, 1963.
- [49] A. Klivans, K. Stavropoulos, A. Vasilyan. Testable learning with distribution shift. arXiv:2311.15142, 2023.
- [50] A. Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- [51] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, L. Fei-Fei. ImageNet: A large-scale hierarchical image database. *CVPR*, 2009.
- [52] D. Hendrycks, T. Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. *ICLR*, 2019.
- [53] D. Hendrycks, S. Basart, N. Mu, S. Kadavath, F. Wang, E. Dorundo, R. Desai, W. Zhu, R. Parajuli, M. Guo, D. Song, J. Steinhardt. The many faces of robustness: A critical analysis of out-of-distribution generalization. *ICCV*, 2021.
- [54] P. W. Koh, S. Sagawa, H. Marklund, S. M. Xie, M. Zhang, A. Balsubramani, W. Hu, M. Yasunaga, et al. WILDS: A benchmark of in-the-wild distribution shifts. *ICML*, 2021.
- [55] P. Bándi, O. Geessink, Q. Manson, M. van Dijk, M. Balkenhol, M. Hermesen, et al. From detection of individual metastases to classification of lymph node status at the patient level: The CAMELYON17 challenge. *IEEE Trans. Med. Imaging* 38(2):550–560, 2019.
- [56] J. Taylor, S. L. Earnshaw, B. M. Mabey, M. Victors, M. N. Annis, D. Y. Torigian, et al. The RXXR1 dataset and challenge: Unsupervised classification of cellular morphological changes. arXiv:1907.04758, 2019.
- [57] H. Venkateswara, J. Eusebio, S. Chakraborty, S. Panchanathan. Deep hashing network for unsupervised domain adaptation. *CVPR*, 2017. (Office-Home dataset release).
- [58] I. Gulrajani, D. Lopez-Paz. In search of lost domain generalization. *ICLR*, 2021.
- [59] F. Croce, M. Andriushchenko, V. Schwag, E. Debenedetti, N. Flammarion, M. Chiang, P. Mittal, M. Hein. RobustBench: A standardized adversarial robustness benchmark. *NeurIPS Datasets and Benchmarks*, 2021.
- [60] B. Recht, R. Roelofs, L. Schmidt, V. Shankar. Do CIFAR-10 classifiers generalize to CIFAR-10? *ICML*, 2019.
- [61] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, et al. PyTorch: An imperative style, high-performance deep learning library. *NeurIPS*, 2019.
- [62] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in Python. *JMLR* 12:2825–2830, 2011.
- [63] B. Schölkopf, D. Janzing, J. Peters, E. Sgouritsa, K. Zhang, J. Mooij. On causal and anticausal learning. *ICML*, 2012.
- [64] X. Wu, M. Gong, J. H. Manton, U. Aickelin, J. Zhu. On causality in domain adaptation and semi-supervised learning: an information-theoretic analysis for parametric models. *JMLR* 25(261):1–57, 2024.