

K-Bound: When Is Label-Free Adaptation Knowable? Helpful, Harmful, and Unknowable Regimes Under Distribution Shift

Pratik Niroula

Minnesota State University, Mankato




Fig. 1. **K-Bound previews the decision.** From an unlabeled batch under shift, the certificate estimates the adaptation benefit $\hat{\Delta}$ with a finite-sample radius ε and returns ADAPT (certified helpful), FREEZE (certified harmful), or ABSTAIN (no supported strict commitment). Under interval coverage it controls the marginal events $\mathbb{P}(\text{ADAPT}, \Delta \leq 0)$ and $\mathbb{P}(\text{FREEZE}, \Delta \geq 0)$ at level α ; no conditional-error guarantee is claimed.

Abstract—Test-time adaptation (TTA) may help or harm on unlabeled target data, so a deployment system must choose whether to *adapt*, *freeze*, or *abstain* from committing an update. We define adaptation benefit as $\Delta = R_T(f_0) - R_T(f_a)$, so positive values favor adaptation. Over a declared calibration-drift class, a strict adapt or freeze commitment is uniformly supportable if and only if the population evidence margin satisfies $|M| > \beta$; inside the closed band, matched evidence or zero-benefit ambiguity prevents a uniformly sound strict commitment. Knowability-Guided Adaptation (KGA) addresses the separate finite-sample layer: it estimates Δ directly from label-free evidence and commits only when a calibrated interval $\hat{\Delta} \pm \varepsilon$ excludes zero. On real benchmarks KGA does not numerically apply $|M| > \beta$, and ε is not an estimate of β . Controlled mixed regimes exercise all three actions with low observed harmful commitment and selective coverage; natural one-sided shifts mainly exercise safe adapt or freeze behavior, while weak evidence yields abstention or diagnostic failure. Lower regret than both fixed policies is reported only as a secondary routing-utility check. K-Bound is therefore a safety and validity layer, not a universal accuracy booster.

Index Terms—test-time adaptation, distribution shift, selective prediction, identifiability, label-free risk estimation

I. INTRODUCTION

Test-time adaptation (TTA) updates a model on the unlabeled test stream and, under favorable distribution shift, recovers accuracy. Under unfavorable shift, the same

objective silently degrades it, and with no target labels the system usually cannot tell which case it faces. The decision that matters therefore comes *before* the update: given a frozen baseline f_0 and an adapted candidate f_a , choose **adapt** when adaptation is knowably helpful, **freeze** when knowably harmful, or **abstain** when label-free evidence cannot identify the sign of the benefit. The field’s default assumption is that adaptation should be made robust enough to always run; we replace it with an explicit, certified decision about whether to run it at all.

Some cases are information-theoretically unknowable. It is established that label-free quantities cannot decide adaptation without an assumption on the shift: identifying OOD accuracy is as hard as identifying the optimal predictor [11], and the optimal predict-versus-abstain rule under a shift budget has a known form [12]. We sharpen this landscape into a strict-commitment criterion: a strict adapt or freeze action is uniformly sound over the declared drift class *if and only if* $|M| > \beta$ (Theorem 2). In the interior, evidence-identical worlds can have opposite nonzero benefits (Theorem 1); on the boundary, zero-versus-strict ambiguity still rules out a uniformly sound strict action (Proposition 2). KGA (Algorithm 1), a benefit estimator plus a conformal radius wrapped around any adapter (Tent, EATA, SAR), realizes the three-way decision as a finite-sample certificate—a direct benefit estimate with a calibrated radius, not a numerical $|M| > \beta$ test—controlling marginal false-adapt and false-freeze events under interval coverage (Theorem 3).

The evidence follows the regime map in Table I. *Identifiable mixed* harmful+helpful regimes test whether the controller adapts, freezes, and abstains in the corresponding conditions, including a pre-registered head-to-head over protocol-matched POEM-style and AETTA-style baselines. Natural shifts primarily test safety: a reconciled result may tie the better fixed policy, while a provisional result remains outside the headline until its saved summary and canonical replay agree. K-Bound does *not* claim universal improvement: when adaptation is already safe, a valid certificate can only tie always-adapt. It provides a safety layer for TTA rather than an accuracy boost.

a) *Plain-language takeaway.*: If the population evidence margin dominates declared drift ($|M| > \beta$), a strict

action is uniformly supportable; otherwise the maximal sound three-way rule ABSTAINS. KGA implements the distinct finite-sample interval decision. On single natural-shift datasets this mostly acts as safety or abstention. Improvement over both fixed policies requires a detectable mixed regime but is not the primary objective; a constructed pooled mixture is reported as such, not as a claim of single-dataset natural-shift dominance.

b) Contributions.:

- 1) **Adapt/freeze/abstain formulation.** We make the target risk difference, strict actions, fallback semantics, and target-label-free information boundary explicit.
- 2) **Exact strict-commitment frontier.** A strict adapt/freeze action is uniformly sound iff $|M| > \beta$; the interior and boundary cases are proved separately (Theorem 2).
- 3) **Coverage-based finite-sample certificate.** Under measurable interval coverage, the empirical rule controls marginal false-adapt and false-freeze events (Theorem 3).
- 4) **KGA deployable wrapper.** KGA wraps Tent/EATA/SAR without changing their adaptation objectives and uses shadow state, commit, rollback, and logging (Algorithm 1). It is adapter-interface agnostic but adapter-specific in calibration: the benefit model and radius are refit per dataset and per candidate adapter.
- 5) **Regime-based empirical evaluation.** A uniform panel evaluates action validity, false adaptation, coverage, natural no-harm behavior, constructed routing, and weak-evidence abstention; fixed-policy regret comparisons are secondary.

Table note. The uniform panel below is the sole empirical index; diagnostic variants and historical runs are not additional evidence for the headline claims.

II. RELATED WORK

The closest literature improves test-time updates, monitors adaptation failures, estimates target performance without labels, or characterizes abstention under shift. K-Bound differs in the decision object: it certifies the sign of the frozen-versus-adapted risk difference before the candidate state is committed.

A. Test-Time Adaptation

Tent [1], SHOT [2], TTT [3], EATA [4], SAR [5], CoTTA [6], and MEMO [7] improve adaptation through entropy minimization, filtering, regularization, sharpness-aware updates, or temporal stabilization (survey: [33]). These methods primarily answer *how* to adapt. Even when the update rule is more stable, the deployment system generally still commits to adaptation rather than certifying that the candidate has positive target benefit.

TABLE I

Regime map: expected behavior vs. result. EACH REGIME, AN EXAMPLE DATASET, THE BEHAVIOR THE FRONTIER PREDICTS FOR A VALID CERTIFICATE, AND WHAT WAS MEASURED. “NO-HARM” MEANS MATCHING THE BETTER FIXED POLICY WHILE AVOIDING THE WORSE; COMPARISONS WITH BOTH FIXED POLICIES ARE SECONDARY ROUTING-UTILITY DIAGNOSTICS. THE EVIDENCE TIER IN TABLE XV STATES WHETHER THAT POINT VERDICT IS LOCKED, RECONCILED, PROVISIONAL, OR DIAGNOSTIC.

Regime	Example	Expected	Result
Helpful-dominated	Camelyon17 OOD	tie	always-locked
Harmful-dominated	RxRx1; iWildCam/Office-Home (OOE lock)	tie	always-locked
Domain-gen. natural	PACS (4 LODO splits)	no-harm or abstain	completed 3-seed null diagnostic
Mixed + detectable	CIFAR-10-C Tent/EATA, ImageNet-C SAR	adapt helpful, freeze harmful, abstain near zero	locked; low false-adapt; CI-supported utility
Weak evidence	ImageNet-R, CIFAR-10.1	abstain / conservative	completed diagnostics; no promoted win
Constructed mixture	three-source OOF stream	route per condition	locked routing result, not transfer

B. Guarded and Monitored Adaptation

POEM [10], sequential risk monitors [8], [9], PeTTA [40], and sample-level filters such as DeYO [41] guard an adaptation stream during or after updating. These approaches are complementary to K-Bound. They typically monitor a proxy or intervene in one failure direction, whereas K-Bound poses a pre-commitment three-way decision over the benefit sign and explicitly represents an unknowable region.

C. Label-Free Performance Estimation

ATC [11], agreement-on-the-line [14], [13], AETTA [15], and related confidence-style estimators [32], [31] estimate target accuracy or error without target labels. Estimating $R_T(f)$ is not identical to certifying $\text{sign}(R_T(f_0) - R_T(f_a))$: a deployment controller needs the sign of a paired risk difference and must account for uncertainty around zero. K-Bound uses this difference as the estimand and abstains when the available evidence does not identify its sign.

D. Impossibility, Abstention, and Risk Control

General label-free target-risk identification requires assumptions on the shift [11], [16], [17], [18]. Kalai and Kanade [12] characterize a related abstain-versus-predict decision under a shift budget, while selective prediction [19] abstains on individual labels. K-Bound specializes this landscape to adaptation benefit, localizes the comparison to the disagreement region, and derives an adapt/freeze/abstain certificate. Its empirical uncertainty layer connects to split

conformal prediction, conformal under shift, RCPS, and Learn-then-Test [34], [36], [35], [38], [39].

E. Surviving Gap

Existing work provides stronger adapters, reactive monitors, label-free accuracy estimates, and general abstention theory. The remaining gap is a pre-commitment rule for the sign of adaptation benefit that (i) distinguishes helpful, harmful, and unsupported decisions, (ii) states the assumptions under which the decision is valid, and (iii) controls marginal false adaptation under a calibrated uncertainty model. K-Bound addresses this gap with the population frontier and its finite-sample realization, KGA.

III. PROBLEM FORMULATION AND VALIDITY

A. Deployment Setting and Adaptation Benefit

Source data follow $P_S(X, Y)$ and may be labeled during training and development. Deployment data follow $P_T(X, Y)$, but only target inputs and model outputs are available when the decision is made. A frozen predictor f_0 serves as the safe fallback, and an adapter $\mathcal{A}$ proposes a temporary candidate $f_a = \mathcal{A}(f_0, X_{1:m})$ from an unlabeled target batch. For loss ℓ , the target risk and adaptation benefit are

$$R_T(f) = \mathbb{E}_{P_T}[\ell(f(X), Y)], \quad \Delta = R_T(f_0) - R_T(f_a). \quad (1)$$

Thus $\Delta > 0$ favors committing the candidate, $\Delta < 0$ favors retaining the frozen state, and $\Delta = 0$ makes the two actions risk-equivalent. Target labels reveal Δ only during offline auditing; they are not available to the deployment rule.

B. Target-Label-Free Evidence and Decision Semantics

The controller observes a label-free evidence vector $Z = \phi(X_{1:m}, f_0, f_a)$ containing quantities such as prediction entropy, confidence, class-balance statistics, frozen-versus-adapted disagreement, output-distribution drift, and update magnitudes. The deployment rule returns

$$g(Z) \in \{\text{ADAPT}, \text{FREEZE}, \text{ABSTAIN}\}.$$

ADAPT commits f_a ; FREEZE rejects the proposal and retains f_0 (or the latest certified state in continual operation); and ABSTAIN declines to change the model, serves the frozen prediction, and may request more evidence or human review. Abstention therefore applies to the update decision, not to prediction itself.

C. What “Label-Free” Means

K-Bound is target-label-free rather than supervision-free. Source labels and a held-out development set may be used to fit the benefit estimator and calibrate its uncertainty radius. No target-test label may influence the adapter configuration, evidence map, benefit estimator, radius, action threshold, or model selection. Target-test labels are opened only after all decisions are frozen, for reporting accuracy, regret, and false-adapt events.

TABLE II

DATA ACCESS BY STAGE. TARGET-TEST LABELS ARE FORBIDDEN UNTIL OFFLINE AUDIT.

Stage	Source labels	Dev labels	Target-test labels
Source training	yes	optional	no
Benefit/radius calibration	optional	yes	no
Deployment decision	no	frozen artifacts only	no
Offline evaluation	no	no	yes, after decisions

TABLE III

ASSUMPTIONS, THEIR ROLE, AND THE SAFE RESPONSE WHEN THEY ARE UNSUPPORTED.

Condition	Role	If unsupported
Risk-aligned Z	Benefit sign is recoverable over the declared class	abstain; do not interpret a point estimate as a certificate
Coverage of $\hat{\Delta} \pm \varepsilon$	Controls marginal false-adapt/freeze	recalibrate, widen the radius, or use the declared fallback
Exchangeable/correct residuals	Transfers calibration to deployment	use shift-aware calibration or abstain
Bounded loss/benefit	Supports concentration and finite-sample analysis	clip/redefine the loss or state no bound
Frozen protocol	Prevents target-test leakage and selection bias	rerun with a locked protocol

D. Assumptions and Validity Scope

The theory and the finite-sample certificate rely on distinct assumptions. *Risk alignment* determines whether the evidence can identify or bound the benefit sign over the declared target class. *Coverage* determines whether the empirical interval contains the true benefit at the stated rate. *Exchangeability* (or an explicit shift correction) connects calibration residuals to deployment residuals. The adapter, evidence schema, estimator, and calibration protocol must be fixed before held-out scoring. If these conditions cannot be justified, the theorem does not license either strict action; deployment must use a separately declared operational fallback, such as abstention or serving the frozen model. This is a safety policy, not a theorem that freezing is optimal. Risk alignment is an assumption on the deployment class, not a property certified by fitting a benefit regressor: support and residual diagnostics may reject implausible evidence maps, but they do not verify it universally.

E. Notation and Observability

Table IV lists the core quantities and marks which are observable at deployment.

F. Multiclass and Regression Scope

The binary 0–1 loss case gives the cleanest frontier because, on $D = \{x : f_0(x) \neq f_a(x)\}$, exactly one predictor is correct. For multiclass 0–1 loss,

$$\Delta = P_T(D)(p_a - p_0), \quad p_j = P_T(f_j(X) = Y \mid D),$$

so the decision becomes an ordinal comparison of the two disagreement-region accuracies. For squared-loss regression, the risk difference reduces to a disagreement-local moment.

TABLE IV
CORE QUANTITIES. THE POPULATION DRIFT BUDGET β AND
EMPIRICAL RADIUS ε ARE RELATED BUT NOT INTERCHANGEABLE.

Symbol	Meaning	Available at deployment?
f_0, f_a	frozen and candidate models	yes
Δ	true target adaptation benefit	no
D	disagreement region	yes
M	observable population margin	through label-free evidence
γ	realized calibration drift	no
β	declared bound on $ \gamma $	supplied pre-deployment
Z	label-free evidence vector	yes
$\hat{\Delta}$	estimated adaptation benefit	yes
ε	calibrated residual radius	fixed pre-deployment
α	marginal error level	chosen pre-deployment




Fig. 2. **Population strict-commitment frontier.** The population rule supports adapt or freeze outside the drift band and abstains inside it. At the boundary, zero-benefit ambiguity rules out a uniformly sound strict action; opposite nonzero signs are claimed only in the interior. KGA uses the separate finite-sample interval rule in Fig. 3.

The main theorem uses a split-observable decomposition $\text{sign}(\Delta) = \text{sign}(M + \gamma)$; the multiclass bridge is stated below and the regression derivation appears in Appendix D.

Proposition 1 (Multiclass benefit decomposition). *For multiclass 0/1 loss, $\Delta = \mu_T(D)(p_a - p_0)$ with $p_j = \mathbb{P}_T(f_j(X)=Y \mid D)$. If the disagreement-region accuracy gap admits a split-observable decomposition $p_a - p_0 = M_{\text{mc}} + \gamma_{\text{mc}}$ with $|\gamma_{\text{mc}}| \leq \beta_{\text{mc}}$, then $|M_{\text{mc}}| > \beta_{\text{mc}}$ implies $\text{sign } \Delta = \text{sign } M_{\text{mc}}$. The converse strict-commitment result additionally requires a richness assumption: the declared class must contain augmented-evidence-identical laws whose admissible γ_{mc} values straddle $-M_{\text{mc}}$, together with the zero-benefit boundary law. At $M_{\text{mc}} = \beta_{\text{mc}} = 0$, the sign is identified as zero.*

In the multiclass experiments KGA estimates Δ directly through the benefit proxy $\hat{\Delta} = h_\theta(Z)$ (not M_{mc} itself); the binary frontier is the identifiability backbone, and the multiclass results are evaluated through this benefit-estimation route.

G. Formal Setup

Source law $P_S(X, Y)$ supplies labels at train/validation; target law $P_T(X, Y)$ exposes only $P_T(X)$ at deployment. Fix loss ℓ , frozen predictor f_0 , candidate f_a , and target measure μ_T on inputs. Write target risk $R_T(f) = \mathbb{E}_{P_T}[\ell(f(X), Y)]$ and the adaptation benefit

$$\Delta := R_T(f_0) - R_T(f_a) \quad (\Delta > 0 \Leftrightarrow \text{adapting helps}).$$

The *disagreement region* $D := \{x : f_0(x) \neq f_a(x)\}$ is observable from (f_0, f_a) .

Assumption 1 (Deployment setup). Unless stated otherwise: binary label $Y \in \{0, 1\}$, 0/1 loss ℓ , and $\mu_T(D) > 0$. Fix a source-calibrated score $s : \mathcal{X} \rightarrow [0, 1]$ estimating the candidate’s *pointwise correctness* (e.g. a Platt-scaled confidence that f_a is right at x), and write the target pointwise correctness $\eta_a(x) := \mathbb{P}_T(f_a(X)=Y \mid X=x)$ on D , so that $\mathbb{E}_{\mu_T}[\eta_a \mid D] = \bar{a}$. Define the *observable margin*

$$M := \mathbb{E}_{\mu_T}[s(X) \mid D] - \frac{1}{2},$$

the *calibration drift* (unobserved at deploy)

$$\gamma := \mathbb{E}_{\mu_T}[\eta_a(X) - s(X) \mid D],$$

and label-free evidence $Z = \phi(X_{1:m}, f_0, f_a)$ for any measurable ϕ using only unlabeled batches and model outputs (entropy, confidence, drift, disagreement, update norms). A *label-free rule* is any measurable map $g : \mathcal{Z} \rightarrow \{\text{ADAPT}, \text{FREEZE}, \text{ABSTAIN}\}$ that does not use target labels on the deployment batch.

Definition 1 (Risk alignment). Evidence Z is *risk-aligned on a class $\mathcal{P}$* if no two laws $P_T^1, P_T^2 \in \mathcal{P}$ with $\text{sign } \Delta_1 \neq \text{sign } \Delta_2$ satisfy $\text{Law}(Z \mid P_T^1) = \text{Law}(Z \mid P_T^2)$.

All frontier claims stated in terms of M additionally require M to be determined on each evidence-law fibre; equivalently, take the population evidence to be $\tilde{Z} = (Z, M)$. Without this requirement, $|M| > \beta$ proves sign stability *given* M , not identifiability from $\text{Law}(Z)$ alone. Fitting a benefit regressor does not certify risk alignment: support, dependence, and residual diagnostics may falsify a proposed evidence map but cannot verify the assumption uniformly over the declared deployment class.

Definition 2 (Regimes). At evidence z : *knowably helpful* if the population frontier certifies $\Delta > 0$ (ADAPT); *knowably harmful* if it certifies $\Delta < 0$ (FREEZE); *unsupported for strict commitment* if the same augmented-evidence law is compatible with distinct signs, including zero versus a strict sign (ABSTAIN).

For a supplied drift budget $\beta \geq 0$, the *declared deployment class* is

$$\mathcal{C}_\beta := \{P_T : |\gamma(P_T)| \leq \beta\}.$$

Identifiability statements below are *over* $\mathcal{C}_\beta$, not over arbitrary shifts.

Definition 3 (Strict directional soundness and maximality). Fix an augmented-evidence law z and let $\mathcal{C}_\beta(z)$ be the laws in $\mathcal{C}_\beta$ that induce z . A strict action $a \in \{\text{ADAPT}, \text{FREEZE}\}$ is *uniformly directionally sound at z* when

$$\begin{aligned} a = \text{ADAPT} &\Rightarrow \Delta(P) > 0 \quad \text{for every } P \in \mathcal{C}_\beta(z), \\ a = \text{FREEZE} &\Rightarrow \Delta(P) < 0 \quad \text{for every } P \in \mathcal{C}_\beta(z). \end{aligned}$$

A sound three-way rule is *pointwise maximal* if it commits whenever some strict action is uniformly directionally sound and abstains otherwise. Zero benefit blocks a strict directional claim even though adapt and freeze have equal task risk when $\Delta = 0$; this is an epistemic validity convention, stronger than zero regret at the boundary.

Remark 1 (Four quantities). Do not confuse M (observable margin, deploy-time), γ (realized drift, latent), β (declared bound $|\gamma| \leq \beta$ on the deployment class), and ε (finite-sample radius from held-out residuals). The radius ε operationalizes coverage; it is *not* an estimate of β .

Symbol	Meaning	At deploy?	Role
M	observable evidence margin	yes	evidence toward adapt/freeze
γ	realized calibration drift	no	can reverse the benefit sign
β	declared bound on $ \gamma $	externally supplied	population frontier
ε	residual quantile radius	estimated pre-deploy	finite-sample certificate

What “label-free” means. No target/test labels at deploy; labeled source and held-out validation may calibrate $\hat{\Delta}$ and ε . Evidence Z uses only unlabeled batches, predictions, entropy, drift, and update statistics. The method is **target-label-free**, not label-free of all supervision: labels may be used only in the source/validation calibration split, never in the deployment batch or held-out test decision.

K-Bound in one sentence. A strict adapt/freeze commitment is uniformly supportable only when the population evidence margin exceeds the declared drift budget; otherwise abstention is the maximal sound three-way action.

IV. THEORY

A. Roadmap

The theoretical spine has three layers. First, the risk difference reduces to the region where the frozen and candidate predictions disagree. Second, matched label-free evidence can support incompatible benefit signs, which establishes an abstention requirement. Third, the decomposition $\text{sign}(\Delta) = \text{sign}(M + \gamma)$ yields the population frontier, while a calibrated interval around $\hat{\Delta}$ yields the finite-sample decision rule.

Lemma 1 (Disagreement-region reduction). *Under Assumption 1, with $\bar{a} := \mathbb{P}_T(f_a = Y \mid D)$,*

$$\Delta = 2\mu_T(D)(\bar{a} - \tfrac{1}{2}),$$

$$\text{sign } \Delta = \text{sign}(\bar{a} - \tfrac{1}{2}) = \text{sign}(M + \gamma).$$

The identity $\text{sign } \Delta = \text{sign}(M + \gamma)$ holds for every target law; it does not assume $|\gamma| \leq \beta$.

Proof. On D exactly one of f_0, f_a is correct under 0/1 loss, so the benefit equals $2\mu_T(D)(\bar{a} - \frac{1}{2})$. By the definitions of M and γ ,

$$\begin{aligned} M + \gamma &= \left(\mathbb{E}_{\mu_T|D}[s] - \tfrac{1}{2} \right) + \mathbb{E}_{\mu_T|D}[\eta_a - s] \\ &= \mathbb{E}_{\mu_T|D}[\eta_a] - \tfrac{1}{2} = \bar{a} - \tfrac{1}{2}, \end{aligned}$$

so $\text{sign } \Delta = \text{sign}(\bar{a} - \frac{1}{2}) = \text{sign}(M + \gamma)$. $\square$

Theorem 1 (Interior matched-evidence impossibility). *Fix $\beta > 0$. For any margin with $|M| < \beta$ there exist $P_T^1, P_T^2 \in \mathcal{C}_\beta$ such that $\text{Law}(\tilde{Z} \mid P_T^1) = \text{Law}(\tilde{Z} \mid P_T^2)$ for $\tilde{Z} = (Z, M)$, $M(P_T^1) = M(P_T^2) = M$, and $\text{sign } \Delta_1 = -\text{sign } \Delta_2 \neq 0$. (For $\beta = 0$ the class $\mathcal{C}_0$ forces $\gamma = 0$, so $\text{sign } \Delta = \text{sign } M$ is determined and no such pair exists: matched-evidence ambiguity is a strictly-positive-drift phenomenon.)*

Proof. Choose $0 < \delta < \min\{\beta - |M|, \frac{1}{2}\}$ and construct the two label laws explicitly. For $\theta \in \{+1, -1\}$ define P_T^θ on $\mathcal{X} \times \{0, 1\}$ by $P_T^\theta(dx, y) = \mu_T(dx) K^\theta(x, y)$ with label kernel

$$K^\theta(x, y) = \begin{cases} \frac{1}{2} + \theta\delta, & x \in D, y = f_a(x), \\ \frac{1}{2} - \theta\delta, & x \in D, y = f_0(x), \\ K^\circ(x, y), & x \notin D, \end{cases}$$

where K° is any fixed conditional label law used by *both* worlds (e.g. the source conditional). On D the binary labels satisfy $\{f_0(x), f_a(x)\} = \{0, 1\}$, so $K^\theta(x, \cdot)$ is a valid Bernoulli law ($\delta < \frac{1}{2}$ keeps both masses in $(0, 1)$); off D the two predictors agree, so those labels contribute identically to both risks and cancel in Δ . *Class membership:* $\eta_a^\theta(x) = \frac{1}{2} + \theta\delta$ on D , hence $\bar{a}_\theta = \frac{1}{2} + \theta\delta$ and $\gamma_\theta = \bar{a}_\theta - (M + \frac{1}{2}) = \theta\delta - M$, with $|\gamma_\theta| \leq \delta + |M| < \beta$; both worlds lie in $\mathcal{C}_\beta$. *Matched evidence:* the input marginal μ_T and (f_0, f_a, s) are unchanged, and $Z = \phi(X_{1:m}, f_0, f_a)$ uses no labels, so $\text{Law}(\tilde{Z} \mid P_T^+) = \text{Law}(\tilde{Z} \mid P_T^-)$ for $\tilde{Z} = (Z, M)$. *Opposite signs:* Lemma 1 gives $\Delta_\theta = 2\mu_T(D)\theta\delta \neq 0$ with $\text{sign } \Delta_\theta = \theta$. This completes the construction. An explicit Gaussian witness appears in Appendix C (Theorem 4). $\square$

Corollary 1 (Matched-evidence abstention lower bound). *Fix $\alpha \in (0, \frac{1}{2})$. On the shared evidence law constructed in Theorem 1, any possibly randomized label-free rule g satisfying $\mathbb{P}[g = \text{ADAPT}, \Delta \leq 0] \leq \alpha$ in every admissible world and $\mathbb{P}[g = \text{FREEZE}, \Delta \geq 0] \leq \alpha$ in every admissible world must satisfy $\mathbb{P}[g = \text{ABSTAIN}] \geq 1 - 2\alpha$.*

Proof. The rule has the same action probabilities in the two evidence-identical worlds. Its adapt probability is at most α in the negative world and its freeze probability is at most α in the positive world. The three action probabilities sum to one. $\square$

Proposition 2 (Boundary case and closed-band abstention). *Fix $\alpha \in (0, \frac{1}{2})$. If $|M| = \beta > 0$, the declared class contains an augmented-evidence-identical zero-benefit world ($\gamma = -M$) and a strict-benefit world of sign M ($\gamma = 0$). In the zero-benefit world both strict actions are errors under the directional events used in Corollary 1, so any rule controlling both marginal errors at level α abstains with probability at least $1 - 2\alpha$. If $M = \beta = 0$, every admissible world has $\Delta = 0$ and the same conclusion holds directly. Thus, under Definition 3, abstention is the pointwise maximal sound three-way action throughout the closed band $|M| \leq \beta$,*

although opposite nonzero signs are asserted only in the interior. This directional convention is stricter than task regret: at $\Delta = 0$, adapt and freeze are risk-equivalent but neither certifies a strict sign.

Theorem 2 (Exact strict-commitment frontier). *Under Assumption 1, fix the observables $(\mu_T, P_S, f_0, f_a, s)$ so M is determined, and let $\beta \geq 0$. The necessity and maximality clauses use the full class $\mathcal{C}_\beta$ and therefore its richness: it contains the evidence-identical label kernels constructed in Theorem 1 and the zero-benefit boundary law. For a narrower declared subclass not closed under these constructions, only the sufficiency clause follows without an additional richness assumption.*

- (i) **Reduction.** For every P_T , $\text{sign } \Delta = \text{sign}(M + \gamma)$ (Lemma 1).
- (ii) **Sufficiency.** If $|M| > \beta$ and $P_T \in \mathcal{C}_\beta$, then $\text{sign } \Delta = \text{sign } M$.
- (iii) **Necessity.** If $|M| < \beta$ (which requires $\beta > 0$), there exist $P_T^1, P_T^2 \in \mathcal{C}_\beta$ with the same augmented-evidence law and opposite nonzero benefit signs; at $|M| = \beta > 0$ the admissible drift $\gamma = -\text{sign}(M)\beta$ yields $\Delta = 0$, and in the degenerate case $M = \beta = 0$ the class $\mathcal{C}_0$ forces $\Delta \equiv 0$. Hence no strict benefit direction is uniformly supported by $\text{Law}(\tilde{Z})$ over $\mathcal{C}_\beta$ whenever $|M| \leq \beta$.
- (iv) **Strict-commitment iff.** At the fixed augmented-evidence law, a strict action in $\{\text{ADAPT}, \text{FREEZE}\}$ is uniformly directionally sound in the sense of Definition 3 if and only if $|M| > \beta$; on the closed band $|M| \leq \beta$ (including $M = \beta = 0$, where $\Delta \equiv 0$) abstention is the pointwise maximal sound action.

The population rule $M > \beta \Rightarrow \text{ADAPT}$, $M < -\beta \Rightarrow \text{FREEZE}$, $|M| \leq \beta \Rightarrow \text{ABSTAIN}$ is the pointwise maximal sound certificate on $\mathcal{C}_\beta$.

Proof. Part (i) is Lemma 1. For (ii), if $M > \beta$ and $|\gamma| \leq \beta$, then $M + \gamma > 0$, so $\text{sign } \Delta = \text{sign}(M + \gamma) = +1$; the negative case is symmetric. For (iii), Theorem 1 supplies two laws in the same augmented-evidence fibre with opposite nonzero signs whenever $|M| < \beta$. At $|M| = \beta$ the boundary law with $\gamma = -\text{sign}(M)\beta$ gives $\Delta = 0$, still precluding a determined nonzero sign. Part (iv) combines (ii), (iii), and Proposition 2. Pointwise maximality follows because (ii) certifies the unique strict direction outside the band, while Theorem 1 and Proposition 2 rule out both strict directions inside and on the boundary. $\square$

Region	Meaning	Action
$M > \beta$	allowed drift cannot flip the positive sign	ADAPT
$M < -\beta$	allowed drift cannot flip the negative sign	FREEZE
$ M < \beta$	drift can reverse the strict sign	ABSTAIN
$ M = \beta > 0$	drift can reduce benefit to zero	ABSTAIN
$M = \beta = 0$	benefit is identically zero	ABSTAIN

Theorem 3 (Finite-sample adapt/freeze/abstain certificate). *Suppose Δ , $\hat{\Delta}$, and the nonnegative random radius ε are measurable on a common probability space containing*

all calibration, benefit-model fitting, and deployment-unit randomness, and that the benefit interval satisfies marginal coverage

$$\mathbb{P}[|\hat{\Delta} - \Delta| \leq \varepsilon] \geq 1 - \alpha.$$

Define

$$g(Z) = \begin{cases} \text{ADAPT} & \hat{\Delta} - \varepsilon > 0, \\ \text{FREEZE} & \hat{\Delta} + \varepsilon < 0, \\ \text{ABSTAIN} & \text{otherwise.} \end{cases}$$

Then, marginally over that common probability space,

$$\begin{aligned} \mathbb{P}[g(Z) = \text{ADAPT}, \Delta \leq 0] &\leq \alpha, \\ \mathbb{P}[g(Z) = \text{FREEZE}, \Delta \geq 0] &\leq \alpha. \end{aligned}$$

Proof. On the event $\{|\hat{\Delta} - \Delta| \leq \varepsilon\}$, the condition $\hat{\Delta} - \varepsilon > 0$ implies $\Delta > 0$, so a false adapt requires the complement of this event, which has probability at most α . The freeze case is symmetric. $\square$

Exchangeability, or an explicit shift correction, is one route to this coverage assumption. Risk alignment is not required for the conditional-on-coverage implication.

Remark 2 (Marginal vs. conditional error). Theorem 3 bounds the *marginal, unconditional* false-adapt and false-freeze probabilities (FA_u). It does *not* bound the conditional rate $\text{FA}_c = \mathbb{P}[\Delta \leq 0 \mid \text{ADAPT}]$, which we report empirically. Risk alignment (Definition 1) is required for the *frontier* identifiability claim; coverage alone certifies the *interval* rule above.

Remark 3 (Fallback when assumptions are unsupported). If interval coverage—or the calibration assumptions supporting it (exchangeability, or an explicit shift correction)—cannot be justified on the deployment batch, Theorem 3 supplies no level- α protection for either strict action. A deployment must then use a separately declared operational fallback, such as abstention or serving the frozen model; this is a policy recommendation, not a conclusion that either fallback is optimal. If instead *risk alignment* (Definition 1) cannot be justified, the population-frontier interpretation and the transferability of the benefit estimator are unsupported, even though the interval implication of Theorem 3 remains valid *conditional on coverage* (cf. Theorem 1).

B. How the Results Fit Together

disagreement-region reduction $\rightarrow$ matched-evidence ambiguity
 $\rightarrow$ **Thm. 1** (opposite-world construction) $\rightarrow$ **Cor. 1**
 (abstention lower bound)

$\rightarrow$ split-observable decomposition $\text{sign } \Delta = \text{sign}(M + \gamma) \rightarrow$
Thm. 2 ($|M| > \beta$ frontier)

$\rightarrow$ coverage condition $\rightarrow$ **Thm. 3** (finite-sample certificate)

C. Extensions and Deferred Results

Additional minimax, anytime-valid, multicandidate, and one-bit results are retained in the long manuscript (`kbound.tex`) and are not used by the short-paper claims; the auditability of the drift budget itself is stated in Appendix L, with proofs in the long manuscript. The saved

iWildCam streaming script computes its betting increments from frozen and adapted window accuracies, which require target labels. We therefore retain that run only as a label-informed offline stress diagnostic (on the native 35,360-image iWildCam test stream, continual TENT collapses to a cumulative macro-F1 of 0.02 versus 0.26 frozen, with predictions degenerating to a handful of classes); it is not a label-free KGA deployment result and no streaming empirical claim is made in this paper.

D. Claim Scope

The population theorem is conditional on the declared class $\mathcal{C}_\beta$, and the finite-sample guarantee is conditional on valid interval coverage. The theorem controls the marginal probabilities $\mathbb{P}(\text{ADAPT}, \Delta \leq 0)$ and $\mathbb{P}(\text{FREEZE}, \Delta \geq 0)$; it does not automatically control the conditional error among committed actions. The framework does not claim that β is identifiable from unlabeled deployment data, that the empirical radius ε estimates β , or that calibration coverage transfers to arbitrary unseen domains without an assumption or correction.

V. KGA METHOD

A. System Overview

KGA is the paper’s practical finite-sample method—a decision wrapper around a fixed adapter, not a new adaptation objective and not a numerical implementation of the population $|M| > \beta$ rule. The wrapper is *adapter-interface agnostic*: it requires only that a candidate can be proposed, evaluated on the batch, and rolled back. Its calibration is *adapter-specific*: the benefit model and radius are refit per dataset and per candidate adapter, and cross-adapter transfer of a fitted calibration is disallowed because it violates the false-adapt bound (Table XIX). We fit the benefit estimator and its radius *out of fold*, so the data used to fit $\hat{\Delta}$ and the data used to calibrate its residual are disjoint for every point. On the per-cell stress grids (CIFAR-10-C and the synthetic grids), for each cell i a gradient-boosted $\hat{\Delta}_{-i}$ is trained on the other $N-1$ cells and evaluated on cell i , and ε is an upper empirical quantile of the leave-one-out residuals $\{|\hat{\Delta}_{-i}(Z_i) - \Delta_i|\}_{i=1}^N$ —leave-one-condition-out cross-fitted empirical residual calibration, so no cell’s residual uses an estimator that saw it. On the natural-shift protocols (Camelyon17, iWildCam, Office-Home) the split is across domains: $\hat{\Delta}$ and ε are fit once on a dev split and the untouched held-out target-domain conditions are scored a single time. At deploy we extract Z from an unlabeled batch and return adapt, freeze, or abstain (Algorithm 1). Labels enter only through the calibration split; the deployment batch and held-out test scorer never use target labels for the adapter, ε , evidence, or rule selection. The coverage guarantee is split-dependent. The current released scorer uses the exact split-conformal rank $\varepsilon = r_{(k)}$, $k = \min\{n_{\text{cal}}, \lceil (n_{\text{cal}}+1)(1-\alpha) \rceil\}$. Under exchangeability this supports finite-sample marginal coverage on a clean split. The archived benchmark JSONs, however,

were generated by the earlier interpolated empirical quantile; they are therefore reported as empirical benchmark evidence rather than retroactively relabeled exact split-conformal runs. The per-cell grids set ε to the $(1-\alpha)$ empirical quantile of the leave-one-out residuals, so the *calibration-set* coverage is approximately nominal (realized 0.898 at nominal 0.90, $\alpha=0.10$)—empirical calibration evidence, not an exact finite-sample statement; this is the standard jackknife radius, which is distribution-free only asymptotically under estimator stability; its formal finite-sample counterpart is the jackknife+ of Barber et al. (2021) at level $1-2\alpha$, which we note but do not implement. The decision rule’s false-adapt guarantee is stated under the coverage hypothesis (§V), which the empirical coverage approximately meets. This construction sits in a known landscape: weighted conformal corrects coverage under known covariate shift [34], adaptive conformal tracks coverage online [36], and beyond exchangeability the coverage gap is bounded by the total-variation drift of the calibration pairs [35], which is conceptually related to but mathematically distinct from our disagreement-region calibration budget β . The certificate’s marginal false-adapt control is risk control in the RCPS/Learn-then-Test lineage [38], [39], specialized to the adapt/freeze/abstain loss; PAC prediction sets give set-valued analogues under covariate shift [37].

B. Evidence Map and Benefit Estimator

The evidence map uses only unlabeled batches and model outputs, and is instantiated as two label-free panels matched to the two calibration regimes below. On the controlled stress grids (CIFAR-10-C, ImageNet-C, PACS) the base panel is an 11-dimensional vector computed from the frozen softmax p_0 and the adapted softmax p_a on the adaptation batch: the mean predictive entropy and mean confidence of f_0 and of f_a ; the normalized marginal (predicted-class) entropy of each; the entropy and marginal-balance drops between them; the fraction of adapted predictions above confidence 0.9 (a collapse indicator); the KL divergence between the adapted and frozen predicted-class marginals; and the ℓ_2 norm of the adapter’s parameter update. On the natural-shift protocols (Camelyon17, iWildCam, Office-Home, RxRx1) a 16-dimensional panel adds distribution- and disagreement-oriented statistics computed from the frozen logits: entropy and confidence quantiles; the frozen/adapted prediction-disagreement rate; the entropy gap; a free-energy source→target shift; a batch-vs-running BatchNorm-statistic KL; an ATC-style threshold-accuracy estimate (its threshold fixed on *source* labels only, never target labels); and a source→target confidence drop. The source-calibrated correctness score s of §IV is operationalized by the adapted-confidence and ATC-style features. KGA does not explicitly estimate the theoretical margin M ; it learns the benefit $\hat{\Delta} = h_\theta(Z)$ directly from the full evidence vector Z .

TABLE V

LABEL-FREE EVIDENCE ACTUALLY COMPUTED BY KGA. THE 11-DIM BASE PANEL DRIVES THE STRESS GRIDS; THE NATURAL-SHIFT PROTOCOLS USE A 16-DIM PANEL THAT ADDS THE ROWS MARKED (NAT.). EXACT FORMULAS AND PER-TRACK SUBSETS ARE IN THE CONFIGURATION FILES.

Family	Statistics
Frozen output Z_0	mean entropy, mean confidence, normalized marginal-class entropy; entropy/confidence quantiles (<i>nat.</i>)
Adapted output Z_a	mean entropy, mean confidence, normalized marginal entropy, frac. confidence > 0.9
Change Z_{change}	entropy drop, marginal-balance drop, entropy gap
Disagreement	frozen/adapted prediction-disagreement rate (<i>nat.</i>)
Distribution drift	adapted-vs-frozen marginal KL; free-energy shift, BN-statistic KL, ATC/confidence drop (<i>nat.</i>)
Update behavior	ℓ_2 norm of the adapter parameter update

The benefit model predicts $\hat{\Delta} = h_\theta(Z)$ from labeled development conditions; h_θ is a gradient-boosted regression tree (scikit-learn `GradientBoostingRegressor`, squared error, 250 trees, depth 2, learning rate 0.05, subsample 0.8, seed 0), refit *per dataset and per candidate adapter* rather than shared, so the estimator itself implies no cross-dataset transfer. The complete schema for both panels—exact per-feature formulas and families—is given in Appendix B (Tables XXII and XXIII), with per-adapter update hyperparameters in Table XXIV.

C. Out-of-Fold Uncertainty Calibration

For calibration condition i , an estimator that did not train on i produces $\hat{\Delta}_{-i}(Z_i)$ and residual

$$r_i = |\hat{\Delta}_{-i}(Z_i) - \Delta_i|.$$

The radius is a pre-specified upper residual quantile. On clean development/test splits, ordinary split conformal provides finite-sample marginal coverage under exchangeability. On stress grids, the reported radius uses leave-one-condition-out cross-fitted empirical residual calibration; its observed coverage does not imply the same finite-sample guarantee as split conformal, and jackknife+ is not claimed. We therefore report the calibration design and guarantee level separately for each track.

On the good event $|\hat{\Delta} - \Delta| \leq \varepsilon$, committing adapt only when $L > 0$ implies $\Delta > 0$ (Theorem 3). The implication requires measurable marginal coverage; exchangeability or shift correction is used to justify coverage, while risk alignment is relevant to the population interpretation and estimator transfer. No target labels are used at deployment.

D. Population Frontier versus Empirical Certificate

KGA’s commit boundary is not a swept hyperparameter, but the population frontier and the empirical radius must be kept distinct. The *population* frontier of Theorem 2 says a strict action is uniformly sound over $\mathcal{C}_\beta$ iff the population margin $|M|$ exceeds the declared budget β . The realized drift satisfies $|\gamma| \leq \beta$ but is unobserved. KGA instead

Algorithm 1 KGA: Knowability-Guided Adaptation

Require: Frozen f_0 ; adapter $\mathcal{A}$; batch $X_{1:m}$; $\hat{\Delta}$, ε , α
Ensure: Decision $g \in \{\text{ADAPT}, \text{FREEZE}, \text{ABSTAIN}\}$

```

1:  $f_a \leftarrow \mathcal{A}(f_0, X_{1:m})$ 
2:  $Z \leftarrow \phi(X_{1:m}, f_0, f_a)$ 
3:  $\hat{\Delta} \leftarrow \hat{\Delta}(Z)$ ;  $L \leftarrow \hat{\Delta} - \varepsilon$ ;  $U \leftarrow \hat{\Delta} + \varepsilon$ 
4: if  $L > 0$  then
5:    $g \leftarrow \text{ADAPT}$ 
6: else if  $U < 0$  then
7:    $g \leftarrow \text{FREEZE}$ 
8: else
9:    $g \leftarrow \text{ABSTAIN}$ 
10: end if
11: return  $g$ 
```




Fig. 3. **The certificate rule in one picture.** KGA estimates the label-free benefit $\hat{\Delta}(Z)$ and forms an interval $\hat{\Delta} \pm \varepsilon$: if the interval is entirely positive it adapts, if entirely negative it freezes, and otherwise it abstains. The vertical axis here uses true benefit only for offline audit; deployment decisions use only unlabeled evidence and the pre-calibrated radius.

learns $\hat{\Delta}$ from Z and uses a finite-sample residual-coverage radius ε , fixed before target-test scoring. The radius neither estimates nor substitutes for β ; conditional on valid coverage it controls marginal decision errors without requiring risk alignment. Confidence-thresholding (ATC) and agreement-trend selection (AGL) fit a scalar against a labeled or held-out signal and predict *accuracy*; KGA certifies a strict benefit direction with a calibrated interval and, unlike a committal thresholded estimator, ABSTAINS where its calibrated interval crosses zero—consistent with the structural region studied by Theorem 1, though finite-sample abstentions may also reflect sampling uncertainty, estimator inadequacy, or calibration conservatism. The impossibility construction, not a tuned cutoff, is what separates KGA from “estimate accuracy, then threshold.”

E. Selecting and Auditing the Drift Budget

The drift budget β is part of the declared deployment class, not a target-test-tuned hyperparameter. It may be justified from historical domain calibration gaps, a documented domain assumption, or an explicit stress envelope.

Sensitivity analysis should report how decisions change over a pre-specified range of budgets. A smaller β asserts stronger calibration transfer and increases commitment; a larger β widens the abstention region. When no credible upper bound is available, the population frontier does not provide a robust sign certificate. Reporting $\beta = 0$ is then only a zero-drift plug-in analysis, not an assumption-free or conservative guarantee. Appendix L states exactly when a *computed* $\hat{\beta}$ can replace the declared budget, and the limits of any such audit.

F. Deployment Semantics and Failure-Safe Behavior

The shadow-state workflow is: (1) checkpoint the active model; (2) propose adaptation on a temporary copy; (3) evaluate frozen and candidate outputs; (4) extract Z ; (5) compute $\hat{\Delta} \pm \varepsilon$; (6) commit or roll back; and (7) log the configuration, evidence, interval, and action. Abstention applies to the adaptation update, not to prediction. ADAPT commits the candidate; FREEZE rolls it back; and ABSTAIN retains the frozen state while recording that evidence was insufficient rather than certifiably harmful. The implementation should automatically freeze or abstain when evidence is missing, the batch is too small, a checkpoint or adapter configuration differs from calibration, the evidence lies outside calibration support, numerical values are invalid, or no valid radius is available. In the reference implementation (`kbound_pkg`), the gate is applied per batch (`KGA.decide_from_batch`); on ABSTAIN or FREEZE the adaptation optimizer scales the candidate update by `abstain_scale=0`, i.e. the update is skipped and the frozen state f_0 is served—this is the rollback mechanism. The failure-safe defaults instantiate the diagnostics of Table XX: abstain when the batch is smaller than the calibration batch size, when Z lies outside calibration support, when a checkpoint or adapter-configuration hash differs from the calibrated one, or when the radius or evidence is missing or numerically invalid. In continual operation the last certified state replaces f_0 as the fallback.

G. Computational Cost

KGA cannot avoid the cost of proposing an update because candidate adaptation includes forward and backward computation before commitment. The controller’s *added* cost is measured by a controller-only microbenchmark (Table VI): evidence extraction 0.074 ms and benefit-model evaluation 0.127 ms per decision (0.20 ms total), plus one model copy (44.8 MB, ResNet-18/CIFAR) for rollback; the benefit model itself is 195 KB. Candidate adaptation and inference are unchanged by KGA— f_a is evaluated regardless—and are expected to dominate end-to-end latency, but this microbenchmark does not measure them: a complete end-to-end component profile (frozen inference, candidate adaptation and inference, rollback time, total window latency) remains pending (App. H).

TABLE VI
KGA CONTROLLER *ADDED* COST (RESNET-18/CIFAR;
CONTROLLER-ONLY MICROBENCHMARK). CANDIDATE
ADAPTATION/INFERENCE ARE UNCHANGED BY KGA AND ARE NOT
MEASURED HERE; THE END-TO-END COMPONENT PROFILE IS PENDING
(APP. H).

Component	Cost
Evidence extraction (per batch)	0.074 ms
Benefit-model evaluation (per decision)	0.127 ms
Controller added latency	0.20 ms
Rollback model copy (memory)	44.8 MB
Benefit model size	195 KB

VI. EXPERIMENTAL SETUP

A. Research Questions

The experiments address seven questions: **RQ1** whether the measured commit transition follows the theoretical frontier; **RQ2** whether the interval controls marginal false adaptation; **RQ3** whether KGA assigns ADAPT, FREEZE, and ABSTAIN consistently with helpful, harmful, and low-margin conditions; **RQ4** whether it prevents damage on held-out natural shifts; **RQ5** whether it improves decision validity over simple and neighboring label-free gates; **RQ6** which evidence and calibration choices matter; and **RQ7** what computational overhead the controller adds.

B. Datasets and Shift Taxonomy

The evaluation separates controlled corruption/composition stress, large-scale corruption, natural institutional or geographic shifts, domain generalization, and low-margin diagnostic transfers. This taxonomy matters because mixed conditions test action discrimination, one-sided shifts test safe commitment, and low-margin transfers test abstention. The nine tracks draw on CIFAR-10-C [21], ImageNet [20], the WILDS Camelyon17, iWildCam, and RxRx1 shifts [22], [23], [24], Office-Home [25], PACS via DomainBed [26], and CIFAR-10.1 [28]; models build on RobustBench checkpoints [27] with PyTorch [29] and scikit-learn [30].

C. Models, Adapters, and Leakage Control

For every track, the source checkpoint, candidate adapter, update mode, learning rate, number of steps, batch construction, adapted parameter subset, and random seeds must be declared before held-out scoring. Adapter selection, evidence selection, benefit-model fitting, and residual calibration use development data only. All choices are frozen before the target test cells are scored once; target labels are opened afterward solely for offline audit. The calibration and scoring configuration is uniform across tracks except where noted in Table VIII.

Exact per-adapter update settings (optimizer, steps, learning rate, and the mild/aggressive operating points) are tabulated per track in Appendix B, Table XXIV.

TABLE VII

DATASET ROLES IN THE EVALUATION. THE EXPECTED REGIME IS A HYPOTHESIS FIXED BY THE PROTOCOL, NOT A CONCLUSION INFERRED FROM THE FINAL SCORE.

Track	Shift type	Candidate role	Evaluation purpose
CIFAR-10-C stress	corruption/composition/update stress	Tent/EATA (SAR withheld)	mixed detectable decision benchmark
ImageNet-C noise	large-scale corruption	Tent/EATA/SAR	scale and candidate dependence
Camelyon17	hospital/institution shift	EATA	helpful-dominated natural safety
iWildCam	geographic camera shift	Tent	harmful-dominated natural safety
Office-Home	visual domain/style shift	SAR	held-out domain safety
RxRx1	experimental/laboratory shift	online adapter	harmful-dominated natural safety
PACS	domain generalization	fixed candidate	breadth and null behavior
ImageNet-R	rendition shift	candidate panel	weak-evidence diagnostic
CIFAR-10.1	natural resampling	TTA candidates	low-margin transfer diagnostic

TABLE VIII

PER-TRACK CALIBRATION AND SCORING CONFIGURATION ($\alpha=0.10$ THROUGHOUT; BENEFIT ESTIMATOR FAMILY AND CALIBRATION RECIPE ARE USED ACROSS TRACKS, WITH PARAMETERS REFIT PER DATASET AND CANDIDATE ADAPTER, §V). BACKBONES: RESNET-18 (CIFAR), RESNET-50 (IMAGENET-C); THE NATURAL-SHIFT BACKBONES FOLLOW THE WILDS/DOMAINBED DEFAULTS PER TRACK.

Track	Calibration unit	Radius rule	Test unit	
CIFAR-10-C stress	grid cell (leave-one-cell-out)	LOO jackknife $q_{0.9}$	5×432 cells	
ImageNet-C	grid cell (leave-one-cell-out)	LOO jackknife $q_{0.9}$	27 cells	
PACS (LODO)	source-domain cells	LOO jackknife $q_{0.9}$	18 cells/target/seed; 3 seeds	test {2, 3, 4}
Camelyon17	dev {0, 1}	archived OOF empirical radius	held-out {1}	
iWildCam H v2	calibration {0}	archived OOF empirical radius	target {0, 1}	
Office-Home M v2	calibration {0, 1}	archived OOF empirical radius	test {5, ..., 9}	
RxRx1 J	development {0, ..., 4}	archived OOF empirical radius		

D. Baselines

We distinguish fixed policies (always-freeze and always-adapt), an offline per-condition oracle, simple label-free gates (confidence, entropy, drift/KL, and ATC-style), KGA without its radius, and neighboring methods or ports inspired by POEM and AETTA. Each comparison must state whether the baseline is an official implementation, a faithful reimplement, or a protocol-matched port.

E. Metrics and Statistical Testing

The primary decision metrics are regret-to-oracle, marginal false-adapt $\text{FA}_u = \mathbb{P}(\text{ADAPT}, \Delta \leq 0)$, conditional false-adapt $\text{FA}_c = \mathbb{P}(\Delta \leq 0 \mid \text{ADAPT})$, adapt rate, action composition, and decision coverage $\mathbb{P}(\text{ADAPT OR FREEZE})$. Accuracy or AUROC is reported as a task metric, not as the controller’s objective. The theorem controls FA_u under its assumptions; FA_c is descriptive. Lower regret than both fixed policies is a secondary routing-utility check and does not establish interval coverage or risk alignment. Paired bootstrap confidence intervals, their resampling unit,

seed aggregation, and Holm correction are declared per comparison family.

F. Evidence-Status Policy

Locked results trace to raw cells/seeds and a frozen aggregate; **reconciled** results correct a prior interpretation using raw records; **provisional** summaries require canonical replay; and **diagnostic** runs are complete but deliberately not promoted. Only locked and reconciled rows support the empirical headline. Every evidence-tier cell in the result tables begins with one of these four labels; qualifiers (seed counts, lock type, known caveats) follow in parentheses.

G. Hardware and Software

Runs use Python 3.12, PyTorch 2.5.1, and scikit-learn for the benefit estimator. Accelerator and checkpoint details are recorded per run manifest rather than pooled into one hardware claim: CIFAR experiments used the saved ResNet-18 setup; the authoritative ImageNet-C result used ResNet-50 on seeds 0–4. Hardware affects runtime but not the decision definitions or replayed metrics.

VII. RESULTS

The results are organized by regime rather than by dataset prestige. Controlled mixed regimes test action discrimination and selective coverage, one-sided natural shifts test no-harm behavior, and low-margin transfers test whether the method declines unsupported claims.

A. Synthetic validation of the frontier

Before the real-data tracks we validate the theory directly on synthetic ground truth, where the observable margin M , the unobserved drift $\gamma \sim \text{U}(-\beta, \beta)$, and hence the true benefit sign $\text{sign } \Delta = \text{sign}(M + \gamma)$ are all controlled. In the real-data tracks that follow, β is *not* numerically supplied to KGA: the empirical decision uses $\hat{\Delta} \pm \varepsilon$, and the β -frontier is evaluated directly only in this synthetic study. Evidence is fed to the reported cross-fitted rule (leave-one-out gradient-boosted $\hat{\Delta}$ with an empirical residual radius, §V); nothing is hand-tuned. Three predictions hold (Fig. 5; $\alpha=0.10$). **(i) Recovery:** $\hat{\Delta}$ tracks Δ at 90.0% empirical coverage, matching the nominal target in this simulation. In this construction unresolved drift dominates





Fig. 4. **Regime geometry for K-Bound.** The important axis is not dataset name but where the deployment sits relative to the frontier. Helpful- or harmful-dominated regimes usually produce no-harm ties with the better fixed policy; low-margin regimes force abstention; mixed and detectable regimes exercise all three actions and make routing utility observable. Fixed-policy regret is secondary. This is a *conceptual schematic*: axis positions are illustrative regime categories, not measured coordinates (the actual per-track verdicts are in Table I).

the benefit-estimation residual, causing the calibrated radius to approximately track β . **(ii) Frontier** (Thm 2): the commit rate jumps from 6.5% for $|M| < \beta$ to 96.7% for $|M| > \beta$ —a transition *at* the predicted boundary—and the committed sign is correct 99.3% of the time. **(iii) Certificate** (Thm 3): the unconditional false-adapt rate stays under budget for every $\beta \in \{0.05, \dots, 0.20\}$ ($\text{FA}_u \leq 0.014 \ll \alpha$). The frontier is not imposed; it appears because the synthetic residual scale is controlled by the deliberately hidden drift.



B. CIFAR-10-C stress grid (5 seeds)

The per-corruption suite is helpful-dominated; the *stress* grid crosses severity, batch size, composition, and update aggressiveness (432 conditions/method, ResNet-18; pre-registered as Protocol A; seeds 0–4). Harmful base rates are 34% (Tent) and 27% (EATA). The primary result is selective routing at 0% false-adapt (Table IX). Among harmful cells KGA adapts 0 and freezes them; among helpful cells it adapts almost all; on marginal cells it abstains. For both Tent and EATA, the regret gaps to *both* fixed policies are positive with 95% bootstrap CIs excluding zero, Holm-corrected over the comparison family, at false-

Fig. 5. **Measured benefit-sign frontier** (synthetic ground truth, archived cross-fitted rule). The commit rate transitions sharply from abstention to adapt/freeze near $|M|=\beta$; the calibrated radius ε (dotted) approximately tracks β in this construction. Companion panels—estimator recovery and $\text{FA}_u \leq \alpha$ across β —in `fig_frontier_recovery.png`, `fig_frontier_fa_coverage.png`.

adapt 0/2160.¹ The previously mismatched SAR aggregate

¹Design-based intervals, resampling the pre-registered mixing-ratio distribution of the stress-grid design (Protocol A; the 11-point Pareto over harmful fraction, $N_{\text{boot}}=5000$): Tent +0.019 [0.011, 0.026] vs. always-adapt, +0.080 [0.052, 0.107] vs. always-freeze; EATA +0.015 [0.009, 0.021] and +0.075 [0.049, 0.100]. A conventional paired per-condition bootstrap over the 432 realized conditions returns the same Tent/EATA verdict.

candidate	policy	mean acc	regret↓	FA _u	coverage	lower than both?
Tent	always-adapt	0.786	0.0079	—	1.00	
	always-freeze	0.671	0.1241	—	1.00	
	K-Bound	0.793	0.0016	0	0.67	yes
EATA	always-adapt	0.798	0.0033	—	1.00	
	always-freeze	0.671	0.1314	—	1.00	
	K-Bound	0.801	0.0013	0	0.63	yes

TABLE IX

CIFAR-10-C STRESS GRID (432 CONDITIONS/METHOD, FIVE SEEDS, $\alpha=0.1$, PROTOCOL A). “COVERAGE” IS DECISION COVERAGE $\mathbb{P}[\text{ADAPT} \vee \text{FREEZE}]$; BOTH FIXED POLICIES COMMIT ON EVERY CELL. SAR IS OMITTED FROM THIS FROZEN DISPLAY; ITS RECONCILED FIVE-SEED CURRENT-TREE RESULT IS REPORTED IN THE TEXT TOGETHER WITH ITS LARGER ACROSS-SEED RADIUS VARIATION.

was rebuilt from all five saved per-condition seed files using the locked Protocol A analysis. The rebuild yields regret 0.0015/0.0112/0.1286 for KGA/always-adapt/always-freeze and false-adapt 0/2160; paired condition-bootstrap intervals exclude zero against both fixed policies. This is retained as a current-tree reconciliation rather than evidence that every historical archived SAR summary was equivalent: the SAR conformal radius is substantially less stable across seeds (CV 0.390) than Tent or EATA.

C. Decision-baseline comparison: empirical safety of the calibrated certificate

KGA’s claim goes beyond “better than always-adapt/always-freeze”: it is better than *simple label-free decision rules* under the same locked protocol. On the CIFAR-10-C stress grid (Tent candidate, identical cells), we score six rules from the same evidence (Table X): a confidence gate (adapt iff adaptation raised mean confidence), an entropy gate (adapt iff it lowered entropy—Tent’s own objective), a drift/KL threshold, an ATC-style accuracy-prediction gate, KGA *without* its conformal radius (adapt iff $\hat{\Delta} > 0$), and the full certificate. Only KGA is derived from a coverage-based false-adapt certificate; on this stress grid the leave-one-cell-out radius provides strong *empirical* coverage, not the exact distribution-free validity of clean split conformal or jackknife+. It attains zero observed FA_u on this grid while staying near-oracle on regret (the radius-free variant keeps FA_u=0.049< α empirically, but with no such guarantee). The confidence/entropy gates false-adapt on ~74% of harmful cells; the drift threshold, tuned to the budget, degenerates to always-freeze. The KGA-without-radius row isolates the radius: removing it lowers regret slightly (0.0017→0.0004) but raises FA_u from 0 to 0.049 (to 0.141 on harmful cells): the radius is what turns a good benefit *estimate* into a false-adapt *guarantee*. Reproduced from the locked per-cell dump (artifact manifest, App. J).

D. Mixed head-to-head vs. POEM-/AETTA-style baselines (pre-registered)

The stress-grid wins above beat the *trivial* policies always-adapt and always-freeze. A harder question is whether KGA improves over strong protocol-matched label-free decision baselines inspired by POEM [10] and AETTA [15], on a *mixed* harmful+helpful

TABLE X

Decision-baseline comparison, CIFAR-10-C stress grid (TENT CANDIDATE, $n=432$ CELLS, 149 HARMFUL, $\alpha=0.1$). LOWER REGRET AND LOWER FA_u ARE BETTER. ONLY THE CONFORMAL CERTIFICATE KEEPS FA_u=0 ON THE FULL GRID AND ON HARMFUL CELLS WHILE STAYING NEAR-ORACLE.

Decision rule	regret↓	FA _u	FA _c	adapt	cov.	FA _u (harm)
confidence gate	0.0084	0.257	0.301	0.85	1.00	0.745
entropy gate	0.0086	0.255	0.304	0.84	1.00	0.738
drift/KL gate	0.1232	0.000	0.000	0.00	1.00	0.000 [†]
ATC-style gate	0.0045	0.116	0.172	0.67	1.00	0.336
KGA (no radius)	0.0004	0.049	0.071	0.68	1.00	0.141
KGA (certificate)	0.0017	0.000	0.000	0.51	0.68	0.000

[†]Degenerate: the budget-tuned drift gate never adapts (adapt-rate 0), so its “0 false-adapt” is just always-freeze (regret 0.123).

TABLE XI


PRE-REGISTERED MIXED HEAD-TO-HEAD ON THE LOCKED CIFAR-10-C STRESS STREAM. POEM/AETTA ROWS ARE PROTOCOL-MATCHED STYLE PORTS, NOT OFFICIAL REPRODUCTIONS. ALL POLICIES USE THE SAME CROSS-FITTED EVALUATION CELLS, CANDIDATE ADAPTER, LOGGED LABEL-FREE EVIDENCE, REGRET-TO-ORACLE METRIC, AND UNCONDITIONAL FALSE-ADAPT METRIC.

policy	mean regret↓	FA _u	decisive rate	KGA gap	Holm beats?
always-adapt	0.0079	0.329	1.00	−0.0064	yes
always-freeze	0.1241	0.000	1.00	−0.1225	yes
POEM-style	0.0088	0.245	1.00	−0.0072 [−0.0088, −0.0057]	yes
AETTA-style	0.0073	0.240	1.00	−0.0058 [−0.0071, −0.0044]	yes
KGA	0.0016	0.0000	0.6847	—	WIN
oracle	0.0000	0.000	1.00	—	ceiling


stream. We pre-registered the benchmark, metrics, and WIN/TIE/LOSE criterion *before* computing any number (MIXED_BENCHMARK_PROTOCOL.md). All six policies consume the *same* logged per-condition signals Z and benefit B from the locked CIFAR-10-C stress grid (Tent, 432 conditions/seed, five seeds; harmful fraction $\approx 33\%$, i.e. the always-adapt FA_u of Table XI); only the decision rule differs. The competitors are *protocol-matched POEM-style and AETTA-style baselines*: ports of the published protectors onto this protocol with documented simplifications (batch-summary entropy rather than per-sample streams). Official per-sample POEM and dropout-based AETTA reproduction remains a camera-ready faithfulness check.

Verdict: WIN. KGA lowers mean regret-to-oracle vs. both ported baselines with paired-bootstrap 95% CIs entirely below zero and Holm correction ($\alpha=0.05$), at FA_u=0 $\leq \alpha$ (Table XI). The ported POEM-/AETTA-style baselines false-adapt at ~24–25% on the same stream, consistent with the absence of an explicit marginal false-adapt certificate under this ported protocol. Secondary pooled sets (Tent+EATA, EATA alone) yield the same WIN verdict, and a separate pre-registered recomputation pipeline (paired bootstrap 10^4) reproduces all three configurations internally, with both regret-gap CIs below zero in each and FA_{KGA}=0 versus 11–33% for the competitors; both artifacts are indexed in the manifest (App. J).

a) *Baseline faithfulness.*: Table XII records exactly what each port matches (evidence, candidate, cross-fitted evaluation cells, metrics) and what it does not claim



(a) decisions by true regime



(b) regret as harmful fraction changes

Fig. 6. **CIFAR-10-C stress grid is a decision result, not only an accuracy result.** KGA adapts on helpful cells, freezes harmful cells, abstains on marginal cells, and therefore lies on the regret front: it ties always-adapt at $p=0$ and improves both fixed-policy regrets once harmful mass is nontrivial.

(POEM’s per-sample betting process; AETTA’s dropout estimator); the comparison claim is limited to this locked protocol-matched setting.

E. ImageNet-C SAR (mechanism-faithful, five seeds)

On ImageNet-C noise (ResNet-50, 27 cells: three noise corruptions $\times$ severities $\{1, 3, 5\} \times$ three batch-composition regimes (iid, imbalanced, single-class), seeds 0–4, full-scale), SAR is harmful on 15.6% of cell-seed pairs under a mechanism-faithful re-implementation (SAM, recovery reset, EMA) at the learning rate shared across candidates. Pooled over five seeds (seed-averaged conditions, the same paired design as the CIFAR-10-C rows), always-adapt drops to 0.385 accuracy while KGA reaches 0.427, **beating both** trivial policies (regret 0.0264 vs 0.0529/0.0319) at $\text{FA}_u = 0.000$, using the exact split-conformal radius $\varepsilon = \rho_{(k)}$, $k = \lceil (n+1)(1 - \alpha) \rceil$. The pooled 95% paired-bootstrap gap to always-freeze is $[-0.0088, -0.0027]$ and to always-adapt $[-0.052, -0.003]$; both intervals exclude zero. As a secondary utility check, point estimates improve both fixed-policy regrets on 5/5 seeds; the per-seed 27-cell CIs exclude zero on 2/5 (seeds 0–1) and narrowly include it on the lighter-harmful-mass seeds (App. K)—the win scales with detectable harmful mass, as the regime map predicts. EATA stays helpful-dominated (KGA ties always-adapt); on Tent the harmful regime dominates (55.6% of cell-seed pairs) and KGA abstains to always-freeze (Table XIII). Under the same logged conditions, an ATC-style confidence gate false-adapts on the harmful SAR cells—it adapts wherever confidence is high, which is exactly where SAR collapses—while KGA abstains on them, so the certificate beats not only the two fixed policies but a simple label-free decision gate under the same protocol. The result uses our mechanism-faithful SAR reimplementation at a shared learning rate; SAR’s official gentler schedule is left as a robustness check, so the claim is about this harmful-SAR operating point rather than all SAR configurations.




Fig. 7. **Natural-shift action safety.** Positive values favor KGA. Office-Home and iWildCam beat always-adapt but tie always-freeze under the corrected out-of-fold replay; the intervals are paired held-out-condition bootstrap intervals from the saved OOF bootstrap lock (App. J).

F. Natural Shifts and Consolidated Summary

All benchmark tracks use the same decision target: the adaptation benefit $\Delta = R(f_0) - R(f_a)$, evaluated against the same per-condition oracle, always-adapt, and always-freeze policies. The candidate adapter is locked using development data. Evidence Z , the benefit estimator $\hat{\Delta}$, and radius ε are then frozen before scoring the declared test cells; test labels are used only after decisions for offline regret and false-adapt audit. Stress grids use leave-one-cell-out residuals; semantic domain-shift tracks use a dev/test split. We report $\text{FA}_u = \mathbb{P}[\text{ADAPT}, \Delta \leq 0]$ and do not interpret FA_c as a certificate.

a) *Evidence tiers.*: Tier labels follow the four-level policy of §VI-F; only locked and reconciled rows support the empirical headline.

b) *Multi-seed stability (Camelyon17).*: The one natural-shift track with per-condition logs at multiple seeds is Camelyon17 (4 seeds, 9 composition-stress conditions/seed). Aggregating the deployed per-seed decisions (locked per-seed artifacts, App. J) shows the no-harm

TABLE XII
BASELINE FAITHFULNESS AUDIT FOR POEM/AETTA COMPARISONS.

Baseline	Original method needs	Our locked protocol uses	What is matched	What is not claimed
POEM-style	online entropy/betting monitor over adaptation stream	batch-summary entropy and logged per-condition evidence	same cross-fitted evaluation cells, same candidate, same false-adapt/regret metrics	not official per-sample POEM reproduction
AETTA-style	dropout/disagreement-based label-free accuracy estimation	protocol-matched accuracy/benefit from logged evidence	same evidence budget, same decision task, same cross-fitted evaluation stream	not official dropout-AETTA reproduction
KGA	benefit estimator + conformal radius	out-of-fold $\hat{\Delta}$ and ε	same stream, oracle/regret/FA _u metric	official method in this paper

candidate	policy	mean acc	regret↓	harmful cells	lower than both?
Tent	always-adapt	0.402	0.0191	56%	
	always-freeze	0.406	0.0145		
	K-Bound	0.407	0.0139		no ($\approx$ ties freeze)
EATA	always-adapt	0.440	0.0001	2%	no (ties adapt)
	always-freeze	0.406	0.0342		
	K-Bound	0.440	0.0003		
SAR	always-adapt	0.385	0.0529	16%	
	always-freeze	0.406	0.0319		
	K-Bound	0.427	0.0264		yes

TABLE XIII
IMAGENET-C, MECHANISM-FAITHFUL SAR (27 CELLS: 3 NOISE CORRUPTIONS, RESNET-50, $\alpha=0.1$); POOLED OVER SEEDS 0–4 (SEED-AVERAGED CONDITIONS), SUPERSEDING THE EARLIER SINGLE-SEED AND 36-CELL PROVISIONAL CONFIGURATIONS. PER-SEED DETAIL IN APP. K.

TABLE XIV
Multi-seed no-harm (Camelyon17, seeds 0–3). REGRET
MEAN $\pm$ STD ACROSS SEEDS; FA_u IS THE PER-SEED MAXIMUM.
REPRODUCED FROM THE COMMITTED PER-SEED LOGS.

candidate	KGA regret	adapt	freeze	FA _u	verdict
Tent	0.020 $\pm$ 0.023	0.138	0.020	0.00	stable no-harm
EATA	0.039 $\pm$ 0.025	0.042	0.042	0.00	inconclusive
SAR	0.041 $\pm$ 0.017	0.000	0.065	0.11	over-freezes (helpful-dom.)

outcome is *candidate-dependent* and stable only for the milder adapter (Table XIV): with Tent, KGA is **stable no-harm** across all four seeds—it ties always-freeze, beats always-adapt, at FA_u=0 every seed; with EATA the margin is inconclusive; and with the aggressive helpful-dominated SAR arm KGA over-freezes (worse than always-adapt, FA_u up to 0.11), consistent with the candidate dependence seen elsewhere. Per-seed multi-seed replays for iWildCam, Office-Home, and RxRx1 are single-run in the current artifacts and remain future work.

G. Weak-Evidence and Negative Regimes

ImageNet-R and CIFAR-10.1 remain diagnostics. Their outcomes are consistent with weak evidence, low margin, estimator inadequacy, or calibration-transfer failure; they do not by themselves prove structural non-identifiability. PACS is complete at three seeds but remains a null diagnostic.

H. Physical-Study Pre-registration

We preregister a two-phone physical package-inspection study with distinct source, calibration, held-out, and

replication sessions. The protocol, leakage checks, and outcome-free result templates are provided in Appendix F. No physical result is claimed until fresh sessions pass the repository’s fail-closed publication gate; mock or reconstructed clips are never admissible evidence.

I. Constructed Heterogeneous Routing

The three-source OOF mixture combines separately calibrated Office-Home, iWildCam, and Camelyon17 conditions. It demonstrates routing value under a researcher-constructed heterogeneous stream, not single-dataset natural-shift dominance or transfer to unseen shift types.

J. Consolidated Claim and Guarantee Accounting

Statement	Status	Requirement
Abstention needed in matched-evidence opposite-benefit worlds	theorem	stated model class
FA _u $\leq \alpha$ (unconditional false-adapt)	theorem	measurable marginal coverage; exchangeability/shift correction only justifies coverage
FA _c $\leq \alpha$ (conditional false-adapt)	not claimed	report FA _c descriptively only
Nine-track panel	accounting	sealed lock inventory (NINE_TRACK_LOCK_SEAL_v1); incomplete tracks stay diagnostic
Stress-grid action validity and routing utility (Tent/EATA)	locked	locked Protocol A v1 stress grid
ImageNet-C SAR action validity and paired-bootstrap utility	locked (5 seeds)	27-cell noise grid $\times$ 5 seeds; pooled paired CIs exclude 0; point-level 5/5
Natural safety (Camelyon17, RxRx1)	locked	OOD reconciliation seal / held-out test artifacts
Office-Home and iWildCam OOF summaries	locked	OOF bootstrap seal (no-harm only; OH LOO BB not promoted)
Three-source mixed routing aggregate	locked	constructed OOF mixture, not transfer
Universal improvement	not claimed	impossible in general
ε estimates drift budget β	false	β is declared; ε is empirical

VIII. ABLATIONS AND SENSITIVITY

A. Value of the Uncertainty Radius

Table X isolates the radius: KGA without the radius has lower point regret but a higher empirical false-adapt rate, whereas the full certificate sacrifices coverage to eliminate observed harmful commitments on the locked

TABLE XV

Uniform nine-track empirical panel. REGRETS ARE ORDERED AS KGA / ALWAYS-ADAPT / ALWAYS-FREEZE; LOWER IS BETTER. FIXED-POLICY COMPARISONS ARE SECONDARY. THE FINAL COLUMN IS PART OF THE RESULT: IT PREVENTS A POINT ESTIMATE, A RECONSTRUCTED SUMMARY, AND A LOCKED MULTI-SEED RESULT FROM BEING PRESENTED AS EQUIVALENT EVIDENCE.

Track	Evaluation unit	Result	Decision interpretation	Evidence tier
CIFAR-10-C stress	5×432 per candidate	Tent: 0.0016/0.0079/0.1241; EATA: 0.0013/0.0033/0.1314; SAR withheld; $FA_u = 0$	selective routing; paired-CI utility	locked (Tent/EATA; five raw seeds); SAR withheld (seed-0 aggregate non-reproducing)
ImageNet-C SAR	27 cells $\times$ 5 seeds	0.0264/0.0529/0.0319 pooled; $FA_u = 0.000$ (exact split-conformal)	selective routing; pooled CI-supported utility; point-level 5/5 seeds	locked (5 seeds; per-seed CIs 2/5)
Camelyon17 OOD	genuine OOD test $n = 18$ (support $n = 36$)	0.0000/0.0000/0.1381; $FA_u = 0$	selects adapt-equivalent behavior	locked (OOD reconciliation seal)
iWildCam H v2	held-out test $n = 72$	0.0041/0.1028/0.0041; $FA_u = 0$	no-harm, ties freeze	locked (OOF bootstrap; single-run)
Office-Home M v2	target test $n = 35$	0.0157/0.0468/0.0158; $FA_u = 0$	no-harm, tiny point edge	locked (OOF no-harm only; LOO BB not promoted)
RxRx1 J	OOD test $n = 60$	0.0000/0.2531/0.0000; $FA_u = 0$	selects freeze-equivalent behavior	locked (real ckpt; single-run)
PACS	3 seeds $\times$ 4 LODO targets $\times$ 18 cells	mean regret 0.0431/0.0176/0.0446; reported mean $FA_u = 0.0093$ (pooled count not retained)	conservative/null action-safety diagnostic	locked diagnostic null; no beats-both claim
ImageNet-R D	4 seeds $\times$ 10 backbones $\times$ 12 conditions	mean-across-backbone regret 0.0112/0.0064/0.0325; $FA_u = 1/480$	candidate-dependent fixed-policy tracking; 0/10 CI beats-both	locked diagnostic null
CIFAR-10.1 K	cross-seed test $n = 48$	0.0021/0.0190/0.0017; $FA_u = 0.167$, $FA_c = 0.444$	fails transfer bar; no claim	locked diagnostic fail

TABLE XVI

Primary numeric evidence used for empirical claims. THE THREE-SOURCE ROW IS A CONSTRUCTED ROUTING MIXTURE, NOT A NATURAL-TRANSFER RESULT.

Track	Unit	KGA	Adapt	Freeze	Status
CIFAR-10-C Tent	5×432	.0016	.0079	.1241	locked; selective routing, $FA_u = 0$; CI utility
CIFAR-10-C EATA	5×432	.0013	.0033	.1314	locked; selective routing, $FA_u = 0$; CI utility
ImageNet-C SAR	27 cells $\times$ 5 seeds	.0264	.0529	.0319	locked (5 seeds); pooled routing utility
Camelyon17 OOD	$n = 18$	.0000	.0000	.1381	locked; no-harm
RxRx1 J	$n = 60$	.0000	.2531	.0000	locked; no-harm
Three-source OOF mixture	$n = 143$	.0059	.0632	.0342	locked; heterogeneous routing (constructed)

grid. This comparison separates benefit-prediction quality from certification.

B. Alpha, Evidence, Estimator, and Transfer Sensitivity

All ablations reuse the logged CIFAR-10-C stress evidence (per-condition Z and true benefit B ; 432 conditions/candidate, seed 0) and the deployed rule with an *exact-rank* empirical residual radius $\varepsilon = r_{(k)}$, $k = \lceil (n+1)(1-\alpha) \rceil$ (locked exact-rank ablation artifact with input hashes, App. J). At the deployed operating point (Tent, $\alpha=0.10$) the harness reproduces the locked gate row

TABLE XVII

ABLATION (I): SAFETY-LEVEL α SWEEP (EXACT-RANK RADIUS), CIFAR-10-C TENT ($n=432$). $FA_u=0$ AT EVERY CERTIFIED α ; ONLY THE RADIUS-FREE ROW BREAKS IT.

α	regret $\downarrow$	FA_u	adapt	coverage
0.01	0.0036	0.000	0.45	0.48
0.05	0.0021	0.000	0.49	0.63
0.10	0.0017	0.000	0.51	0.69
0.20	0.0011	0.000	0.53	0.74
no radius	0.0004	0.051	0.68	1.00

of Table X (regret 0.0017, $FA_u=0$, adapt 0.51, coverage 0.69).

(i) **Safety level α .** Sweeping α (Table XVII) trades decision coverage for regret while holding $FA_u=0$ at every certified level; only the radius-free variant breaks the guarantee ($FA_u = 0.051$). (ii) **Estimator.** Swapping the benefit regressor (Table XVIII) leaves the false-adapt guarantee intact: gradient boosting, random forest, and ridge all keep $FA_u \leq 0.002$; a small MLP over-abstains on this small calibration set. (iii) **Evidence family.** Dropping any single evidence family (frozen, adapted, change, drift, or update) keeps regret in $[0.0015, 0.0019]$ at $FA_u=0$: no single family is load-bearing. (iv) **Cross-adaptor transfer.** Fitting the benefit model on one adaptor and applying it to another (Table XIX) *breaks* the false-adapt bound in most directions (SAR $\rightarrow$ Tent $FA_u = 0.25$), which is exactly why the deployed method refits per adaptor. (v) **Drift budget β .** The population-frontier β sweep is the synthetic ground-truth validation of §VII-A, kept distinct from the empirical radius ε .

TABLE XVIII

ABLATION (II): BENEFIT ESTIMATOR (EXACT-RANK RADIUS), CIFAR-10-C TENT ($\alpha=0.10$). THE RADIUS KEEPS $FA_u \approx 0$ REGARDLESS OF ESTIMATOR.

estimator	regret $\downarrow$	FA_u	adapt	coverage
gradient boosting	0.0017	0.000	0.51	0.69
random forest	0.0017	0.002	0.51	0.65
ridge (linear)	0.0043	0.000	0.45	0.47
small MLP	0.0143	0.000	0.36	0.36

TABLE XIX

ABLATION (IV): CROSS-ADAPTER TRANSFER (EXACT-RANK RADIUS, $\alpha=0.10$). TRANSFERRING ONE BENEFIT MODEL ACROSS ADAPTERS VIOLATES $FA_u \leq \alpha$ IN MOST DIRECTIONS—PER-ADAPTER CALIBRATION IS REQUIRED.

fit $\rightarrow$ apply	regret $\downarrow$	FA_u	adapt	coverage
Tent $\rightarrow$ EATA	0.0020	0.000	0.54	0.63
EATA $\rightarrow$ Tent	0.0010	0.007	0.55	0.69
Tent $\rightarrow$ SAR	0.0173	0.125	0.50	0.72
EATA $\rightarrow$ SAR	0.0236	0.192	0.57	0.62
SAR $\rightarrow$ Tent	0.0085	0.255	0.84	0.84
SAR $\rightarrow$ EATA	0.0027	0.116	0.75	0.75

IX. DISCUSSION AND FAILURE MODES

A. Why Natural Shifts Mainly Produce No-Harm

A single natural benchmark is often one-sided: adaptation helps almost everywhere or harms almost everywhere. In the first case, always-adapt is already near-oracle; in the second, always-freeze is near-oracle. A three-way router gains a strict margin over both fixed policies only when helpful and harmful conditions coexist and the evidence distinguishes them. Natural no-harm results and mixed-regime wins are therefore different operating points of the same decision geometry.

B. Four Sources of Abstention

Not every abstention has the same interpretation. *Structural non-identifiability* means opposite-benefit worlds share the same admissible evidence law. *Finite-sample uncertainty* means the sign is identifiable in principle but the interval still crosses zero. *Estimator inadequacy* means the chosen model fails to extract an available signal. *Calibration failure* means a development relationship does not transfer to deployment. We distinguish these causes throughout the evaluation rather than treating every null result as evidence of structural unknowability.

C. Operational Failure Modes

Table XX maps each deployment failure to a diagnostic and a safe response.

D. When to Deploy K-Bound

K-Bound is most valuable when harmful updates are plausible, the cost of a bad commit is high, a frozen fallback is available, and historical or structural information supports calibration. It is less useful when adaptation is

TABLE XX

DEPLOYMENT FAILURES, DIAGNOSTICS, AND SAFE RESPONSES.

Failure	Diagnostic	Safe response
Calibration transfer fails	support/residual audit	recalibrate, widen radius, or abstain
True drift exceeds β	external domain monitoring	revise class/budget; freeze meanwhile
Evidence out of support	distance or density check in Z-space	abstain
Adapter/checkpoint changed	configuration hash mismatch	require recalibration
Batch too small	minimum-size rule	wait for more data or freeze
Continual state accumulates error	checkpoint/stream monitor	rollback to last certified state
Missing/invalid evidence	schema and numerical validation	freeze

uniformly benign, when no credible drift class or calibration set exists, or when the candidate update cannot be evaluated and rolled back safely.

X. LIMITATIONS

Scope of empirical decision evidence. On the stress grid KGA exercises selective adapt/freeze/abstain behavior for Tent/EATA (Table IX); CIFAR-10-C SAR is withheld. On ImageNet-C the pattern flips by candidate—Tent/EATA are robust while SAR collapses and KGA retains low regret under a mechanism-faithful re-run (Table XIII). Per-corruption CIFAR-10-C and helpful natural shifts are insurance regimes, not policy wins.

Real-shift scope. The uniform panel deliberately distinguishes three outcomes. Camelyon17 genuine OOD and RxRx1 are locked no-harm results: KGA ties the better fixed policy at zero observed false-adapt. Office-Home and iWildCam are locked from their out-of-fold bootstrap seal (App. J; Office-Home promotes OOF no-harm only; LOO beats-both is not claimed). PACS and ImageNet-R are completed, locked null diagnostics; CIFAR-10.1 is a locked diagnostic fail. The three-source mixture is a constructed routing aggregate; it is not a transfer result and does not change the per-dataset conclusion.

Theory and calibration. The conformal radius assumes exchangeable calibration pairs; sign bracketing without structure is impossible (Theorems 1 and 2; Appendix C). KGA’s value scales with how often catastrophic, detectable harm occurs, not with universal accuracy gains. On helpful-dominated shifts the resulting near-tie with always-adapt reflects the certificate abstaining rather than committing harmful updates: KGA approximately matches the better policy, at most paying a small abstention cost when it declines a helpful but uncertifiable update, so the value proposition is conditional, like insurance against detectable harm. *Setting β .* The budget is *declared*, not estimated from unlabeled data—and provably so: label-free audits of β are vacuous over the unrestricted class (Appendix L). It can, however, be *audited*: a labeled audit budget, a Lipschitz-transport assumption, or an unknown-

direction index-drift assumption yields a computed upper bound with finite-sample validity, at the price of the named side information. Absent an audit, β is fixed from domain knowledge, historical dev-to-deployment gaps, or an explicit stress envelope. The choice $\beta=0$ is the strongest zero-drift assumption and may be reported only as a plug-in baseline; it is not an assumption-free or conservative default. When no credible bound is justifiable, the population frontier cannot certify robustness to calibration drift, and the safe operational response is to abstain or freeze. Larger β widens the abstain band monotonically, trading decision coverage for robustness.

Detectability and transfer. ImageNet-R is a weak, split-specific evidence regime: no candidate survives both comparisons across four seeds. CIFAR-10.1 fails the cross-seed certificate bar ($\text{FA}_u = 0.167$ and descriptive $\text{FA}_c = 0.444$). These are useful boundary cases, not failed headline experiments: they show that a low-margin signal should not be promoted merely because one split or one candidate looks favorable.

Theory limits versus empirical limits. Some abstention is compatible with the population frontier, whereas baseline faithfulness, seed count, estimator quality, and replay status are ordinary empirical limitations requiring further validation. Tying always-adapt on a helpful-dominated shift is the no-harm guarantee *holding*: one cannot beat an oracle that is already always-adapt. Declining ImageNet-R is consistent with a weak-evidence, low-margin regime; the null result does not by itself establish structural non-identifiability under Theorem 2 (it could also reflect weak evidence, poor benefit estimation, or calibration-transfer failure). Abstention in the structurally ambiguous closed band prevents a uniformly sound strict commitment over the declared class. Empirical abstention, however, may instead stem from an overly wide radius or a weak estimator and therefore does not establish that a tested shift lies in that structural band. Richer label-free evidence may improve benefit estimation and reduce avoidable abstention, while better-justified or more data-efficient calibration can narrow finite-sample uncertainty; neither removes the population limitation without strengthening the declared assumptions or evidence.

Baseline scope and recomputation. On the ImageNet-C SAR configuration, the committal confidence gate adapts on harmful cells while KGA abstains on the low-margin band. The POEM/AETTA comparison is complete only for protocol-matched ports; official implementations remain future baseline work. The Office-Home/iWildCam summaries have a separate saved recomputation lock, but no external-lab replication. PACS and ImageNet-R are completed registered diagnostics and are retained as null/weak-evidence results rather than tuned into wins. The frozen external procedure is documented in `INDEPENDENT_REPLICATION_PROTOCOL.md`. Abstention lowers *decision* coverage but never withholds a prediction: it retains the frozen fallback f_0 .

What we do not claim. Universal accuracy gains; adapter selection by test regret; multicandidate routes with violated false-adapt bounds; or wins on helpful-dominated shifts where always-adapt is already oracle-optimal.

XI. REPRODUCIBILITY

A. Artifacts and Result Lineage

All locked/reconciled rows in Tables XV–XVI link to saved raw cells or seed aggregates in the anonymized repository. Provisional rows are retained only to document completed evidence and the required replay; they are not reused as headline claims. The core certificate module is packaged as an installable module in the anonymized repository.

B. Formal Verification Scope

The repository contains a kernel-checked Lean 4 formalization of indexed algebraic, finite-decision, measure-containment, and conditional uniform-rank results, with no `sorry`, `admit`, or project-defined axioms. It checks frontier sufficiency, all three deterministic branches, algebraic witnesses for the closed and open bands, both zero-versus-strict boundary witnesses, and the marginal false-adapt/false-freeze implication once coverage is assumed. The current development does not formalize the target-law witness construction, equality of its induced evidence laws, membership of those laws in $\mathcal{C}_\beta$, realization of the algebraic witnesses by measurable target laws, distributional frontier necessity or pointwise maximality, risk alignment, multiclass identifiability, calibration transfer, or a general exchangeable-process lift. These clauses are proved only in the manuscript; compilation of the mapped helper declarations is not described as kernel verification of the complete paper theorem. The implementation theorem uses the same exact rank rule as the current clean-split scorer; historical benchmark artifacts produced with an interpolated empirical quantile are not presented as kernel-certified split-conformal experiments. To our knowledge, machine-checked conformal-coverage results remain uncommon, but no priority claim is made.

C. Claim-to-Support Map

XII. CONCLUSION

K-Bound reframes test-time adaptation as a validity decision rather than an unconditional update. The central object is the target benefit $\Delta = R_T(f_0) - R_T(f_a)$. Over a declared drift-budget class, a strict adapt or freeze commitment is uniformly supportable if and only if the observable disagreement-region margin satisfies $|M| > \beta$; on $|M| \leq \beta$, abstention is the maximal sound three-way action under the paper’s strict-decision semantics.

KGA operationalizes this principle with a label-free benefit estimate and a pre-calibrated uncertainty radius. The resulting controller commits the candidate only when the interval excludes zero, retains the frozen model when harm is certified, and otherwise abstains from changing

TABLE XXI

Claim-to-support map. “THEOREM” MEANS PROVED IN THE MANUSCRIPT; LEAN COVERAGE IS LIMITED TO THE INVENTORY DESCRIBED ABOVE. “EMPIRICAL” AND “DIAGNOSTIC” CLAIMS ARE RESTRICTED TO THE NAMED SAVED ARTIFACTS AND PROTOCOLS.

Claim	Type	Exact requirement	Evidence location	Caveat
Interior matched-evidence impossibility	theorem	$\beta > 0, M < \beta$, declared class contains the constructed worlds	Thm. 1; App. C	Opposite nonzero benefits are an interior result only.
Closed-band abstention	theorem	Strict three-way soundness over $ \gamma \leq \beta$	Prop. 2	Boundary is zero-versus-strict ambiguity, not opposite signs.
Strict-commitment frontier	theorem	Declared drift class and risk alignment	Thm. 2	Uniform strict commitment, not universal sign recovery.
Marginal FA _u certificate	theorem	$\mathbb{P}(\hat{\Delta} - \Delta \leq \varepsilon) \geq 1 - \alpha$	Thm. 3; Lean inventory	Does not control conditional FA _c .
CIFAR-10-C mixed-regime action validity	empirical	Five-seed Tent/EATA raw decisions, action rates, false-adapt, and paired CIs	Table IX; empirical audit manifest	SAR is withheld after a replay mismatch and contributes no evidence.
ImageNet-C authoritative result	empirical	27 cells $\times$ 5 seeds, SAR candidate, pooled seed-averaged paired bootstrap	Table XIII; pooled bootstrap JSON	Per-seed CIs exclude 0 on 2/5 at $n=27$; point-level 5/5.
Natural no-harm results	empirical	Reconciled OOF or locked held-out artifacts	Tables XV–XVI	One-sided regimes mainly exercise the adapt or freeze branch.
Controlled multimodal routing	empirical	Pre-registered 13 $\times$ 10 injected-corruption grid	App. G; KB-CLAIM-027	Mechanism confirmation, not a natural multimodal benchmark.
POEM/AETTA comparison	empirical	Same logged stress cells and protocol-matched ports	Table XI	Not official implementations.
Constructed heterogeneous mixture	empirical	Researcher-constructed three-source routing aggregate	Sec. VII-I	Not unseen-shift transfer.
Universal improvement	diagnostic	No supporting requirement is declared	Sec. X	Explicitly not claimed.
Universal accuracy improvement	diagnostic	Would require a different objective and task-specific held-out evidence	Natural-shift tables	Explicitly not claimed.
Real-camera study	pending	Fresh preregistered held-out sessions and raw logs	pre-registration templates	Templates are not evidence.

state. Across the benchmark panel, the observed behavior follows the regime map: controlled mixed regimes—including the pre-registered two-view multimodal check of Appendix G—exercise all three actions with low observed false adaptation and reduced regret, one-sided natural shifts primarily exercise the safe adapt or freeze branch, and weak or non-transferable evidence produces abstention or remains diagnostic rather than being promoted as a win.

The broader implication is that an adaptation system should expose two components: an update mechanism and a validity layer that decides whether the update is supportable under the available evidence and assumptions. Strengthening K-Bound therefore requires not only better adapters, but also richer risk-aligned evidence, more defensible calibration across shifts, explicit drift-budget auditing, and reproducible deployment semantics.

APPENDIX A

CALIBRATION AND HELD-OUT EVALUATION

Algorithm 2 gives the evaluation protocol behind the uniform nine-track panel (Table XV) and the primary numeric evidence (Table XVI). Its two branches are *not* interchangeable: stress grids are scored cell-by-cell with leave-one-cell-out cross-fitting (no single deployment estimator exists), whereas natural-shift tracks fit one estimator and radius on the development split and score the untouched held-out domains exactly once.

APPENDIX B

EVIDENCE FEATURE SCHEMA AND ADAPTER HYPERPARAMETERS

Two label-free evidence panels are used, matched to the two calibration regimes of §V. Notation: p_0, p_a are the frozen and adapted softmax rows on the adaptation batch; $\bar{m}_0 = \text{mean}_n p_{0,n}$ and $\bar{m}_a$ are the predicted-class marginals; $H(p) = -\sum_c p_c \log p_c$; C is the number of classes. The stress-grid tracks (CIFAR-10-C, ImageNet-C, PACS) use the 11-dimensional base panel of Table XXII (kbound_pkg/kbound/evidence.py). The natural-shift tracks (Camelyon17, iWildCam, Office-Home, RxRx1) use the 16-dimensional panel of Table XXIII: a 10-statistic frozen-logit summary plus a 6-feature rich panel (scripts/run_wilds_camelyon17.py). All features are target-label-free; the ATC feature calibrates its threshold on *source* labels only.

APPENDIX C

COMPLETE CORE PROOFS

This appendix is not required for the main empirical claim; it preserves the full theorem stack and validator mapping behind the three main-paper theorems: matched-evidence impossibility (Theorem 1), the benefit-sign frontier (Theorem 2), and the finite-sample certificate (Theorem 3). The main text contains the complete algebraic proofs; the material below records only the supporting matched-evidence form and zero-drift interpretation used by those proofs.

Algorithm 2 Calibration and evaluation for K-Bound (grid vs. natural-shift branches)

Require: Conditions $\mathcal{C}$; frozen f_0 ; adapters $\mathcal{A}_1, \dots, \mathcal{A}_K$; track type (stress grid or natural shift); level α

Ensure: Regret-to-oracle, false-adapt rate, coverage, and action composition

- 1: **for all** $c \in \mathcal{C}$, adapters $\mathcal{A}_k$ **do** ▷ shared evidence logging
- 2: $f_a^{(k,c)} \leftarrow \mathcal{A}_k(f_0, X_c)$; $Z_{k,c} \leftarrow \phi(X_c, f_0, f_a^{(k,c)})$;
 $\Delta_{k,c} \leftarrow R_c(f_0) - R_c(f_a^{(k,c)})$
- 3: **end for**
- 4: **if stress grid** (CIFAR-10-C, ImageNet-C, PACS): $\mathcal{C}$ is one grid of cells **then**
- 5: **for all** cells $i \in \mathcal{C}$ **do**
- 6: Fit $\hat{\Delta}_{-i}$ on the other $N-1$ cells; $\rho_i \leftarrow |\hat{\Delta}_{-i}(Z_i) - \Delta_i|$ ▷ leave-one-cell-out
- 7: **end for**
- 8: $\varepsilon \leftarrow \rho_{(k)}$, $k = \lceil (N+1)(1-\alpha) \rceil$ ▷ exact split-conformal rank quantile
- 9: **for all** cells $i \in \mathcal{C}$ **do**
- 10: Run Algorithm 1 with $(f_0, \mathcal{A}_k, X_i, \hat{\Delta}_{-i}, \varepsilon)$ ▷ each cell scored by the estimator that never saw it (cross-fitted)
- 11: **end for**
- 12: **else natural shift** (Camelyon17, iWildCam, Office-Home, RxRx1): split $\mathcal{C} = \mathcal{C}_{\text{cal}} \dot{\cup} \mathcal{C}_{\text{test}}$ across domains
- 13: Form out-of-fold residuals on $\mathcal{C}_{\text{cal}}$; $\varepsilon \leftarrow \rho_{(k)}$, $k = \lceil (N+1)(1-\alpha) \rceil$ (exact split-conformal rank quantile)
- 14: Fit deployment estimator $\hat{\Delta}(Z)$ on all of $\mathcal{C}_{\text{cal}}$; freeze $(\hat{\Delta}, \varepsilon)$
- 15: **for all** $c \in \mathcal{C}_{\text{test}}$ **do**
- 16: Run Algorithm 1 with $(f_0, \mathcal{A}_k, X_c, \hat{\Delta}, \varepsilon)$ ▷ untouched held-out domains, scored once
- 17: **end for**
- 18: **end if**
- 19: Record decision, regret-to-oracle, coverage, false-adapt indicator per condition
- 20: Aggregate metrics vs. always-adapt, always-freeze, per-condition oracle

a) *Core-theory inventory.*: The main text gives complete proofs of disagreement reduction, interior impossibility, boundary abstention, the strict-commitment frontier, and the coverage-based certificate. This appendix records the matched-evidence full form and the zero-drift interpretation. Minimax, one-bit, capacity, and rate extensions are kept in the long manuscript and are not part of the short-paper claim stack.

Theorem 4 (Matched evidence impossibility (full form)). *Theorem 1 is parts (i)–(ii) of the following. Fix $\alpha \in (0, \frac{1}{2})$.*

- (i) *Witness worlds with $\text{Law}(Z \mid P_T^1) = \text{Law}(Z \mid P_T^2)$ and $\text{sign } \Delta_1 = -\text{sign } \Delta_2 \neq 0$.*
- (ii) *Le Cam identity: $\inf_g \text{err}(g) = 1 - \text{TV}$ between the evidence laws, where $\text{err}(g)$ is the total test error*

TABLE XXII
BASE 11-DIM EVIDENCE PANEL (STRESS GRIDS).

Feature	Definition	Family
pre_entropy	$\text{mean}_n H(p_{0,n})$	frozen output
pre_conf	$\text{mean}_n \max_c p_{0,n}$	frozen output
pre_pbal	$H(\bar{m}_0) / \log C$	frozen marginal
post_entropy	$\text{mean}_n H(p_{a,n})$	adapted output
post_conf	$\text{mean}_n \max_c p_{a,n}$	adapted output
post_pbal	$H(\bar{m}_a) / \log C$	adapted marginal
pbal_drop	$\text{pre_pbal} - \text{post_pbal}$	change
entropy_drop	$\text{pre_entropy} - \text{post_entropy}$	change
frac_highconf	$\text{frac}_n(\max_c p_{a,n} > 0.9)$	collapse
marginal_KL	$\sum \bar{m}_{a,c} \log(\bar{m}_{a,c} / \bar{m}_{0,c})$	drift
update_norm	$\ \Delta \theta_{\text{aff}}\ _2$	update

TABLE XXIII
RICH 16-DIM PANEL (NATURAL SHIFT): 10 FROZEN-LOGIT STATISTICS + 6 RICH FEATURES. H IS SOFTMAX ENTROPY, $s = \max_c \text{softmax}$ THE CONFIDENCE, $E(x) = -\log \sum_c e^{z_c}$ THE FREE ENERGY; SUBSCRIPTS s, t DENOTE SOURCE/TARGET POOLS.

Feature	Definition	Family
<i>Frozen-logit summary (10)</i>		
ent mean/std	mean/std H	frozen
ent p25/p75	25/75 pct. of H	frozen
conf mean/std	mean/std s	frozen
frac hi-conf	$\text{frac}(s > 0.9)$	frozen
frac hi-ent	$\text{frac}(H > \log 2)$	frozen
marg. max	$\max_c \text{mean}_n p$	frozen marginal
marg. std	$\text{std}_c \text{mean}_n p$	frozen marginal
<i>Rich features (6)</i>		
disagree	$\text{frac}(\arg \max f_0 \neq \arg \max f_a)$	disagreement
ent_gap	$\text{mean } H(f_0) - \text{mean } H(f_a)$	change
en_shift	$\text{mean } E_t - \text{mean } E_s$	drift
bn_kl	$\text{mean}_{\text{ch}} \text{KL}(\mathcal{N}_{\text{run}} \parallel \mathcal{N}_{\text{batch}})$	drift
atc	$\text{frac}_t(s \geq t)$, t src-calib.	accuracy proxy
conf_drop	$\text{mean } s_s - \text{mean } s_t$	drift

(distinct from the observable margin M).

- (iii) *Forced abstention at rate $\geq 1 - 2\alpha$ under dual error control.*

Corollary 2 (Closed-band abstention under dual error control). *In the $|M| \leq \beta$ wedge of Theorem 2, Theorem 4(iii) implies any label-free rule with false-adapt and false-freeze $\leq \alpha$ must abstain with probability at least $1 - 2\alpha$ on matched evidence. Opposite nonzero signs are required only for $|M| < \beta$; at $|M| = \beta > 0$ zero-versus-strict ambiguity suffices, and at $M = \beta = 0$ the benefit is identically zero so neither strict action is sound.*

Remark 4 ($\beta = 0$ face). On the $\beta = 0$ face the class $\mathcal{C}_0$ forces $\gamma = 0$, so Lemma 1 gives $\text{sign } \Delta = \text{sign } M$ exactly, and the frontier of Theorem 2 specializes as follows. If $M \neq 0$, the strict action of $\text{sign } M$ (ADAPT for $M > 0$, FREEZE for $M < 0$) is uniformly sound over $\mathcal{C}_0$. If $M = 0$, then $\Delta = 0$ identically, so under the strict-decision semantics neither strict action is sound and the maximal sound action is ABSTAIN. Existing label-free estimators may be interpreted as $\beta = 0$ plug-ins only when their estimand realizes the stated disagreement-score decomposition; assuming $\gamma = 0$ alone does not establish their soundness.

TABLE XXIV

PER-TRACK ADAPTER HYPERPARAMETERS (FROM THE RUN CONFIGS). ADAPTED PARAMETERS ARE BN/LAYER-NORM-AFFINE ONLY; TTA OPTIMIZER IS ADAM. “AGGRESSIVENESS” CELLS USE MILD = (10 STEPS, LR 10^{-3}) OR AGGRESSIVE = (50 STEPS, LR 2.5×10^{-3}). EATA USES AN ENTROPY FILTER $e_{\text{thr}} = 0.4 \ln K$ (≈ 0.92 FOR CIFAR-10, 2.76 FOR IMAGENET); SAR ADDS SAM + ENTROPY RESET. BACKBONES FOR CIFAR-10-C, IMAGENET-C, PACS, AND CAMELYON17 ARE VERIFIED FROM THE RUN CODE; THE REMAINING WILDS/DOMAINBED TRACKS USE THE STANDARD BASELINE.

Track	Backbone	Adapter(s)	Steps / lr
CIFAR-10-C stress	ResNet-18	Tent/EATA/SAR	mild & aggressive (per cell)
ImageNet-C	ResNet-50	Tent/EATA/SAR (SAM+reset)	mild & aggressive (per cell)
PACS (LODO)	ResNet-18	mild/aggressive	$10/10^{-3}$ & $50/2.5 \times 10^{-3}$
Camelyon17	DenseNet-121	EATA online	$10 / 10^{-3}$
iWildCam	ResNet-50	Tent episodic	$10 / 10^{-3}$
Office-Home	ResNet (DomainBed)	SAR online aggr.	$50 / 2.5 \times 10^{-3}$
RxRx1	ResNet-50	episodic	5 / (batch 64)

b) Asymmetric fixed-policy regret characterization (secondary).: For completeness, a router strictly improves both fixed-policy regrets only when the optimal action varies *and* is label-free detectable. The following proposition makes the condition exact and, crucially, *asymmetric*, which explains why natural-shift tracks tie (rather than beat) always-adapt. Let δ^* be the strict-commitment frontier (Theorem 2) with abstention resolved to FREEZE, and let $R(\pi) = \mathbb{E}_Q[|\Delta| \mathbf{1}\{\pi \neq \text{oracle}\}]$ be regret-to-oracle over a deployment mixture Q .

Proposition 3 (Asymmetric fixed-policy regret). *Since $\Delta > 0$ on $\{M > \beta\}$ and $\Delta < 0$ on $\{M < -\beta\}$ for every admissible drift (Theorem 2). The excess regret of always-freeze is exactly the detectable-helpful mass $\mathbb{E}_Q[|\Delta| \mathbf{1}\{M > \beta\}]$. For always-adapt it is the detectable-harmful mass plus harmful-minus-helpful regret inside $\{|M| \leq \beta\}$. Hence δ^* weakly dominates always-freeze, strictly iff detectable-helpful mass $Q(M > \beta) > 0$; but δ^* beats always-adapt iff the certified harmful-mass win exceeds the helpful adaptation forgone in the abstention region. Strict improvement over both fixed policies thus requires both detectable-helpful and net-detectable-harmful mass.*

Because abstention defaults to FREEZE, the freeze side is free but the adapt side is conditional: a helpful-dominated shift with negligible detectable-harmful mass ties or slightly loses to always-adapt. This predicts the candidate-dependent Camelyon17 outcome (Table XIV): the harmful-mass-bearing arm wins narrowly while the helpful-dominated arms tie. Full partition proof and a 25-seed numerical check are in the repository note [theory_v2/minimax_frontier/](#). This sharpens, and is consistent with, the frontier maximality already established in Theorem 2.

APPENDIX D

MULTICLASS AND REGRESSION DERIVATIONS

For multiclass 0/1 loss, errors cancel outside $D = \{f_0 \neq f_a\}$, giving $\Delta = \mu_T(D)(p_a - p_0)$ with $p_j = \mathbb{P}_T(f_j = Y | D)$. In contrast to binary classification, p_0 need not equal $1 - p_a$. If $p_a - p_0 = M_{\text{mc}} + \gamma_{\text{mc}}$ and $|\gamma_{\text{mc}}| \leq \beta_{\text{mc}}$, then $|M_{\text{mc}}| > \beta_{\text{mc}}$

is sufficient for a strict commitment; necessity additionally needs the declared class to realize the admissible evidence-identical drifts.

For squared-loss regression, let $m_T(x) = \mathbb{E}_T[Y | X = x]$. Direct expansion gives

$$\Delta = \mathbb{E}_T[(f_a - f_0)(2m_T - f_0 - f_a)].$$

The integrand vanishes where $f_a = f_0$, so the comparison is again disagreement-local. For an observable proxy $u(x)$, define

$$M_{\text{reg}} = \mathbb{E}_T[(f_a - f_0)(2u - f_0 - f_a)],$$

$$\gamma_{\text{reg}} = 2\mathbb{E}_T[(f_a - f_0)(m_T - u)].$$

Then $\Delta = M_{\text{reg}} + \gamma_{\text{reg}}$. A declared bound $|\gamma_{\text{reg}}| \leq \beta_{\text{reg}}$ yields the same sufficient strict-commitment rule $|M_{\text{reg}}| > \beta_{\text{reg}}$; a converse again requires a richness assumption on the target class.

APPENDIX E

HISTORICAL PROTOCOL RECONCILIATION

The main text contains the complete empirical accounting in the uniform nine-track panel. Historical route studies and superseded per-dataset tables are deliberately excluded from this short-paper build. Their raw logs remain in the repository for audit, but they are not independent empirical claims. The declared evaluation unit is a condition or evaluation cell (cross-fitted on the stress grids, held-out on the natural-shift tracks), never an arbitrary pool of frames or domains; all policies replay the same unit stream, and labels are revealed only after the KGA decision for offline scoring.

APPENDIX F

PHYSICAL-STUDY PRE-REGISTRATION

The two-phone package-inspection protocol is locked before fresh capture. Table XXV records the split and deployment design; Table XXVI fixes the primary endpoint layout. Every result cell remains blank until the exact physical inventory, source-model gate, development seal, held-out chronology, replication replay, and anti-leakage audit pass the machine-readable publication gate. Browser previews, mock clips, and reconstructed media are software diagnostics, not evidence.

APPENDIX G

CONTROLLED MULTIMODAL MECHANISM CHECK

Protocol D33 is a pre-registered controlled two-view MNIST experiment: each digit is split into left- and right-half modalities, and the right modality is corrupted by Gaussian noise over 13 severities and 10 batches, yielding 130 conditions. This construction deliberately creates the named regime in which fusion changes from helpful to harmful and the change is visible in label-free reliability evidence. Table XXVII reports the locked result.

The result confirms that the controller can exercise all three actions and improve on both fixed policies when

TABLE XXV

R1 – Pre-registered real-world deployment protocol. THE SPLIT AND ONLINE-LABEL CONSTRAINTS ARE FIXED BEFORE PHYSICAL HELD-OUT CAPTURE. SESSION IDENTIFIERS ARE POPULATED FROM THE SIGNED RECORDING MANIFEST.

Item	Locked design
Task	Four-class package/label inspection: correct, missing label, misaligned label, damaged label
Camera	Two phone cameras and real physical objects; the second phone is reserved for external-device replication
Base model	MobileNetV3-Small
Candidate adaptation	Episodic Tent; one update step per decision window
Decision window	32 frames
Evidence	14 label-free frozen/candidate disagreement and shift features
Calibration-fit sessions	— (distinct recorded session; IDs locked in manifest)
Conformal calibration sessions	— (different day/session from calibration-fit)
Held-out physical test sessions	— (different day/session; inaccessible during tuning)
Physical shifts	Lighting, shadow, blur, background, viewpoint, distance, and batch composition
Policies compared	Always-freeze, always-adapt, confidence gate, entropy gate, KGA without radius, KGA certificate
Deployment labels used live?	No; labels are joined only during offline evaluation
Miscoverage level	$\alpha = 0.10$
Official live model	Frozen fallback unless the candidate is certified for adaptation

TABLE XXVI

R2 – Primary held-out real-world result (RESULT PENDING). ALL POLICIES WILL BE SCORED ON THE IDENTICAL LOCKED PHYSICAL STREAM AFTER REAL CAPTURES PASS THE SOURCE-MODEL GATE (≥ 0.80 BALANCED ACCURACY AND MACRO-F1 ON S02). UNTIL THEN EVERY METRIC CELL READS *RESULT PENDING* — *NO MEASURED DATA AVAILABLE*; DEVELOPMENT REPLAYS ON MOCK CLIPS ARE AUDIT-ONLY AND MUST NOT BE CITED.

Policy	Balanced acc. $\uparrow$	Macro-F1 $\uparrow$	Regret $\downarrow$	FA_u $\downarrow$	FA_c $\downarrow$	Adapt rate	Abstain rate	Mean latency (ms) $\downarrow$
Always-freeze	—	—	—	—	—	—	—	—
Always-adapt	—	—	—	—	—	—	—	—
Confidence gate	—	—	—	—	—	—	—	—
Entropy gate	—	—	—	—	—	—	—	—
KGA, no radius	—	—	—	—	—	—	—	—
KGA certificate	—	—	—	—	—	—	—	—

$FA_u = \mathbb{P}(\text{adapt} \wedge \Delta \leq 0)$ is the certificate-aligned unconditional false-adapt rate; $FA_c = \mathbb{P}(\Delta \leq 0 \mid \text{adapt})$ is descriptive. Ground-truth labels are unavailable during live decisions and are revealed only for offline evaluation. Latency reporting includes candidate adaptation, the KGA decision, and full end-to-end window latency.

TABLE XXVII

Controlled multimodal Protocol D33 (130 conditions).

ACCURACY IS AVERAGED OVER THE FIXED CORRUPTION GRID. THE SUPERIORITY PROBABILITIES ARE PAIRED-BOOTSTRAP QUANTITIES SAVED BY THE LOCKED ANALYSIS. THIS IS A CONTROLLED MECHANISM CONFIRMATION, NOT A NATURAL-SHIFT RESULT.

Policy or diagnostic	Result
Always fuse (adapt)	0.5832
Always single-A (freeze)	0.8536
KGA routed	0.8568
Oracle	0.8573
Actions (adapt / freeze / abstain)	9/119/2
Observed false-adapt	0/130
$P(\text{KGA} > \text{fuse}) / P(\text{KGA} > \text{single-A})$	1.000/1.000

the helpful-to-harmful transition is present and detectable. Its scientific scope is narrow: the corruption is injected, the gain over the safe single-modality policy is 0.0032, and the experiment does not establish a natural multimodal win or universal accuracy improvement. The protocol was pre-registered in `research_lock/CONTROLLED_MULTIMODAL_PROTOCOL_D33_v1.yaml` (sealed 2026-06-15, before any result was computed); the authoritative

artifact is `experiments/kbound/results/controlled_multimodal_d33/results.json`, and the claim ledger records this result as KB-CLAIM-027 (supported).

APPENDIX H RUNTIME DETAILS

The adapter update requires forward and backward computation before the controller can decide. The saved controller microbenchmark reports evidence extraction, benefit-model evaluation, rollback copy memory, and benefit-model size; these are the added-cost quantities reported in Table VI. A complete end-to-end profile separating frozen inference, candidate adaptation, candidate inference, rollback time, and total window latency remains pending and is not inferred from the controller-only microbenchmark.

APPENDIX I FORMALIZATION INVENTORY

Lean checks the indexed algebraic, finite-decision, measure-containment, and conditional uniform-rank results listed in `formal/KBound/TheoremMap.lean`. These declarations verify the named helper results and clauses;

TABLE XXVIII

PER-SEED SAR DETAIL (IMAGENET-C, 27 CELLS/SEED, $\alpha=0.1$). GAP COLUMNS ARE 95% PAIRED-BOOTSTRAP UPPER BOUNDS (NEGATIVE = CI EXCLUDES ZERO).

seed	harmful	regret KGA	gapadapt hi	gapfreeze hi	FA _u
0	0.148	0.0108	−0.0011	−0.0111	0.000
1	0.222	0.0091	−0.0047	−0.0105	0.000
2	0.111	0.0128	+0.0007	+0.0148	0.037
3	0.148	0.0056	+0.0041	−0.0143	0.000
4	0.148	0.0154	+0.0063	−0.0115	0.000
pooled	0.156	0.0107	−0.0030	−0.0093	0.007

successful compilation does not by itself certify a whole paper theorem that also contains unformalized clauses. In particular, Lean does not formalize the explicit target-law witness construction, equality of the induced evidence laws, membership of those laws in $\mathcal{C}_\beta$, the zero-versus-strict boundary construction, frontier necessity/maximality, risk alignment, calibration transfer, or the lift from a conditional uniform-index model to arbitrary exchangeable deployment processes.

APPENDIX J

CLAIM-TO-ARTIFACT MAP

The machine-readable manifest `paper/generated/kbound_result_manifest.json` is the authoritative index for every promoted number, seed count, protocol, quantile convention, and known replay caveat. The derived audit `paper/generated/empirical_audit/decision_metrics.json` adds `adapt/freeze/abstain` counts and rates, 95% Wilson intervals for action rates and marginal false adaptation, empirical interval-hit coverage where the raw fields were retained, the exact source artifact, and a reproduction command. Observed interval-hit coverage is explicitly marked *empirical only*; it neither verifies exchangeability nor establishes the theorem’s coverage premise or risk alignment. Historical natural-shift summaries that did not retain interval-hit records are marked `not_retained`, not assigned a synthetic coverage value. The compact claim matrix `paper/generated/empirical_audit/claim_matrix.md` records claim, dataset, seeds, artifact, paper consumer, and closure status. Table XXI gives the reader-facing summary.

APPENDIX K

IMAGENET-C MULTI-SEED DETAIL

Table XXVIII reports the SAR row of Table XIII per seed. Point estimates improve both fixed-policy regrets on every seed; the per-seed 27-cell paired bootstrap excludes zero on both gaps for seeds 0–1 and narrowly includes zero on the remaining seeds, whose realized harmful mass is lighter—consistent with the regime map: the certified win scales with detectable harmful mass. The pooled seed-averaged bootstrap (the same design as the CIFAR-10-C rows) excludes zero on both gaps. $\text{FA}_u \leq \alpha$ at every seed.

APPENDIX L

AUDITABLE DRIFT BUDGETS: STATEMENTS

The drift budget β of the population frontier is declared in the main text. The following results (proofs in the long manuscript, appendix “Auditable drift budgets”; synthetic validation in the repository artifact `audit_validation/results_audited_drift.json`) settle when a *computed* budget can replace it. Write $u = \eta_a - s \in [-1, 1]$ on D and $\gamma(P) = \mathbb{E}_{\mu_P}[u \mid D]$.

Theorem 5 (Vacuity of label-free audits). *Fix the observables (μ_T, f_0, f_a, s) , hence M . Over the class of all target laws with these observables, any label-free audit $\hat{\beta} = g(Z)$ satisfying $\mathbb{P}(|\gamma| \leq \hat{\beta}) \geq 1 - \delta$ uniformly must satisfy $\hat{\beta} \geq \frac{1}{2} + |M|$ with probability $\geq 1 - \delta$; the audited frontier then abstains in every world. (Corollary of Theorem 1’s label-kernel construction.)*

Theorem 6 (Computed budgets under purchasable structure). *Each of the following side informations yields a computable $\hat{\beta}$ with $\mathbb{P}(|\gamma(P_T)| \leq \hat{\beta}) \geq 1 - \delta$: (i) n deployment labels (Hoeffding; width $\sqrt{2 \ln(2/\delta)/n}$); (ii) an L -Lipschitz drift on a bounded 1-d representation (Kantorovich–Rubinstein + DKW, using only unlabeled batches and labeled calibration data); (iii) unknown-direction index drift $u = g(\theta_*^\top \psi)$ on $\|\psi\| \leq B$ (rotation-uniform adaptive max-sliced audit at rate $O(\sqrt{(d + \ln(1/\delta))/n})$ via halfspace-VC uniformity; sharper constants from known max-sliced concentration). More generally, any symmetric witness class of Rademacher/VC complexity $O(\sqrt{v/n})$ yields a useful root- n integral-probability-metric audit whenever the unknown drift lies in that class (long-manuscript Theorem Aud-H); full Lipschitz W_1 witnesses are excluded from that rate. The structural assumptions are declared and falsifiable, not verifiable—some such purchase is unavoidable by Theorem 5.*

Theorem 7 (Composition; fully empirical rule). *For any audit with $\mathbb{P}(|\gamma| \leq \beta) \geq 1 - \delta$: the audited rule (adapt iff $M > \hat{\beta}$, freeze iff $M < -\hat{\beta}$, else abstain) falsely strict-commits with probability $\leq \delta$; replacing M by the batch estimate $\bar{M} \pm t_M(m, \delta')$ gives a rule computable entirely from data with false-commit probability $\leq \delta + \delta'$. On the synthetic sweep: 0/12 false commits (empirical rule) and 0/13 (population rule) versus 3/20 for the $\beta=0$ plug-in.*

Theorem 8 (Domain-level verifiability with an exact floor). *With K calibration domains exchangeable with deployment, common per-domain labeled size, and a drift predictor independent of all $K+1$ domains, a domain-level conformal bound certifies $|\gamma_{K+1}| \leq |\gamma_{\text{pred}}(Z_{K+1})| + s_{(j)} + \epsilon$ at level $1 - \alpha$, feasible iff $\alpha \geq \frac{1}{K+1} + \delta$. Alignment is thus verifiable exactly at the rate one can afford exchangeable domains ($K=5$: nothing below $\alpha \approx 0.17 + \delta$); per-deployment verification remains impossible.*

Open (long manuscript). (1) *Maximal auditable class:* useful $O(\sqrt{d/n})$ audits exist for rotation-closed classes whose

dual witnesses have root- n complexity (Aud-F/Aud-H), while full Lipschitz/ W_1 is rate-excluded; whether anything strictly larger than adaptive MSW_1 -ridge admits a useful root- n audit—and whether one-sided audits can beat the spiked-transport estimation line—remains open (estimation lower bounds alone do not finish the question: Wasserstein testing can be easier than estimation). (2) *Certified poly-time supremum*: whether a poly(d, n) algorithm can output a sound and null-complete upper bound on linear max-sliced W_1 is open; only net certificates (not gradient-ascent surrogates) are admissible in Aud-F until settled.

REFERENCES

- [1] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, T. Darrell. Tent: Fully test-time adaptation by entropy minimization. *ICLR*, 2021.
- [2] J. Liang, D. Hu, J. Feng. Do we really need to access the source data? Source hypothesis transfer for unsupervised domain adaptation. *ICML*, 2020.
- [3] Y. Sun, X. Wang, Z. Liu, J. Miller, A. Efros, M. Hardt. Test-time training with self-supervision for generalization under distribution shifts. *ICML*, 2020.
- [4] S. Niu, J. Wu, Y. Zhang, Y. Chen, S. Zheng, P. Zhao, M. Tan. Efficient test-time model adaptation without forgetting. *ICML*, 2022.
- [5] S. Niu, J. Wu, Y. Zhang, Z. Wen, Y. Chen, P. Zhao, M. Tan. Towards stable test-time adaptation in dynamic wild world. *ICLR*, 2023.
- [6] Q. Wang, O. Fink, L. Van Gool, D. Dai. Continual test-time domain adaptation. *CVPR*, 2022.
- [7] M. Zhang, S. Levine, C. Finn. MEMO: Test time robustness via adaptation and augmentation. *NeurIPS*, 2022.
- [8] M. Schirmer, M. Jazbec, C. A. Naesseth, E. Nalisnick. Monitoring risks in test-time adaptation. *NeurIPS*, 2025. arXiv:2507.08721.
- [9] A. Podkopaev, A. Ramdas. Tracking the risk of a deployed model and detecting harmful distribution shifts. *ICLR*, 2022. arXiv:2110.06177.
- [10] Y. Bar, S. Shaer, Y. Romano. Protected test-time adaptation via online entropy matching: A betting approach. *NeurIPS*, 2024. arXiv:2408.07511.
- [11] S. Garg, S. Balakrishnan, Z. Lipton, B. Neyshabur, H. Sedghi. Leveraging unlabeled data to predict out-of-distribution performance. *ICLR*, 2022. arXiv:2201.04234.
- [12] A. T. Kalai, V. Kanade. Towards optimally abstaining from prediction with OOD test examples. *NeurIPS*, 2021. arXiv:2105.14119.
- [13] J. Miller, R. Taori, A. Raghunathan, S. Sagawa, P. W. Koh, V. Shankar, P. Liang, Y. Carmon, L. Schmidt. Accuracy on the line: On the strong correlation between out-of-distribution and in-distribution generalization. *ICML*, 2021. arXiv:2107.04649.
- [14] C. Baek, Y. Jiang, A. Raghunathan, J. Z. Kolter. Agreement-on-the-line: Predicting the performance of neural networks under distribution shift. *NeurIPS*, 2022. arXiv:2206.13089.
- [15] T. Lee, S. Chottananurak, T. Gong, S.-J. Lee. AETTA: Label-free accuracy estimation for test-time adaptation. *CVPR*, 2024. arXiv:2404.01351.
- [16] J. Steinhardt, P. Liang. Unsupervised risk estimation using only conditional independence structure. *NeurIPS*, 2016. arXiv:1606.05313.
- [17] S. Ben-David, J. Blitzer, K. Crammer, A. Kulesza, F. Pereira, J. Wortman Vaughan. A theory of learning from different domains. *Machine Learning* 79(1–2):151–175, 2010.
- [18] E. Rosenfeld, S. Garg. (Almost) provable error bounds under distribution shift via disagreement discrepancy. *NeurIPS*, 2023. arXiv:2306.00312.
- [19] Y. Geifman, R. El-Yaniv. Selective classification for deep neural networks. *NeurIPS*, 2017.
- [20] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, L. Fei-Fei. ImageNet: A large-scale hierarchical image database. *CVPR*, 2009.
- [21] D. Hendrycks, T. Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. *ICLR*, 2019.
- [22] P. W. Koh, S. Sagawa, H. Marklund, S. M. Xie, M. Zhang, A. Balsubramani, W. Hu, M. Yasunaga, et al. WILDS: A benchmark of in-the-wild distribution shifts. *ICML*, 2021.
- [23] P. Bándi, O. Geessink, Q. Manson, M. van Dijk, M. Balkenhol, M. Hermesen, et al. From detection of individual metastases to classification of lymph node status at the patient level: The CAMELYON17 challenge. *IEEE Trans. Med. Imaging* 38(2):550–560, 2019.
- [24] J. Taylor, S. L. Earnshaw, B. M. Mabey, M. Victors, M. N. Annis, D. Y. Torigian, et al. The RXRX1 dataset and challenge: Unsupervised classification of cellular morphological changes. *arXiv:1907.04758*, 2019.
- [25] H. Venkateswara, J. Eusebio, S. Chakraborty, S. Panchanathan. Deep hashing network for unsupervised domain adaptation. *CVPR*, 2017. (Office-Home dataset release).
- [26] I. Gulrajani, D. Lopez-Paz. In search of lost domain generalization. *ICLR*, 2021.
- [27] F. Croce, M. Andriushchenko, V. Schwag, E. Debenedetti, N. Flammarion, M. Chiang, P. Mittal, M. Hein. RobustBench: A standardized adversarial robustness benchmark. *NeurIPS Datasets and Benchmarks*, 2021.
- [28] B. Recht, R. Roelofs, L. Schmidt, V. Shankar. Do CIFAR-10 classifiers generalize to CIFAR-10? *ICML*, 2019.
- [29] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, et al. PyTorch: An imperative style, high-performance deep learning library. *NeurIPS*, 2019.
- [30] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, et al. Scikit-learn: Machine learning in Python. *JMLR* 12:2825–2830, 2011.
- [31] R. Xie, A. Odonnat, V. Feofanov, W. Deng, J. Zhang, B. An. MaNo: exploiting matrix norm for unsupervised accuracy estimation under distribution shifts. *NeurIPS*, 2024. arXiv:2405.18979.
- [32] W. Deng, Y. Suh, S. Gould, L. Zheng. Confidence and dispersity speak: characterizing prediction matrix for unsupervised accuracy estimation. *ICML*, 2023. arXiv:2302.01094.
- [33] J. Liang, R. He, T. Tan. A comprehensive survey on test-time adaptation under distribution shifts. *IJCV* 133(1):31–64, 2025. arXiv:2303.15361.
- [34] R. J. Tibshirani, R. F. Barber, E. Candès, A. Ramdas. Conformal prediction under covariate shift. *NeurIPS*, 2019. arXiv:1904.06019.
- [35] R. F. Barber, E. J. Candès, A. Ramdas, R. J. Tibshirani. Conformal prediction beyond exchangeability. *Annals of Statistics* 51(2):816–845, 2023. arXiv:2202.13415.
- [36] I. Gibbs, E. Candès. Adaptive conformal inference under distribution shift. *NeurIPS*, 2021. arXiv:2106.00170.
- [37] S. Park, E. Dobriban, I. Lee, O. Bastani. PAC prediction sets under covariate shift. *ICLR*, 2022. arXiv:2106.09848.
- [38] S. Bates, A. Angelopoulos, L. Lei, J. Malik, M. I. Jordan. Distribution-free, risk-controlling prediction sets. *JACM* 68(6), 2021. arXiv:2101.02703.
- [39] A. N. Angelopoulos, S. Bates, E. J. Candès, M. I. Jordan, L. Lei. Learn then test: calibrating predictive algorithms to achieve risk control. *Annals of Applied Statistics*, 2025. arXiv:2110.01052.
- [40] T. Hoang, D. M. Vo, M. N. Do. Persistent test-time adaptation in recurring testing scenarios. *NeurIPS*, 2024. arXiv:2311.18193.
- [41] J. Lee, D. Jung, S. Lee, J. Park, J. Shin, U. Hwang, S. Yoon. Entropy is not enough for test-time adaptation: from the perspective of disentangled factors. *ICLR*, 2024. arXiv:2403.07366.
- [42] S. Sreeram, Y. D. Kwon, C. Mascolo. Tempora: Characterising the time-contingent utility of online test-time adaptation. *ICML*, 2026. arXiv:2602.06136.
- [43] A. Corrada-Emmanuel. The logic of NTQR evaluations of noisy AI agents: Complete postulates and logically consistent error correlations. 2023. arXiv:2312.05392.
- [44] C. Gupta, A. Ramdas. Top-label calibration and multiclass-to-binary reductions. *ICLR*, 2022. arXiv:2107.08353.