

ELARA: When Reliability Gates Misfire — A Cross-Benchmark Study of Drift-Signal Generalization in Multimodal Anomaly Fusion

Pratik Niroula
Independent Researcher

Abstract—We report a contrastive negative result for validation-derived drift gates in multimodal anomaly fusion: the same Kolmogorov–Smirnov drift signal that improves all-domain coherent score-collapse robustness on a label-aligned secondary benchmark *hurts* clean and adversarial ROC-AUC on naturally paired MVTec 3D-AD, with mechanism isolation identifying KS-drift mis-firing on legitimate inter-category variation as the cause. The cross-benchmark asymmetry is the central scientific finding.

The result is produced inside ELARA (Evidence-Layered Reliability for Anomaly fusion), a research prototype for score-level multimodal anomaly fusion. The mechanism under study is Reliability-Gated Attention (RGA), which couples masked cross-domain attention with a post-hoc reliability estimator (validation expected calibration error, test-time KS drift, prediction sharpness) and a conservative gate that preserves the static attention model unless mean reliability falls below a threshold. The system bundle includes a reproducible long-format fusion schema, strong score-level baselines (random forest, early-fusion MLP, late-fusion ensemble, Tent and TTT score adapters, static attention), missing-domain ablations, adversarial perturbation stress tests, calibration analysis, and counterfactual domain attribution.

The empirical study spans two benchmarks. The primary benchmark is naturally paired MVTec 3D-AD across eight categories (3,226 paired RGB/depth samples, 22.4% positive). The secondary benchmark, ELARA-Bench-LA, combines four real local datasets (Credit Card Fraud, UNSW-NB15, On-line Shoppers Intention, labeled news text) with explicit label alignment. Both benchmarks honor their underlying dataset’s predefined train/validation/test split so the fusion test fold is disjoint from the rows the per-domain scorers fit on. On MVTec 3D-AD, random forest fusion dominates (clean ROC-AUC 0.959), static attention is mid-pack (0.650), and RGA regresses by -0.040 clean and -0.117 to -0.124 under all-domain coherent attacks. On ELARA-Bench-LA the same gate improves all-domain **zero_attack** and **max_attack** ROC-AUC by $+0.0506$ and $+0.0319$ respectively. Component ablation shows the KS-drift signal alone controls whether the gate fires; removing it eliminates both the gain and the harm. We report both directions of the result with confidence intervals and explicit mechanism isolation, not as a blanket superiority claim.

Index Terms—anomaly detection, multimodal fusion, reliability, calibration, attention, score-level fusion, counterfactual explanation, fraud detection, cybersecurity

I. INTRODUCTION

Operational anomaly detection is usually a system problem before it is a single-model problem. A high-risk case may appear weakly in transaction records, network telemetry, user behavior, text evidence, device signals, or document checks. Each domain expert may be accurate inside its own data distribution, but the final decision still depends on a fusion layer

that decides how much to trust each signal. This trust problem becomes harder when domains are missing, score distributions drift, calibration degrades, or an adversary corrupts one or more sources.

ELARA is a research system for this setting. The system is not designed as a new fraud detector, network detector, or text detector in isolation. Instead, it treats anomaly detection as a verification pipeline: domain experts generate scores and compact evidence, those scores are converted into a unified fusion schema, and a reliability-aware fusion module produces a final risk estimate with a domain-level explanation. This architecture matches how many practical anomaly systems are deployed: detectors are built and maintained by different teams, and the fusion layer must remain robust as their data quality changes.

The paper treats ELARA as the system boundary and RGA as one mechanism inside that boundary. This matters because the contribution is not a new domain detector. It is a reliability-aware fusion layer with explicit input schema, stress tests, and limitations.

A. Research Question

The paper asks:

Can a multimodal anomaly-verification system preserve clean score-fusion accuracy while using calibration and drift evidence to reduce brittle behavior under domain failure?

This question is intentionally narrower than a claim of general anomaly detection superiority. The present codebase already supports domain experts, fusion schemas, attention fusion, reliability scoring, and robustness tests. The research target is the fusion and verification layer that combines those experts under imperfect operating conditions.

B. Thesis and Evidence Standard

The working thesis is that reliability-aware fusion is most useful when the operating condition changes. A strong paper must therefore show three things:

- 1) clean performance remains competitive with strong baselines,
- 2) reliability adaptation improves behavior under meaningful stress, and
- 3) the data and experimental protocol justify the scope of the claim.

The current results support the first two conditions for score-level fusion but do not yet fully satisfy the third condition for co-observed multimodal incidents. This limitation is explicit throughout the manuscript.

C. Contributions

We make four scoped contributions:

- A contrastive negative result for validation-derived drift gates. The same Kolmogorov–Smirnov drift signal carries the entire all-domain stress-test gain on the synthetically paired secondary benchmark and drives the clean-and-adversarial regression on RGB/3D paired MVTEC 3D-AD; component ablation identifies the signal as the cause.
- RGA couples masked cross-modal attention with a three-signal reliability estimator (validation ECE, test-time KS drift, score sharpness) and a conservative gate that preserves the static path unless mean reliability falls below τ .
- ELARA-Bench-LA is a reproducible the secondary benchmark real-domain score-fusion benchmark over Credit Card Fraud, UNSW-NB15, Online Shoppers Intention, and labeled news text. A naturally paired MVTEC 3D-AD benchmark path covers eight categories with predefined-split discipline.
- ELARA packages the schema, baselines (random forest, early-fusion MLP, late-fusion ensemble, Tent, TTT, static attention, RGA), missing-domain ablation, drift curves, adversarial perturbation stress tests, calibration analysis, and counterfactual domain attribution with an explicit construct-validity boundary.

The evidence is strongest as a reliability-stress study over score-level fusion. It is not a general claim that RGA is always superior to static attention, nor as evidence that the label-aligned benchmark captures naturally co-observed multimodal incidents.

Figure 1 shows the system boundary. Table I summarizes the implementation requirements used to keep the paper aligned with the codebase.

II. RELATED WORK

A. Anomaly Detection Systems

Deep anomaly detection has been studied across tabular, sequential, text, graph, and visual data. Pang et al. [3] emphasize that anomaly detection is tightly coupled to assumptions about labels, imbalance, and application context. ELARA follows this systems view: it does not assume a single representation can absorb every domain cleanly. Instead, it treats each domain expert as a separate source of evidence.

B. Multimodal and Score-Level Fusion

Multimodal learning surveys distinguish early fusion, late fusion, hybrid fusion, and interaction-based fusion [2]. Early fusion concatenates features before prediction; late fusion combines model outputs; stacking learns a meta-model over domain predictions. Transformer attention [1] gives a natural mechanism for modeling interactions across available

TABLE I
SYSTEM REQUIREMENTS FOR ELARA AND HOW THE CURRENT IMPLEMENTATION ADDRESSES THEM.

Requirement	Current implementation
Heterogeneous evidence	Fraud, cyber, behavior, and text domains are mapped to a shared score-level schema.
Missing-domain support	The attention model uses explicit masks; baselines receive missingness-aware fallbacks.
Reliability awareness	Domain weights combine calibration error, KS drift evidence, and score sharpness.
Strong baselines	Random forest, early fusion MLP, late fusion ensemble, static attention, and confidence weighting are included.
Explanation	Counterfactual domain attribution masks one domain at a time and measures risk change.
Reproducibility	Data generation, benchmark execution, asset generation, and PDF compilation are scriptable.

domains. The limitation is that ordinary attention learns a trust policy from training data and can remain overconfident when the inference condition changes.

C. RGB–3D Industrial Anomaly Detection

MVTEC 3D-AD provides paired RGB and high-resolution 3D scans for industrial anomaly detection and localization [15]. Recent RGB–3D methods show that multimodal anomaly detection is not solved by direct feature concatenation alone. M3DM uses hybrid feature and decision-layer fusion with memory banks [16]; crossmodal feature mapping detects anomalies by checking consistency between RGB and point-cloud features [17]; Real3D-AD expands high-precision point-cloud benchmarking [18]; and noise-resistant multimodal work studies robustness when training data are not clean [19]. These papers motivate the paired benchmark path in ELARA: the fusion layer should be judged on naturally paired modalities and on reliability stress, not only on clean score fusion.

D. Calibration, Drift, and Uncertainty

Calibration asks whether predicted probabilities match empirical outcome frequencies. Post-hoc calibration methods [4], [5], [6] improve probability quality, but calibration alone does not solve the problem of domain trust under drift. Uncertainty and calibration can degrade under dataset shift [7]. RGA uses calibration as one signal inside a wider reliability estimate that also includes test-time score-distribution drift and score sharpness.

E. Test-Time Adaptation

Test-time training and test-time adaptation update models at inference time to handle distribution shift [8], [9]. ELARA now includes two score-level adaptation comparators: an entropy-minimizing Tent adapter over the fusion head and a high-confidence pseudo-label TTT adapter. They are intentionally conservative because the fusion setting receives domain scores rather than raw modality tensors. Their role is




Fig. 1. ELARA system architecture. Domain experts produce scores, confidence values, and compact embeddings; a unified fusion schema feeds static attention and the RGA reliability estimator; the system emits a risk score and counterfactual domain attribution.

to test whether RGA adds value beyond generic score-level adaptation.

F. Ensembles and Strong Baselines

Ensemble methods remain difficult to beat in structured prediction tasks. Random forests [10] and gradient boosting [11] are strong baselines for tabular and score-level settings, and outlier-ensemble theory provides a related foundation for combining heterogeneous anomaly scorers [20]. Classical classifier-combiner theory [21] further characterises when sum, product, and weighted rules differ in behaviour under noise. The benchmark therefore includes random forest fusion, early fusion MLP, late fusion ensemble, confidence-weighted mean, and static attention. A credible systems paper must compare against these practical alternatives, not only against weak averages. The reliability gate of RGA is itself an inference-time switching rule, conceptually adjacent to selective-prediction work that abstains rather than commits when confidence is low [23].

G. Explanation for Fusion Systems

Feature attribution methods such as SHAP [12] explain individual model predictions, but fusion systems also need domain-level explanations. Analysts may need to know whether fraud evidence, network evidence, behavior evidence, or text evidence drove the risk score. Counterfactual-explanation work [22] provides the conceptual basis for the masking-based domain attribution used here: we ask what the risk score would have been had a single domain not been observed. ELARA therefore includes counterfactual domain

attribution: it masks one domain and reports the change in predicted risk.

III. SYSTEM AND METHOD

Each sample i contains up to D domain observations. Domain d contributes a feature vector $x_{i,d} \in \mathbb{R}^F$ containing a score, optional confidence, and compact embedding features. A binary mask $m_{i,d}$ indicates whether domain d is missing. The target label is $y_i \in \{0, 1\}$ when supervised labels are available.

The system estimates

$$\hat{p}_i = P(y_i = 1 \mid \{x_{i,d}, m_{i,d}\}_{d=1}^D), \quad (1)$$

where $\hat{p}_i$ is the final anomaly risk. In static fusion, the learned relationship among domains is fixed after training. In RGA, the system also computes a batch-level reliability vector $r \in [0, 1]^D$ at inference time and applies the sample mask to form effective per-sample weights $\tilde{r}_i \in [0, 1]^D$. The reliability vector should reduce trust in domains that appear miscalibrated, shifted, uninformative, or missing.

The evaluation problem is not only to maximize ROC-AUC. A useful verification system must also satisfy:

- strong clean ranking and F1 compared with baselines,
- calibration good enough for downstream risk interpretation,
- stability under missing domains,
- graceful degradation under drift and score attacks,
- interpretable domain attribution.

A. ELARA System Architecture

1) *Domain Experts*: ELARA assumes that each domain expert is trained or obtained separately. For the current benchmark, tabular scorers are trained for fraud, cyber, and behavior data, while a text scorer is trained for labeled text. The fusion layer does not require these experts to share the same raw feature space. It only requires score-level outputs and optional compact evidence features.

2) *Fusion Schema*: The unified schema is long-format: each row contains a sample identifier, a domain name, a binary label when available, a score, a confidence value, and embedding columns. This design allows different samples to have different available domains. Missing domains are represented by masks rather than by fabricated raw observations.

3) *Static Attention Path*: The static attention path encodes each available domain into a shared latent space and applies multi-head attention across domains. Domain embeddings allow the model to distinguish evidence sources. Missing domains are ignored through key-padding masks. This is the high-reliability default path.

4) *Reliability-Aware Path*: The reliability-aware path is activated when the average reliability of present domains falls below a configured threshold. This avoids the common failure mode where reliability weights disturb clean predictions even though no domain appears degraded. The reliability path is therefore a stress-response mechanism, not a universal replacement for the static model.

5) *Explanation Path*: For an analyst-facing explanation, the system masks one available domain at a time and recomputes risk. The resulting probability change gives a domain-level counterfactual attribution. This produces statements such as: if text evidence were absent, the risk would decrease by a given amount.

B. RGA Method

1) *Domain Encoding and Attention*: Each domain encoder maps the domain input vector into a shared latent space:

$$h_{i,d} = g_d(x_{i,d}) + e_d + p_d, \quad (2)$$

where e_d is a learned domain embedding and p_d is an optional learned position embedding used by the implementation. A multi-head self-attention block operates across available domain embeddings:

$$z_i = f_\theta(\text{Pool}(\text{Attention}(H_i, m_i))), \quad \hat{p}_i = \sigma(z_i). \quad (3)$$

2) *Reliability Estimation*: After model training, the system fits a reliability estimator on the validation split. For domain d , it stores validation score distributions, fits an isotonic calibrator when enough labels are available, and computes expected calibration error. At inference time:

$$r_d = \alpha(1 - \text{ECE}_d) + \beta p_{\text{KS},d} + \gamma S_d, \quad (4)$$

where ECE_d is validation calibration error, $p_{\text{KS},d}$ is the Kolmogorov-Smirnov p-value comparing current scores with validation scores, and S_d is score sharpness from squared distance to 0.5. The current benchmark uses $(\alpha, \beta, \gamma) = (0.45, 0.35, 0.20)$.

Algorithm 1 RGA inference

Require: Inputs $x_{i,1\dots D}$ and mask $M_i \in \{0,1\}^D$ for sample i

Require: Fitted attention model f_{att} with encoders $\{g_d\}$

Require: Fitted reliability estimator producing per-sample weights $r_{i,d}$

Require: Gate threshold $\tau \in (0,1)$, default $\tau = 0.66$

Ensure: Risk score $\hat{p}_i \in [0,1]$

```

1:  $h_{i,d} \leftarrow g_d(x_{i,d}) + e_d + p_d$  for each  $d$  with  $M_{i,d} = 0$ 
2:  $r_{i,d} \leftarrow \alpha(1 - \text{ECE}_d) + \beta p_{\text{KS},d}(x_{i,d}) + \gamma S_{i,d}$ 
3:  $\bar{r}_i \leftarrow \frac{1}{|D_i^{\text{obs}}|} \sum_{d: M_{i,d}=0} r_{i,d}$ 
4: if  $\bar{r}_i \geq \tau$  then
5:    $\hat{p}_i \leftarrow \sigma(f_{\text{att}}(h_i, \cdot; M_i))$  ▷ static path
6: else
7:    $\hat{p}_i \leftarrow \sigma(f_{\text{att}}(h_i, \cdot; M_i, w_{i,\cdot} = r_{i,\cdot} \odot \neg M_i))$ 
8:   ▷ reliability-injected fusion
9: end if
10: return  $\hat{p}_i$  and (optional) counterfactual attributions  $\{\Delta_{i,d}\}_{d \in D_i^{\text{obs}}}$ 

```

3) *Reliability Gate*: Let $\bar{r}$ be the mean reliability over present domains and let $\tau = 0.66$. If $\bar{r} \geq \tau$, ELARA uses the static attention path. If $\bar{r} < \tau$, reliability weights are injected into the fusion block:

$$\tilde{r}_{i,d} = (1 - m_{i,d})r_{i,d}. \quad (5)$$

This makes RGA different from confidence weighting. Confidence weighting uses the local score confidence directly. RGA uses validation calibration, test-time drift evidence, and sharpness to decide when a domain should be trusted less.

4) *Counterfactual Domain Attribution*: For sample i , let $\hat{p}_i$ be the original prediction and $\hat{p}_{i,-d}$ be the prediction when domain d is masked. The attribution is

$$\Delta_{i,d} = \hat{p}_i - \hat{p}_{i,-d}. \quad (6)$$

A positive $\Delta_{i,d}$ indicates that domain d increased anomaly risk. A negative value indicates that domain d reduced risk.

5) *Algorithm Summary*: The full inference procedure is given in Algorithm 1. Steps 1–3 are deterministic feature operations; steps 4–6 are the reliability-gated switch. The reported runs use the batch-level reliability vector described above, applied through each sample’s observed-domain mask. A per-sample gating mode is implemented as an ablation path, but it is not the default for the reported tables.

IV. BENCHMARK CONSTRUCTION

A. Primary Benchmark: Naturally Paired MVTEC 3D-AD

The primary benchmark of this manuscript is MVTEC 3D-AD [15]. Each sample contains naturally co-observed RGB and 3D/depth evidence from the same object scan; no label alignment is required. The current run covers all eight extracted categories — bagel, cable_gland, cookie, dowel, foam, peach, rope, and tire — for a total of 3,226 paired samples with a positive fraction of 0.224 and full RGB and depth

TABLE II
NATURALLY PAIRED MVTEC 3D-AD MULTI-CATEGORY METADATA.

Property	Value
Benchmark type	naturally_paired_mvtec3d_score_fusion
Natural pairing	yes
Categories	bagel, cable_gland, cookie, dowel, foam, peach, rope,
Samples	3226
Positive fraction	0.224
Score fit split	train/good
Score fit samples	2070
Score normalization	train_good_distance_minmax_clipped
Domains	rgb, depth_or_xyz

coverage. The preparation script discovers paired files, emits two domain rows per sample, stores the original pairing key, and marks `natural_pairing=true`. The initial scorer is deliberately lightweight: normal-reference image statistics produce reproducible RGB and depth anomaly scores without requiring a heavy specialized 3D detector. This keeps the first paired benchmark auditable; stronger RGB-3D scorers can replace these domain experts later while preserving the fusion interface.

Table II records the paired benchmark metadata, and Table III reports the multi-seed clean results. Random forest fusion is the strongest baseline by a clear margin (ROC-AUC 0.959); static attention is mid-pack (0.650); the reliability gate of RGA *hurts* both clean and adversarial conditions on this data. The paper reports this outcome directly as a headline result rather than framing it as a smoke run. The RGB/depth anomaly scores are fit from original MVTEC train-good observations only; the supervised fusion evaluation uses a held-out fusion split over paired observations and should not be read as the canonical MVTEC train/test anomaly-detection protocol.

Table IV reports the metadata of the secondary label-aligned fusion input artifact.

B. Real-Domain Sources

ELARA-Bench-LA is built from four local datasets: Credit Card Fraud, UNSW-NB15, Online Shoppers Intention, and labeled news text. Domain-specific out-of-fold scorers are trained first. Their scores, confidences, and compact embeddings are then converted into the common fusion schema.

The generated fusion table contains 8,000 composite samples, four domains, eight embedding features per domain, a 0.307 positive rate, and 11.8% nominal missingness. Domain coverage is 0.883 for fraud, 0.885 for cyber, 0.885 for behavior, and 0.877 for text.

C. Label Alignment Limitation

The datasets do not share natural entity identifiers or timestamps. Therefore, composite samples are the secondary benchmark: a positive composite draws positive records from available domains, and a negative composite draws negative records. This is a meaningful real-domain score-fusion benchmark, but it is not evidence that the system has learned co-observed multimodal incident structure. That distinction is

central to the claim boundary of the paper. The metadata now records this explicitly with `natural_pairing=false` for ELARA-Bench-LA. Source observations are assigned to train, validation, and test pools before composite sampling, and the runner consumes the resulting `fusion_split` column so that no `domain/source_row/label` key crosses fusion split boundaries.

D. Training Protocol

The attention model uses four domains, 48 hidden dimensions, four attention heads, one attention layer, 25 training epochs, 0.15 domain dropout during training, and seeds 42 through 46. Clean benchmark results are reported as mean±standard deviation across five seeds. The experiment runner also repeats missing-domain, drift, adversarial, calibration, and CDA analyses for every configured seed, then stores both raw per-seed rows and aggregated mean, standard deviation, and 95% confidence intervals.

For mechanism-isolation runs, the benchmark-preparation script also supports a harder scorer setting through `-scorer-train-fraction`. Values such as 0.05 or 0.10 train each out-of-fold domain scorer on only a stratified fraction of its training fold before fusion rows are built. This is intended to reduce clean-score saturation and expose whether the fusion mechanism helps when domain experts are useful but imperfect.

As a post-audit check, we ran the 0.05 setting and archived the results in `experiments/fusion/craf_real_results_hard.json`. The individual domain scorers weakened substantially for fraud and behavior (OOF ROC-AUC 0.913 and 0.798, respectively), but the label-aligned fusion task remained near-saturated: static attention and RGA both reached mean ROC-AUC 0.99896 across five seeds, random forest reached 0.99908, and RGA did not improve drift degradation area on any domain. This hard run is therefore reported as a construct-validity warning rather than a new performance claim.

E. Baselines and Metrics

Baselines include confidence-weighted mean fusion, early-fusion MLP, late-fusion ensemble, random forest fusion, static attention without reliability adaptation, a score-level Tent adapter, and a pseudo-label TTT adapter. Metrics include ROC-AUC, PR-AUC, F1, expected calibration error, Brier score, and accuracy. Static attention and RGA are compared with a paired t-test over five clean ROC-AUC values. ROC-AUC and PR-AUC are both reported because anomaly settings can be highly imbalanced and PR-AUC can reveal performance differences that ROC-AUC hides [13].

F. Confidence Intervals

Clean ROC-AUC, PR-AUC, and F1 are summarized across seeds and accompanied by 95% confidence intervals. Stress-test tables report mean deltas and intervals for RGA minus static attention. The result JSON stores raw per-seed stress

TABLE III

NATURALLY PAIRED MVTEC 3D-AD CLEAN HELD-OUT PERFORMANCE ACROSS EIGHT CATEGORIES. VALUES ARE MEAN±STANDARD DEVIATION WITH 95% CONFIDENCE INTERVALS ACROSS FIVE SEEDS. RANDOM FOREST FUSION IS THE STRONGEST BASELINE; RGA ATTENTION DEGRADES STATIC-ATTENTION PERFORMANCE ON THIS PAIRED BENCHMARK.

Method	ROC-AUC	PR-AUC	F1	ECE	Brier
Random forest	0.9592±0.0076 [0.9508, 0.9717]	0.8938±0.0216 [0.8682, 0.9286]	0.8146±0.0063 [0.8068, 0.8201]	0.0867±0.0089 [0.0777, 0.1008]	0.0721±0.0041 [0.0657, 0.0775]
Conf.-weighted mean	0.5092±0.0185 [0.4802, 0.5294]	0.3106±0.0144 [0.2971, 0.3348]	0.2427±0.0192 [0.2093, 0.2582]	0.2172±0.0143 [0.1974, 0.2350]	0.2520±0.0081 [0.2418, 0.2621]
Tent score adapter	0.6700±0.0193 [0.6404, 0.6946]	0.4075±0.0241 [0.3801, 0.4445]	0.3761±0.0129 [0.3586, 0.3932]	0.1283±0.0116 [0.1137, 0.1447]	0.1848±0.0071 [0.1765, 0.1941]
TTT pseudo-label	0.5692±0.0390 [0.5228, 0.6246]	0.3257±0.0283 [0.3062, 0.3751]	0.3394±0.0245 [0.3048, 0.3657]	0.2847±0.0224 [0.2515, 0.3129]	0.2693±0.0192 [0.2412, 0.2931]
Early fusion MLP	0.8278±0.0211 [0.7949, 0.8505]	0.6760±0.0277 [0.6403, 0.7184]	0.5784±0.0347 [0.5247, 0.6214]	0.2050±0.0078 [0.1935, 0.2159]	0.1684±0.0083 [0.1587, 0.1787]
Late fusion ensemble	0.6019±0.0235 [0.5692, 0.6321]	0.3209±0.0147 [0.2981, 0.3400]	0.3736±0.0229 [0.3387, 0.4012]	0.2678±0.0020 [0.2649, 0.2700]	0.2415±0.0027 [0.2378, 0.2445]
Static attention	0.6497±0.0384 [0.6071, 0.7086]	0.3913±0.0659 [0.2943, 0.4765]	0.0447±0.0894 [0.0000, 0.2012]	0.0354±0.0166 [0.0120, 0.0587]	0.1648±0.0065 [0.1539, 0.1704]
RGA attention	0.6098±0.0312 [0.5766, 0.6607]	0.3507±0.0382 [0.3051, 0.4117]	0.0106±0.0212 [0.0000, 0.0477]	0.0402±0.0229 [0.0106, 0.0751]	0.1698±0.0041 [0.1632, 0.1748]

TABLE IV

BENCHMARK METADATA FOR THE SECONDARY SYNTHETICALLY PAIRED FUSION INPUT ARTIFACT.

Property	Value
Benchmark type	label_aligned_real_domain_score_fusion
Natural pairing	no
Pairing unit	binary-label-aligned composite sample
Samples	8000
Positive fraction	0.301
Scorer train fraction	1.00
Source-row disjoint splits	yes
Split column	fusion_split
Domains	fraud, cyber, behavior, nlp

TABLE V

OUT-OF-FOLD DOMAIN SCORER QUALITY USED TO CONSTRUCT ELARA-BENCH-LA.

Domain	Rows	Pos. rate	OOF ROC	OOF PR
fraud	20000	0.025	0.9686	0.7826
cyber	20000	0.500	0.9684	0.9696
behavior	12330	0.155	0.8711	0.6042
nlp	20000	0.500	0.9973	0.9973

rows so that the interval computation is auditable rather than embedded only in the manuscript.

G. Threat Model

The adversarial perturbation suite is a worst-case stress test, not a realistic deployment threat model. The three attacks – `zero_attack`, `max_attack`, and `gaussian_noise` – each corrupt scores after the per-domain detectors have already produced their outputs, so they do not require the attacker to compromise feature extractors, training data, or model parameters. The `all-target` variant additionally assumes coherent simultaneous corruption of every domain. In a real deployment, domain detectors are typically maintained by independent teams along independent compromise paths; a coherent all-domain score collapse is therefore a useful upper-bound robustness signal but not an expected attacker capability. Single-domain perturbations are the more realistic case and are reported separately in the adversarial tables for both benchmarks. We do not claim defenses against feature- or training-time attacks; those threat models are out of scope for a fusion layer that consumes score-level evidence from upstream detectors.

H. Secondary Clean Sanity Check

Table VI and Fig. 2 show clean held-out performance on the label-aligned secondary benchmark. This table is intentionally treated as a sanity check, not as the headline result, because the synthetically paired split is near saturated. Bold values indicate the highest rounded mean in each column; multiple bold values are rounded ties.

Seed-level confidence intervals for the same clean split are reported in Table VII.

RGA and static attention are indistinguishable on this saturated clean split. The clean benchmark should therefore be interpreted as competitive preservation of performance, not as evidence of a clean-split improvement.

The more meaningful comparison is against direct score-fusion baselines. RGA improves over confidence-weighted mean fusion by about 0.0015 ROC-AUC, 0.0039 PR-AUC, 0.0100 F1, and reduces ECE from 0.0834 to 0.0042. Random forest fusion obtains the best F1 at 0.9906, but its ECE is 0.0287, substantially worse than attention-based fusion. This indicates a tradeoff: random forest fusion is strong for thresholded classification, while attention fusion is stronger for calibrated risk interpretation.

V. PRIMARY HEADLINE RESULTS: MVTEC 3D-AD

This section reports the primary multi-seed evidence on RGB/3D paired MVTEC 3D-AD across all eight extracted object categories (3,226 paired samples, 22.4% positive). The same suite is repeated on the contrastive the secondary benchmark in §VI; the asymmetry between the two outcomes is the central finding.

A. Clean Performance

Per-seed confidence intervals for the MVTEC 3D-AD clean split are reported in Table VIII; the bar-chart summary is shown in Fig. 3.

Random forest fusion is the strongest baseline (clean ROC-AUC 0.959 ± 0.008); early-fusion MLP is the second strongest tabular baseline (0.829). The Tent score adapter [9] reaches 0.670 and outperforms the TTT pseudo-label adapter [8] at 0.569. Static attention is competitive at 0.650 but loses substantial ground to the random-forest ceiling. RGA attention reaches 0.610, -0.040 below static attention, with bootstrap intervals that overlap zero on the lower bound. This is a negative result for the reliability mechanism on naturally

TABLE VI

SATURATED CLEAN HELD-OUT PERFORMANCE ON ELARA-BENCH-LA. VALUES ARE MEAN $\pm$ STANDARD DEVIATION ACROSS FIVE SEEDS; THIS TABLE IS NOT USED AS THE HEADLINE EVIDENCE.

Method	ROC-AUC	PR-AUC	F1	ECE	Brier
Random forest	0.9999 [0.9999, 1.0000]	0.9999 [0.9998, 0.9999]	0.9946 $\pm$ 0.0004 [0.9939, 0.9948]	0.0297 $\pm$ 0.0004 [0.0294, 0.0304]	0.0061 $\pm$ 0.0001 [0.0060, 0.0062]
Conf.-weighted mean	0.9987 [0.9987, 0.9987]	0.9968 [0.9968, 0.9968]	0.9746 [0.9746, 0.9746]	0.0799 [0.0799, 0.0799]	0.0232 [0.0232, 0.0232]
Tent score adapter	0.9999 [0.9999, 0.9999]	0.9999 [0.9999, 0.9999]	0.9959 [0.9959, 0.9959]	0.0028 [0.0028, 0.0028]	0.0023 [0.0023, 0.0023]
TTT pseudo-label	0.9999 [0.9999, 0.9999]	0.9998 [0.9998, 0.9998]	0.9938 [0.9938, 0.9938]	0.0032 [0.0032, 0.0032]	0.0025 [0.0025, 0.0025]
Early fusion MLP	0.9999 [0.9998, 0.9999]	0.9998 $\pm$ 0.0001 [0.9996, 0.9999]	0.9919 $\pm$ 0.0004 [0.9917, 0.9927]	0.0043 $\pm$ 0.0010 [0.0030, 0.0058]	0.0036 $\pm$ 0.0006 [0.0029, 0.0045]
Late fusion ensemble	1.0000 [1.0000, 1.0000]	0.9999 [0.9999, 0.9999]	0.9938 [0.9938, 0.9938]	0.0123 [0.0123, 0.0123]	0.0036 [0.0036, 0.0036]
Static attention	0.9999 [0.9999, 0.9999]	0.9998 [0.9998, 0.9998]	0.9927 $\pm$ 0.0009 [0.9917, 0.9938]	0.0038 $\pm$ 0.0008 [0.0024, 0.0046]	0.0032 $\pm$ 0.0002 [0.0029, 0.0034]
RGA attention	0.9999 [0.9999, 0.9999]	0.9998 [0.9998, 0.9998]	0.9927 $\pm$ 0.0009 [0.9917, 0.9938]	0.0038 $\pm$ 0.0008 [0.0024, 0.0046]	0.0032 $\pm$ 0.0002 [0.0029, 0.0034]

TABLE VII

CLEAN RANKING AND F1 CONFIDENCE INTERVALS ACROSS CONFIGURED SEEDS.

Method	ROC-AUC 95% CI	PR-AUC 95% CI	F1 95% CI
Random forest	[0.9999, 1.0000]	[0.9998, 0.9999]	[0.9939, 0.9948]
Conf.-weighted mean	[0.9987, 0.9987]	[0.9968, 0.9968]	[0.9746, 0.9746]
Tent score adapter	[0.9999, 0.9999]	[0.9999, 0.9999]	[0.9959, 0.9959]
TTT pseudo-label	[0.9999, 0.9999]	[0.9998, 0.9998]	[0.9938, 0.9938]
Early fusion MLP	[0.9998, 0.9999]	[0.9996, 0.9999]	[0.9917, 0.9927]
Late fusion ensemble	[1.0000, 1.0000]	[0.9999, 0.9999]	[0.9938, 0.9938]
Static attention	[0.9999, 0.9999]	[0.9998, 0.9998]	[0.9917, 0.9938]
RGA attention	[0.9999, 0.9999]	[0.9998, 0.9998]	[0.9917, 0.9938]




Fig. 2. Clean benchmark ROC-AUC and PR-AUC across score-level and attention fusion methods.

paired data: the gate hurts even before any adversarial stress is applied.

B. Adversarial Score Perturbations

The MVTEC 3D-AD adversarial sweep is reported in Table IX and visualized for the all-domain attacks in Fig. 4. Across every attack type and every target subset (RGB-only, depth-only, all), RGA produces a *negative* ROC-AUC delta relative to static attention. The harshest case is `zero_attack:all` at -0.124 (95% CI $[-0.182, -0.078]$); the smallest all-domain negative is `max_attack:all` at -0.117 . There is no attack scenario in which RGA improves over static attention on this benchmark.

C. Missing-Domain Ablation

Inference-time domain dropout is reported in Table X and Fig. 5. RGA is below static attention by -0.040 at zero extra dropout and remains negative or near-zero through the 0.5 dropout sweep. The single-domain MVTEC benchmark, like the adversarial sweep, provides no missing-domain advantage for the current gate.

D. Score Drift

The MVTEC 3D-AD drift summary is shown in Table XI and Fig. 6. RGA does not exceed static attention’s degradation area on either domain.

TABLE VIII
NATURALLY PAIRED MVTEC 3D-AD: CLEAN ROC-AUC, PR-AUC, AND F1 95% CONFIDENCE INTERVALS ACROSS CONFIGURED SEEDS.

Method	ROC-AUC 95% CI	PR-AUC 95% CI	F1 95% CI
Random forest	[0.9508, 0.9717]	[0.8682, 0.9286]	[0.8068, 0.8201]
Conf.-weighted mean	[0.4802, 0.5294]	[0.2971, 0.3348]	[0.2093, 0.2582]
Tent score adapter	[0.6404, 0.6946]	[0.3801, 0.4445]	[0.3586, 0.3932]
TTT pseudo-label	[0.5228, 0.6246]	[0.3062, 0.3751]	[0.3048, 0.3657]
Early fusion MLP	[0.7949, 0.8505]	[0.6403, 0.7184]	[0.5247, 0.6214]
Late fusion ensemble	[0.5692, 0.6321]	[0.2981, 0.3400]	[0.3387, 0.4012]
Static attention	[0.6071, 0.7086]	[0.2943, 0.4765]	[0.0000, 0.2012]
RGA attention	[0.5766, 0.6607]	[0.3051, 0.4117]	[0.0000, 0.0477]




Fig. 3. Clean benchmark ROC-AUC and PR-AUC across score-level and attention fusion methods on naturally paired MVTEC 3D-AD.

E. Mechanism Isolation on MVTEC 3D-AD

The mechanism-isolation tables exposed why the gate hurts on this benchmark. Table XII reports the gate threshold sweep:

On MVTEC 3D-AD, the gate is dormant at $\tau \leq 0.6$ on every condition, because the mean reliability is high enough that the static path runs. At $\tau = 0.66$ the gate begins to fire on clean data ($\sim 20\%$ adaptation rate) and clean ROC-AUC immediately drops by -0.040 ; raising τ further both raises the adaptation rate to 100% and increases the clean penalty to -0.127 at $\tau = 0.90$. Under all-domain adversarial attacks the gate fires for 100% of samples from $\tau = 0.66$ onward, and the delta saturates between roughly -0.117 and -0.124 depending on the attack. Every substantial adaptation rate on this benchmark is paired with a negative ROC-AUC delta; the clean $\tau = 0.66$ interval touches zero, while the stronger clean thresholds and all-domain attack intervals are strictly negative.

The reliability-component ablation in Table XIII shows where the failure originates.

Only the `no_ks` variant avoids harm: when the KS-drift sig-

nal is removed, the adaptation rate stays at 0% and the deltas stay at exactly zero. Every other variant (`full`, `no_ece`, `no_gate`, `no_sharpness`) reaches a 100% adaptation rate and a strongly negative delta. The reading is unambiguous: the KS-drift component is the entire reason RGA fires on this data, and once it fires the adapted weighting is worse than the static attention solution. ECE contributes a small amount in either direction; sharpness is near-neutral; the gate cannot help because the KS signal already places mean reliability below τ on most batches.

F. Calibration and Counterfactual Attribution

Static attention retains the better MVTEC 3D-AD calibration (Table XIV and Fig. 7). The reliability diagram for RGA is more dispersed, consistent with the clean ROC-AUC penalty. The counterfactual domain-attribution impacts (Fig. 8) remain meaningfully non-zero on both RGB and depth, so the explanation pathway continues to function even when the overall mechanism is harmful.

TABLE IX
MVTec 3D-AD ADVERSARIAL PERTURBATION BENCHMARK. DELTA IS RGA MINUS STATIC ATTENTION; ALL DELTAS ARE NEGATIVE OR ZERO.

Attack	Target	Static ROC	RGA ROC	Delta	Delta 95% CI
gaussian_noise	all	0.6498	0.5329	-0.1170	[-0.1823, -0.0603]
gaussian_noise	depth_or_xyz	0.6496	0.5231	-0.1265	[-0.1740, -0.0947]
gaussian_noise	rgb	0.6483	0.5763	-0.0720	[-0.1177, -0.0389]
max_attack	all	0.6330	0.5164	-0.1166	[-0.1831, -0.0394]
max_attack	depth_or_xyz	0.6426	0.5851	-0.0575	[-0.1262, -0.0019]
max_attack	rgb	0.6420	0.6165	-0.0255	[-0.0684, 0.0000]
zero_attack	all	0.6442	0.5199	-0.1243	[-0.1819, -0.0784]
zero_attack	depth_or_xyz	0.6447	0.5826	-0.0620	[-0.1209, -0.0023]
zero_attack	rgb	0.6494	0.6203	-0.0291	[-0.0741, 0.0000]




Fig. 4. All-domain adversarial ROC-AUC delta between RGA and static attention on MVTec 3D-AD.




Fig. 5. MVTec 3D-AD ROC-AUC under additional missing-domain dropout.

TABLE X
MVTec 3D-AD MISSING-DOMAIN ABLATION. DELTA IS RGA MINUS STATIC ATTENTION.

Dropout	Static ROC	RGA ROC	Delta	Delta 95% CI
0.0	0.6497	0.6098	-0.0398	[-0.1206, 0.0000]
0.1	0.6187	0.6159	-0.0028	[-0.0101, 0.0000]
0.2	0.5953	0.5706	-0.0246	[-0.0477, 0.0000]
0.3	0.5756	0.5566	-0.0190	[-0.0516, -0.0011]
0.5	0.5369	0.5170	-0.0199	[-0.0446, -0.0010]

G. Failure Cases

Representative MVTec 3D-AD contrast cases are reported in Table XV.

VI. CONTRASTIVE SECONDARY BENCHMARK (LABEL-ALIGNED)

The same evaluation suite is now applied to ELARA-Bench-LA, where four real-domain datasets are joined by explicit label alignment rather than co-observation. The aim of this

TABLE XI
MVTec 3D-AD DRIFT ROBUSTNESS SUMMARY. HIGHER DEGRADATION AREA IS BETTER.

Domain	Static	RGA	Delta	Seeds
depth_or_xyz	0.1948±0.0114 [0.1822, 0.2123]	0.1589±0.0072 [0.1512, 0.1696]	-0.0359	5
rgb	0.1948±0.0113 [0.1822, 0.2121]	0.1771±0.0066 [0.1726, 0.1886]	-0.0176	5

section is to read each result alongside its MVTec 3D-AD counterpart in §V-E, so the cross-benchmark asymmetry is visible per condition rather than only in the aggregate. The saturated clean split is discussed in §IV-H; what follows are the robustness, drift, adversarial, and mechanism-isolation tables.

A. Missing-Domain Ablation

Table XVI and Fig. 9 report additional inference-time domain dropout. At zero extra dropout, RGA is equal to static attention, and the same equality holds through the configured dropout sweep. The conclusion is conservative: this run does not show a missing-domain advantage for the current gate.




Fig. 6. Per-domain score-drift curves on the MVTec 3D-AD benchmark.




Fig. 7. Reliability diagram on MVTec 3D-AD for static attention and RGA.

B. Score Drift

Table XVII summarizes degradation area under per-domain score-noise sweeps. Higher area means slower degradation. All deltas round to zero or near zero because clean performance is near saturated, so this table is best read as a robustness check rather than a win. The full curves are shown in Fig. 10.

C. Adversarial Score Perturbations

Table XVIII evaluates all score-level perturbations, while Fig. 11 isolates the all-domain attacks so the meaningful deltas are visible. Single-domain attacks are mostly ties. The important stress case is coherent all-domain score collapse. Under the all-domain zero attack, RGA improves ROC-AUC from 0.7130 to 0.7770. Under the all-domain max attack, it improves ROC-AUC from 0.7114 to 0.7418. Under all-domain Gaussian noise, it is slightly worse by 0.0003 ROC-AUC.




Fig. 8. Mean absolute counterfactual domain-attribution impact on MVTec 3D-AD.

D. Mechanism Isolation Protocol

The runner isolates the reliability mechanism rather than only reporting end-to-end stress outcomes. It emits two diagnostics when the mechanism-isolation block is enabled. First, a τ sweep evaluates $\tau \in \{0.4, 0.5, 0.6, 0.66, 0.7, 0.8, 0.9\}$ on clean data and all-domain stress conditions, recording both ROC-AUC deltas and the adaptation rate. Second, a reliability-component ablation compares the full estimator with no-ECE, no-KS, no-sharpness, and no-gate variants on all-domain score attacks. These diagnostics show whether the gate and reliability terms are active contributors rather than incidental implementation details.

Threshold-sweep outcomes and adaptation rates are reported in Table XIX.

At the configured $\tau = 0.66$, clean adaptation is rare: the mean adaptation rate is 0.032 with a confidence interval that includes zero. This explains why the missing-domain deltas are effectively zero: the gate usually keeps the static path

TABLE XII
MVTec 3D-AD RELIABILITY-GATE THRESHOLD SWEEP. DELTA IS RGA MINUS STATIC ATTENTION; ADAPTATION RATE IS THE FRACTION OF EVALUATED SAMPLES USING THE RELIABILITY-AWARE PATH.

Condition	τ	Static ROC	RGA ROC	Delta	Delta 95% CI	Adapt rate
clean	0.40	0.6497	0.6497	0.0000	[0.0000, 0.0000]	0.000
clean	0.50	0.6497	0.6497	0.0000	[0.0000, 0.0000]	0.000
clean	0.60	0.6497	0.6497	0.0000	[0.0000, 0.0000]	0.000
clean	0.66	0.6497	0.6098	-0.0398	[-0.1206, 0.0000]	0.204
clean	0.70	0.6497	0.5812	-0.0684	[-0.1226, -0.0063]	0.404
clean	0.80	0.6497	0.5317	-0.1180	[-0.1506, -0.0901]	0.879
clean	0.90	0.6497	0.5225	-0.1271	[-0.1757, -0.0845]	1.000
gaussian_noise:all	0.40	0.6502	0.6502	0.0000	[0.0000, 0.0000]	0.000
gaussian_noise:all	0.50	0.6502	0.6502	0.0000	[0.0000, 0.0000]	0.000
gaussian_noise:all	0.60	0.6502	0.5478	-0.1025	[-0.1380, -0.0698]	0.800
gaussian_noise:all	0.66	0.6502	0.5267	-0.1235	[-0.1882, -0.0576]	1.000
gaussian_noise:all	0.70	0.6502	0.5267	-0.1235	[-0.1882, -0.0576]	1.000
gaussian_noise:all	0.80	0.6502	0.5267	-0.1235	[-0.1882, -0.0576]	1.000
gaussian_noise:all	0.90	0.6502	0.5267	-0.1235	[-0.1882, -0.0576]	1.000
max_attack:all	0.40	0.6330	0.6330	0.0000	[0.0000, 0.0000]	0.000
max_attack:all	0.50	0.6330	0.6330	0.0000	[0.0000, 0.0000]	0.000
max_attack:all	0.60	0.6330	0.6330	0.0000	[0.0000, 0.0000]	0.000
max_attack:all	0.66	0.6330	0.5164	-0.1166	[-0.1831, -0.0394]	1.000
max_attack:all	0.70	0.6330	0.5164	-0.1166	[-0.1831, -0.0394]	1.000
max_attack:all	0.80	0.6330	0.5164	-0.1166	[-0.1831, -0.0394]	1.000
max_attack:all	0.90	0.6330	0.5164	-0.1166	[-0.1831, -0.0394]	1.000
zero_attack:all	0.40	0.6442	0.6442	0.0000	[0.0000, 0.0000]	0.000
zero_attack:all	0.50	0.6442	0.6442	0.0000	[0.0000, 0.0000]	0.000
zero_attack:all	0.60	0.6442	0.6442	0.0000	[0.0000, 0.0000]	0.000
zero_attack:all	0.66	0.6442	0.5199	-0.1243	[-0.1819, -0.0784]	1.000
zero_attack:all	0.70	0.6442	0.5199	-0.1243	[-0.1819, -0.0784]	1.000
zero_attack:all	0.80	0.6442	0.5199	-0.1243	[-0.1819, -0.0784]	1.000
zero_attack:all	0.90	0.6442	0.5199	-0.1243	[-0.1819, -0.0784]	1.000

active. Under all-domain zero and max score collapse, the same threshold activates the reliability path for every evaluated sample and yields the measured gains. Lower thresholds through $\tau = 0.60$ suppress adaptation and lose those gains; higher thresholds adapt on clean data more often and slightly hurt clean ranking.

Per-component contributions on all-domain score attacks are reported in Table XX.

The component ablation produces a useful negative result. Removing the ECE term does not hurt the all-domain collapse cases; it improves the max-attack delta from 0.0319 to 0.0457 and the zero-attack delta from 0.0506 to 0.0629. In contrast, removing the KS drift term eliminates adaptation and returns exactly to static attention. This suggests that, in the current stress benchmark, the distribution-shift signal is the necessary trigger while validation ECE is not the driver of the robustness gain.

E. Cross-Benchmark Contrast

The same mechanism-isolation protocol applied to the primary co-observed MVTec 3D-AD benchmark (§V-E) gives the opposite verdict. On MVTec 3D-AD the KS-drift term is again the only component that controls whether the gate fires — but every firing condition produces a *negative* ROC-AUC delta of about -0.10 to -0.13 relative to static attention, with 95% confidence intervals that exclude zero. The asymmetry is sharp: the KS-drift signal that delivers the entire $+0.0506$ *zero_attack* gain on the label-aligned secondary benchmark misfires on legitimate inter-category variation in RGB/3D paired data, places the mean reliability below τ on most batches, and the resulting adapted weighting consistently underperforms the static attention solution. The cleanest reading of the two results together is that a validation-derived drift detector cannot tell adversarial coherent score collapse apart from natural cross-category heterogeneity on this scale of paired data; the mechanism therefore needs a benchmark-

TABLE XIII
MVTEC 3D-AD RELIABILITY-COMPONENT ABLATION ON ALL-DOMAIN SCORE ATTACKS.

Variant	Attack	Static ROC	RGA ROC	Delta	Delta 95% CI	Adapt rate
full	gaussian_noise	0.6495	0.5188	-0.1307	[-0.1870, -0.0825]	1.000
full	max_attack	0.6330	0.5164	-0.1166	[-0.1831, -0.0394]	1.000
full	zero_attack	0.6442	0.5199	-0.1243	[-0.1819, -0.0784]	1.000
no_ece	gaussian_noise	0.6495	0.5261	-0.1234	[-0.1877, -0.0809]	1.000
no_ece	max_attack	0.6330	0.5315	-0.1015	[-0.1612, -0.0402]	1.000
no_ece	zero_attack	0.6442	0.5299	-0.1143	[-0.1612, -0.0704]	1.000
no_gate	gaussian_noise	0.6495	0.5188	-0.1307	[-0.1870, -0.0825]	1.000
no_gate	max_attack	0.6330	0.5164	-0.1166	[-0.1831, -0.0394]	1.000
no_gate	zero_attack	0.6442	0.5199	-0.1243	[-0.1819, -0.0784]	1.000
no_ks	gaussian_noise	0.6495	0.6495	0.0000	[0.0000, 0.0000]	0.000
no_ks	max_attack	0.6330	0.6330	0.0000	[0.0000, 0.0000]	0.000
no_ks	zero_attack	0.6442	0.6442	0.0000	[0.0000, 0.0000]	0.000
no_sharpness	gaussian_noise	0.6495	0.5314	-0.1181	[-0.1633, -0.0796]	0.921
no_sharpness	max_attack	0.6330	0.5180	-0.1150	[-0.1803, -0.0383]	1.000
no_sharpness	zero_attack	0.6442	0.5188	-0.1253	[-0.1822, -0.0775]	1.000

TABLE XIV
MVTEC 3D-AD CALIBRATION AND COUNTERFACTUAL
DOMAIN-ATTRIBUTION SUMMARY.

Metric	Static	RGA
ECE	0.0354	0.0402
Brier	0.1648	0.1698
CDA samples	–	400
CDA/ECE Spearman	–	–
CDA/ECE Spearman status	–	n/a (fewer than 3 domains)

TABLE XV
MVTEC 3D-AD REPRESENTATIVE FAILURE AND CONTRAST CASES.

Case	Label	Static p	RGA p	Sample
biggest_elara_win	0	0.5951	0.0520	mvtec3d_foam_train_good_150
biggest_elara_loss	1	0.7030	0.0329	mvtec3d_cable_gland_test_cut_0
elara_correct_static_wrong	0	0.5601	0.0505	mvtec3d_foam_train_good_151
static_correct_elara_wrong	1	0.6612	0.0319	mvtec3d_cable_gland_test_hole_
high_confidence_elara_failure	1	0.5319	0.0338	mvtec3d_cable_gland_test_threa

conditioned drift threshold or a different drift signal before it can be deployed on natural paired evidence streams.

VII. CALIBRATION AND EXPLANATION

A. Calibration

Table XXI reports calibration for the final configured seed. Static attention is slightly better than RGA on ECE and Brier score in that seed. Both attention variants remain much better calibrated than confidence-weighted mean fusion in the five-seed clean table. The reliability diagram is shown in Fig. 12.

B. Counterfactual Domain Attribution

Counterfactual domain attribution was computed for 400 test samples across configured seeds. Mean absolute probability changes were 0.0138 for fraud, 0.0172 for cyber, 0.0059 for behavior, and 0.0353 for text. Text has the largest average

TABLE XVI
MISSING-DOMAIN ABLATION. DELTA IS RGA MINUS STATIC ATTENTION.

Dropout	Static ROC	RGA ROC	Delta	Delta 95% CI
0.0	0.9999	0.9999	0.0000	[0.0000, 0.0000]
0.1	0.9995	0.9995	0.0000	[0.0000, 0.0000]
0.2	0.9981	0.9981	-0.0000	[-0.0000, 0.0000]
0.3	0.9964	0.9964	-0.0001	[-0.0002, 0.0000]
0.5	0.9784	0.9783	-0.0001	[-0.0004, 0.0000]




Fig. 9. ROC-AUC under additional missing-domain dropout.

counterfactual impact, consistent with the strong text scorer in Table V; the per-domain magnitudes are visualized in Fig. 13.




Fig. 10. Per-domain score-drift curves on the real-domain benchmark.

TABLE XVII
DRIFT ROBUSTNESS SUMMARY. HIGHER AREA IS BETTER.

Domain	Static	RGA	Delta	Seeds
behavior	0.3000 [0.3000, 0.3000]	0.3000 [0.3000, 0.3000]	0.0000	5
cyber	0.3000 [0.3000, 0.3000]	0.3000 [0.3000, 0.3000]	0.0000	5
fraud	0.3000 [0.3000, 0.3000]	0.3000 [0.3000, 0.3000]	-0.0000	5
nlp	0.2999 [0.2999, 0.2999]	0.2999 [0.2999, 0.2999]	-0.0000	5




Fig. 11. All-domain adversarial ROC-AUC delta between RGA and static attention. Complete single-domain scenarios are reported in Table XVIII.

VIII. DISCUSSION AND LIMITATIONS

A. Scope of the Empirical Claim

The evidence supports three scoped claims and rules out one broad one. On ELARA-Bench-LA the reliability gate improves all-domain coherent score-collapse robustness (+0.0506 zero_attack, +0.0319 max_attack) while preserving the saturated clean ranking. On naturally paired MVTEC 3D-AD the same gate is strictly dominated: random forest fusion leads (0.959), static attention sits mid-pack (0.650), and RGA regresses to 0.610, with every adversarial delta also negative. The mechanism-isolation results together




Fig. 12. Reliability diagram for static attention and RGA.

identify a specific failure mode of validation-derived drift detectors — a KS-drift signal calibrated on validation can deliver gains under coherent adversarial collapse and simultaneously misfire on legitimate inter-category heterogeneity — which is the central claim of the paper. What the paper does not support is a broad claim that RGA improves multimodal fusion in general: the label-aligned gain is bounded to coherent all-domain attacks on a near-saturated split, and no attention variant matches random forest on naturally paired MVTEC 3D-AD at the current scorer level. Both directions are bounded by benchmark design and the lightweight image-statistic features used for MVTEC 3D-AD domain scoring.

B. Failure Analysis

The result artifact exports representative cases for the largest RGA win, largest RGA loss, cases where RGA fixes static

TABLE XVIII
ADVERSARIAL PERTURBATION BENCHMARK. DELTA IS RGA MINUS STATIC ATTENTION.

Attack	Target	Static ROC	RGA ROC	Delta	Delta 95% CI
gaussian_noise	all	0.9999	0.9998	-0.0001	[-0.0001, 0.0000]
gaussian_noise	behavior	0.9999	0.9999	-0.0000	[-0.0000, 0.0000]
gaussian_noise	cyber	0.9999	0.9999	0.0000	[0.0000, 0.0000]
gaussian_noise	fraud	0.9999	0.9999	0.0000	[0.0000, 0.0000]
gaussian_noise	nlp	0.9999	0.9999	-0.0000	[-0.0000, 0.0000]
max_attack	all	0.7401	0.7720	0.0319	[0.0050, 0.0617]
max_attack	behavior	0.9997	0.9997	-0.0000	[-0.0000, -0.0000]
max_attack	cyber	0.9972	0.9972	0.0000	[0.0000, 0.0000]
max_attack	fraud	0.9996	0.9996	0.0000	[0.0000, 0.0000]
max_attack	nlp	0.9840	0.9843	0.0003	[-0.0009, 0.0016]
zero_attack	all	0.7212	0.7718	0.0506	[0.0315, 0.0681]
zero_attack	behavior	0.9997	0.9997	0.0000	[-0.0000, 0.0000]
zero_attack	cyber	0.9980	0.9980	0.0000	[0.0000, 0.0000]
zero_attack	fraud	0.9995	0.9995	0.0000	[0.0000, 0.0000]
zero_attack	nlp	0.9713	0.9716	0.0003	[-0.0004, 0.0009]



Fig. 13. Mean absolute counterfactual domain-attribution impact.

attention, cases where static attention fixes RGA, and high-confidence RGA failures. Each case includes the sample identifier, label, static and RGA probabilities, per-domain scores, missing-domain mask, and reliability weights.

Representative cases are listed in Table XXII.

Across both benchmarks, the experiments expose four failure modes:

- the synthetically paired construction makes the secondary clean split highly separable and obscures method differences,
- even the 5% scorer-training hard-mode rerun weakens

individual domain scorers without fully removing the secondary benchmark saturation,

- on co-observed MVTec 3D-AD the reliability gate misfires on legitimate cross-category variation and degrades both clean and adversarial ROC-AUC,
- random forest fusion remains the strongest baseline on the paired benchmark by a clear margin,
- the lightweight image-statistic domain scorer caps the achievable ceiling on MVTec 3D-AD and is the natural next replacement target.

C. Internal Validity

The benchmark scripts are deterministic for configured seeds, and clean results are reported across five seeds. The runner repeats robustness ablations across configured seeds and stores confidence intervals. Residual risk remains if the stress distributions are too synthetic relative to deployment failures.

D. Construct Validity

ELARA-Bench-LA measures real-domain score fusion, not RGB/3D paired multimodal incident detection. The MVTec 3D-AD path addresses natural RGB/3D pairing, but the current local run is a small bagel-subset smoke benchmark with lightweight features. This remains the largest construct-validity threat. The paper therefore avoids claiming that the model has learned causal or temporal relationships among fraud, cyber, behavior, and text domains, or that the paired path is competitive with specialized RGB-3D anomaly detectors.

TABLE XIX
RELIABILITY-GATE THRESHOLD SWEEP. DELTA IS RGA MINUS STATIC ATTENTION; ADAPTATION RATE IS THE FRACTION OF EVALUATED SAMPLES USING THE RELIABILITY-AWARE PATH.

Condition	τ	Static ROC	RGA ROC	Delta	Delta 95% CI	Adapt rate
clean	0.40	0.9999	0.9999	0.0000	[0.0000, 0.0000]	0.000
clean	0.50	0.9999	0.9999	0.0000	[0.0000, 0.0000]	0.000
clean	0.60	0.9999	0.9999	0.0000	[0.0000, 0.0000]	0.000
clean	0.66	0.9999	0.9999	0.0000	[0.0000, 0.0000]	0.000
clean	0.70	0.9999	0.9999	-0.0000	[-0.0000, 0.0000]	0.320
clean	0.80	0.9999	0.9999	-0.0000	[-0.0001, 0.0000]	0.960
clean	0.90	0.9999	0.9999	-0.0000	[-0.0001, 0.0000]	1.000
gaussian_noise:all	0.40	0.9999	0.9999	0.0000	[0.0000, 0.0000]	0.000
gaussian_noise:all	0.50	0.9999	0.9999	0.0000	[0.0000, 0.0000]	0.000
gaussian_noise:all	0.60	0.9999	0.9998	-0.0000	[-0.0001, -0.0000]	0.680
gaussian_noise:all	0.66	0.9999	0.9998	-0.0001	[-0.0002, -0.0000]	1.000
gaussian_noise:all	0.70	0.9999	0.9998	-0.0001	[-0.0002, -0.0000]	1.000
gaussian_noise:all	0.80	0.9999	0.9998	-0.0001	[-0.0002, -0.0000]	1.000
gaussian_noise:all	0.90	0.9999	0.9998	-0.0001	[-0.0002, -0.0000]	1.000
max_attack:all	0.40	0.7401	0.7401	0.0000	[0.0000, 0.0000]	0.000
max_attack:all	0.50	0.7401	0.7401	0.0000	[0.0000, 0.0000]	0.000
max_attack:all	0.60	0.7401	0.7401	0.0000	[0.0000, 0.0000]	0.000
max_attack:all	0.66	0.7401	0.7720	0.0319	[0.0050, 0.0617]	1.000
max_attack:all	0.70	0.7401	0.7720	0.0319	[0.0050, 0.0617]	1.000
max_attack:all	0.80	0.7401	0.7720	0.0319	[0.0050, 0.0617]	1.000
max_attack:all	0.90	0.7401	0.7720	0.0319	[0.0050, 0.0617]	1.000
zero_attack:all	0.40	0.7212	0.7212	0.0000	[0.0000, 0.0000]	0.000
zero_attack:all	0.50	0.7212	0.7212	0.0000	[0.0000, 0.0000]	0.000
zero_attack:all	0.60	0.7212	0.7212	0.0000	[0.0000, 0.0000]	0.000
zero_attack:all	0.66	0.7212	0.7718	0.0506	[0.0315, 0.0681]	1.000
zero_attack:all	0.70	0.7212	0.7718	0.0506	[0.0315, 0.0681]	1.000
zero_attack:all	0.80	0.7212	0.7718	0.0506	[0.0315, 0.0681]	1.000
zero_attack:all	0.90	0.7212	0.7718	0.0506	[0.0315, 0.0681]	1.000

E. External Validity

The datasets are useful but not sufficient for deployment claims. A production system would need domain-specific label definitions, temporal splits, deployment-time calibration monitoring, privacy review, and dataset license review.

F. Statistical Validity

Five clean seeds provide only limited statistical support, especially on a near-saturated split. This draft therefore reports seed-level confidence intervals for clean PR-AUC and F1 and treats stress-test deltas as the primary diagnostic. DeLong tests are still used where assumptions hold [14], but the paper avoids treating a single significance test as proof of broad superiority.

G. Anticipated Reviewer Questions

We anticipate five lines of reviewer pushback and address each here explicitly:

(Q1) Why does the gate hurt on naturally paired data? The mechanism-isolation results in §V-E are unambiguous: removing the KS-drift component returns RGA to static-attention performance, and every condition that fires the gate produces a negative ROC-AUC delta with confidence intervals that exclude zero. The KS-drift signal misfires on legitimate inter-category variation in MVTec 3D-AD’s eight object categories and pushes the model into a degraded adapted mode.

(Q2) Why is random forest fusion so much stronger than attention on MVTec 3D-AD? The MVTec domain scorers in this study are deliberately lightweight (normal-reference image statistics with eight-dim embeddings), which caps the achievable ceiling for any fusion model that consumes them. Random forest exploits the residual non-linear structure better than a 32-dimensional attention head at this scorer level. The paper does not claim that attention is generally worse than RF for paired RGB+3D fusion; it reports the contrast at the current feature quality. Replacing the scorer with stronger RGB–3D

TABLE XX
RELIABILITY-COMPONENT ABLATION ON ALL-DOMAIN SCORE ATTACKS. DELTA IS RGA MINUS STATIC ATTENTION.

Variant	Attack	Static ROC	RGA ROC	Delta	Delta 95% CI	Adapt rate
full	gaussian_noise	0.9999	0.9997	-0.0002	[-0.0005, -0.0000]	1.000
full	max_attack	0.7401	0.7720	0.0319	[0.0050, 0.0617]	1.000
full	zero_attack	0.7212	0.7718	0.0506	[0.0315, 0.0681]	1.000
no_ece	gaussian_noise	0.9999	0.9996	-0.0003	[-0.0005, -0.0001]	1.000
no_ece	max_attack	0.7401	0.7858	0.0457	[0.0129, 0.1067]	1.000
no_ece	zero_attack	0.7212	0.7841	0.0629	[0.0493, 0.0808]	1.000
no_gate	gaussian_noise	0.9999	0.9997	-0.0002	[-0.0005, -0.0000]	1.000
no_gate	max_attack	0.7401	0.7720	0.0319	[0.0050, 0.0617]	1.000
no_gate	zero_attack	0.7212	0.7718	0.0506	[0.0315, 0.0681]	1.000
no_ks	gaussian_noise	0.9999	0.9999	0.0000	[0.0000, 0.0000]	0.000
no_ks	max_attack	0.7401	0.7401	0.0000	[0.0000, 0.0000]	0.000
no_ks	zero_attack	0.7212	0.7212	0.0000	[0.0000, 0.0000]	0.000
no_sharpness	gaussian_noise	0.9999	0.9997	-0.0002	[-0.0005, -0.0000]	1.000
no_sharpness	max_attack	0.7401	0.7808	0.0408	[0.0123, 0.0811]	1.000
no_sharpness	zero_attack	0.7212	0.7759	0.0547	[0.0382, 0.0716]	1.000

TABLE XXI
CALIBRATION AND COUNTERFACTUAL DOMAIN-ATTRIBUTION SUMMARY.

Metric	Static	RGA
ECE	0.0038	0.0038
Brier	0.0032	0.0032
CDA samples	–	400
CDA/ECE Spearman	–	0.052
CDA/ECE Spearman status	–	computed

TABLE XXII
REPRESENTATIVE FAILURE AND CONTRAST CASES FROM THE LATEST CONFIGURED SEED.

Case	Label	Static p	RGA p	Sample
biggest_elara_win	0	0.0001	0.0001	real_fusion_test_000000
biggest_elara_loss	1	0.9998	0.9998	real_fusion_test_001072
high_confidence_elara_failure	1	0.0470	0.0470	real_fusion_test_001517

features [16], [17], [18] is the natural next step.

(Q3) Why $\tau = 0.66$ and not a learned gate? The threshold sweep in §V-E and the matching sweep on the secondary benchmark show that the conservative threshold preserves clean accuracy on the label-aligned data while activating under coherent stress. A learned-gate prototype exists in the codebase (src/uais/fusion/attention/learned_gate.py) and is the planned next experimental step. On the naturally paired benchmark, however, a learned gate would still have to overcome the same KS-drift misfire that the component ablation identifies.

(Q4) Why is the all-domain attack model meaningful? It is not, in the strong realistic sense, and the threat model in §IV-G says so explicitly. Single-domain perturbations are the realistic case and are reported separately. The all-domain coherent attack is a worst-case stress test that bounds the mechanism’s contribution; it is not an expected attacker capability.

(Q5) Why no top-tier conference submission for this result?

The cross-benchmark contrast is real but the negative-result direction needs at least one additional co-observed dataset before the misfire pattern can be claimed as a property of validation-derived drift signals rather than a property of MVtec 3D-AD specifically. The submission target for this manuscript is therefore an arXiv preprint plus a workshop submission (e.g. a NeurIPS workshop on distribution shifts or reliable ML); extension toward a mid-tier conference paper is described in the accompanying publication roadmap.

IX. REPRODUCIBILITY

The reproducibility flow has three scripted steps:

- 1) build the real-domain fusion input table,
- 2) run the real fusion benchmark with the YAML configuration,
- 3) generate tables, figures, and the PDF manuscript.

The implementation uses the local scripts under src/scripts/, the configurations configs/attention_real_fusion.yaml and configs/attention_mvtec3d_fusion.yaml, result artifacts under experiments/fusion/, and generated paper assets under docs/research/. The paired benchmark entry point is src/scripts/prepare_mvtec3d_fusion_benchmark.py. Targeted regression tests cover attention masking, statistics, baselines, reliability estimation, counterfactual explanation, paired-benchmark metadata, and result aggregation.

X. RELATIONSHIP TO COMPANION THESIS CHAPTER

A longer, thesis-style treatment of this work is maintained alongside the manuscript as docs/research/THESIS_CHAPTER_v1.tex, compiled to output/pdf/THESIS_CHAPTER_v1.pdf.

The conference manuscript and the thesis chapter share the same code base, configurations, and result artifacts; what differs is the audience and the level of reflection. The thesis chapter is longer, contains a full Background section, expanded construct-validity and statistical-validity discussion, an explicit “what did not work” subsection, and explicit framing of failure cases as evidence rather than appendix material. The conference manuscript prioritizes the contrastive headline result – gate-helps-on-label-aligned versus gate-hurts-on-natural-paired – and uses the thesis chapter as the place to develop those threads in more detail.

The two documents are kept consistent through a single asset pipeline. The benchmark JSONs (`experiments/fusion/craf_real_results.json` and `experiments/fusion/mvtec3d_results.json`) generate the same `rga_*.tex` and `mvtec3d_*.tex` tables and `*.png` figures, which both manuscripts input or includegraphics. Changes to the experimental evidence therefore propagate to both documents without divergence. Readers who want the systems-paper framing should stay with this manuscript; readers who want a thesis-style discussion (including the extended background, threats-to-validity, and forward-looking research agenda) should read the companion chapter.

XI. CONCLUSION

ELARA presents a reliability-aware anomaly-fusion pipeline with domain experts, unified score schema, masked attention, reliability-gated adaptation, counterfactual explanation, benchmark construction, baselines, ablations, calibration, and failure analysis. The strongest measured claim is that RGA preserves clean attention performance, improves over confidence-weighted mean fusion, and helps under severe all-domain score collapse. The main remaining research gap is full RGB/3D paired multimodal validation with stronger RGB–3D domain experts.

REFERENCES

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems*, 2017.
- [2] T. Baltrušaitis, C. Ahuja, and L.-P. Morency, “Multimodal machine learning: A survey and taxonomy,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 41, no. 2, pp. 423–443, 2019.
- [3] G. Pang, C. Shen, L. Cao, and A. van den Hengel, “Deep learning for anomaly detection: A review,” *ACM Computing Surveys*, vol. 54, no. 2, 2021.
- [4] B. Zadrozny and C. Elkan, “Transforming classifier scores into accurate multiclass probability estimates,” in *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2002.
- [5] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proceedings of the International Conference on Machine Learning*, 2017.
- [6] J. Platt, “Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods,” in *Advances in Large Margin Classifiers*, 1999.
- [7] Y. Ovadia et al., “Can you trust your model’s uncertainty? Evaluating predictive uncertainty under dataset shift,” in *Advances in Neural Information Processing Systems*, 2019.

- [8] Y. Sun, X. Wang, Z. Liu, J. Miller, A. Efros, and M. Hardt, “Test-time training with self-supervision for generalization under distribution shifts,” in *Proceedings of the International Conference on Machine Learning*, 2020.
- [9] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, and T. Darrell, “Tent: Fully test-time adaptation by entropy minimization,” in *International Conference on Learning Representations*, 2021.
- [10] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, pp. 5–32, 2001.
- [11] J. H. Friedman, “Greedy function approximation: A gradient boosting machine,” *The Annals of Statistics*, vol. 29, no. 5, pp. 1189–1232, 2001.
- [12] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems*, 2017.
- [13] J. Davis and M. Goadrich, “The relationship between precision-recall and ROC curves,” in *Proceedings of the International Conference on Machine Learning*, 2006.
- [14] E. R. DeLong, D. M. DeLong, and D. L. Clarke-Pearson, “Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach,” *Biometrics*, vol. 44, no. 3, pp. 837–845, 1988.
- [15] P. Bergmann, X. Jin, D. Sattlegger, and C. Steger, “The MVTEC 3D-AD Dataset for Unsupervised 3D Anomaly Detection and Localization,” in *Proceedings of VISAPP*, 2022.
- [16] Y. Wang, J. Peng, J. Zhang, R. Yi, Y. Wang, and C. Wang, “Multimodal Industrial Anomaly Detection via Hybrid Fusion,” in *Proceedings of CVPR*, 2023.
- [17] A. Costanzino, P. Zama Ramirez, G. Lisanti, and L. Di Stefano, “Multimodal Industrial Anomaly Detection by Crossmodal Feature Mapping,” *arXiv preprint arXiv:2312.04521*, 2023.
- [18] J. Liu, G. Xie, R. Chen, X. Li, J. Wang, Y. Liu, C. Wang, and F. Zheng, “Real3D-AD: A Dataset of Point Cloud Anomaly Detection,” *arXiv preprint arXiv:2309.13226*, 2023.
- [19] C. Wang, H. Zhu, J. Peng, Y. Wang, R. Yi, Y. Wu, L. Ma, and J. Zhang, “M3DM-NR: RGB-3D Noisy-Resistant Industrial Anomaly Detection via Multimodal Denoising,” *arXiv preprint arXiv:2406.02263*, 2024.
- [20] C. C. Aggarwal and S. Sathe, *Outlier Ensembles: An Introduction*. Springer, 2017.
- [21] J. Kittler, M. Hatef, R. P. W. Duin, and J. Matas, “On combining classifiers,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 20, no. 3, pp. 226–239, 1998.
- [22] S. Wachter, B. Mittelstadt, and C. Russell, “Counterfactual explanations without opening the black box: Automated decisions and the GDPR,” *Harvard Journal of Law & Technology*, vol. 31, no. 2, 2017.
- [23] Y. Geifman and R. El-Yaniv, “Selective classification for deep neural networks,” in *Advances in Neural Information Processing Systems*, 2017.